



BITS ZG553: Real Time Systems

BITS Pilani
Pilani|Doha|Goa|Hyderabad

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

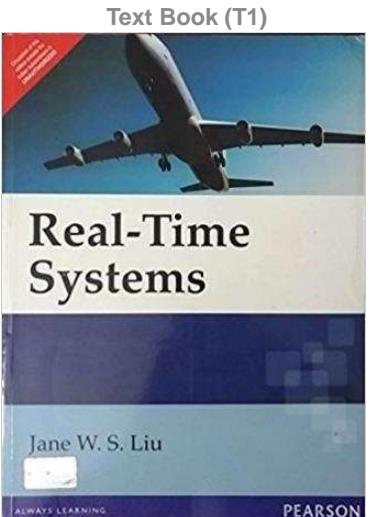
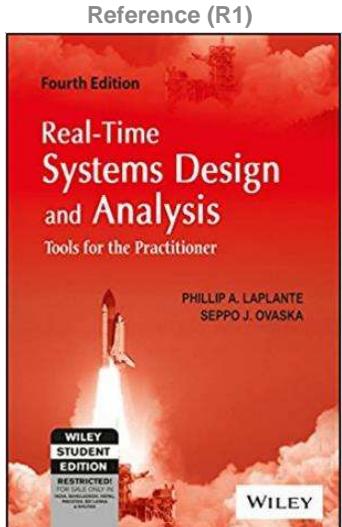


Welcome To BITS ZG553: Real Time Systems -KG Krishna

Manufacturer	Raytheon
Introduced	August 1966; 52 years ago
Discontinued	July 1975; 44 years ago
Type	Avionics Guidance Computer
Processor	Discrete IC RTL based
Frequency	2.048 MHz
Memory	15-bit wordlength + 1-bit parity, 2048 words RAM (magnetic-core memory) ^[1] , 36,864 words ROM (core rope memory) ^[1]
Ports	DSKY, IMU, Hand Controller, Rendezvous Radar (CM), Landing Radar (LM), Telemetry Receiver, Engine Command, Reaction Control System
Power consumption	55 W ^{[2]:120}
Weight	70 lb (32 kg)
Dimensions	24x12.5x6.5 inches(61x32x17 cm)



Text Book / References



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

Excellent MOOCs Videos (Coursera, edX,...)



HONOR CODE CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

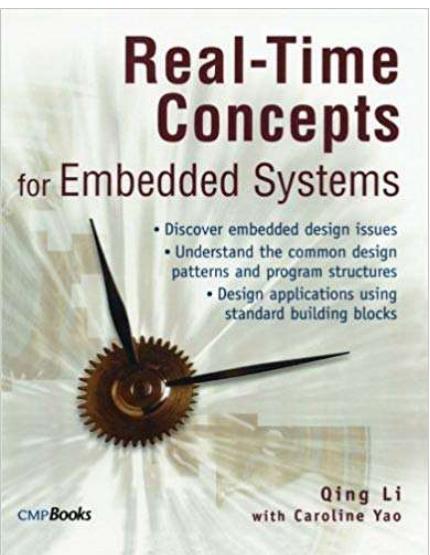


Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eefb47d4e54

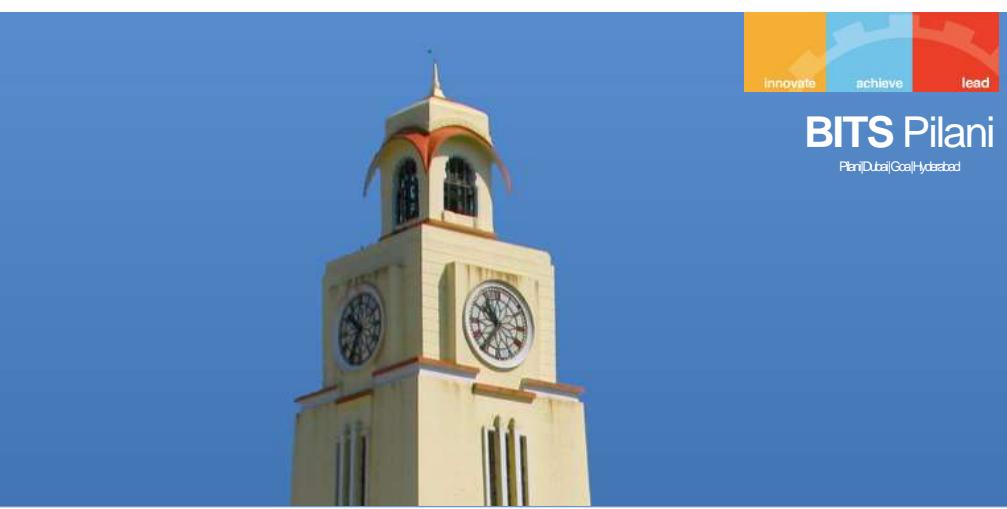
RTS Primer – For Light Reading



5 BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Real Time Computing

- Correctness = Logical Result of Computation + Time at which result is produced
- Usually part of a larger cyber-physical system
- Response Time – Time between presentation of inputs and the realization of the required behavior.
- How Punctual ?? / How Fast ?? – System Specifications



L-1: Real Time Systems – An Introduction

Ref: R1/ [Informal Discussion]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;
PLEASE DO NOT PRINT PPTs, Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations by Prof K R Anupa, BITS-Pilani, Goa

6

RTS

- System satisfies bounded **Response Time** constraint
- Severe Consequence
- **Is this always failure ???**
- Safe/ Reliable
- RTS are mostly embedded/ reactive
- Concurrency
- Distributed
- **Every Practical System is real time**

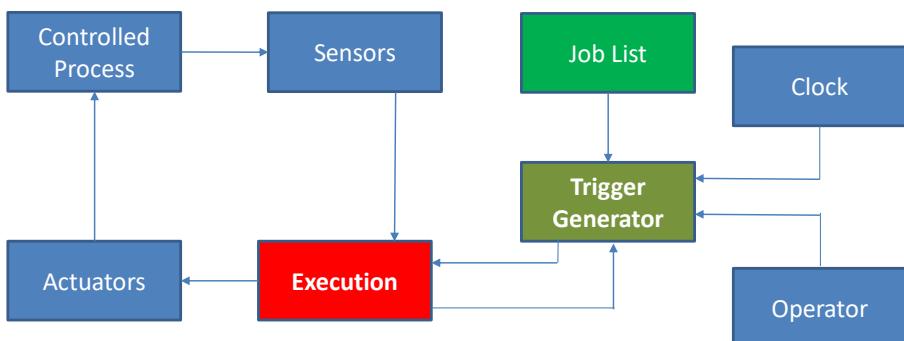
Hard Vs Soft

- Soft RTS - Performance is **degraded but not destroyed** by failure to meet deadlines (response time requirements)
 - Hard RTS – Even a **single deadline missed** will be **catastrophic or complete failure**
- ? Missile Launching System
? Video Games
? Weed killing robot
- Firm RTS – **> N deadlines missed** – complete failure.

Where deadlines come from?

- Animated Display – 30 frames/sec
- System Specification
 - Specification Languages
 - Predict execution times of programs
 - Model reliability of s/w and h/w
 - Assign tasks to processors
 - Failure Recovery

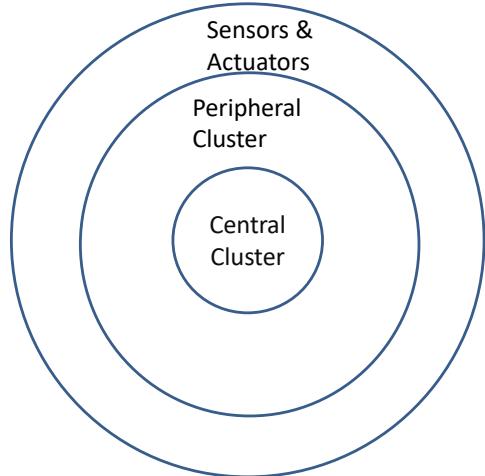
Structure of a RTS



Tasks

- Periodic/ Non Periodic
- Critical / Non Critical

In terms of Speed



Where do deadlines come from:

- Elevator Door Control

	Contributing parameters	Min	Avg	Max
1	Sensor Response Time	5ms	9ms	15ms
2	Hardware Response Time	1µs	1.2µs	2µs
3	System Software Response Time	16µs	37µs	48µs
4	Appln Software Response Time	0.5µs	0.5µs	0.5µs
5	Door Actuator Response Time	300ms	400	500

Real – Time Punctuality

- Every response System has an average t_R
- Upper Bound $t_R + \varepsilon_U$
- Lower Bound $t_R - \varepsilon_L$
- $\varepsilon_U \varepsilon_L \rightarrow 0^+$
- Practical Systems – non zero
 - Cumulative latency/ propagation delay in software & hardware

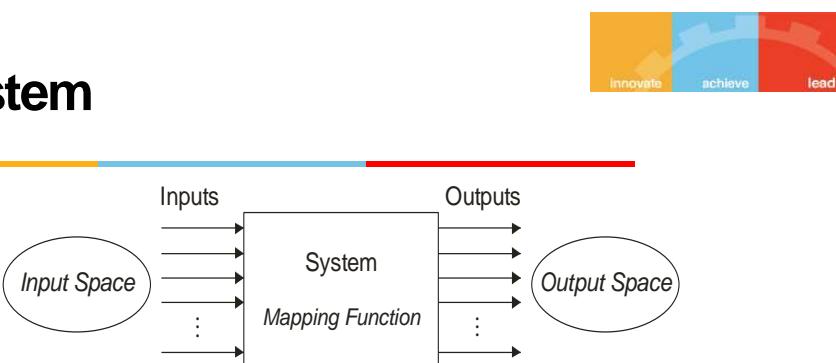
Issues in RTS

- Architectural Issues
 - Processor Architecture
 - Network Architecture
 - Clock Sync
 - Fault Tolerance and Reliability
- OS Issues
 - Task Assignment & Scheduling
 - Commn Protocols
 - Failure Management
 - Clock Sync
- Other Issues
 - Programming Languages
 - Databases
 - Performance Measures

Hard Vs Soft

- Soft RTS - Performance is **degraded but not destroyed** by failure to meet deadlines (response time requirements)
 - Hard RTS – Even a **single deadline missed** will be **catastrophic or complete failure**
- ? Missile Launching System
? Video Games
? Weed killing robot
- Firm RTS – > N deadlines missed – complete failure.

A System



(Source: P. A. Laplante & S. J. Ovaska, *Real-Time Systems Design and Analysis: Tools for the Practitioner*, Wiley, 4th edition)

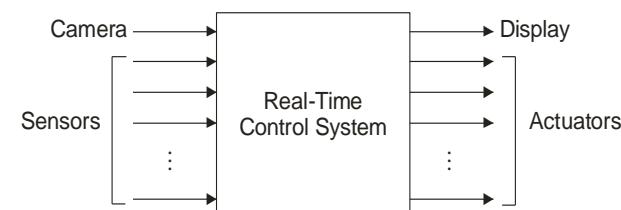
- A system is an assembly of components connected together in an organized way
- A system is fundamentally altered if a component joins or leaves it
- It has a purpose
- It has a degree of permanence
- It has been defined as being of particular interest

Where deadlines come from?

- Animated Display – 30 frames/sec
- System Specification
 - Specification Languages
 - Predict execution times of programs
 - Model reliability of s/w and h/w
 - Assign tasks to processors
 - Failure Recovery

Example: A Real-Time Control System

- Inputs are *excitations* and outputs are corresponding *responses*
- Inputs and outputs may be digital or analog
- Inputs are associated with sensors, cameras, etc.
- Outputs with actuators, displays, etc.



(Source: P. A. Laplante & S. J. Ovaska, *Real-Time Systems Design and Analysis: Tools for the Practitioner*, Wiley, 4th edition)

Sample Data System

Feedback Control Loop pseudocode:

```

set timer to interrupt periodically with period 'T'
at each timer interrupt do
{
    do analogue-to-digital conversion to get 'y'
    compute control output 'u'
    do digital-to-analogue conversion
}

```

Sample Data System

Determining the Sampling Rate

Sampling time interval = T

In general, the faster a system can respond to changes, the faster the input to the actuator varies, hence the shorter the sampling period should be.

We can measure the responsiveness of the system by its *rise time R*.

R is the time it takes to get close to its final state after the input changes.

Typically the ratio R/T of rise time to sampling period should be between 10 and 20 in order to give a smooth response.

One can also consider this in terms of bandwidth ω which is approximately $1/2R$ Hz.

The Nyquist sampling theorem says that any time-continuous signal of bandwidth ω can be reproduced faithfully from its sampled values only if the sampling rate is at least 2ω .

Recommended sampling rate is **20 to 40 times the bandwidth ω** .

21

BITS Pilani, Pilani Campus

Sample Data System

Multirate System

- ❑ A system typically has state defined by multiple state variables. e.g. rotation speed, temperature, fuel consumption, etc. of an engine.
- ❑ The state is monitored by multiple sensors and controlled by multiple actuators.
- ❑ Hence there are multiple sampling rates, one for each state variable.
- ❑ **A system with multiple sampling rates is called a multirate system.**

Determining Sampling Rate for a Multirate System

- ❑ Fastest sample rate is the choice ?
- ❑ Usually longer sampling period is an integer multiple of every shorter sampling period
- ❑ **Design is done in iterative way**
 - Select the fastest sampling rate controller assuming it is independent of other controller. Integrate it with the plant
 - Select the next fastest one and integrate. Continue till the slowest sampling rate controller is integrated
 - Correct the error, aroused due to the assumption that the first one is independent of others

22

BITS Pilani, Pilani Campus

Sample Data System

Determining the Sampling Rate

Sampling time interval = T

In general, the faster a system can respond to changes, the faster the input to the actuator varies, hence the shorter the sampling period should be.

We can measure the responsiveness of the system by its *rise time R*.

R is the time it takes to get close to its final state after the input changes.

Typically the ratio R/T of rise time to sampling period should be between 10 and 20 in order to give a smooth response.

One can also consider this in terms of bandwidth ω which is approximately $1/2R$ Hz.

The Nyquist sampling theorem says that any time-continuous signal of bandwidth ω can be reproduced faithfully from its sampled values only if the sampling rate is at least 2ω .

Recommended sampling rate is **20 to 40 times the bandwidth ω** .

22

BITS Pilani, Pilani Campus

A More Complex Control Law Computation

```

Set timer to interrupt periodically with period 'T'
At each timer interrupt do
{
    Sample and digitize sensor readings;
    Compute control output from measured and state-variable values;
    Convert control output to analogue form;
    Estimate and update system parameters;
    Compute and update state variables;
}

```

23

BITS Pilani, Pilani Campus

Controller Hierarchy

Controllers in a complex systems are typically organised hierarchically.

Example – Air Traffic Control System

- The air-traffic control system is at the highest level.
 - Regulates the flow of flights to the destination
 - Does by assigning each aircraft an arrival time at each metering fix (a geographical point) en route to the destination
 - Assigned arrival time to the next metering fix is the reference for the on board flight management system
- The on-board flight management system chooses the flight paths etc. and sets parameters for the lower level controllers.
 - It has to optimize the flight path based on certain constraints
 - Constraints include the ones imposed by the flight characteristics such as maximum and minimum allowed cruise speeds and decent/ascend rates.
 - Another constraint is the cost i.e. Fuel consumption. A most desirable path is the most fuel efficient path.
 - When the flight is late, it should try to bring the aircraft to the next metering fix in the shortest time.
- The flight controller at the lowest level handles cruise speed, turn radius, ascend/descend rates etc.

Real Time Application – Real Time Database

- Refers to diverse spectrum from stock price information system to track record databases
- Real-time data is perishable (in contrast with non-real time data such as pay-roll database)
- The **age of an object** measures how up-to-date the information is. It is the **length of the time since the instant of its last update**.
- The age of an object whose value is computed from other objects is equal to that of the oldest of those objects.
- **Absolute Temporal Consistency**
 - A set of data objects is said to be absolutely temporally consistent if the maximum age in the set is no greater than a certain threshold.

Examples:

- A fighter jet and the targets it tracks move at supersonic speeds, so the information on where they are must be less than 50 ms.
- An air traffic control system monitors commercial aircrafts at subsonic speed, hence the absolute temporal consistency threshold for it is much larger

➤ **Relative Temporal Consistency**

- A set of data objects is said to be relatively temporally consistent if the maximum difference in ages of the objects in the set is no greater than the relative consistency threshold used by the application.

Controller Hierarchy

Air Traffic Control System

➤ Real Time Command and Control Operation

- ATC gathers information on the state of each aircraft via one or more active radars.
- Such radars interrogates the aircrafts periodically and the aircraft responds with its state variables such as identifiers, position, altitude, heading etc.
- ATC stores the information received from the aircrafts in a database.
- This information is processed by the processors.
- At the same time the surveillance system continuously analyzes the scenario and alerts the operators whenever there is a possibility of collision.

Real Time Application – Multimedia Systems

□ **MPEG Compression / Decompression**

- Motion Estimation
- Discrete Cosine Transform / Inverse Discrete Cosine Transform
- Huffman Encoding / Decoding

□ **Real Time Characteristics**

- SD Video requires 30 frames per sec, HD video requires 60 frames per sec
- Lower frame-rates are fine for applications like video conferencing
- DCT/IDCT, Motion Estimation and Compensation and Huffman Encoding / Decoding are computation intensive.
- Computation need to be done per frame basis.
- Lip-sync between Audio and Video also need to be performed.
- Muxing / Demuxing also add to the computation.

Types of Real Time Systems

Purely cyclic

- Every task executes periodically. Even I/O operations are polled.
- Demands on resources do not vary significantly from period to period.
- Example: [Real-time monitors](#)

Mostly cyclic

- Most tasks execute periodically.
- The system can also respond to some external events asynchronously.
- Example: [Process control systems](#)

Asynchronous and somewhat predictable

- Most tasks are not periodic.
- Duration between executions of a task may vary considerably or the resource requirements may vary. Variations have bounded ranges or known statistics.
- Example: [Multimedia systems](#), [Radar signal processing](#)

Asynchronous and unpredictable

- Systems which react to asynchronous events and have tasks with high run-time complexity.
- Example: [Intelligent real-time control system](#)

39

BITS Pilani, Pilani Campus

Jobs and Tasks

Job: Each unit of work that is scheduled and executed by the system is a job.

Task: A set of related jobs which jointly provide some system function is a task.

Example:

Computation of FFT of sensor data at a particular time instance, sending a data packet are examples of job.

Whereas the function of sending data packets is an example of a task.

Soft & Hard Real Time Systems

Soft Real Time System

- A soft real-time system is one in which performance is degraded but not destroyed by failure to meet response-time constraints.
- User requires only a demonstration that the system always meet some *statistical constraints*.
- Best effort service
- Example: Multimedia system

Hard Real Time System

- A hard real-time system is one in which failure to meet a single deadline may lead to complete and catastrophic system failure
- User requires validation that the system always meet the *timing constraint*. (validation means demonstration by a provably correct and efficient procedure)
- Guaranteed service
- Example: Avionics weapons delivery system

Firm Real Time System

- A firm real-time system is one in which failure to meet a single deadline will not lead to total system failure, but missing more than few may lead to complete and catastrophic system failure
- Example: Navigation controller for an autonomous weed-killer robot
(In most of the texts, Firm Realtime Systems are not mentioned)

30

BITS Pilani, Pilani Campus

Processors, Servers, Resources

The jobs are executed by [operating systems](#).

Each job requires some resource to execute such as a CPU, a Network, a Disk etc.

In Queuing Theory, these are called [Servers](#).

In Operating Systems literature, these are called [Active Resources](#).

Here we will call them [Processors](#).

31

BITS Pilani, Pilani Campus

32

BITS Pilani, Pilani Campus

Release Time and Deadline

Release time of a job is the instance of time at which the job becomes available for execution.

Deadline of a job is the instant of time by which its execution is required to be completed.

It is also called **Absolute Deadline**.

Relative Deadline = **Absolute Deadline** – **Release time**.

Response Time

The time between the presentation of a set of inputs to a system and the realization of the required behavior, including the availability of all associated outputs, is called the **response time of the system**

In other words, **Response time** is the length of the time from the release time of the job and the time instance when it completes.

So **Relative Deadline** is the maximum allowable Response Time.

Example

A system which monitors and controls several furnaces

At the beginning, the system is initialized and starts execution (time 0 ms).
The system samples and reads each temperature sensor every 100 ms.
The system computes the control law of each furnace every 100 ms.

Each of the control law computation is a **job**, say J_k $k = 0, 1, 2, \dots$
The function of control law computation all the time is a **task**.

The first control law computation started at 20 ms.

So, release time of first job, $r_0 = 20 \text{ ms}$

Subsequent release times: $120 \text{ ms}, 220 \text{ ms}, \dots$

So release time of k th job $r_k = (20 + 100 * k) \text{ ms}$

If each job needs to be completed within 70 ms , then the absolute deadlines of the jobs: $90 \text{ ms}, 190 \text{ ms}, 290 \text{ ms}, \dots$

Relative deadline $D_k = 70 \text{ ms}$

Real-Time Punctuality

Real Time Punctuality means every response time has an average value t_R with upper and lower bounds of $(t_R + \varepsilon_U)$ and $(t_R - \varepsilon_L)$ with $\varepsilon_U, \varepsilon_L \rightarrow 0^+$.

In all practical situations, $\varepsilon_U, \varepsilon_L \neq 0$.

It is due to the propagation delays and cumulative latencies in both hardware and software components.

So the response time will have jitter within the interval $[-\varepsilon_L, +\varepsilon_U]$.

Response Time - Example

For the door reopening operation of an elevator system, individual response times of the component involved are following.

Components	Min Response Time	Max Response Time
Sensor	5 ms	15 ms
Hardware	1 μ s	2 μ s
System Software	16 μ s	48 μ s
Application Software	0.5 μ s	0.5 μ s
Door drive	300 ms	500 ms
Total	305 ms	515 ms

Please note that overall response time is dominated by the response time of the door drive.

Some Other Definitions

Hard and Soft Deadlines

- A timing constraint or deadline is **hard**, if the failure to meet it is considered to be a fatal fault.
- A late completion of a job that has a **soft deadline** is undesirable. However, few misses of soft deadline don't cause serious harm.

Tardiness

- The **tardiness** of a job measures how late it completes respective to its deadline.
- Tardiness is 0, if it completes at or before the deadline
- Tardiness is positive and is equal to the difference between its completion time and the deadline, if the job completes after the deadline.

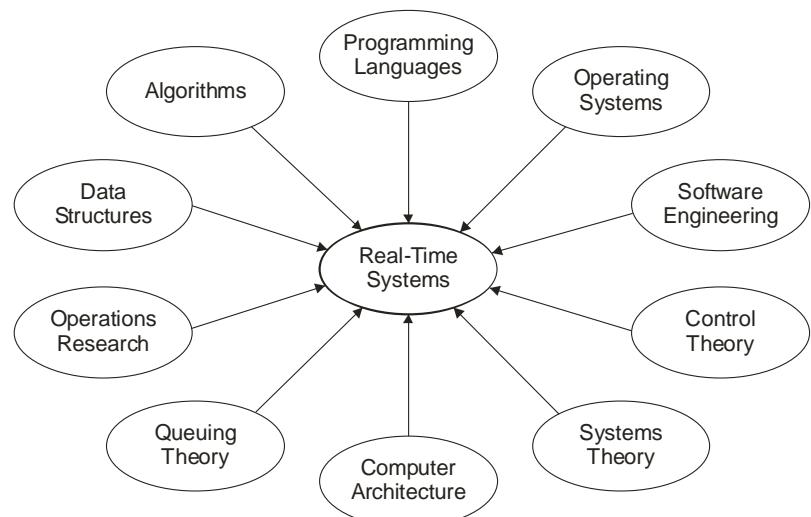
Usefulness

- **Usefulness** of a hard real-time job falls off abruptly and becomes negative when its tardiness becomes positive.
- Usefulness of a soft real-time job decreases as the tardiness increases.
- The deadline of a job is softer if its usefulness decreases at slower rate as its tardiness increases.

Real Time Systems: Usual Misconceptions

- Real-time systems are synonymous with “fast” systems
 - Many (but not all) hard real-time systems deal with deadlines in the tens of milliseconds
- There are universal, widely accepted methodologies for real-time systems specification and design
 - There is still no methodology available that answers all of the challenges of real-time specification and design all the time and for all applications
- There is no more a need to build a real-time operating system, because many commercial products exist
 - Commercial solutions have certainly their place, but choosing when to use an off-the-shelf solution and choosing the right one are continuing challenges
- Rate-monotonic analysis has solved “the real-time problem”
 - While rate-monotonic systems provide guidance in the design of real-time systems, and while there is theory surrounding them, they are not a panacea
- The study of real-time systems is mostly about scheduling theory
 - While it is scholarly to study scheduling theory, most published results require impractical simplifications and clairvoyance in order to make the theory work

Supporting Disciplines for “Real-Time Systems”



Thank You.

Any Questions?



BITS ZG553: Real Time Systems

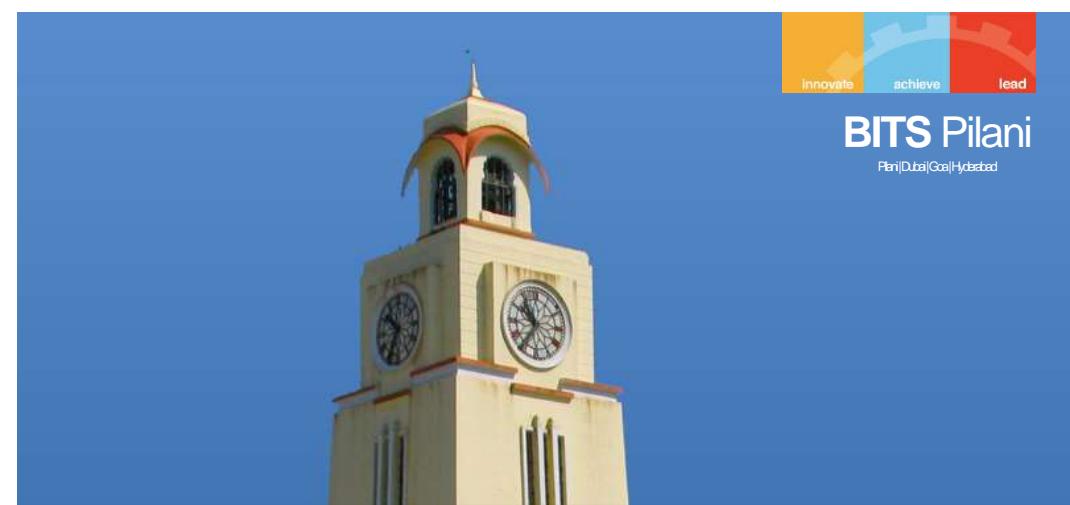
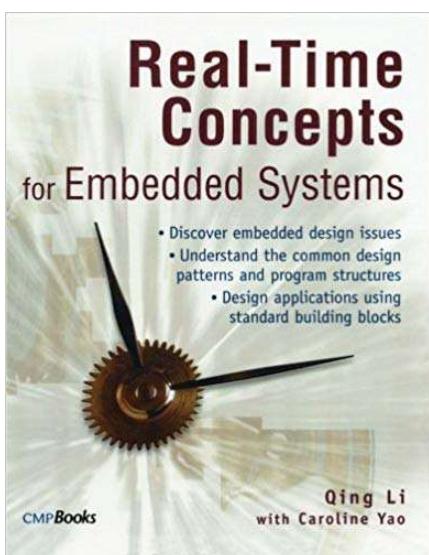
BITS Pilani
PanjDoba|Goa|Hyderabad

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

41

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

RTS Primer – For Light Reading



L-1a: Real Time Systems -
Overview/Review of OS/RTS Concepts

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Performance Criteria for RTS

RT

Timeliness
Simultaneity
Predictability
Dependability

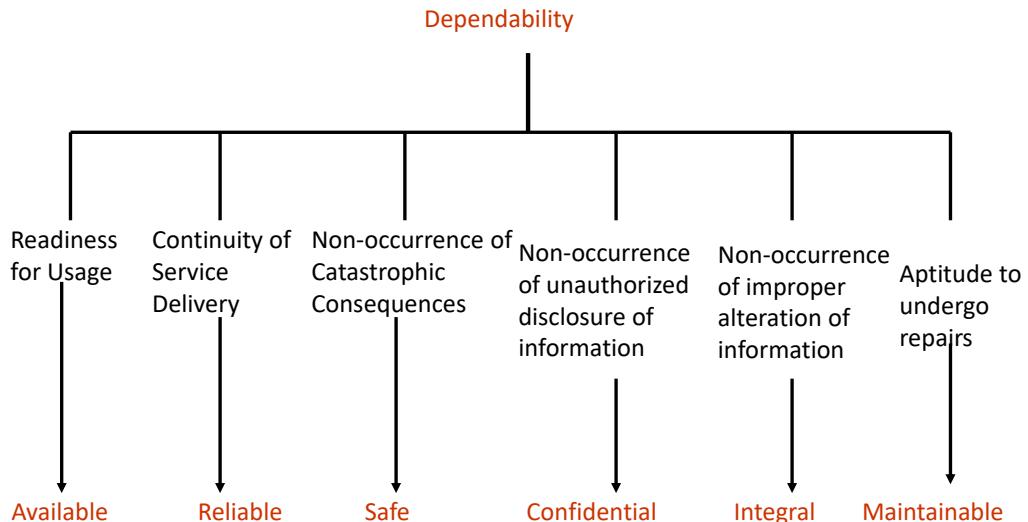
NRT

Throughput
CPU Utilization Factor

Prediction of execution

Worst Case Execution Time

Aspects of Dependability



Factors that affect execution time

Source Code

Compiler

Hardware

- Processor
- Memory
- I/O
- Interconnections
- Interrupt Priorities and Latency

Cache

OS

Source Code

```
L1: a = b * c
L2: g = d + e
L3: h = a - f
```

Load c
Load b
Multiply
Store into a

 $\sum_1^4 t_{\text{exec}}(L_{1,i})$

Loops - while

```
while (p) do
    Q1
    Q2
    Q3
end while
```

Loops if then else

```
if B1 then
    S1
else if B2 then
    S2
else if B3 then
    S3
else
    S4
endif
```

$T(B_1) + T(S_1) + T(\text{JMP})$
 $T(B_2) + T(B_1) + T(S_2) + T(\text{JMP})$
 $T(B_3) + T(B_2) + T(B_1) + T(S_3) + T(\text{JMP})$

A Quick Review

The Basics

What's an Operating System?

Provides environment for executing programs
Process abstraction for multitasking/concurrency

- Scheduling

Hardware abstraction layer (device drivers)

Filesystems

Communication

We will focus on concurrent, real-time issues

Do I Really Need An OS?

Not always

Simplest approach: **cyclic executive**

loop

do part of task 1

do part of task 2

do part of task 3

end loop

Cyclic Executive

Advantages

- Simple implementation
- Low overhead
- Very predictable

Disadvantages

- Can't handle sporadic events
- Everything must operate in lockstep
- Code must be scheduled manually

Interrupts

Some events can't wait for next loop iteration

- Communication channels
- Transient events

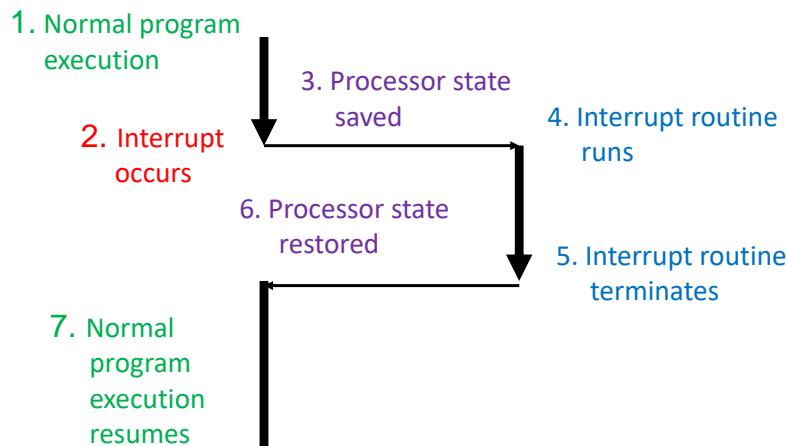
A solution: Cyclic executive plus interrupt routines

Interrupt: environmental event that demands attention

- Example: “byte arrived” interrupt on serial channel

Interrupt routine: piece of code executed in response to an interrupt

Handling an Interrupt



Drawbacks of CE + Interrupts

Main loop still running in lockstep
 Programmer responsible for scheduling
 Scheduling static
 Sporadic events handled slowly

Interrupt Service Routines

Most interrupt routines:

- Copy peripheral data into a buffer
- Indicate to other code that data has arrived
- Acknowledge the interrupt (tell hardware)
- Longer reaction to interrupt performed outside interrupt routine
- E.g., causes a process to start or resume running

Cooperative Multitasking

A cheap alternative

Non-preemptive

Processes responsible for relinquishing control

Examples: Original Windows, Macintosh

A process had to periodically call `get_next_event()` to let other processes proceed

Drawbacks:

- Programmer had to ensure this was called frequently
- An errant program would lock up the whole system

Alternative: preemptive multitasking

Concurrency Provided by OS

Basic philosophy:

- Let the operating system handle scheduling, and let the programmer handle function

Scheduling and function usually orthogonal

Changing the algorithm would require a change in scheduling.

Timesharing Operating Systems

Solution

- Store multiple batch jobs in memory at once
- When one is waiting for the tape, run the other one

Basic idea of timesharing systems

Fairness primary goal of timesharing schedulers

- Let no one process consume all the resources
- Make sure every process gets “equal” running time

Batch Operating Systems

Original computers ran in batch mode:

- Submit job & its input
- Job runs to completion
- Collect output
- Submit next job

Processor cycles very expensive at the time

Jobs involved reading, writing data to/from tapes

Cycles were being spent waiting for the tape!

Real-Time Is Not Fair



Main goal of an RTOS scheduler: meeting deadlines

If you have five homework assignments and only one is due in an hour, you work on that one

Fairness does not help you meet deadlines

Role of an OS in Real Time Systems

Standalone Applications

- Often no OS involved
- Micro controller based Embedded Systems

Some Real Time Applications are huge & complex

- Multiple threads
- Complicated Synchronization Requirements
- Filesystem / Network / Windowing support
- OS primitives reduce the software design time

Features of RTOS's

Scheduling.

Resource Allocation.

Interrupt Handling.

Other issues like kernel size.

Scheduling in RTOS

More information about the tasks are known

- No of tasks
- Resource Requirements
- Release Time
- Execution time
- Deadlines

Being a more deterministic system better scheduling algorithms can be devised.

Characteristics

RTS - Tasks

Process

A program in execution

An instance of a program running on a computer

The entity that can be assigned to and executed on a processor

A unit of activity characterized by

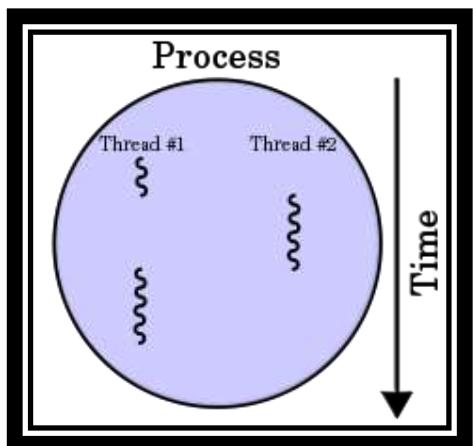
- the execution of a sequence of instructions
- a current state
- an associated set of system resources

Process Concept

Process Includes

- Program Counter
- Code
- Data
- Stack

Process & Threads



Multiprogramming

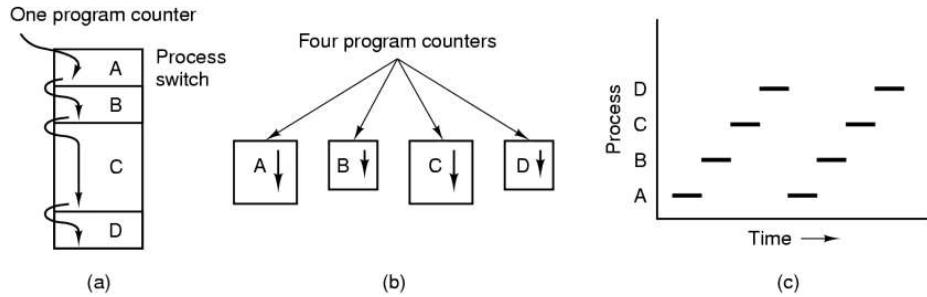
The interleaved execution of two or more computer programs by a single processor

An important technique that

- enables a time-sharing system
- allows the OS to overlap I/O and computation, creating an efficient system

Processes

The Process Model

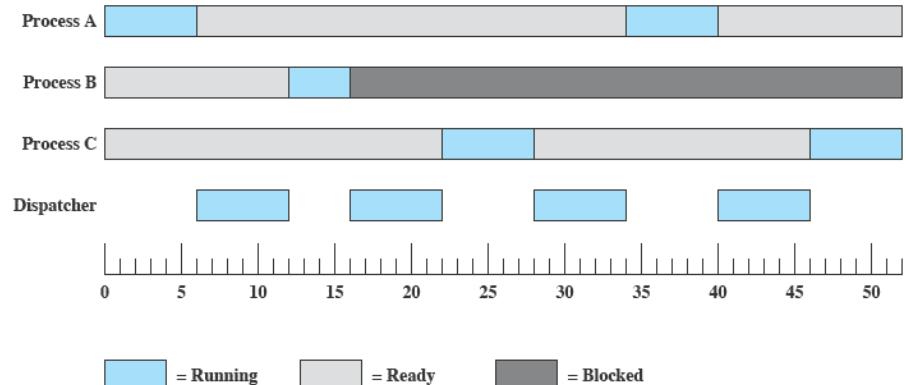


Multiprogramming of four programs

Conceptual model of 4 independent, sequential processes

Only one program active at any instant

Multiprogramming



Cooperating Processes

Sequential programs consist of a single process

Concurrent applications consist of multiple cooperating processes that execute concurrently

Advantages

- Can exploit multiple CPUs (hardware concurrency) for speeding up application

Cooperating Processes

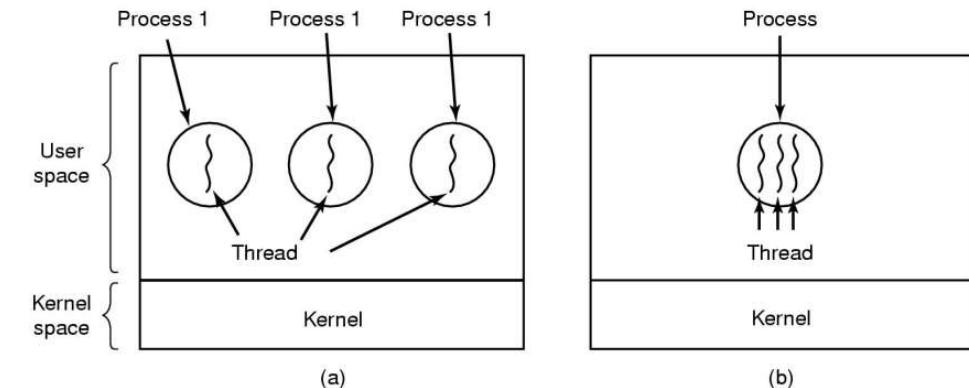
- Cooperating processes need to share information
- Since each process has its own address space, OS mechanisms are needed to let process exchange information
- Two paradigms for cooperating processes
 - Shared Memory
 - OS enables two independent processes to have a shared memory segment in their address spaces
 - Message-passing
 - OS provides mechanisms for processes to send and receive messages

Threads: Motivation

Process created and managed by the OS kernel

- Process creation expensive
- Context switching expensive
- IPC requires kernel intervention - expensive
- Cooperating processes – no need for memory protection, i.e., separate address spaces

Threads- The Thread Model



(a) Three processes each with one thread
 (b) One process with three threads

The Thread Model

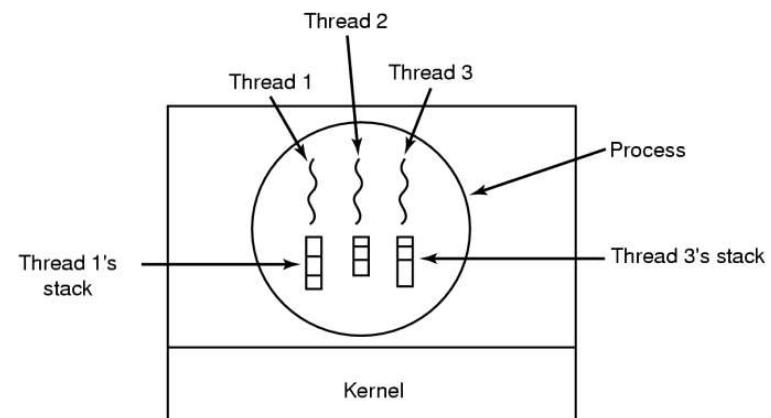
Items shared by all threads in a process

Items private to each thread

Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

The Thread Model

Each thread has its own stack



Threads/ Process

Similarities

- Share CPU - only one thread active (running) at a time
- Threads within a processes execute sequentially.
- Can create children.
- If one thread is blocked- another thread can run.

Differences

- Threads are not independent of one another
- All threads can access every address in the task
- Thread are designed to assist one other- processes might/not – as they originate from different users

Why Threads

Process with multiple threads make a great server - printer server

Threads can share common data - do not need to use inter-process commn

Threads can take advantage of multiprocessors.

Threads are cheap in the sense that

- They only need a stack and storage for registers - cheap to create.
- Threads use very little resources of OS - they do not need new address space, global data, program code or OS resources.

Context switching fast - only have to save and/or restore PC, SP & regs

But this cheapness does not come free - the biggest drawback is that there is no protection between threads.

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Thank You.

Any Questions?



[AMD Official Use Only - General]



BITS ZG553: Real Time Systems

BITS Pilani

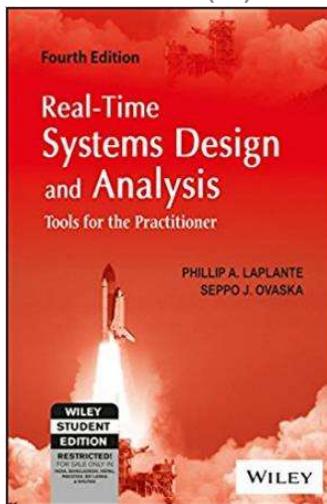
Pilani|Duba|Goa|Hyderabad

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

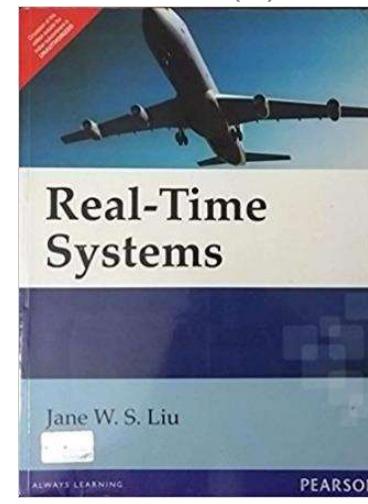


Text Book / References

Reference (R1)



Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

Excellent MOOCs Videos (Coursera, edX,...)



HONOR CODE CERTIFICATE



Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.



HONOR CODE CERTIFICATE
Issued November 28, 2015

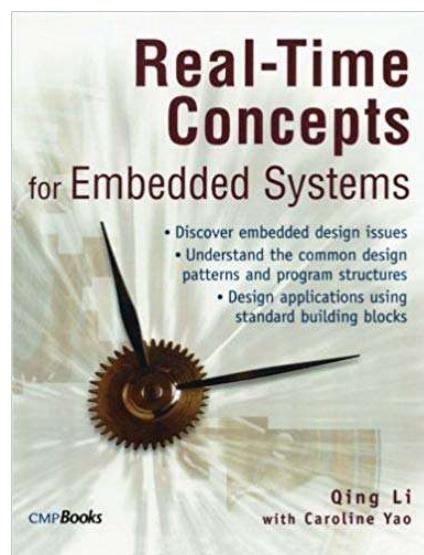
VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



BITS Pilani
Pilani|Duba|Gwalior|Hyderabad

RTS Primer – For Light Reading



L-2: Real Time Systems – Definitions/Notations, Task Model, Scheduling Approaches

Ref: T1/[C3, Lecture PPT/Notes]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations by Prof B Mishra / Prof K R Anupa, BITS-Pilani WILP Faculty



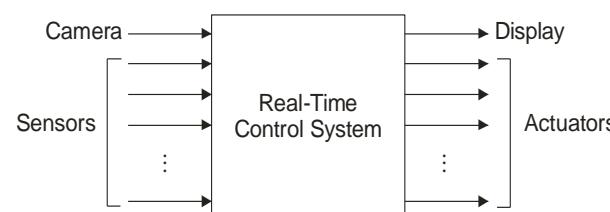
Definitions

1. A real-time system is a computer system that must satisfy bounded response-time constraints or risk severe consequences, including failure.
2. A real-time system is any information processing system which has to respond to externally generated input stimuli within a finite and specified period



A Real-Time System

- Inputs are *excitations* and outputs are corresponding *responses*
- Inputs and outputs may be digital or analog
- Inputs are associated with sensors, cameras, etc.
- Outputs with actuators, displays, etc.



Source: P. A. Laplante & S. J. Ovaska, *Real-Time Systems Design and Analysis: Tools for the Practitioner*, Wiley, 4th edition



What Is Special about Real-Time Systems?

- More number of embedded/real-time systems than desktops
- Stricter requirements: More focus on
 - Timeliness
 - Robustness
 - Safety



Soft and Hard Real Time Systems

- Soft Real Time System
 - A soft real-time system is one in which performance is degraded but not destroyed by failure to meet response-time constraints.
 - User requires only a demonstration that the system always meet some *statistical constraints*.
 - Best effort service
 - Example: Multimedia system
- Hard Real Time System
 - A hard real-time system is one in which failure to meet a single deadline may lead to complete and catastrophic system failure
 - User requires validation that the system always meet the *timing constraint*. (validation means demonstration by a provably correct and efficient procedure)
 - Guaranteed service
 - Example: Avionics weapons delivery system



Types of Real Time Tasks

➤ Periodic Task

- The specified task is executed at a **fixed period**.
- Example: A process control system sampling temperature from a temperature sensor every 5 sec and processing it.

➤ Aperiodic Task

- A task is said to be Aperiodic, when it is asynchronously arrived and it has a soft or no deadline.
- Example: Low priority interrupts

➤ Sporadic Task

- A task is said to be Sporadic, when it is asynchronously arrived and it has a hard deadline.
- Example: High priority interrupts



Processors and Resources

All system resources are divided into two categories:

Processors

- Active resources e.g. Computers, Transmission links, Disks etc
- Denoted by **P**

Resources

- Passive resources e.g. Memory, Sequence numbers, Mutexes, Database locks etc
- Denoted by **R**



RTS: Reference Model

Each system is characterised by three elements

- **Workload model:**
 - Describes applications supported by the system
- **Resource model**
 - Resources available to the system
- **Algorithms**
 - Algorithm that defines how the system utilizes the resources at all times



Assumptions

- Many parameters of hard real-time jobs are known at all time
 - For example, number of tasks in a hard-real time system is known at the beginning
- Rate of progress of a job depends upon the speed of the processor on which it executes
- Rate of progress of a job doesn't depend upon the resources it uses during execution
- Plentiful resources (resources for which no job is prevented from execution) are omitted from the consideration
- Memory is considered as a plentiful resource, hence omitted from the consideration



Parameters associated with a Job

Job: Work done by a system e.g. computation of a control law, FFT computation etc

Each job J_i is characterized by its

- Temporal Parameters
- Functional Parameters
- Resource Parameters
- Interconnection Parameters

14 BITS Pilani, Pilani Campus

Example



A system which monitors and controls several furnaces

At the beginning, the system is initialized and starts execution (time 0 ms). The system samples and reads each temperature sensor every 100 ms. The system computes the control law of each furnace every 100 ms.

Each of the control law computation is a job, say J_k , $k = 0, 1, 2, \dots$

The first control law computation started at 20 ms.

So, release time of first job, $r_0 = 20 \text{ ms}$

Subsequent release times: $120 \text{ ms}, 220 \text{ ms}, \dots$

So release time of k th job $r_k = 20 + 100 * k \text{ ms}$

Period $p_k = 100 \text{ ms}$

If each job needs to be completed within 70 ms , then the absolute deadlines of the jobs: $90 \text{ ms}, 190 \text{ ms}, 290 \text{ ms}, \dots$

Relative deadline $D_k = 70 \text{ ms}$

16 BITS Pilani, Pilani Campus



Timing Specifications

- **Release Time (r_i):** The instance of time at which the job becomes available for execution
- **Deadline :** Deadline of a job is the instant of time by which its execution is required to be completed
- **Execution time (e_i):** The amount of time required to complete the execution of the job. The actual amount of time required by a job may vary due to various reasons like conditional branches, performance enhancing features like caches and pipelines etc.
- **Response Time:** The length of time from the release time of the job to the instance when it completes
- **Relative Deadline (D_i):** Maximum allowable response time
- **Absolute Deadline (d_i)** = Release time + Relative deadline
- **Timing Constraint** is specified in terms of its release time and relative or absolute deadlines
- **Feasible interval:** The time between the release time and the absolute deadline i.e. (r_i, d_i)

Release Time, Relative Deadline and Absolute Deadline are the Temporal Parameters of Job J_i

15 BITS Pilani, Pilani Campus

Periodic Task Model



Well-known deterministic workload model

Each computation or data transmission is executed repeatedly at regular or semi regular intervals.

Period p_i of a periodic task T_i is the minimum length of all time intervals between release time of all jobs in it.

Execution time e_i is the maximum amount of time required to complete the execution of a task T_i when it executes alone and has all the resources it requires.

The accuracy of the periodic model decreases with increasing jitter in release time and variations in execution times.

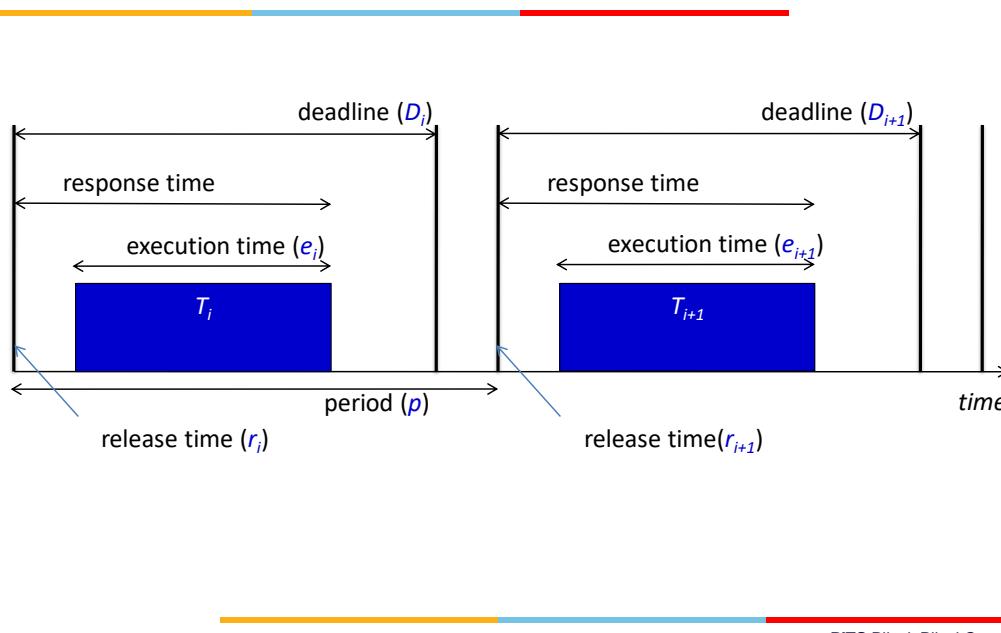
Tasks in the system are denoted as: $T_1, T_2, T_3, \dots, T_n$

Jobs in Task T_i are denoted as: $J_{i,1}, J_{i,2}, J_{i,3}, \dots, J_{i,k}$

The **release times** of the jobs in Task T_i are denoted as: $r_{i,1}, r_{i,2}, r_{i,3}, \dots, r_{i,k}$

17 BITS Pilani, Pilani Campus

Timing Specifications – Periodic Task



The release time of the first job $J_{i,1}$ of Task T_i is called **Phase ϕ** of the Task T_i i.e.

$$\phi_i = r_{i,1}$$

Two tasks are said to be in phase, when they have same phase. → starting period of both the task are same; first instance of each task gets released at the same time.

Hyperperiod H of a periodic task T_i is LCM of all periods of that task i.e.

$$H = \text{LCM}(p_1, p_2, \dots, p_k)$$

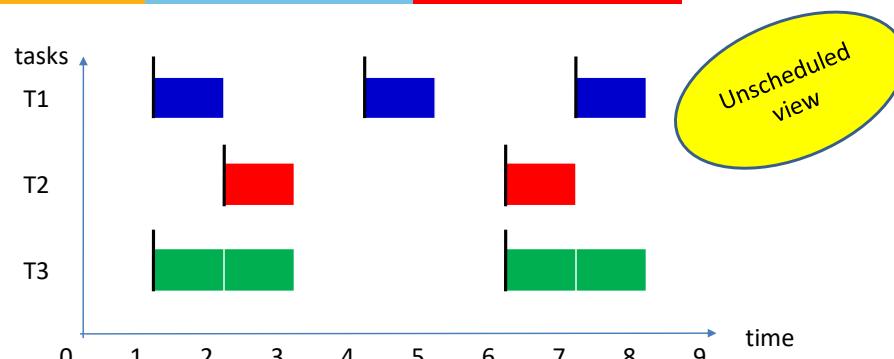
Relative deadline (or simply **deadline**) D_i of a task T_i is the maximum allowable time for each job of the task to complete i.e. maximum response time for each job of the task.

We will often assume that for every task, a job is released and become ready at the beginning of each period and must complete by the end of the period. Hence D_i will become equal to p_i . However, in general D_i can have an arbitrary value less than p_i .

Absolute deadline of the job $J_{i,j}$ $d_{i,j} = r_{i,j} + D_i$

Feasible Interval is $(r_{i,j}, d_{i,j}]$

Periodic Task Model



Tasks	ϕ_i	p_i	e_i	D_i	u_i
T_1	1	3	1	3	0.33
T_2	2	4	1	4	0.25
T_3	1	5	2	5	0.4



Periodic Task Model

The release time of the first job $J_{i,1}$ of Task T_i is called **Phase ϕ** of the Task T_i i.e.

$$\phi_i = r_{i,1}$$

Two tasks are said to be in phase, when they have same phase. → starting period of both the task are same; first instance of each task gets released at the same time.

Hyperperiod H of a periodic task T_i is LCM of all periods of that task i.e.

$$H = \text{LCM}(p_1, p_2, \dots, p_k)$$

Relative deadline (or simply **deadline**) D_i of a task T_i is the maximum allowable time for each job of the task to complete i.e. maximum response time for each job of the task.

We will often assume that for every task, a job is released and become ready at the beginning of each period and must complete by the end of the period. Hence D_i will become equal to p_i . However, in general D_i can have an arbitrary value less than p_i .

Absolute deadline of the job $J_{i,j}$ $d_{i,j} = r_{i,j} + D_i$

Feasible Interval is $(r_{i,j}, d_{i,j}]$



Periodic Task Model

Utilization of task T_i

$$u_i = \frac{e_i}{p_i}$$

Total Utilization

$$U = \sum_{i=1}^n \frac{e_i}{p_i}$$



CPU Utilization Zones

Utilization %	Zone Type	Typical Application
< 26	Unnecessarily safe	Various
26 – 50	Very safe	Various
51 – 68	Safe	Various
69	Theoretical limit (RM Scheduling)	Embedded systems
70 – 82	Questionable	Embedded systems
83 – 99	Dangerous	Embedded systems
100	Critical	M marginally stressed system
> 100	Overloaded	Stressed system



Parameters Summary

Precedence constraints: specifies if any task(s) needs to precede other tasks.

Release or arrival time r_{ij} : the release time of the j^{th} instance of task i .

Phase: Φ_i the release time of the first instant of task i .

Response time: Time span between the task activation and its completion.

Absolute deadline d_i : is the instant of time by which the task must complete.

Relative deadline D_i : is the maximum allowable response time of the task.

Laxity type: Notion of urgency or leeway in a task's execution.

Period p_i : is the minimum length of intervals between the release times of consecutive tasks.

Execution time e_i : is the (maximum) amount of time required to complete the execution of a task i when it executes alone and has all the resources it requires.

Aperiodic and Sporadic jobs



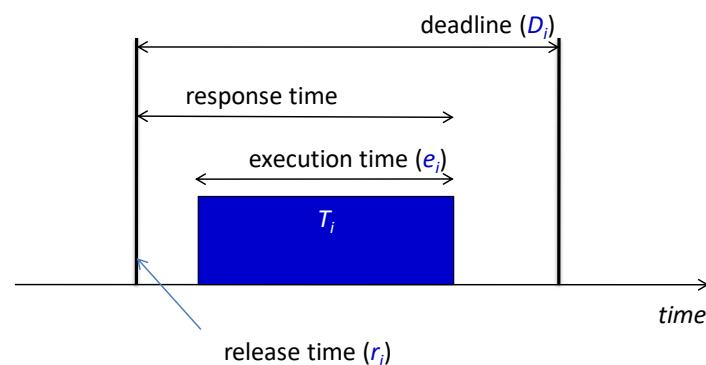
- We call a task to be **aperiodic** if the jobs in it have either soft deadlines or no deadline.

Example: Task to adjust radar's sensitivity

- We call a task to be **sporadic** when the jobs are released at random time instants, but have hard deadlines.

Example: A command and control system (of airplane control) must process sporadic data messages in order to continuous voice and video traffic.

Timing Specifications – Aperiodic & Sporadic Task





Precedence Constraints, Predecessors and Successors

- Jobs are said to have **precedence constraints** if they are constrained to execute in some order.
- If jobs can execute in any order, they are called **independent**.
- A job J_i is a **predecessor** of another job J_k if J_k can't begin execution until the execution of J_i completes. It is denoted by

$$J_i < J_k$$

Here J_k is a **successor** of J_i .

Let us consider the relation $J_i < J_j < J_k$.

J_i is a **immediate predecessor** of J_j and predecessors of J_j and J_k

J_k is a **immediate successor** of J_j and successor of J_i and J_j

- A job with predecessors is ready for execution when the time is at or after its release time and all of its predecessors are completed.



Clock-driven Approach

- Scheduling decisions are made at specific time instants, which are chosen a priori before the system begins execution
- Typically this type of scheduling is **suitable for hard real-time systems** and **smaller systems**, where the parameters are fixed and known.
- Scheduling decisions are computed off-line and stored for use at run-time, thus scheduling overhead is minimal.
- Generally a hardware time is set to expire periodically.
- After the system gets initialized, the scheduler selects and schedules jobs which execute till the next scheduling decision is made. Then the scheduler blocks itself waiting for the expiration of the timer.
- When the timer expires, the scheduler wakes up, does necessary scheduling and sleeps again. This process repeats.



Three Approaches for Scheduling Real Time System Tasks

- Clock-driven
- Round-robin / Weighted round-robin
- Priority-driven

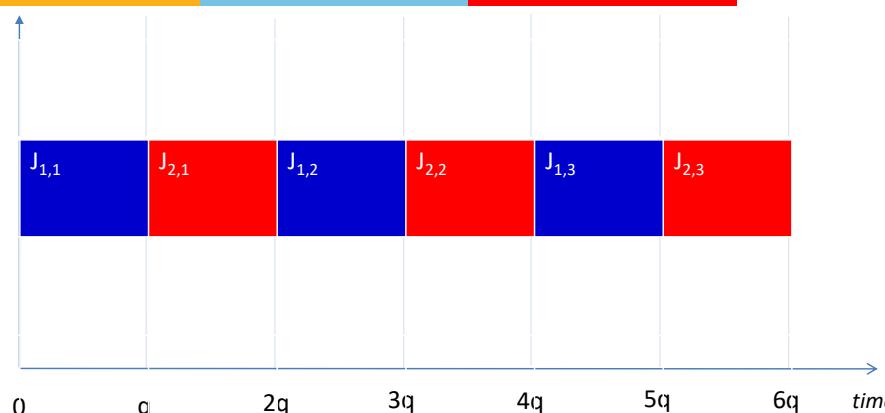


Round-robin Approach

- Also known as **time-sharing**
- Every job joins a FIFO (First-in-first-out) queue when it becomes ready for execution
- The entire time period is divided into several **time-slices**
- The job at the **head of the queue** executes for one time-slice.
- If the job doesn't complete at the end of the time-slice, it gets pre-empted and placed at the end of the queue to waits for its next turn.
- If there are ' n ' jobs ready for execution, each job gets $1/n$ th share of the processor.



Round-robin Approach - Example



(Round-robin execution of two tasks on a single processor)
(Time quantum = q)

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Priority Driven Approach



- Priorities are assigned to the jobs based on their criticality
- Jobs ready for execution are placed in one or more queues ordered by priorities of the jobs.
- At any scheduling decision time, the jobs with the highest priorities are scheduled and executed on the available processors.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Weighted round-robin Approach

- This approach is a round robin approach with different weights assigned to different jobs.
- If a job has weight ' wt ', then it will get ' wt ' time slices every round for execution.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Priority-Driven Approach - Example



- 3 Jobs to be scheduled based on priority-driven approach:
 - J1: priority 1, release time 15, execution time 10
 - J2: priority 2, release time 0, execution time 30
 - J3: priority 3, release time 18, execution time 20

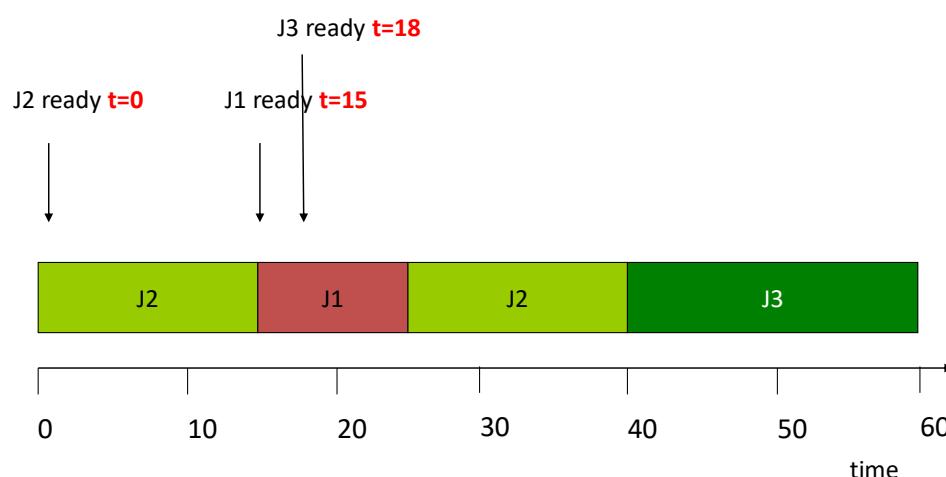
Rules:

- Each process has a fixed priority (1 highest);
- Highest-priority ready process gets CPU;
- Process continues until done.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Priority-Driven Approach - Example



Priority-Driven VS Clock-Driven Approaches

- Priority driven approaches have many advantages compared to clock driven approach:
 - They **don't have to have the information on the release time, execution time etc** (in contrast with clock driven approach, where these parameters are required to be known *a priori*)
 - It is best suited for applications with **varying time and resource requirements**
 - Many well-known priority –driven algorithms use very simple priority assignments **reducing the overhead** of maintaining multiple queues.
- Despite all these advantages, **Clock-driven approaches** are used for hard real-time systems, especially in safety-critical systems.
 - **The major reason is that the timing behaviour of a priority-driven system is nondeterministic when job parameters vary.**



What to do if CPU is getting overloaded ?

CPU is said to be overloaded, when the **Utilization > 1**.

What to do in this case?

Reduce Execution Time !

How to do it ?

- Replace the CPU with a **faster CPU**
 - CPU with higher frequency
 - CPU with more cores
 - CPU with a deeper pipeline and more cache
- Add **more RAM** (it will reduce the page fault overhead)
- Have **dedicated Hardware** (e.g. custom ASICs, DSP Processors) for specific functions
- **Optimize** the code
- ...

Thank You.

Any Questions?



BITS ZG553: Real Time Systems



BITS Pilani
Pilani|D.b.a|Goa|Hyderabad

K G Krishna
WILP Division, BITS-Pilani, Hyderabad



Excellent MOOCs Videos (Coursera, edX,...)



HONOR CODE CERTIFICATE



Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

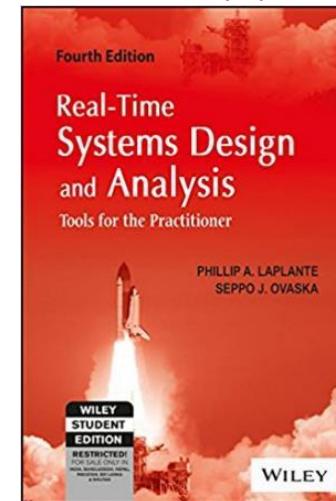
a course of study offered by IEEEex, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

HONOR CODE CERTIFICATE
Issued November 28, 2015

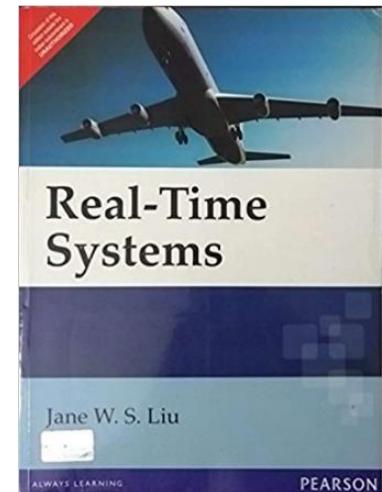
VALID CERTIFICATE ID
06ae78da4df44a218f5a1eebf47d4e54

Text Book / References

Reference (R1)



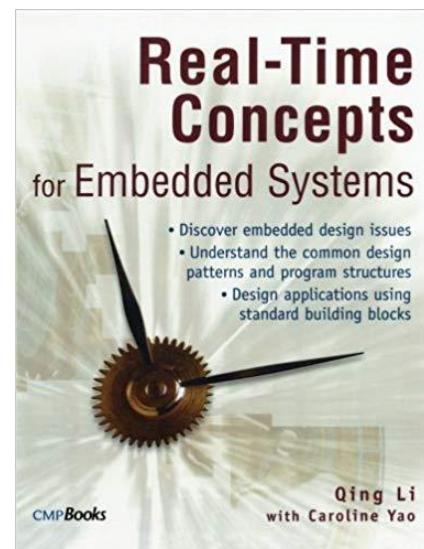
Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

RTS Primer – For Light Reading





L-3: Introduction to Scheduling - Approaches/Algorithms

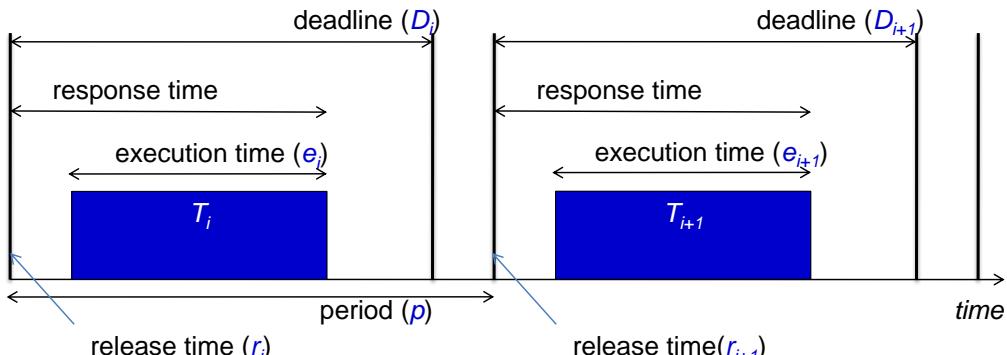
Ref: [Lecture PPT/Notes]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

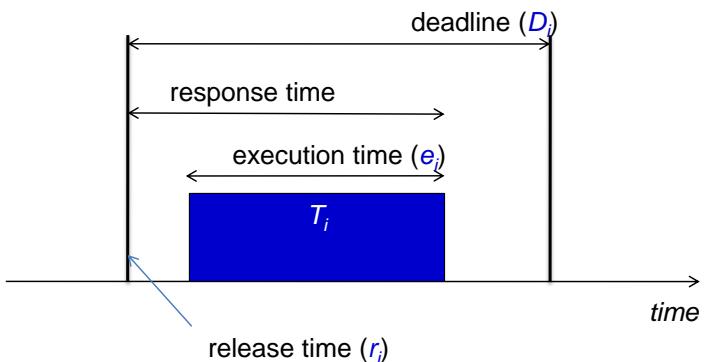
PLEASE DO NOT PRINT PPTs, Save the Environment!

5

Timing Specifications – Periodic Task



Timing Specifications – Aperiodic & Sporadic Task



Three Approaches for Scheduling Real Time System Tasks

- Clock-driven
- Round-robin / Weighted round-robin
- Priority-driven



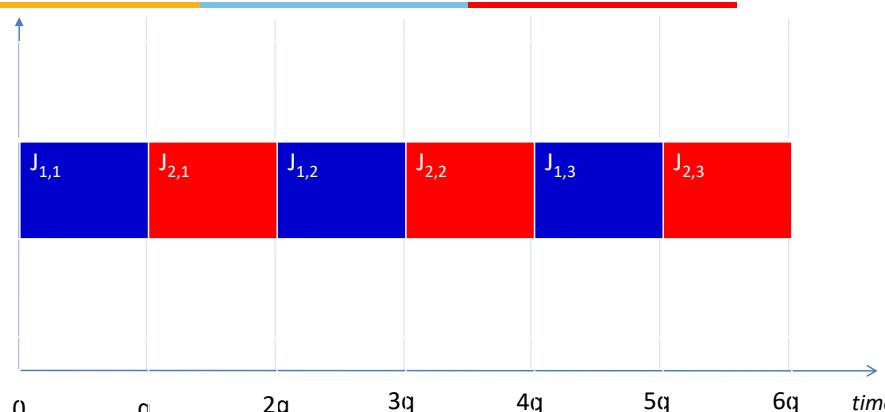
Clock-driven Approach

- Scheduling decisions are made at specific time instants, which are chosen a priori before the system begins execution
- Typically this type of scheduling is **suitable for hard real-time systems** and **smaller systems**, where the parameters are fixed and known.
- Scheduling decisions are **computed off-line** and stored for use at run-time, thus scheduling overhead is minimal.
- Generally a hardware time is set to expire periodically.
- After the system gets initialized, the scheduler selects and schedules jobs which execute till the next scheduling decision is made. Then the scheduler blocks itself waiting for the expiration of the timer.
- When the timer expires, the scheduler wakes up, does necessary scheduling and sleeps again. This process repeats.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Round-robin Approach - Example



(Round-robin execution of two tasks on a single processor)
(Time quantum = q)

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Round-robin Approach

- Also known as **time-sharing**
- Every job joins a FIFO (First-in-first-out) queue when it becomes ready for execution
- The entire time period is divided into several **time-slices**
- The job at the **head of the queue** executes for one time-slice.
- If the job doesn't complete at the end of the time-slice, it gets pre-empted and placed at the end of the queue to waits for its next turn.
- If there are ' n ' jobs ready for execution, each job gets $1/n$ th share of the processor.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Weighted round-robin Approach

- This approach is a round robin approach with different weights assigned to different jobs.
- If a job has weight ' wt ', then it will get ' wt ' time slices every round for execution.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



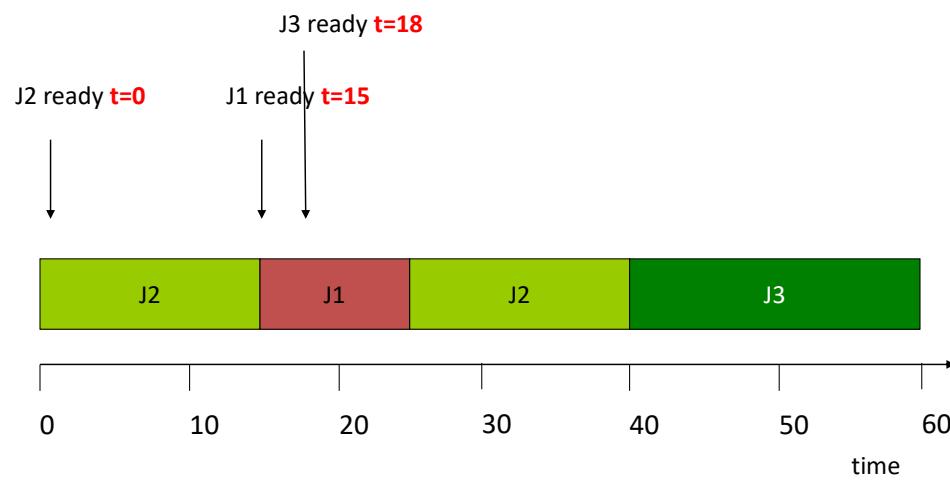
Priority Driven Approach

- Priorities are assigned to the jobs based on their **criticality**
- Jobs ready for execution are placed in one or more queues ordered by priorities of the jobs.
- At any scheduling decision time, the jobs with the highest priorities are scheduled and executed on the available processors.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Priority-Driven Approach - Example



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Priority-Driven Approach - Example

- 3 Jobs to be scheduled based on priority-driven approach:

- J1: priority 1, release time 15, execution time 10
- J2: priority 2, release time 0, execution time 30
- J3: priority 3, release time 18, execution time 20

Rules:

- Each process has a fixed priority (1 highest);
- Highest-priority ready process gets CPU;
- Process continues until done.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Priority-Driven vs Clock-Driven Approaches

- Priority driven approaches have many advantages compared to clock driven approach:

- They **don't have to have the information on the release time, execution time etc** (in contrast with clock driven approach, where these parameters are required to be known *a priori*)
- It is best suited for applications with **varying time and resource requirements**
- Many well-known priority –driven algorithms use very simple priority assignments **reducing the overhead** of maintaining multiple queues.

- Despite all these advantages, **Clock-driven approaches are used for hard real-time systems, especially in safety-critical systems.**

- **The major reason is that the timing behaviour of a priority-driven system is nondeterministic when job parameters vary.**

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



What to do if CPU is getting overloaded ?

CPU is said to be overloaded, when the [Utilization > 1](#).

What to do in this case?

[Reduce Execution Time !](#)

How to do it ?

- Replace the CPU with a [faster CPU](#)
 - CPU with higher frequency
 - CPU with more cores
 - CPU with a deeper pipeline and more cache
- Add [more RAM](#) (it will reduce the page fault overhead)
- Have [dedicated Hardware](#) (e.g. custom ASICs, DSP Processors) for specific functions
- [Optimize](#) the code
- ...

Let's Examine each of the [Scheduling Algorithms](#) in little more detail...

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



How to Assign Priority to Jobs?

Based on [Criticality](#) ?

But on what is the basis of determining criticality?

Scheduling Algorithms comes for rescue.

[Priorities](#) of the jobs are determined by applying [Scheduling algorithms](#).

Two types of Scheduling algorithms:

➤ [Fixed Priority](#) Scheduling Algorithms:

Priorities of all the jobs of a task remain constant
(of course applicable for periodic tasks only)

➤ [Dynamic Priority](#) Scheduling Algorithms

Priorities of the jobs of a task may vary

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Rate Monotonic Scheduling Algorithm

- Fixed priority algorithm
- [Shorter the period, higher the priority](#)
- Rate is inverse of the period. Hence higher the rate, higher the priority. – so the name '[rate monotonic](#)'

Example: Consider 3 periodic tasks with following timing parameters

T1: Release time 0, period 4, execution time 1, relative deadline 4

T2: Release time 0, period 5, execution time 2, relative deadline 5

T3: Release time 0, period 20, execution time 5, relative deadline 20

Schedule them applying RM Algorithm.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

RM Scheduling Algorithm - Example (contd.)

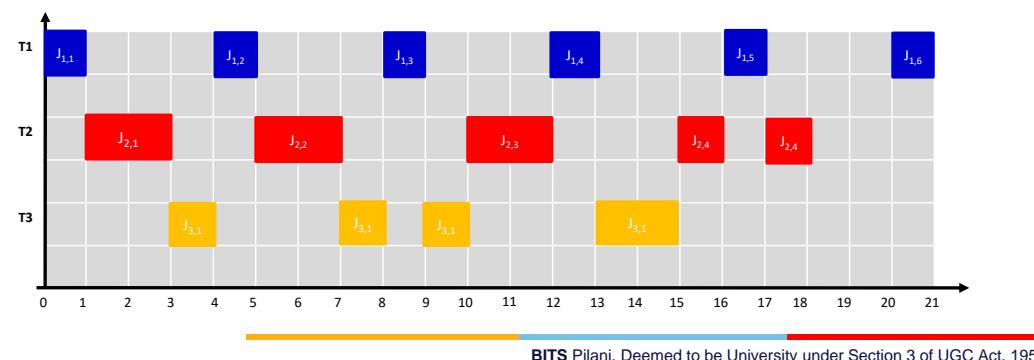
Solution

T1: Release time 0, period 4, execution time 1, relative deadline 4

T2: Release time 0, period 5, execution time 2, relative deadline 5

T3: Release time 0, period 20, execution time 5, relative deadline 20

T1 has shortest period (i.e. 4), so should have higher priority, followed by T2 and T3.

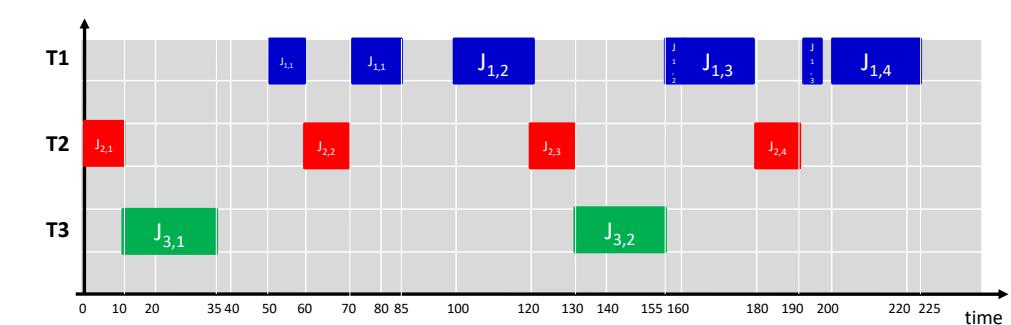


Deadline Monotonic (DM) – Example (contd.)

T1: Release time 50, period 50, execution time 25, relative deadline 100

T2: Release time 0, period 60, execution time 10, relative deadline 20

T3: Release time 0, period 125, execution time 25, relative deadline 50



Deadline Monotonic (DM) Scheduling Algorithm

➤ Fixed priority algorithm

➤ Shorter the relative deadline, higher the priority

➤ When the relative deadlines of every task is proportional to their period, the schedule produced by RM and DM algorithms are identical.

Example: 3 periodic tasks with following timing parameters

T1: Release time 50, period 50, execution time 25, relative deadline 100

T2: Release time 0, period 60, execution time 10, relative deadline 20

T3: Release time 0, period 125, execution time 25, relative deadline 50

Schedule them applying DM Algorithm.

Solution

According to DM algorithm, T2 has highest priority because its relative deadline is 20.

Similarly T1 has lowest priority and T3 has priority in between.

Earliest Deadline Fast (EDF) Scheduling Algorithm

➤ Dynamic priority algorithm

➤ Job with earliest (absolute) deadline has highest priority

Example: Consider two tasks

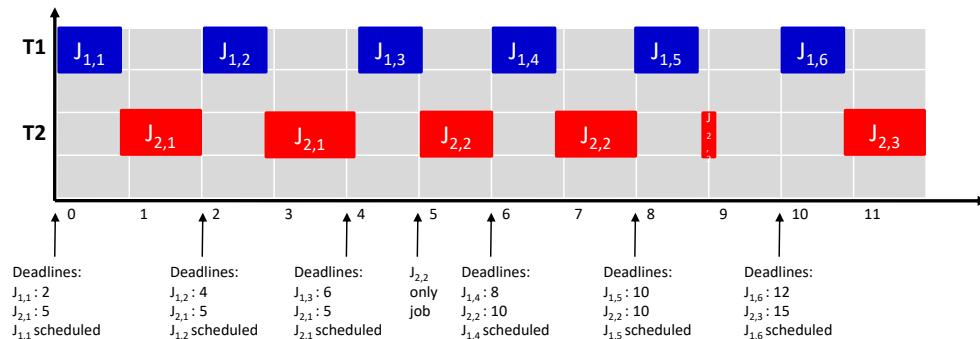
T1: Release time 0, period 2, execution time 0.9, relative deadline 2

T2: Release time 0, period 5, execution time 2.3, relative deadline 5

Schedule them applying EDF Algorithm.

EDF Algorithm – Example (contd.)

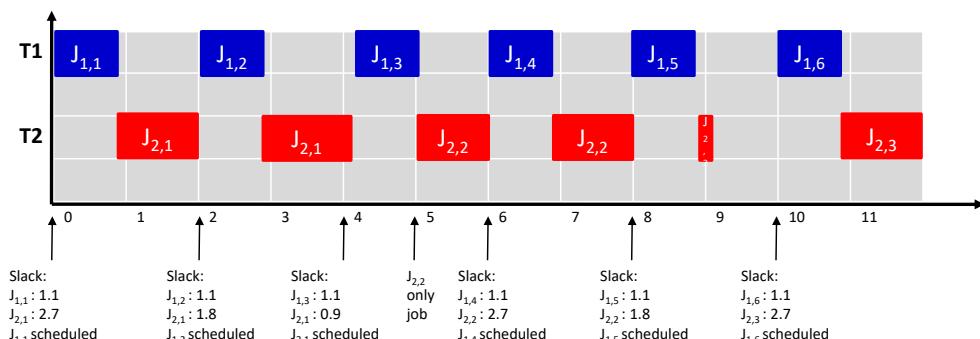
T1: Release time 0, period 2, execution time 0.9, deadline 2
T2: Release time 0, period 5, execution time 2.3, deadline 5



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

LST Algorithm – Example (contd.)

T1: Release time 0, period 2, execution time 0.9, relative deadline 2
T2: Release time 0, period 5, execution time 2.3, relative deadline 5



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

LST (Least-Slack-Time-First) Scheduling Algorithm

- Dynamic priority algorithm
- Job with smallest slack has highest priority
- At time t , the slack of a job whose remaining execution time is x and whose deadline is $d = d - t - x$

Note: Absolute deadline (d) = Release time + Relative deadline

Example: Consider two tasks

T1: Release time 0, period 2, execution time 0.9, relative deadline 2
T2: Release time 0, period 5, execution time 2.3, relative deadline 5
Schedule them applying LST Algorithm.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Any Questions?

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Text Book / References

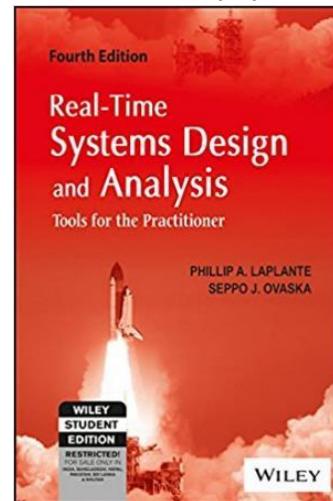


BITS Pilani
Pilani|D.b.a|Gwalior|Hyderabad

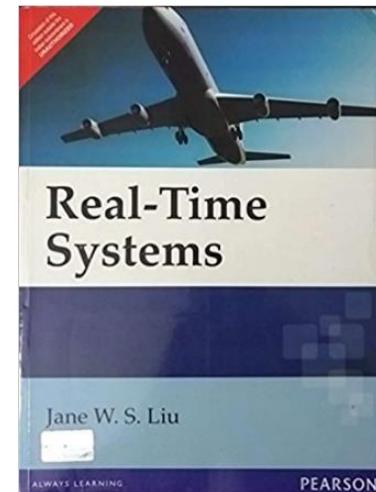
BITS ZG553: Real Time Systems [L-3a: Clock-driven Scheduler Examples]

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Reference (R1)



Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Excellent MOOCs Videos (Coursera, edX,...)



**HONOR CODE
CERTIFICATE**



This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEex, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

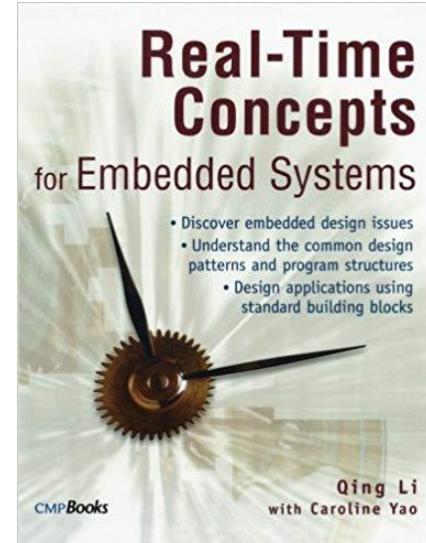
HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44a218f5a1eebf47d4e54

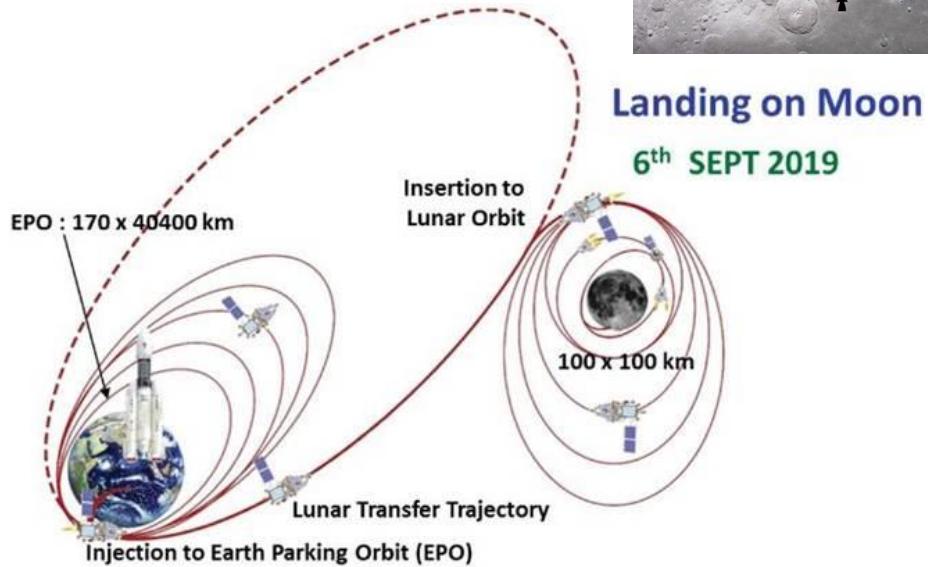
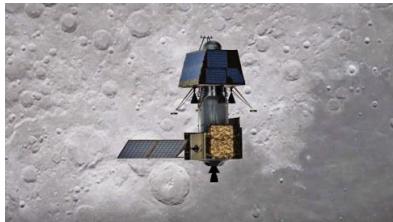
Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

RTS Primer – For Light Reading



What Really Happened? Lessons from 'Unknown-Unknowns'!



5

6



BITS Pilani

Pilani|Doha|Gwalior|Hyderabad



L-3a: Introduction to Scheduling - Clock-driven Scheduler, Examples

Ref: [T1]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs, Save the Environment!

Assumptions

- ❑ There are ' n ' periodic tasks in the system and ' n ' is fixed
- ❑ The parameters of all periodic tasks are known a priori.
- ❑ Variation in the inter-release time is negligibly small. For all practical purpose each job in T_i is released p_i units of time after the previous job in T_i .
- ❑ Each job $J_{i,k}$ is ready for execution at release time $r_{i,k}$

A periodic task T_i is characterised by the 4-tuple (Φ_i, p_i, e_i, D_i) .

Example: (1, 10, 3, 6) periodic task

- 1st job is released at time 1. It must complete by time 7 ($=1 + 6$). It executes for 3 units of time.
- 2nd job is released at time 11 ($=1 + 10$). It must complete by time 17 ($=11 + 6$). It executes for 3 units of time.
- ...

- ❑ Tuples having default values are omitted.
- ❑ By default, the phase of each task is 0 and relative deadline is equal to the period.

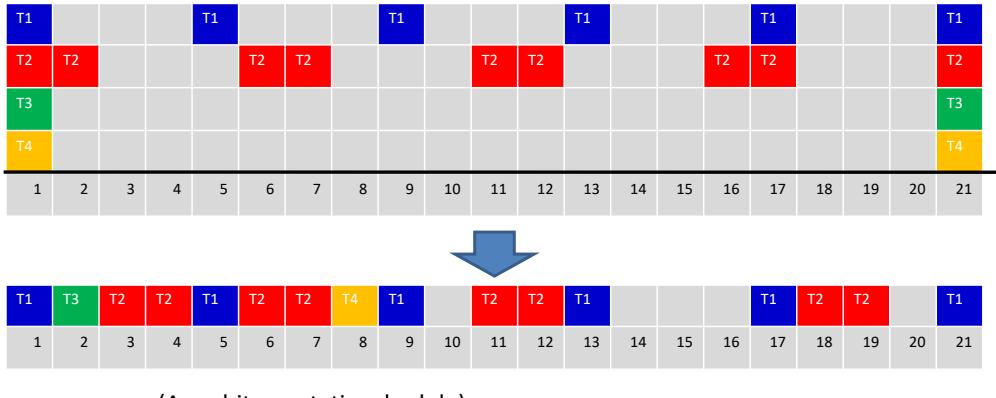
Example: (10, 3, 6) and (10, 3) have zero phases and relative deadlines of 6 and 10 respectively.

A Clock-Driven Scheduler

- Whenever the parameters of the jobs are known a priori (As in case of Hard Real-time Systems), it is better to construct a **static schedule offline**.
- We call a periodic static schedule a **cyclic schedule** as well.
- Store the pre-computed schedule in a table
- Each entry $(t_k, T(t_k))$ in the table gives a time t_k , and the task T_k or I to be scheduled at that time.
- T_k is a periodic job and I is idle interval.
- There should be a **timer**, which expires at the decision times t_k and raises an interrupt as it expires.
- This interrupt wakes up the scheduler, which schedules the appropriate job from the off-line schedule.
- If the job to be scheduled is I , then it is the idle time (no periodic job is available to be scheduled). Hence the job at the head of the **aperiodic queue** is scheduled.

Example

- ❑ Consider 4 tasks: $T_1 = (4, 1); T_2 = (5, 2); T_3 = (20, 1); T_4 = (20, 1)$
- ❑ Hyperperiod = LCM (4, 5, 20, 20) = 20 \rightarrow Schedules can repeat every 20 time units
- ❑ Max no of jobs in each hyperperiod (No of entries in the schedule) = $(20 / 4) + (20 / 5) + (20 / 20) + (20 / 20) = 5 + 4 + 1 + 1 = 11$



General Structure of Cyclic Schedules

- A cyclic scheduler repeats a pre-computed schedule.
- The pre-computed schedule needs to be stored only for one major cycle.
- Usually the major cycle is the **Hyperperiod**.
- Each task in the task set to be scheduled repeats identically in every major cycle.
- The **major cycle** is divided into one or more **frames**
- The **scheduling points of a cyclic scheduler occur at frame boundaries**. This means that a task can start executing only at the beginning of a frame.
- The frame boundaries are defined through the interrupts generated by a periodic timer. Each task is assigned to run in one or more frames.

Tasks	Frame
T_1	f_1
T_3	f_2
T_5	f_3
T_4	f_4
I	f_5

A Clock-Driven Scheduler

H = Length of the hyperperiod

N = Number of entries in the schedule of each hyperperiod

I = Idle Interval

Input: Stored schedule $(t_k, T(t_k))$, for $k = 0, 1, \dots, N-1$

Task: SCHEDULER:

```

set the next decision point i and table entry k to 0;
set the timer to expire at  $t_k$ 
do {
    accept the timer interrupt;
    if(an aperiodic job is executing) preempt it;
    current task  $T = T(t_k)$ 
     $i = i++$ ;
    compute the next table entry  $k = i \bmod(N)$ 
    set the timer to expire at  $\lfloor i/N \rfloor H + t_k$ 
    if(the current task  $T == i$ )
        schedule the job at the head of the aperiodic queue
    else
        schedule the task T
}
End SCHEDULER
  
```

Floor and Ceiling Functions

Floor

$\text{floor}(x) = \lfloor x \rfloor$ is the largest integer not greater than x

Ceiling

$\text{ceiling}(x) = \lceil x \rceil$ is the smallest integer not less than x

Example:

x	$\lfloor x \rfloor$	$\lceil x \rceil$
2.4	2	3
5.5	5	6
-2.1	-3	-2
-2	-2	-2

General Structure of Cyclic Schedules

Frames

- Scheduling time is divided into multiple frames of size '*f*'.
- Scheduling is done at starting of each frame, not at any other time.
- Hence there is no preemption within a frame.
- The phase of each periodic task is a nonnegative integer multiple of a frame. In other words, the first job of every task is released at the beginning of some frame.
- Apart from scheduling decision, schedule carries out monitoring and enforcement actions at the beginning of each frame.

Size of a Frame

A selected frame size should satisfy the following three constraints.

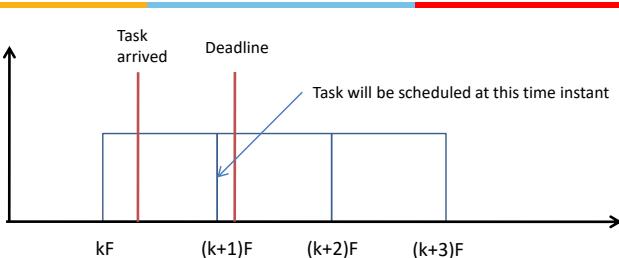
1. Minimum Context Switching

- To avoid unnecessary context switches, the selected frame size should be larger than the execution time of each task, so that when a task starts at a frame boundary it should be able to complete within the same frame.
- It means, the frame size should be larger than maximum execution time.

2. Minimization of Table Size

- This constraint requires that the number of entries in the schedule table should be minimum
- Unless a frame divides the major cycle, the number of entries in a major cycle will not be sufficient for the schedule.
- Hence the frame should divide the major cycle i.e. Hyper period of the tasks.

Size of a Frame

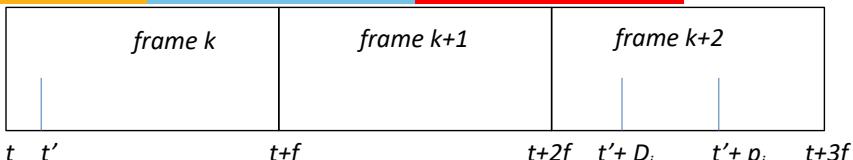


3. Satisfaction of Task Deadline

- At least one frame should be there between the time job is released and its deadline.
- Otherwise it will miss the deadline, just because it gets scheduled late in the next frame boundary (as shown in the diagram)

These three conditions are formulated into three equations in the next slides.

Frame Size Constraints



t = Beginning of the frame in which a job in T_i is released

t' = Release time of the job

If $(t' > t)$, we need frame $k+1$ to be in t' and $t'+D_i$.

For this to happen,

$$\begin{aligned} t + 2f &\leq t' + D_i \\ \Rightarrow 2f - (t' - t) &\leq D_i \end{aligned}$$

t' is an integer multiple of period p_i (say lp_i) and t is an integer multiple of frame f (say kf).

So, $t' - t = lp_i - kf$, where k, l are positive integers.

$kf - lp_i$ will be at least be equal to GCD of f and p_i .

Hence $t' - t \geq \text{gcd}(f, p_i)$

Then the above inequality becomes

$$2f - \text{gcd}(f, p_i) \leq D_i \quad \dots(3)$$

Frame Size Constraints

Minimum Context Switching

- Frame size should be larger than maximum execution time, i.e.
- $$f \geq \max_{1 \leq i \leq n} (e_i) \quad \dots(1)$$

Minimization of Table Size

- Length of a major cycle is equal to the Hyperperiod H .
- The frame size should be as short as possible so that there are integer number of frames in an Hyperperiod H . For this it needs to divide the Hyperperiod, which in turn implies that it should divide the period p_i of at least one task T_i i.e.

$$\lfloor p_i / f \rfloor - p_i / f = 0 \quad \dots(2)$$

- Let F = No of frames in an Hyperperiod (should be an integer, as per the previous equation)
- The Hyperperiod, which begins at the beginning of $(kF + 1)$ st frame, is called a '**major cycle**', for $k = 0, 1, \dots$

Satisfaction of Task Deadline

- Hence frame size should be small enough so that between the release time and deadline of a job, there should be at least one frame.

Frame Size Constraints

For $t = t'$, we need $f \leq D_i$ for a frame to be included between t' and D_i

If eq. (3) is satisfied, then above condition will automatically be satisfied, since $\text{gcd}(f, p_i) < f$.

So eq. (3) is sufficient for both the conditions $t = t'$ and $t < t'$.

Equations (1), (2) and (3) are referred as '**frame size constraints**'.

Multiple frame sizes satisfying the Frame size constraints



- For a given task set it is possible that more than one frame size satisfies all the three constraints.
- In such cases, **it is better to choose the shortest frame size**.
- This is because of the fact that the schedulability of a task set increases as more frames become available over a major cycle.

Job Slices

Example

Find out an appropriate frame size for the following system of periodic tasks to be scheduled executed according to a cyclic schedule.
 $\{(4,1), (5,2,7) \text{ and } (20,5)\}$

Solution:

Max execution time = 5.
So frame size $f \geq 5$... (1)

Frame size should divide at least one period.
So, $f = 1, 2, 4, 5, 10, 20 \dots (2)$

$2f - gcd(f, p_i) \leq D_i$
So for T1, $2f - gcd(f, 4) \leq 4$
 $\Rightarrow f = 1, 2, 4$

For T2, $2f - gcd(f, 5) \leq 7$
 $\Rightarrow f = 1, 2, 3, 4, 5$

For T3, $2f - gcd(f, 20) \leq 20$
 $\Rightarrow f = 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 20$

Hence values of f , which will satisfy $2f - gcd(f, p_i) \leq D_i$
 $f = 1, 2, 4$... (3)

Satisfying eq (1) and (3) is not possible.



So it is better, if we can slice the job with larger execution time T3 to 3 subtasks (a job of computation can be broken into multiple procedures, a message transmission job can be broken into sending multiple smaller messages etc) so that all frame size constraints gets satisfied.

In this case, if we slice T3 as, $T3 = \{(20,1), (20, 3), (20, 1)\}$, then eq (1) will become $f \geq 3$. So frame size constraints will be satisfied e.g. $f = 4$.

This process is called **job slicing**.

Frame Size Constraints - Example



Exercise 5.1 (a)

Find out an appropriate frame size for the following system of periodic tasks to be scheduled executed according to a cyclic schedule.

$\{(6,1), (10,2) \text{ and } (18,2)\}$.

Solution:

Max execution time = 2.

So frame size $f \geq 2$

... (1)

Frame size should divide at least one period.

So, $f = 2, 3, 5, 6, 10, 18$

... (2)

$2f - gcd(f, p_i) \leq D_i$

So for T1, $2f - gcd(f, 6) \leq 6$

$\Rightarrow f = 1, 2, 3, 4, 6$

For T2, $2f - gcd(f, 10) \leq 10$

$\Rightarrow f = 1, 2, 3, 4, 5, 6, 10$

For T3, $2f - gcd(f, 18) \leq 18$

$\Rightarrow f = 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12$

Hence values of f , which will satisfy $2f - gcd(f, p_i) \leq D_i$

$f = 1, 2, 3, 4, 6$

... (3)

Tasks	Period / Deadline p_i / D_i	Execution time e_i
T1	6	1
T2	10	2
T3	18	2

The values of 'f', which will satisfy eq (1), (2) and (3) are: $f = 2, 3, 6$

Hence, possible values for frame size are 2, 3 and 6.

So choose 2 (since it is smallest) as the frame size.

22

BITS Pilani, Pilani Campus

Process of Constructing Cyclic Schedules



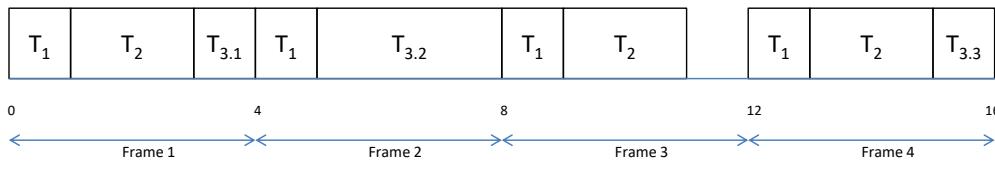
Process

- Choosing frame size
- Partitioning the jobs into slices
- Placing the slices in the frames

Tasks	Period p_i	Deadline D_i	Execution time e_i
T1	4	4	1
T2	5	7	2
T3.1	20	20	1
T3.2	20	20	3
T3.3	20	20	1

For the previous example, we have

- Partitioned T3 into 3 job slices.
- Fixed the frame size $f = 4$
- Now we have to create the schedules by putting the jobs and job slices into frames.



23

BITS Pilani, Pilani Campus

Cyclic Executive

- “Cyclic Executive” refers to a scheduler that deterministically interleaves and sequentializes the execution of periodic tasks on a CPU according to a given cyclic schedule.
- Each job slice is a ‘procedure’
- A cyclic executive executes in a single do loop, beginning of each frame and executes the slices scheduled within a frame.
- The stored table that gives the precomputed cyclic schedule has F entries, where F is the number of frames per major cycle (please note that the schedule will repeat in every major cycle, hence storing the schedule for one major cycle is enough)
- Each entry is called a ‘scheduling block’. A scheduling block for k th frame is denoted by $L(k)$ and it stores the job slices to be scheduled in k th frame.
- Usually a ‘periodic task server’ executes the job slices of a frame.
- Upon completion of executing the periodic job slices scheduled in the current frame, if there is time remaining in the current frame, the cyclic executive schedules aperiodic jobs.

Cyclic Executive

F = Number of frames in a major cycle
 $L(k)$ = Names of the job slices that are scheduled to be executed in frame k , called ‘scheduling block’

Input: Stored schedule: $L(k)$, $k = 0, 1, \dots, F-1$
 Aperiodic job queue
 Task: SCHEDULER:
 current time $t = 0$;
 current frame $k = 0$;
 do {
 accept the timer interrupt at time t_i ;
 current block = $L(k)$;
 $t = t + 1$;
 $k = k \bmod F$;
 if last job is not completed take appropriate action;
 if any of the job slices are not released, take appropriate action;
 execute the slices in the current block by waking of the periodic task server;
 sleep until the periodic task server completes;
 while (aperiodic job queue is not empty) {
 wake up the job at the head of the aperiodic job queue and schedule it;
 sleep until the aperiodic job completes;
 remove the aperiodic job from the queue;
 }
 sleep until next timer interrupt arrives;
 }
 End SCHEDULER

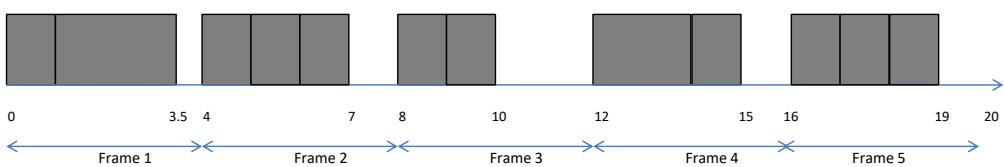
Scheduling Aperiodic Jobs

- Response time of aperiodic jobs is an important criterion
- Sooner the aperiodic job completes, the more responsive the system is.
- Hence improving average response time of all aperiodic jobs is a design goal
- One of the way is ‘Slack Stealing’

Slack Stealing

- There is no advantage completing a periodic task much before the deadline.
- Hence execute the aperiodic jobs ahead of periodic jobs whenever possible
- Every periodic job slice is scheduled in a frame that ends no later than its deadline.
- If the total time allocated to the job slices in a frame (k th frame) is x_k , the slack (time) available in the frame is equal to $(f - x_k)$.
- If the aperiodic job queue is not empty, then aperiodic jobs can be scheduled at the beginning of the frame for the slack i.e. $(f - x_k)$.
- If the aperiodic job queue is empty, then job slices from the next frame can be executed in the current block.

Example – Slack Stealing



- Above is the major cycle of a cyclic executive schedule.
- Three aperiodic jobs A1, A2, A3 arrived as shown in the table.

Aperiodic jobs	Release time	Execution time
A1	4	1.5
A2	9.5	0.5
A3	10.5	2

Example – Slack Stealing (contd.)



Cyclic Executive Schedule with Slack Stealing

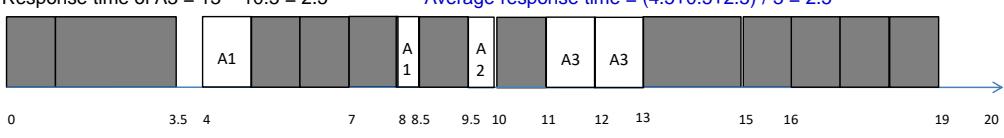
- Time 4: A1 is released, scheduler finds 1 slot slack in frame 2, so schedules A1
- Time 5: A1 is preempted and periodic job slices are scheduled
- Time 8: Next frame started, scheduler finds 2 slots of slack, so schedules A1
- Time 8.5: A1 is completed, so periodic job slices of frame 3 are scheduled
- Time 9.5: A2 is released, the scheduler calculates remaining slack 1.5, so schedules A2.
- Time 10: A2 is completed, so remaining job slices of frame 3 are scheduled
- Time 10.5: A3 is released
- Time 11: Job slices for frame 3 are completed, the scheduler calculates the remaining slack as 1 time slot, so A3 is scheduled
- Time 12: Next frame started, the scheduler calculates the slack as 1 time slot, so A3 continues to execute
- Time 13: Job slices of frame 4 should be scheduled. A3 also completes at this time

Response time of A1 = $8.5 - 4 = 4.5$

Response time of A2 = $10 - 9.5 = 0.5$

Response time of A3 = $13 - 10.5 = 2.5$

$$\text{Average response time} = (4.5+0.5+2.5) / 3 = 2.5$$



Example – Slack Stealing (contd.)

Cyclic Executive Schedule without Slack Stealing

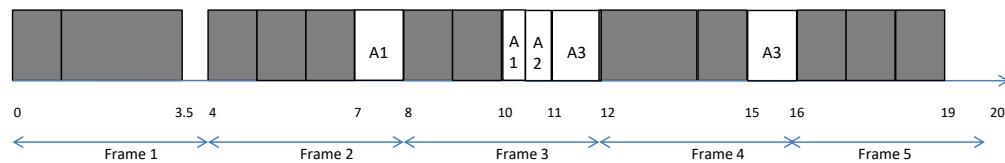
- Time 4: A1 is released
- Time 7: Slot is available, so A1 is scheduled
- Time 8: Next frame started, so A1 is preempted
- Time 9.5: A2 is released
- Time 10: A1 is scheduled
- Time 10.5: A1 completes, so A2 is scheduled
- Time 10.5: A3 is released
- Time 11: A2 completes, so A3 is scheduled
- Time 12: Next frame started, so A3 is preempted
- Time 15: A3 is scheduled
- Time 16: A3 completes.

Response time of A1 = $10.5 - 4 = 6.5$

Response time of A2 = $11 - 9.5 = 1.5$

Response time of A3 = $16 - 10.5 = 5.5$

$$\text{Average response time} = (6.5 + 1.5 + 5.5) / 3 = 4.5$$



Scheduling Sporadic Jobs

- Scheduler performs an acceptance test for the sporadic jobs before accepting them.
- Acceptance test is done at the frame boundaries.
- During Acceptance Test, the scheduler checks if the newly released sporadic job can be feasibly scheduled with all the jobs in the system (all the jobs means either the periodic jobs for which time has been allocated or sporadic jobs scheduled but not yet completed).
- If there is sufficient amount of time in the frames so that the newly released sporadic job can be scheduled and completes before its deadline without causing all the jobs in the system miss their deadlines, then the scheduler accepts the job. Otherwise it rejects it.
- We assume that maximum execution time of the sporadic jobs are known upon their releases.
- We also assume that all sporadic jobs are preemptable. Therefore a sporadic job can execute in more than one frames if no frame has sufficient time to accommodate the entire job.

Acceptance Test of Sporadic Jobs



- Let us suppose that at the beginning of frame ' t ', an acceptance test is done on a sporadic job ' S ' with deadline ' d ' and execution time ' e '.
- Suppose that the deadline ' d ' of S is in the frame ' $t+1$ ', and $t > t$.
- Then the job must be scheduled in the t^h or earlier frames.
- Hence the job can complete on time only if the current (total) amount of **slack time** in frames $t, t+1, \dots, l$

$$\sigma_c(t, l) \geq e$$

- The scheduler accepts the job if above condition is satisfied.

EDF Scheduling of Accepted Jobs

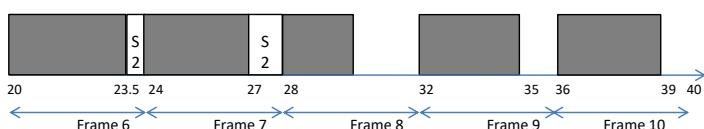
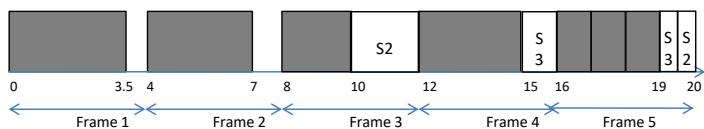


- EDF algorithm is good way to schedule sporadic jobs.
- Scheduler maintains a queue of accepted sporadic jobs in non decreasing order of their deadlines.
- Inserts newly accepted jobs into this queue in this order.
- Whenever all slices of the periodic tasks scheduled in each frame is completed, the scheduler schedules the **sporadic** job from head of the queue
- The scheduler allows **aperiodic** jobs to execute only if the accepted sporadic job queue is empty.

EDF Scheduling of Sporadic Jobs - Example



Sporadic jobs	Release time	Deadline	Execution time
S1	3	17	4.5
S2	5	29	4
S3	11	22	1.5



23
BITS Pilani, Pilani Campus

EDF Scheduling of Sporadic Jobs – Example (Contd)



- Time 3: S1 is released. Its deadline falls in frame 5. Acceptance test is done at starting of frame 2, since it is released at time 3, which falls inside frame 1. Total slack time between frame 2 and frame 4 is 4. Hence this job was rejected, since its execution time is 4.5, which is greater than total available slack.
- Time 5: S2 was released. Acceptance test is done at start of frame 3. Frames 3 through 7 ends before its deadline, having total slack time 5.5. So S2 was accepted.
- Time 10: S2 was scheduled, since there was a slack time.
- Time 11: S3 was released. Acceptance test is done at start of frame 4. Frame 4 to 5 ends before its deadline, having total slack time 2. We need to make sure the already accepted job S2 should also be able to meet its deadline. At start of Frame 3, S2 needs another 2 time slots (it has already completed 2 time slots. The slack available between frame 3 and frame 7 is 3.5. If we subtract S3's execution time 1.5 from it, remaining slack will be 2, which is sufficient for S2 to complete. So acceptance of S3 will not affect S2 meeting its deadline. So the scheduler accepts S3.

25
BITS Pilani, Pilani Campus

24
BITS Pilani, Pilani Campus

26
BITS Pilani, Pilani Campus

EDF Scheduling of Sporadic Jobs – Example (Contd)



- Time 12: S2 is preempted because new frame starts and job slices gets scheduled. Since S3 has earlier deadline than S2, S2 is put after S3 in the sporadic job queue.
- Time 15: S3 is scheduled, since job slices of frame 4 are done.
- Time 16: S3 is preempted because next frame started and its job slices are scheduled.
- Time 19: S3 is scheduled, since job slices of frame 4 are done.
- Time 19.5: S3 is completed, so S2 is scheduled.
- Time 20: S2 is preempted because next frame started and its job slices are scheduled.
- Time 23.5: S2 is scheduled, since job slices of frame 6 are done.
- Time 24: S2 is preempted because next frame started and its job slices are scheduled.
- Time 27: S2 is scheduled, since job slices of frame 7 are done.
- Time 28: S2 is completed

Practical Considerations



➤ Mode Changes

- During mode change, the system is reconfigured, some of the old tasks may be deleted, some new tasks may come.
- Assume that the timing parameters are known a priori for the tasks in the new mode as well.
- Mode change can be accomplished in two ways:
 - Consider it as an [Aperiodic job](#)
 - Mode change task is executed at highest priority among all aperiodic jobs.
 - If there are other aperiodic jobs, their execution is delayed after the mode change.
 - If there are sporadic jobs, delay the switching over to new schedule till the sporadic jobs are completed.
 - Acceptance test is suspended and the mode changer is scheduled during the time allocated for the deleted periodic tasks.
 - Once the mode changer task is executed, the acceptance test resumes at the start of the next frame.
 - Consider it as an [Sporadic job](#)
 - Mode change task is executed as any other Sporadic job.

➤ General Workloads and Multiprocessor Scheduling

- Clock driven approach can be used to do scheduling among multiple processors in a multi-processor system, whenever the workload parameters are known a priori.

Practical Considerations



➤ Handling Frame Overruns

- Unexpected reasons like a transient h/w or a s/w bug may cause a job to execute longer than expected and cross the frame boundary.
- A way to handle it is to simply abort the overrun job at the beginning of the next frame. Such a fault can be handled by a recovery mechanism.
- Another way ([which is used most frequently](#))
 - To preempt the overrun job immediately if it is not in critical section or once it exists from the critical section.
 - The unfinished job is scheduled later as an aperiodic job.

Pros & Cons of Clock-Driven Scheduling



Advantages

- Conceptual Simplicity: Since the schedule is statically determined,
 - Deadlocks can be avoided during scheduling
 - There will be no unpredictable delays
 - No need of concurrency control
 - No need for any synchronization mechanism
- Suitable for hard real-time systems, where all the parameters are known a priori

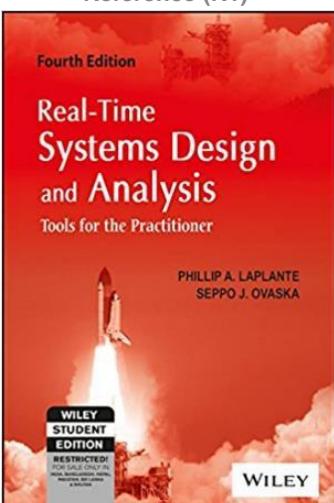
Disadvantages

- Release times of all jobs must be fixed
- All combinations of periodic tasks that might execute at the same time must be known a priori, so a schedule for the combination can be computed.
- Not suitable for systems which are combination of both hard real-time and soft real-time systems.

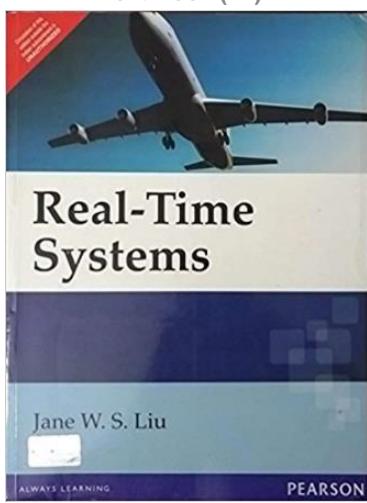
Any Questions?

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Text Book / References



Reference (R1)



Text Book (T1)

Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

BITS ZG553: Real Time Systems

BITS Pilani
Panaji|Doha|Gurgaon|Hyderabad

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Excellent MOOCs Videos (Coursera, edX,...)

HONOR CODE CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIx: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

HONOR CODE CERTIFICATE
Issued November 28, 2015

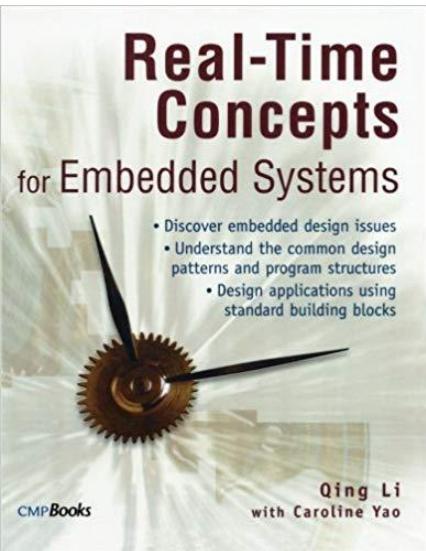
VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54



Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

RTS Primer – For Light Reading



- Discover embedded design issues
- Understand the common design patterns and program structures
- Design applications using standard building blocks

with Caroline Yao

4 BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

J_i - Parameters

Temporal
Functional
Resources
Interconnection



L-5: RTS Scheduling - Functional/Temporal Parameters, Feasibility & Optimality

Ref. T1/[C3, Lecture PPT/Notes]

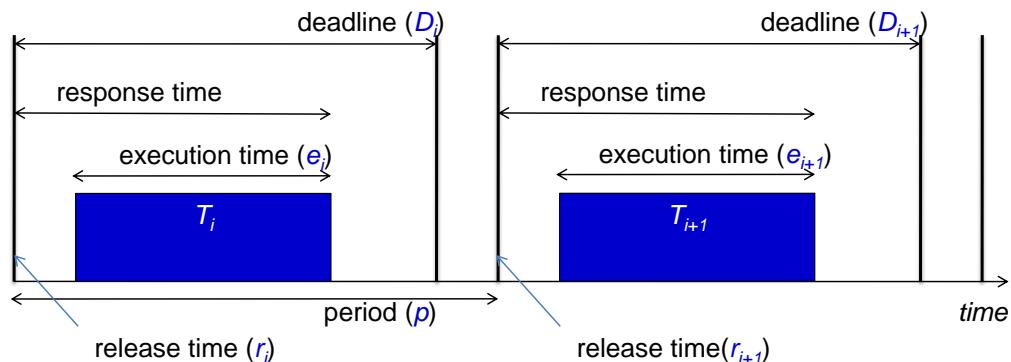
Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs, Save the Environment!

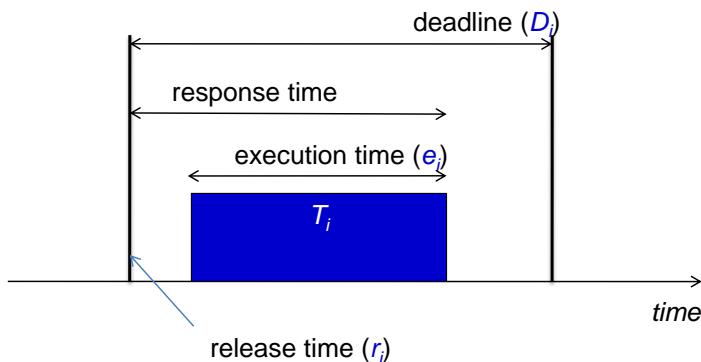
Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupama, BITS-Pilani WILP Faculty



Timing Specifications – Periodic Task



Timing Specifications – Aperiodic & Sporadic Task



Precedence Constraints, Predecessors and Successors



- Jobs are said to have **precedence constraints** if they are constrained to execute in some order.
- If jobs can execute in any order, they are called **independent**.
- A job J_i is a **predecessor** of another job J_k if J_k can't begin execution until the execution of J_i completes. It is denoted by

$$J_i < J_k$$

Here J_k is a **successor** of J_i .

Let us consider the relation $J_i < J_j < J_k$.

J_i is a **immediate predecessor** of J_j and predecessors of J_j and J_k

J_k is a **immediate successor** of J_j and successor of J_i and J_j

- A job with predecessors is ready for execution when the time is at or after its release time and all of its predecessors are completed.

Interconnection - Precedence Constraints

J_i - Parameters

Temporal
Functional
Resources
Interconnection

- Control
- Data

$$J_i < J_k$$

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



Precedence Graph

Precedence Constraints are represented through a graph called '**Precedence Graph**'.

How to draw a precedence graph :

- Each vertex in the graph represents a job.
- Above each vertex, the feasible interval for the job is mentioned in brackets.
- Draw a directed edge (arrowhead) from vertex J_i to J_k if J_i is a immediate predecessor of J_k .

Precedence Graph



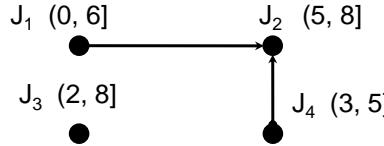
(0,7] (2,9] (4,11] (6,13] (8,15]



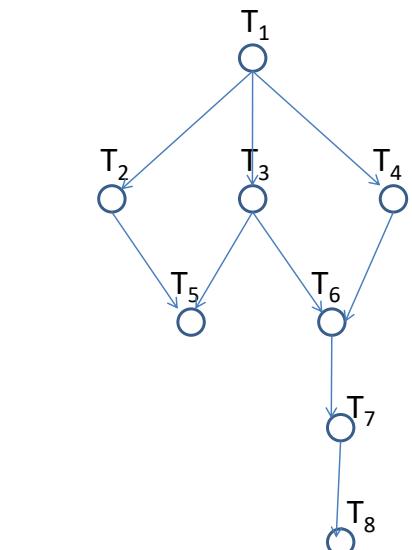
(2,5] (5,8] (8,11] (11,14] (14,17]



Precedence Graph - Example



- J_1, J_2, J_3 and J_4 have feasible intervals $(0, 6]$, $(5, 8]$, $(2, 8]$ and $(3, 5]$ respectively
- J_1 and J_4 are predecessors of J_2
- J_3 is neither a predecessor, nor a successor of any other jobs



Task Graph



Task Graph is like a Precedence Graph.

Each vertex represents a job, which is indicated by circles and squares.

Above each vertex, the feasible interval for the job is mentioned in brackets.

A sub-graph becomes a **chain**, when the jobs in the sub-graph need to be serially executed i.e. either $J_i < J_k$ or $J_i > J_k$.

Unlike precedence graph, the edges are of different types representing different types of dependencies.

The types of an edge connecting two vertices and other parameters of the edge are **interconnection parameters** of the jobs represented by the vertices.

Data Dependency

- In a **producer-consumer** environment, producer puts the data in a shared memory and the consumer consumes the data from the shared memory.
- The precedence graph will show that the producer and consumer are independent, since they are not explicitly constrained for execution.
- In task graphs, these two jobs will be connected via data-dependency edges.

AND / OR Precedence Constraints

- An **AND job** is the one which can begin execution at or after its release time provided all of its predecessors has been completed.
- **AND jobs** are represented by **unfilled circles** in the task graph.
- An **OR job** is the one which can begin execution at or after its release time provided **one or some of its immediate predecessors has been completed**.
- **OR jobs** are represented by **unfilled squares** in the task graph. The fraction of the number of predecessor to be completed is indicated below the vertex e.g. if 2 out of 3 jobs are to be completed, then put $2/3$ there.

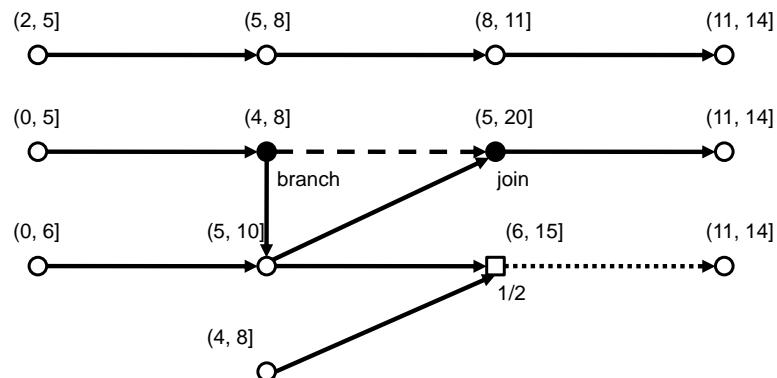
Temporal Dependency

- Difference in completion time of two jobs is called **temporal distance**.

Example: To have lip synchronization, the time between the display of each frame and the generation of corresponding audio must be now more than 160 ms.

- In task graph, temporal distance is represented by **temporal dependency edges**.
- There is a temporal dependency edge from vertex J_i to vertex J_k , if J_k has to be completed after a certain time after J_i completes. The value of this parameter is infinite if there is no temporal dependency between these jobs.

Task Graph - Example

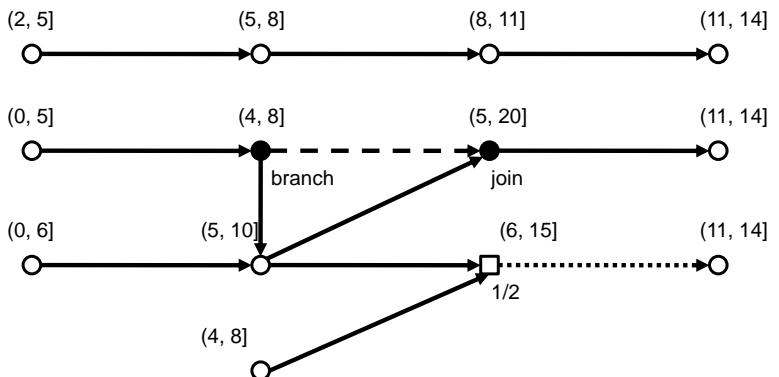


Conditional Branches

- The job from which the conditional branch arises is called a **branch** job.
- The job at which the conditional branches join is called a **join** job.
- Both these jobs are represented as **filled circles** in the task graph.
- The sub-graph that begins from the branch job and ends at the join job is called **conditional block**.

20
BITS Pilani, Pilani Campus

Task Graph - Example



21
BITS Pilani, Pilani Campus

Pipeline Relationship

- Pipelined relationship is used when between a pair of producer-consumer jobs, data can be piped.
- In the task graph, **the vertices are the granules of the producer and the consumer**.
- Each granule of the consumer can begin execution when the previous granule of the job and the corresponding granule of the consumer job is completed.
- This relationship is represented by a **dotted arrow** in the task graph.

21
BITS Pilani, Pilani Campus

Functional Parameters

- Preemptivity
- Criticality
- Optional Interval
- Laxity type

22
BITS Pilani, Pilani Campus

Preemptivity of Jobs

- A job is **preemptable** if its execution can be suspended at any point of time to allow execution of other jobs and later its execution can be resumed from the point of suspension.
- A job is **nonpreemptable** if it must be executed from start to end without any interruption.

Example: Interrupt Handling is non-preemptable.

Context Switching

- During preemption, the system must first save the state of the preempted job, so that it can resume the execution of it later. The state of a job typically includes
 - Program Counter Register
 - Stack Pointer Register
 - Other General Purpose Registers
 - ...
- After doing this, the system must load new processor status register, initialize the program counter, clear the processor pipeline etc.
- These actions are called **context switching**.
- Time taken for context switching is called **context switch time**.

Criticality of Jobs

- The criticality or importance of a job, is a positive number, which indicates **how critical the execution of the job is**, for the system.
- **Priority and Weights** are often used to refer to the criticality.
- During overload, when it is not possible to schedule all the jobs to meet their deadline, less critical jobs are sacrificed.

Optional Execution

- A job or a portion of a job can be called **optional**, if it completes late or not executed at all, the **system performance may degrade** but nevertheless the **function is satisfactory**.
- During overload, when it is not possible to complete all the jobs on time, we may choose to discard the optional jobs (i.e. leave them unexecuted or partially executed).
- The jobs and portion of the jobs which are not optional, are **mandatory**.

Laxity Type

- **Laxity** type of a job indicates whether the time constraints are **hard** or **soft**.
- It is represented by a '**usefulness function**', which gives the usefulness of the result produced by the job as the function of its **tardiness**.
- If the usefulness of a job becomes zero or negative as soon as the job becomes tardy, then it is better not to execute that job rather than completing it late.

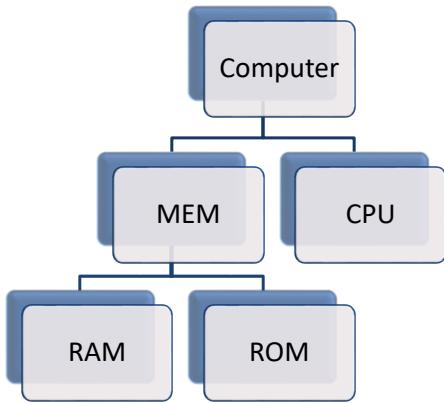
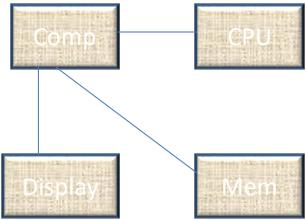
Resource Parameters

Premptivity of Resources

Exclusive/Shared Resources

Resource Graph

- Connectivity
- A part of edge



Feasibility, Optimality and Performance Measures

- A valid schedule is **feasible schedule** if every job completes by its deadline.
- A set of job is **schedulable** according to a scheduling algorithm if when using the algorithm the scheduler always produces a feasible schedule.
- **A hard real-time scheduling algorithm is said to be optimal** if using this algorithm the scheduler always produces a feasible schedule if the given set of jobs has feasible schedule.
- Other than the above other performance measures used are:

Maximum and average

- Tardiness,
- Lateness,
- Response time,
- Miss, Loss and Invalid rates.

Resource Parameter

Preemptivity of Resources

- A resource is **nonpreemptible** if each unit of the resource is constrained to be used serially. It means if a nonpreemptible resource is allocated to a job, other jobs needing it must wait until the current job completes its use.
- If jobs can use every unit of the resource in an interleaved manner, then the resource is **preemptible**.

Feasibility, Optimality and Performance Measures (contd.)

- **Lateness** of a job is the difference between its completion time and the deadline (i.e. completion time – deadline). **It can become negative if the job completes early and positive if it completes late.**
- For a set of jobs in the system, the response time of the job which completes last, is called **makespan** of the schedule.
- **Makespan is used to compare scheduling algorithm performance. The algorithm which produces a schedule with a shorter makespan is better.**
- **Most frequently used performance measure for jobs that have soft deadlines is their average response time.**
- For some **soft real-time systems**, it is acceptable to complete some jobs late or to discard late jobs. For those systems, a suitable performance measure include the **miss rate** (% of jobs those complete late) or **loss rate** (% of jobs those are discarded)
- **Invalid rate = Loss rate + Miss rate** : It should be minimized

What is scheduling ??

- Tasks have timing constraints
- Execution is bounded by max delay
- Ensure that time bound is respected as imperatively as possible
- Objective – allow tasks to fulfill timing constraints when running in nominal mode

Schedule

Valid schedule

A valid schedule satisfies the following requirements

1. every processor is assigned to at most one job at any time
2. every job is assigned to at most one processor at any time
3. No job is scheduled before its release time
4. Depending on the scheduling algo(s) used- total amount of processor time assigned to every job is equal to its max or actual execution time
5. All the precedence & resource usage constraints are satisfied

Some Terminology

- Valid Schedule
- Feasible Schedule
- Proficient Scheduler
- Optimal Scheduler
- Scheduling Pts.
- Pre-emptive Scheduling
- Utilization
- Jitter

Scheduling Algo. Taxonomy

Online/Offline Scheduling (Static/Dynamic)

- Deterministic
- Less complexity
- Flexibility
- Online-used
 - Dynamic variations in user demand/resources
 - Future loads unpredictable
 - More overhead
 - Less precise
 - Optimal schedule not possible
- Online – overloaded – clairvoyant scheduler!

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

36

Feasibility

Valid schedule is feasible if all tasks are completed on time

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

38

Scheduling algo. Taxonomy

Preemptive / Non-preemptive

- Preemptive disadvantages
 - Overhead Involved
- Non-preemptive advantages
 - Sharing of resources easier
- Non-preemptive disadvantages
 - No optimal schedule possible

Best Effort/ Timing Fault Intolerant

Centralized/Distributive

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

37

Performance Measure

Schedulable task set

Optimal

Lateness

Avg absolute lateness

Make span

Avg response Time

Miss Rate

Loss Rate

Invalid Rate

RTS- K.R.Anupama

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

39

Scheduler does not run continuously but at scheduling pts



Scheduler Implementations

Election Table
Priority Queuing list
Constant/ Varying Priority

- Constant
 - Completion time
 - Relative deadline
 - Period
- Variable
 - Pending computation time
 - Residual laxity
 - Absolute deadline /release time (effective)



Effective Release Time

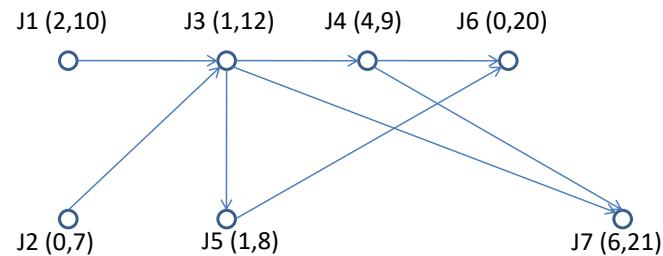
ERT = ART if no predecessor

ERT = max(ART, RT of all its predecessors)

Effective Deadline Time

EDT = ADT if no successor

EDT = min(ADT, DT of all its successors)



Does this satisfy basic requirements ??

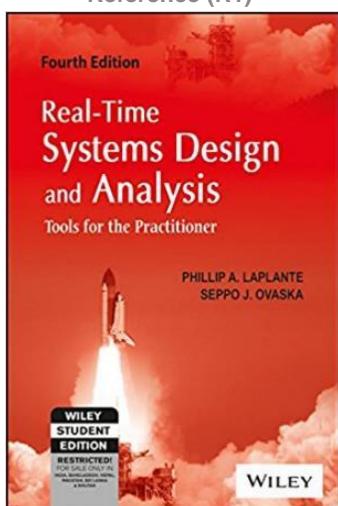
Thank You.

Any Questions?

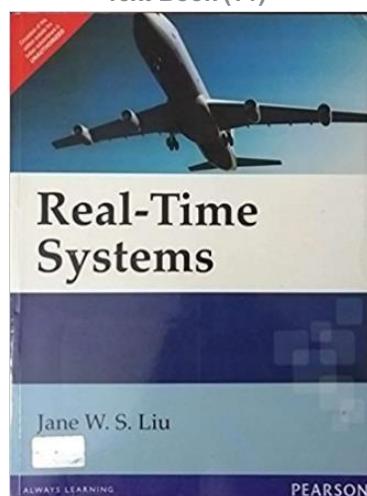
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

44

Text Book / References



Reference (R1)



Text Book (T1)

Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

2

BITS ZG553: Real Time Systems

[L6a – Performance, Optimality, Schedulability]



BITS Pilani
Pilani Campus

K G Krishna
WILP Division, BITS-Pilani, Hyderabad



Excellent MOOCs Videos

(Coursera, edX,...)

HONOR CODE CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.



Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

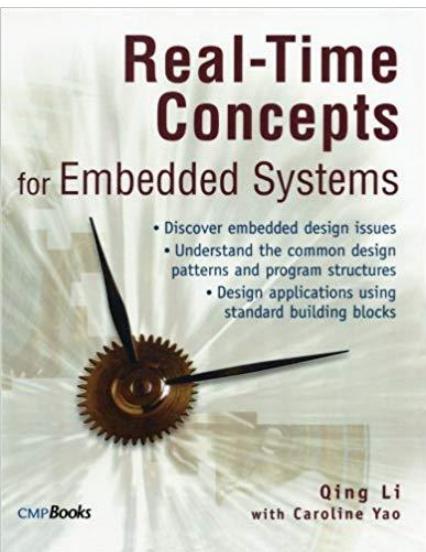
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eefb47d4e54

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

3



4 BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Three Approaches

- Clock-driven
- Round-robin / Weighted round-robin
- Priority-driven

L-6: Review of Common RTS Schedulers -

[/w Functional/Temporal Parameters, Feasibility & Optimality]

Ref.T1]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs, Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

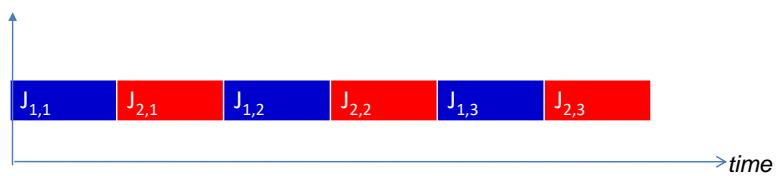
Clock-driven Approach

- Scheduling decisions are made at specific time instants, which are chosen a priori before the system begins execution
- Typically this type of scheduling is **suitable for hard real-time systems**, where the parameters are fixed and known.
- Scheduling decisions are computed off-line and stored for use at run-time, thus scheduling overhead is minimal.
- Generally a hardware time is set to expire periodically.
- After the system gets initialized, the scheduler selects and schedules jobs which execute till the next scheduling decision is made. Then the scheduler blocks itself waiting for the expiration of the timer.
- When the timer expires, the scheduler wakes up, does necessary scheduling and sleeps again. This process repeats.

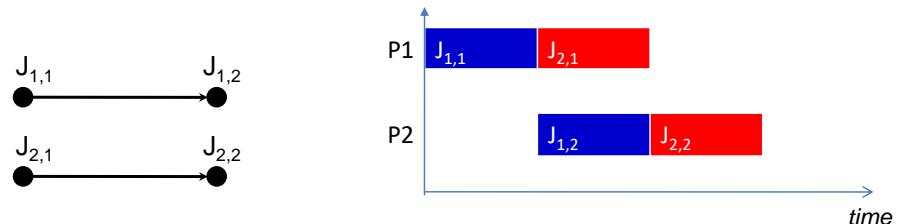
Round-robin Approach

- Also known as **time-sharing**
- Every job joins a FIFO (First-in-first-out) queue when it becomes ready for execution
- The entire time period is divided into several time-slices
- The job at the head of the queue executes for one time-slice.
- If the job doesn't complete at the end of the time-slice, it gets pre-empted and placed at the end of the queue to wait for its next turn.
- If there are ' n ' jobs ready for execution, each job gets $\frac{1}{n}$ th share of the processor.

Round-robin Approach - Example



(Round-robin execution of two tasks on a single processor)



(Round-robin execution of two tasks on two processors)

Weighted round-robin Approach

- This approach is a round robin approach with different weights assigned to different jobs.
- If a job has weight ' wt ', then it will get ' wt ' time slices every round for execution.

Priority Driven Approach

- Also known as **greedy scheduling**, **list scheduling** and **work-conserving scheduling**
- Priorities are assigned to the jobs based on their criticality
- Jobs ready for execution are placed in one or more queues ordered by priorities of the jobs.
- At any scheduling decision time, the jobs with the highest priorities are scheduled and executed on the available processors.
- Most scheduling algorithms used in non-real-time systems are priority driven.
 - FIFO (First In First Out)
 - LIFO (Last In First Out)
 - SETF (Shortest Execution Time First)
 - LETF (Longest Execution Time First)
- For jobs of same priority, round-robin scheduling is used

Priority-driven Scheduling Example



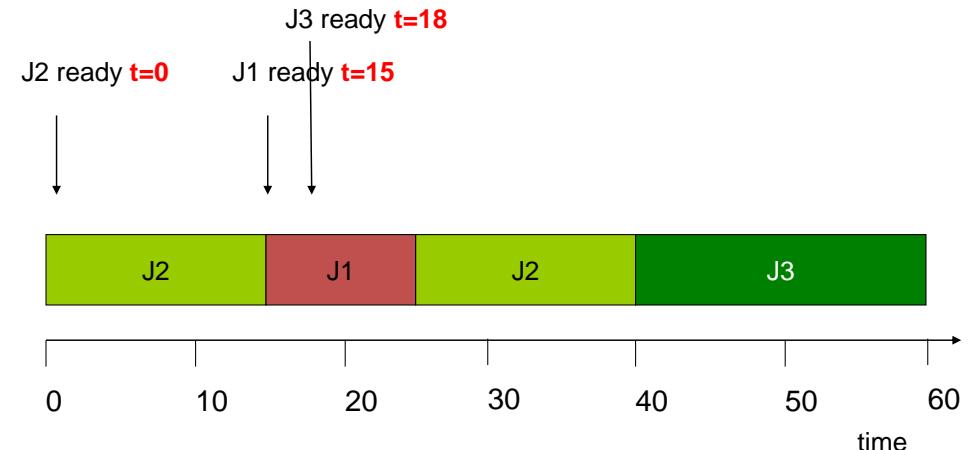
Rules:

- each process has a fixed priority (1 highest);
- highest-priority ready process gets CPU;
- process continues until done.

Processes

- J1: priority 1, release time 15, execution time 10
- J2: priority 2, release time 0, execution time 30
- J3: priority 3, release time 18, execution time 20

Priority-driven Scheduling Example



Effective Release Time and Deadline



- It may happen that a job has release time later than that of its successor and deadline earlier than that of its predecessor.
- Hence it may not be useful to work with individual jobs release time and deadlines.

Effective Release Time

- Effective Release Time of a job without any predecessor is its given release time.
- Effective Release Time of a job with predecessors is **maximum of its given release time and the effective release times of its predecessors**.

Effective Deadline

- Effective Deadline of a job without any successor is its given deadline.
- Effective Deadline of a job with successors is **minimum of its given deadline and the effective deadlines of its successors**.

By release time and deadline, we will mean effective release time and effective deadlines respectively.

Off-line vs On-line scheduling

Off-line scheduling: When the schedule is pre-computed and kept, it is called off-line scheduling.

- Example: [Clock-driven scheduling](#)
- It is possible only [when the system parameters are known a priori](#)
- Advantages: Deterministic timing behaviour, Lesser complexity, Very less scheduling overhead

On-line scheduling: When the scheduler makes each scheduling decision without knowledge about the jobs that will be released in the future

- Example: [Priority driven scheduling](#)
- It is the only option when [future workload is unpredictable](#)
- The price of the flexibility and adaptability is a reduced ability for the scheduler to come up with an optimal schedule making the best use of the system resources

Priority Driven vs Clock Driven approaches



- Priority driven approaches have many advantages compared to clock driven approach:
 - They **don't have to have the information on the release time, execution time etc** (in contrast with clock driven approach, where these parameters are required to be known *a priori*)
 - It is best suited for applications with **varying time and resource requirements**
 - Many well-known priority –driven algorithms use very simple priority assignments **reducing the overhead** of maintaining multiple queues.
- Despite all these advantages, **Clock-driven approaches** are used for hard real-time systems, especially in safety-critical systems.
 - **The major reason is that the timing behaviour of a priority-driven system is nondeterministic when job parameters vary.**
 - Consequently it is difficult to validate the deadlines of all the jobs in a priority-driven approach that they meet the deadline, when job parameters vary.

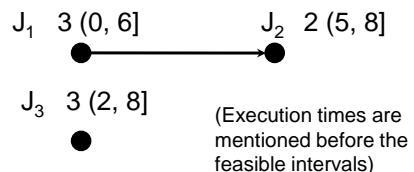
16
BITS Pilani, Pilani Campus

EDF (Earliest Deadline First) Algorithm

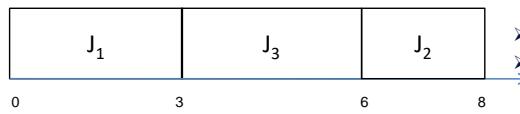


- Job with earliest (absolute) deadline has highest priority (Process closest to its deadline has highest priority)
- Priorities are assigned dynamically, since deadlines of the jobs varied. So it is a **dynamic-priority algorithm**.

Example



- t = 0: J₁ is released, no other job in the system, so gets scheduled
- t = 2: J₃ is released.
 - Deadline of J₁ = 6,
 - Deadline of J₃ = 8.
 - So J₁ has higher priority, hence it continues
- t = 3: J₁ completes, J₃ starts
- t = 5: J₂ is released
 - Deadline of J₂ = 8
 - Deadline of J₃ = 8
 - So both have same priority, let J₃ continue.
- t = 6: J₃ is done, J₂ gets scheduled
- t = 8: J₂ is done



Overload condition



- A system is said to be **overloaded** when the job offered to the scheduler can't be feasibly scheduled by a clairvoyant scheduler.
- When the system is not overloaded, an optimal on-line scheduling algorithm is one that always produces a feasible schedule of all offered job.
- No optimal on-line scheduling algorithm exists when some jobs are non-preemptable.
- During an overload, some jobs must be discarded in order to allow other jobs to complete on time.

18
BITS Pilani, Pilani Campus

Optimality of EDF Algorithm



Theorem:

When preemption is allowed and jobs don't contend for resources, the EDF algorithm can produce a feasible schedule if a set **J** of jobs with arbitrary release times and deadlines on a processor if and only if **J** has feasible schedules.

19
BITS Pilani, Pilani Campus

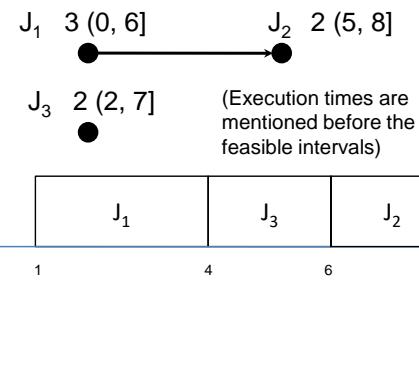
10

20
BITS Pilani, Pilani Campus

LRT (Latest Release Time) Algorithm

- There is no advantage in finishing early in a hard-real time system.
- LRT Algorithm takes the advantage of this fact.
- It treats release time as deadline and deadline as release time and schedule the jobs backward starting from latest deadline in 'priority-driven' manner.

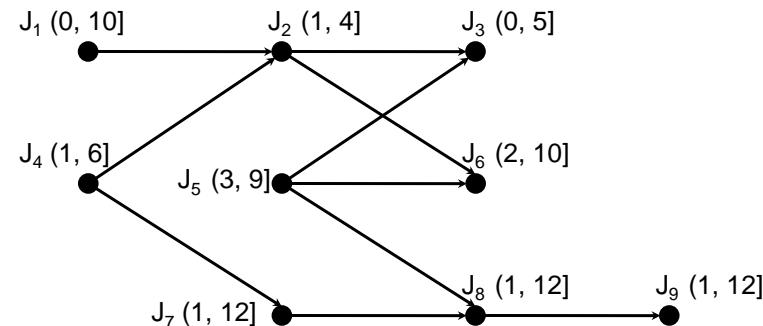
Example



- Latest deadline is 8. So start scheduling backward from time 8.
- J₂ should start at time 6, so that it completes at time 8 (=6 + execution time 2)
- At time 7, J₃ is ready to be scheduled (its deadline is 7). So it gets scheduled at time 6 (after J₂).
- J₃ should start at time 4, so that it completes at time 6 (=4 + execution time 2)
- At time 6, J₁ is ready to be scheduled. So it gets scheduled at time 4 after J₃.
- J₁ should start at time 1, so that it completes at time 4 (=1 + execution time 3)

BITs Pilani, Pilani Campus

Exercise



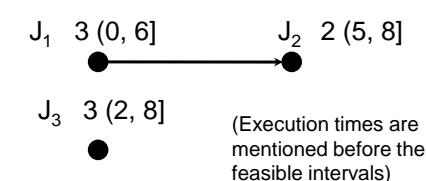
The feasible interval of each job in the precedence graph is given next to its name. The execution time of all jobs are equal to 1.

- Find out the effective release times and deadlines of the jobs
- Find an EDF schedule of the jobs

LST (Least Slack Time First) Algorithm

- Also known as **Minimum-Laxity-First (MLF) Algorithm**
- **Slack (or laxity)** of a job at time 't' = $d - t - \text{time required to complete the remaining portion of the job}$, where 'd' is the deadline
- Smaller the slack, higher the priority.
- As long as a job is executing the slack at time 't' = $d - t - (r + e - t) = d - r - e$, So slack is same as long as the job is executing. It changes after preemption of the job

Example



- t = 0: J₁ is released, no other job in the system, so gets scheduled
- t = 2: J₃ is released
 - Slack of J₁ = $6 - 2 - (3 - 2) = 3$,
 - Slack of J₃ = $8 - 2 - 3 = 3$.
 - So both J₁ and J₃ are of same priority. Let J₁ continue
- t = 3: J₁ completes, J₃ starts.
- t = 5: J₂ is released
 - Slack of J₂ = $8 - 5 - 2 = 1$
 - Slack of J₃ = $8 - 5 - (3 - 2) = 2$
 - So J₂ will have higher priority, hence J₂ will preempt J₃.
- t = 7: J₂ is done, J₃ gets scheduled
 - Slack of J₃ = $8 - 7 - (3 - 2) = 0$
- t = 8: J₃ is done

BITs Pilani, Pilani Campus

Exercise – Answer (a)

Effective Release Times

- J₁, J₄ and J₅ has no predecessors. So effective release times for these jobs are same as their own release times.
- Then traverse the diagram from left.
- Effective release time of other jobs = Max (own release time, effective release times of the predecessors)
 - Effective Rel Time (J₂) = Max (0, 1, 1) = 1
 - Effective Rel Time (J₇) = Max (1, 1) = 1
 - Effective Rel Time (J₃) = Max (0, 1, 3) = 3
 - Effective Rel Time (J₆) = Max (2, 3, 3) = 3
 - Effective Rel Time (J₈) = Max (1, 3) = 3
 - Effective Rel Time (J₉) = Max (1, 3) = 3

Jobs	Effective Release Times
J ₁	0
J ₂	1
J ₃	3
J ₄	1
J ₅	3
J ₆	3
J ₇	1
J ₈	3
J ₉	3

BITs Pilani, Pilani Campus

23

BITs Pilani, Pilani Campus

24

Exercise – Answer (a)

Effective Deadlines

- J3, J6 and J9 has no successors. So effective deadlines for these jobs are same as their own deadlines.
- Then traverse the diagram from right.
- Effective deadline of other jobs = Min (own deadline, effective deadlines of the successors)
- Effective Deadline(J8) = Min (12, 12) = 12
- Effective Deadline(J2) = Min (4, 5, 10) = 4
- Effective Deadline(J5) = Min (9, 5, 10, 12) = 5
- Effective Deadline(J7) = Min (12, 12) = 12
- Effective Deadline(J4) = Min (6, 4, 12) = 4
- Effective Deadline(J1) = Min (10, 4) = 4

Jobs	Effective Deadlines
J1	4
J2	4
J3	5
J4	4
J5	5
J6	10
J7	12
J8	12
J9	12

24

BITS Pilani, Pilani Campus

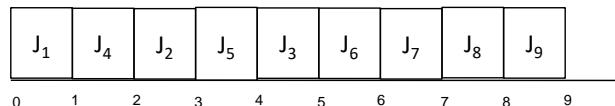
Thank You.

Any Questions?

Exercise – Answer (b)

EDF Schedule

- Time 0: J1 is released and scheduled.
- Time 1: J2, J4, J7 gets released. J2 and J4 have same deadlines, which is earlier than that of J7. But J4 is predecessor of J2. Hence we have to schedule J4.
- Time 2: Out of J2 and J7, J2 has earlier deadline i.e. 4, so gets scheduled.
- Time 3: All other jobs gets released. J3 and J5 have earlier deadlines i.e. 5. But J5 is predecessor of J3. So we have to schedule J5.
- Time 4: J3 gets scheduled, since it has the earliest deadline i.e. 5.
- Time 5: J6 gets scheduled, since it has the earliest deadline i.e. 10.
- Time 6, 7, 8: J7, J8, J9 have same deadlines i.e. 12, so same priority. But J7 is predecessor of J8 and J8 is predecessor of J9. So they have to be scheduled in the order J7, J8 and J9.



25

BITS Pilani, Pilani Campus

BITS ZG553: Real Time Systems

[L6: Overview of RTS Schedulers,
Performance, Optimality, Schedulability]



BITS Pilani
Pilani Campus

K G Krishna
WILP Division, BITS-Pilani, Hyderabad





L-6: Review of RTS Schedulers - (Feasibility, Optimality & Schedulability)

[Ref: T1/C4]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs, Save the Environment!

2

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

Off-line vs On-line scheduling

Off-line scheduling: When the schedule is pre-computed and kept, it is called off-line scheduling.

- Example: Clock-driven scheduling / Table-driven Scheduling / Round-Robin (Time-Slicing) / Weighted Round-Robin
- It is possible only when the system parameters are known a priori
- Advantages: Deterministic timing behaviour, Lesser complexity, Very less scheduling overhead

On-line scheduling: When the scheduler makes each scheduling decision without knowledge about the jobs that will be released in the future

- Example: Priority driven scheduling
- It is the only option when future workload is unpredictable
- The price of the flexibility and adaptability is a reduced ability for the scheduler to come up with an optimal schedule making the best use of the system resources

Task (T1) vs. Jobs (J1,1 → J1,2, → J1,3 → ..)

Remember:

- Tasks are composed of Individual Jobs that get released at intervals
- Each Task has parameters (Period, Execution, Relative Deadline, Phase)
- **Dynamic Priority (Scheduler-driven Priority/On-line Scheduling):** Priorities of individual Jobs for any Tasks can be changing at run-time (not fixed by Design or at the Beginning); Job parameters for a given Task may vary,...
- **Fixed Priority (Static Priority / Off-line Scheduling):** All Jobs of a Task are given the same priority assigned at the beginning (design time) and can't be changed (by scheduler)

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Priority Driven vs Clock Driven approaches

➤ Priority driven approaches have many advantages compared to clock driven approach:

- They don't have to have the information on the release time, execution time etc (in contrast with clock driven approach, where these parameters are required to be known a priori)
- It is best suited for applications with varying time and resource requirements
- Many well-known priority –driven algorithms use very simple priority assignments reducing the overhead of maintaining multiple queues.

➤ Despite all these advantages, Clock-driven approaches are used for hard real-time systems, especially in safety-critical systems.

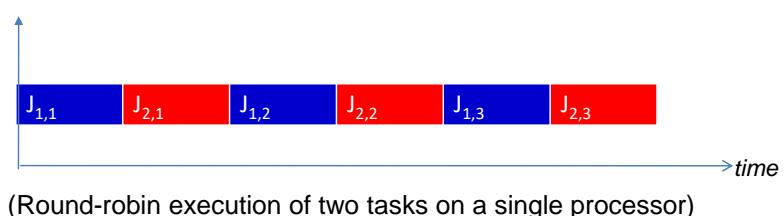
- The major reason is that the timing behaviour of a priority-driven system is nondeterministic when job parameters vary.
- Consequently it is difficult to validate the deadlines of all the jobs in a priority-driven approach that they meet the deadline, when job parameters vary.

Clock-driven Approach

- ❑ Scheduling decisions are made at specific time instants, which are chosen a priori before the system begins execution
- ❑ Typically this type of scheduling is **suitable for hard real-time systems**, where the parameters are fixed and known.
- ❑ Scheduling decisions are computed off-line and stored for use at run-time, thus scheduling overhead is minimal.
- ❑ **Generally a hardware timer is set to expire periodically.**
- ❑ After the system gets initialized, the scheduler selects and schedules jobs which execute till the next scheduling decision is made. Then the scheduler blocks itself waiting for the expiration of the timer.
- ❑ When the timer expires, the scheduler wakes up, does necessary scheduling and sleeps again. This process repeats.

BITS Pilani, Pilani Campus

Round-robin Approach - Example



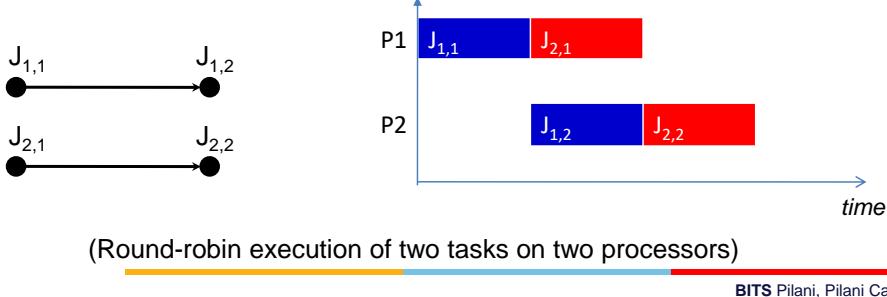
Round-robin Approach

- ❑ Also known as **time-sharing**
- ❑ Every job joins a FIFO (First-in-first-out) queue when it becomes ready for execution
- ❑ The entire time period is divided into several time-slices
- ❑ The job at the head of the queue executes for one time-slice.
- ❑ If the job doesn't complete at the end of the time-slice, it gets pre-empted and placed at the end of the queue to wait for its next turn.
- ❑ If there are ' n ' jobs ready for execution, each job gets $1/n$ th share of the processor.

BITS Pilani, Pilani Campus

Weighted round-robin Approach

- ❑ This approach is a round robin approach with different weights assigned to different jobs.
- ❑ If a job has weight ' wt ', then it will get ' wt ' time slices every round for execution.



BITS Pilani, Pilani Campus

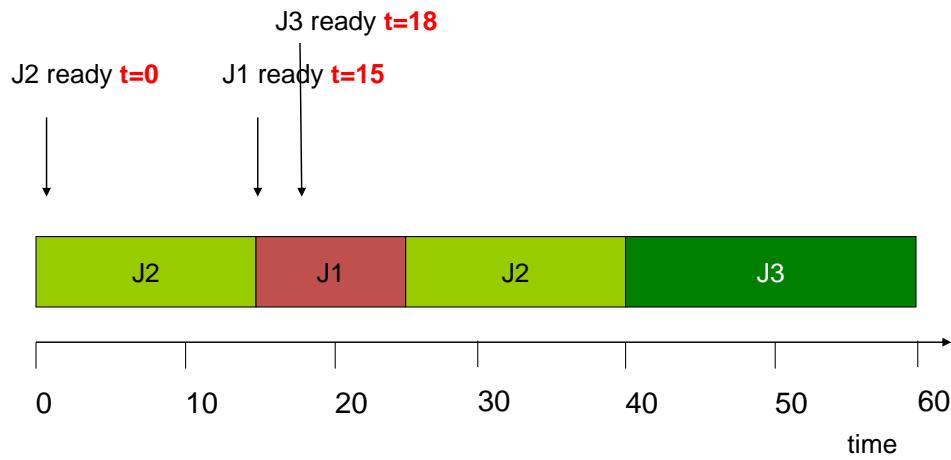
BITS Pilani, Pilani Campus

Priority Driven Approach

- Also known as **greedy scheduling**, **list scheduling** and **work-conserving scheduling**
- Priorities are assigned to the jobs based on their criticality
- Jobs ready for execution are placed in one or more queues ordered by priorities of the jobs.
- At any scheduling decision time, the jobs with the highest priorities are scheduled and executed on the available processors.
- Most scheduling algorithms used in non-real-time systems are priority driven.
 - FIFO (First In First Out)
 - LIFO (Last In First Out)
 - SETF (Shortest Execution Time First)
 - LETF (Longest Execution Time First)
- For jobs of same priority, round-robin scheduling is used

10
BITS Pilani, Pilani Campus

Priority-driven Scheduling Example



12
BITS Pilani, Pilani Campus

Priority-driven Scheduling Example

Rules:

- each process has a fixed priority (1 highest);
- highest-priority ready process gets CPU;
- process continues until done.

Processes

- J1: priority 1, release time 15, execution time 10
- J2: priority 2, release time 0, execution time 30
- J3: priority 3, release time 18, execution time 20

11
BITS Pilani, Pilani Campus

Assumptions

- Tasks are independent
- There are no aperiodic and sporadic tasks
- Every job is ready for execution as soon as released
- A job can be preempted any time
- A job never suspends itself
- Scheduling decisions are made immediately upon the job releases and completions
- Context switch overhead is negligibly small compared with the execution times of the tasks
- Number of priority levels are unlimited
- Number of periodic tasks are fixed

When an application creates a new task,

- The application first requests the scheduler to add a new task by providing the scheduler with relevant parameters of the task including period, execution time and relative deadline.
- Then the scheduler does acceptance test to check if the new task can be feasibly scheduled with all other existing tasks. If it is not feasible, the scheduler rejects the task.

13
BITS Pilani, Pilani Campus

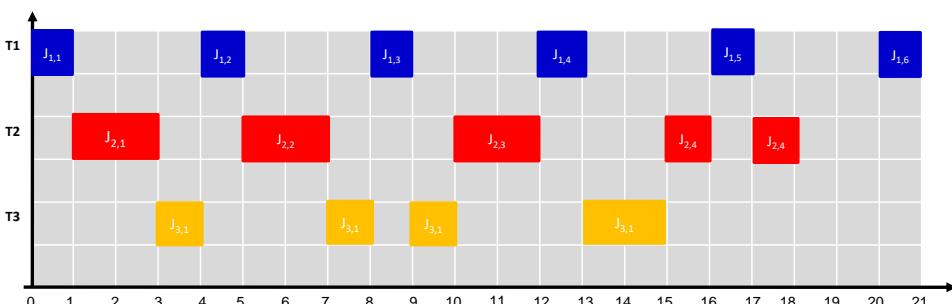
Fixed Priority vs Dynamic Priority Algorithms



- ❑ A priority driven scheduler is an on-line scheduler. It doesn't pre-compute the schedule.
- ❑ It assigns priorities to jobs after they are released, and places them in appropriate queues.
- ❑ **Fixed priority algorithm:** A fixed priority algorithm assigns the same priority to all the jobs in each task.
- ❑ **Dynamic priority algorithm:** A dynamic priority algorithm assigns different priorities to the individual jobs in each task. So priorities of a task can change as the jobs of that task are released.

Generally the real time algorithms of practical interest are fixed priority algorithms.

Rate Monotonic (RM) Algorithm



(A different representation of the schedule)

Rate Monotonic (RM) Algorithm



- ❑ Fixed priority algorithm
- ❑ **Shorter the period, higher the priority**
- ❑ Rate is inverse of the period. Hence higher the rate, higher the priority. – so the name '**rate monotonic**'

Example: 3 tasks T1 = (4,1), T2 = (5, 2), T3 = (20, 5) to be scheduled based on RM algorithm.

- T1 has shortest period (i.e. 4), so should have higher priority, followed by T2 and T3.
- T1 gets scheduled at after 4 time units i.e. at times 0, 4, 8, 12, 16, 20, ...
- T2 gets scheduled at time units 1, 5, 9, 13 for 2 time slots. When T2 gets released at time 15, it is given this slot, preempting T3. But at time 16, T1 gets released. It gets scheduled preempting T2. Once T1 is done, T2 again gets scheduled at time 18.
- T3 gets scheduled in the remaining slots 3, 7, 11. It gets time slot 9, since T1 is done and T2 is not released that time. But it gets preempted in the next slot because T2 gets released. Again it gets scheduled in slot 13 and 14. With this T3 completes for one period 20.



Deadline Monotonic (DM) Algorithm

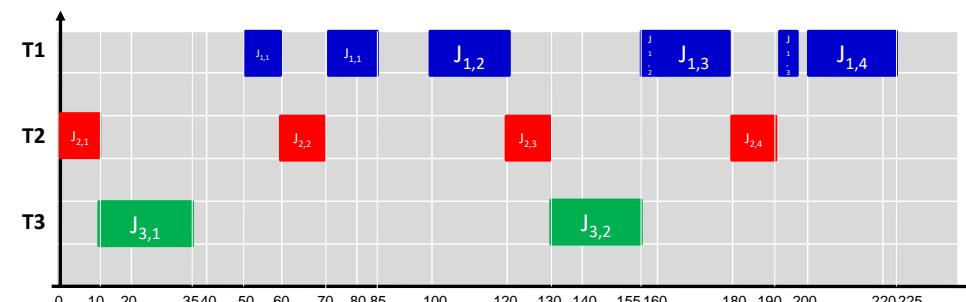


- ❑ Fixed priority algorithm
- ❑ **Shorter the relative deadline, higher the priority**
- ❑ When the relative deadlines of every task is proportional to their period, the schedule produced by RM and DM algorithms are identical.
- ❑ But when the relative deadlines are arbitrary, DM may produce a feasible schedule when RM fails.

Example:

T1 = (50, 50, 25, 100), T2 = (0, 60, 10, 20), T3 = (0, 125, 25, 50)

According to DM algorithm, T2 has highest priority because its relative deadline is 20.
Similarly T1 has lowest priority and T3 has priority in between.

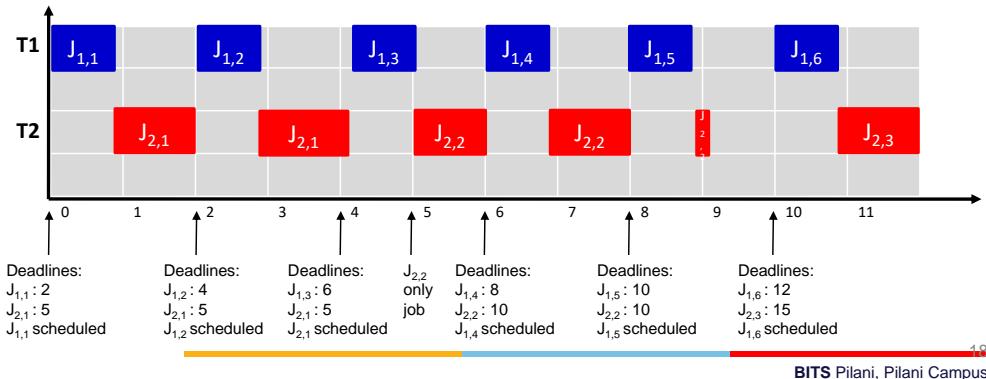


EDF (Earliest Deadline First) Algorithm



- Dynamic priority algorithm
- Job with earliest (absolute) deadline has highest priority

Example: Consider two tasks T1 = (2, 0.9) and T2 = (5, 2.3)



EDF (Earliest Deadline First) Algorithm

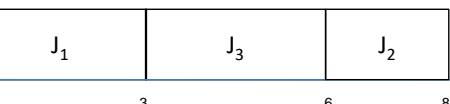
- Job with earliest (absolute) deadline has highest priority (Process closest to its deadline has highest priority)
- Priorities are assigned dynamically, since deadlines of the jobs varied. So it is a *dynamic-priority algorithm*.

Example (Jobs with Precedence)

J₁ 3 (0, 6] J₂ 2 (5, 8]

J₃ 3 (2, 8]

(Execution times are mentioned before the feasible intervals)



- t = 0: J1 is released, no other job in the system, so gets scheduled
- t = 2: J3 is released.
 - Deadline of J1 = 6,
 - Deadline of J3 = 8.
 - So J1 has higher priority, hence it continues
- t = 3: J1 completes, J3 starts.
- t = 5: J2 is released
 - Deadline of J2 = 8
 - Deadline of J3 = 8
 - So both have same priority, let J3 continue.
- t = 6: J3 is done, J2 gets scheduled
- t = 8: J2 is done

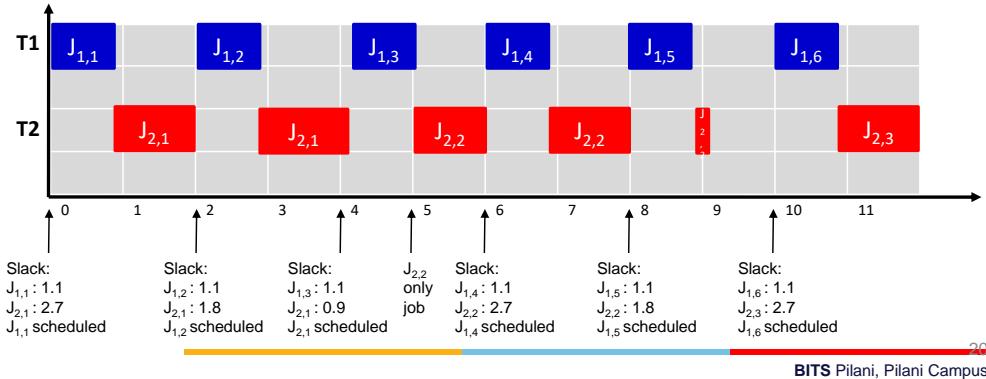
BITs Pilani, Pilani Campus

LST (Least-Slack-Time-First) Algorithm



- Task level dynamic algorithm, and job level dynamic algorithm
- Job with smallest slack has highest priority
- At time t , the slack of a job whose remaining execution time is x and whose deadline is d
 $= d - t - x$

Example: Consider two tasks T1 = (2, 0.9) and T2 = (5, 2.3)



LST (Least Slack Time First) Algorithm (Example)

Jobs with Precedence

J₁ 3 (0, 6] J₂ 2 (5, 8]

J₃ 3 (2, 8]

(Execution times are mentioned before the feasible intervals)



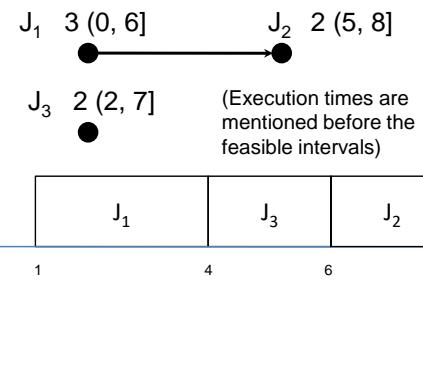
- t = 0: J1 is released, no other job in the system, so gets scheduled
- t = 2: J3 is released.
 - Slack of J1 = 6 - 2 - (3 - 2) = 3,
 - Slack of J3 = 8 - 2 - 3 = 3.
 - So both J1 and J3 are of same priority. Let J1 continue
- t = 3: J1 completes, J3 starts.
- t = 5: J2 is released
 - Slack of J2 = 8 - 5 - 2 = 1
 - Slack of J3 = 8 - 5 - (3 - 2) = 2
 - So J2 will have higher priority, hence J2 will preempt J3.
- t = 7: J2 is done, J3 gets scheduled
 - Slack of J3 = 8 - 7 - (3 - 2) = 0
- t = 8: J3 is done

BITs Pilani, Pilani Campus

LRT (Latest Release Time) Algorithm

- There is no advantage in finishing early in a hard-real time system.
- LRT Algorithm takes the advantage of this fact.
- It treats release time as deadline and deadline as release time and schedule the jobs backward starting from latest deadline in 'priority-driven' manner.

Example (Jobs with Precedence)



- Latest deadline is 8. So start scheduling backward from time 8.
- J₂ should start at time 6, so that it completes at time 8 (=6 + execution time 2)
- At time 7, J₃ is ready to be scheduled (its deadline is 7). So it gets scheduled at time 6 (after J₂).
- J₃ should start at time 4, so that it completes at time 6 (=4 + execution time 2)
- At time 6, J₁ is ready to be scheduled. So it gets scheduled at time 4 after J₃.
- J₁ should start at time 1, so that it completes at time 4 (=1 + execution time 3)

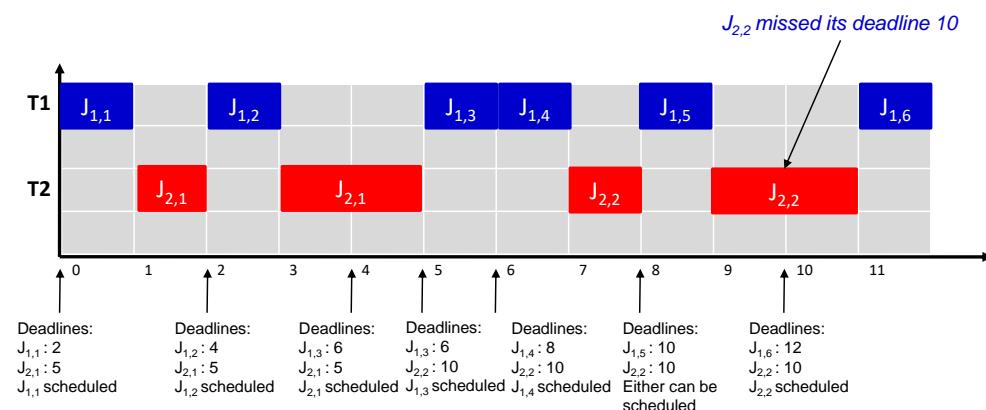
22 BITS Pilani, Pilani Campus

Schedulable Utilization

Example for $U > 1$:

Consider EDF schedule for tasks: T₁ = (2,1), T₂ = (5,3)

Total utilization $U = 1/2 + 3/5 = 0.5 + 0.6 = 1.1$



23 BITS Pilani, Pilani Campus

Schedulable Utilization

A scheduling algorithm can feasibly schedule any set of periodic task on a processor if the total utilization of the tasks is equal to or less than the schedulable utilization of the algorithm.

No algorithm can schedule a set of tasks with a total utilization greater than 1

23 BITS Pilani, Pilani Campus

Relative Merits

- Criterion to measure performance: Schedulable Utilization
- So higher the schedulable utilization, better is the algorithm.
- An algorithm whose schedulable utilization is 1, is an optimal algorithm.
- Optimal dynamic-priority algorithms outperforms fixed-priority algorithms in terms of schedulable utilization.
- But advantage of fixed-priority algorithms is predictability.

25 BITS Pilani, Pilani Campus

Schedulable Utilization of EDF Algorithm



Theorem:

A system T of independent, preemptable tasks with relative deadlines equal to their respective periods can be feasibly scheduled on one processor if and only if its total utilization is equal to or less than 1.

Schedulable utilization of the LST algorithm is also 1.

Schedulability Test of EDF Algorithm

Schedulability condition for EDF algorithm:

$$\sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} \leq 1$$

If above condition is satisfied, the system is schedulable according to EDF algorithm.

Optimality of RM & DM Algorithms



- Since these algorithms assign fixed priorities, they can't be optimal.
- While RM algorithm is not optimal for arbitrary periods, it is optimal in the special case when the periodic tasks in the system are simply periodic.
- A system of periodic tasks is **simply periodic** if for every pair of tasks T_i and T_k in the system and $p_i < p_k$, p_k is an integer multiple of p_i .
- In other words, for simply periodic tasks, the period of all tasks are integer multiple of the shortest period.

Optimality of RM Algorithm



Theorem:

A system of simply periodic, independent, preemptable tasks whose relative deadlines are equal to or larger than their periods is schedulable on one processor according to the RM algorithm if and only if its total utilization is equal to or less than 1.

RM Algorithm is optimal among all fixed-priority algorithms whenever the relative deadlines of the tasks are proportional to their periods.



Optimality of DM Algorithm

Theorem:

A system T of periodic, independent, preemptable tasks that are in phase and have relative deadlines equal to or less than their periods can be feasibly scheduled on one processor according to the DM algorithm whenever it can be feasibly scheduled according to any fixed-priority algorithms.

Sufficient schedulability condition for RM algorithm

Theorem:

A system of ' m ' independent, preemptable periodic tasks with relative deadlines equal to their respective periods can be feasibly scheduled on a processor according to RM algorithm if its total utilization ' U ' is less than or equal to

$$U_{RM} = m(2^m - 1)$$

As number of tasks approaches infinity, schedulable bound

$$U_{RM} = \lim_{m \rightarrow \infty} m(2^m - 1) = \ln 2 = 0.69 = 69\%$$

$U(n) \leq U_{RM}(n)$ is not a necessary condition, a system of tasks may nevertheless be scheduled even when its total utilization exceeds the schedulable bound $U_{RM}(n)$.

Lehoczky's Schedulability Test for Fixed-Priority Algorithms

Used when Sufficient Schedulability Condition fails

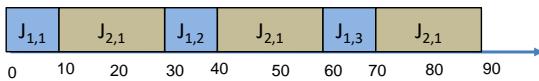
Theorem

A set of periodic real-time tasks is RMA schedulable under any task phasing, if all the tasks meet their respective first deadlines under zero phasing (i.e. when all tasks have phases equal to 0).

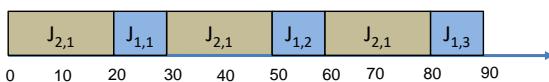
Explanation

Let there are two tasks $T_1 = (30, 10)$ and $T_2 = (120, 60)$ scheduled as per RM Algorithm.

Then T_1 has higher priority than T_2 .



(T1 and T2 have 0 phases – First job of T2 finishes at 90)



(T1 has phase 20 and T2 has phase 0 – First job of T2 finishes at 80)

From this example, it is obvious that worst case response time occurs for a lower priority task, the phase of this task and the phases of all other higher priority task are 0.

Lehoczky's Schedulability Test (Contd.)

As seen in the example, within deadline of the first job of T_2 i.e. 120, higher priority task T_1 can be scheduled $120 / 30 = 4$ times.

So T_2 has to wait for 4×10 (execution time of each job of T_1) time slots during execution of its first job.

Hence in the worst case, the amount of time a low priority task T_i has to wait due to the higher priority tasks $(T_1, T_2, \dots, T_{i-1})$ in the system is

$$\sum_{k=1}^{i-1} \left\lceil \frac{D_i}{P_k} \right\rceil e_k$$

So in worst case, T_i will be in the system for

$$e_i + \sum_{k=1}^{i-1} \left\lceil \frac{D_i}{P_k} \right\rceil e_k$$

Then for all the tasks to be feasibly scheduled, this time should be less than or equal to the respective deadline i.e.

$$e_i + \sum_{k=1}^{i-1} \left\lceil \frac{D_i}{P_k} \right\rceil e_k \leq D_i$$

This is Lehochky's Schedulability Test.

Floor and Ceiling Functions

Floor

$\text{floor}(x) = \lfloor x \rfloor$ is the largest integer not greater than x

Ceiling

$\text{ceiling}(x) = \lceil x \rceil$ is the smallest integer not less than x

Example:

x	$\lfloor x \rfloor$	$\lceil x \rceil$
2.4	2	3
5.5	5	6
-2.1	-3	-2
-2	-2	-2

Example (contd.)

As per RMA, the priorities of these tasks are $T_1 > T_2 > T_3$.

For T_1 , the execution time 3 is less than its deadline 8, so it is schedulable.

For T_2 , the time it will be in the system in worst case scenario =

$$e_i + \sum_{k=1}^{i-1} \left\lceil \frac{D_i}{P_k} \right\rceil e_k = 3 + \left\lceil \frac{9}{8} \right\rceil \times 3 = 3 + 2 \times 3 = 9$$

T_2 is schedulable since this time is equal to its deadline.

For T_3 , the time it will be in the system in worst case scenario =

$$e_i + \sum_{k=1}^{i-1} \left\lceil \frac{D_i}{P_k} \right\rceil e_k = 3 + \left\lceil \frac{15}{8} \right\rceil \times 3 + \left\lceil \frac{15}{9} \right\rceil \times 3 = 3 + 2 \times 3 + 2 \times 3 = 15$$

T_3 is schedulable since this time is equal to its deadline.

Hence all the 3 tasks pass Lehoczky's schedulability test.

Therefore all the 3 tasks are schedulable by RM Algorithm.

Example

Question:

Please check if following sets of tasks can be scheduled by EDF and RM Algorithms.

$$T_1 = (8, 3), T_2 = (9, 3), T_3 = (15, 3)$$

Answer:

$$\text{Utilization } U = (3 / 8) + (3 / 9) + (3 / 15) = 0.375 + 0.333 + 0.2 = 0.9083$$

$U < 1$, so these tasks are schedulable by EDF algorithm.

Let us calculate the sufficient condition for RM schedulability.

$$U_{RM} = m(2^{\frac{1}{m}} - 1) = 3(2^{\frac{1}{3}} - 1) = 3 \times (1.26 - 1) = 0.78$$

So $U > U_{RM}$, hence fails this test. But this doesn't mean that these tasks can't be schedulable by RM algorithm.

Let us perform Lehoczky's test.

Overload condition

- A system is said to be **overloaded** when the job offered to the scheduler can't be feasibly scheduled by a clairvoyant scheduler.
- When the system is not overloaded, an optimal on-line scheduling algorithm is one that always produces a feasible schedule of all offered job.
- **No optimal on-line scheduling algorithm exists when some jobs are non-preemptable.**
- During an overload, some jobs must be discarded in order to allow other jobs to complete on time.

Practical Factors

Nonpreemptability

- So jobs are by nature nonpreemptable e.g. disk scheduling.
- When a low priority job is scheduled and it happens to be nonpreemptable, if a high priority job arrives later (either from blocked state or it gets released), then it has to wait.
- The high priority job has to wait until the nonpreemptable low priority job completes.
- This will increase the response time of the high priority job.
- Hence while considering whether the high priority jobs can meet their deadline or not, we also need to consider **the effect of low priority nonpreemptable jobs** on them.

Practical Factors

Limited-Priority Levels

- In practical systems number of priority levels are limited (e.g. in a token ring network, there are 8 priority levels, in RTOSes, usually there are 256 levels)
- Hence tasks(jobs) have **non-distinct priorities**, which need to be considered during the analysis.

Tick Scheduling

- During our analysis till now, we have assumed that scheduler does the schedulability tests as and when the jobs arrives (i.e. scheduler is event-driven).
- **But practically, there will be a timer running and the scheduler will be waking up at each timer tick.**
- So even if a job is ready, the scheduler may not notice it until the next timer interrupt. This introduces a certain delay in completion of the job.
- Also the ready job which yet to be noticed and to be put in a ready queue, should be placed in some other queue.
- These factors should also be considered during the analysis.

Practical Factors

Self-suspension

- A job may suspend itself during execution due to various reasons like waiting for an I/O or remote procedure call etc.
- As a result O/S removes it from the ready queue and puts it in the suspended queue.
- The **time spent during self suspension** should also be considered during timing analysis of the jobs.

Context Switches

- Context switch is a usual phenomenon in a priority driven system.
- Hence **context switch time** should also be taken into consideration during the timing analysis.

Practical Factors

Varying Priority in Fixed-Priority Systems

- In order to tackle priority inversion problem, sometimes the priorities of the lower priority jobs are raised. Such an operation will have effect on the analysis.

Hierarchical Scheduling

- This scheduling is done, when there are multiple tasks/jobs of same priorities.
- The tasks/jobs having same priorities are put into a cluster/subsystem.
- Two common type of scheduling approaches are used.
- **Priority-driven/round-robin** system: Here The clusters are scheduled in priority driven manner and the tasks/jobs in a clusters are scheduled in a round-robin manner.
- **Fixed-time partitioning** scheme: The clusters/subsystems are scheduled according to a cyclic schedules and the tasks/jobs in the subsystem are scheduled as per the scheduling algorithm chosen for the subsystem.

Assumptions

We will confine our attention to the case where

- Priorities are fixed (e.g. RM Algorithm)
- Response times of the jobs are smaller than or equal to their respective periods.

Schedulability Test for Fixed-Priority Tasks with Short Response Times

Critical Instants

Critical Instant of a task T_i is a time instant which is such that

- The job in T_i released at the instant has the maximum response time of all jobs in T_i , if the response time of every job in T_i is equal or less than the relative deadline D_i of T_i

And

- The response time of the job released at that instant is greater than D_i if the response time of some jobs in T_i exceeds D_i .

The response time of a job in T_i released at **the critical instance** has **maximum response time**.

Theorem

In a fixed-priority system, where every job completes before the next job in the same task is released, a **critical instant** of any task T_i occurs when one of its job $J_{i,c}$ is released at the same time with a job in every higher-priority task, that is $r_{i,c} = r_{k, l_k}$ for some l_k for every $k = 1, 2, \dots, i-1$.

Time Demand Function

Suppose the release time t_0 is the job is a **critical instant** of task T_i .

Then at time $t_0 + t, t > 0$, the total processor time demand $w_i(t)$ of this job and all the higher-priority jobs released in $[t_0, t]$ is given by

$$w_i(t) = e_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{p_k} \right\rceil e_k, 0 < t \leq p_i$$

$w_i(t)$ is called **the time demand function** of task T_i .

If $w_i(t) > t$ for all $0 < t \leq D_i$, then the job **can not complete by its deadline**.

The **maximum possible response time W_i** of all jobs in T_i is equal to the **smallest value of 't'** (because the job will try to complete earliest) that satisfies the equation

$$t = e_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{p_k} \right\rceil e_k, 0 < t \leq p_i$$

Time Demand Function - Example

4 Tasks: $T_1 = (\phi_1, 3, 1)$, $T_2 = (\phi_2, 5, 1.5)$, $T_3 = (\phi_3, 7, 1.25)$, $T_4 = (\phi_4, 9, 0.5)$ are scheduled based on RM algorithm.

So priorities of these tasks are as following.

$T_1 > T_2 > T_3 > T_4$.

$$w_1(t) = e_1 = 1, 0 < t \leq 3$$

$$w_3(t) = e_3 + \left\lceil \frac{t}{p_1} \right\rceil e_1 + \left\lceil \frac{t}{p_2} \right\rceil e_2, 0 < t \leq p_3$$

$$w_2(t) = e_2 + \left\lceil \frac{t}{p_1} \right\rceil e_1, 0 < t \leq p_2 \Rightarrow w_3(t) = 1.25 + \left\lceil \frac{t}{3} \right\rceil \times 1 + \left\lceil \frac{t}{5} \right\rceil \times 1.5, 0 < t \leq 7$$

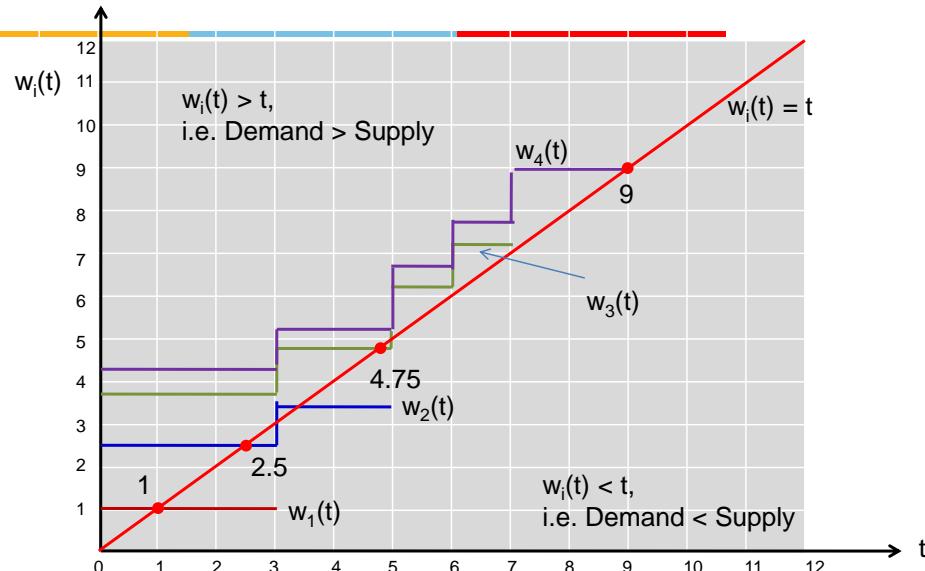
$$\Rightarrow w_2(t) = 1.5 + \left\lceil \frac{t}{3} \right\rceil, 1 < t \leq 5$$

$$\Rightarrow w_2(t) = \begin{cases} 2.5, & 0 < t \leq 3 \\ 3.5, & 3 < t \leq 5 \end{cases}$$

$$\Rightarrow w_3(t) = \begin{cases} 3.75, & 0 < t \leq 3 \\ 4.75, & 3 < t \leq 5 \\ 6.25, & 5 < t \leq 6 \\ 7.25, & 6 < t \leq 7 \end{cases}$$

46
BITS Pilani, Pilani Campus

Time Demand Function



Please note that, for T_2 , the $y=t$ line crosses 2.5 and 3.5. The max response time is 2.5 (the smallest value of t satisfying the equation). 3.5 can't be considered as the max response time, because in one period 5, a job of T_2 can execute at earliest possible instance, which is 2.5.

49
BITS Pilani, Pilani Campus

Time Demand Function – Example (contd.)

$$w_4(t) = e_4 + \left\lceil \frac{t}{p_1} \right\rceil e_1 + \left\lceil \frac{t}{p_2} \right\rceil e_2 + \left\lceil \frac{t}{p_3} \right\rceil e_3, 0 < t \leq p_4$$

$$\Rightarrow w_4(t) = 0.5 + \left\lceil \frac{t}{3} \right\rceil \times 1 + \left\lceil \frac{t}{5} \right\rceil \times 1.5 + \left\lceil \frac{t}{7} \right\rceil \times 1.25, 0 < t \leq 9$$

$$\Rightarrow w_4(t) = \begin{cases} 4.25, & 0 < t \leq 3 \\ 5.25, & 3 < t \leq 5 \\ 6.75, & 5 < t \leq 6 \\ 7.75, & 6 < t \leq 7 \\ 9.0, & 7 < t \leq 9 \end{cases}$$

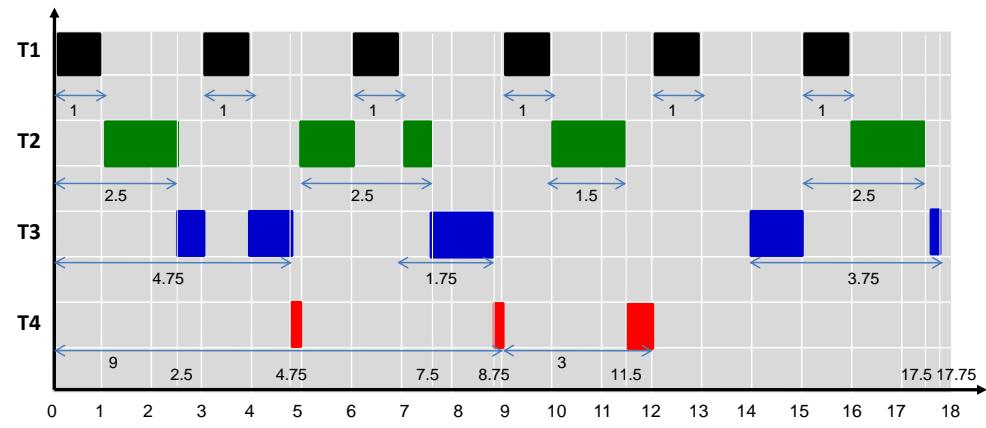
47
BITS Pilani, Pilani Campus

Time Demand Function

4 Tasks: $T_1 = (3, 1)$, $T_2 = (5, 1.5)$, $T_3 = (7, 1.25)$, $T_4 = (9, 0.5)$ are scheduled based on RM algorithm.

So priorities of these tasks are: $T_1 > T_2 > T_3 > T_4$.

Response times are shown as blue arrows in the diagram.



49
BITS Pilani, Pilani Campus

Time Demand Analysis

- Time demand function is a staircase function, with steps at integer multiple of the periods of high priority tasks.
- $w_i(t)$ is the demand of time and t is the supply of time.
- The task is schedulable if at any point of time during the inter-release time of two adjacent jobs of the task, the demand is less than or equal to the supply.
- It happens if $w_i(t) \leq t$ at any point of time during the inter-release time of two adjacent jobs of the task
- It means that $w_i(t)$ must intersect the straight line $y(t) = t$.
- It can only happen if $w_i(t) \leq t$ for some $t =$ the integer multiple the period of any of the high priority tasks or that of the current task
- Intersection of $w_i(t)$ and $y(t)=t$, indicates the maximum response time, since this is the time instance where demand = supply.

Time Demand Analysis

Time Demand Analysis method proposed by Lehoczky

For each task T_i ,

1. Compute the time demand function $w_i(t)$.

$$w_i(t) = e_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{p_k} \right\rceil e_k, 0 < t \leq p_i$$

2. Check whether the inequality

$$w_i(t) \leq t$$

is satisfied for the values of t that are equal to

$$t = jp_k,$$

$$k = 1, 2, 3, \dots, i;$$

$$j = 1, 2, \dots, \lfloor \min(p_i, D_i)/p_k \rfloor$$

If this inequality is satisfied at any of these instants, T_i is schedulable.

Time Demand Analysis - Example

Let us take the old example i.e. $T_1 = (\varphi_1, 3, 1), (\varphi_2, 5, 1.5), (\varphi_3, 7, 1.25), (\varphi_4, 9, 0.5)$.

The time demand functions for all 4 tasks have been calculated.

$$w_1(t) = 1.0, 0 < t \leq 3$$

$$w_2(t) = \begin{cases} 2.5, 0 < t \leq 3 \\ 3.5, 3 < t \leq 5 \end{cases}$$

$$w_3(t) = \begin{cases} 3.75, 0 < t \leq 3 \\ 4.75, 3 < t \leq 5 \\ 6.25, 5 < t \leq 6 \\ 7.25, 6 < t \leq 7 \end{cases}$$

$$w_4(t) = \begin{cases} 4.25, 0 < t \leq 3 \\ 5.25, 3 < t \leq 5 \\ 6.75, 5 < t \leq 6 \\ 7.75, 6 < t \leq 7 \\ 9.0, 7 < t \leq 9 \end{cases}$$

Time Demand Analysis – Example (contd.)

For T1:

$$w_1(t) = 1.0, 0 < t \leq 3$$

i = 1, so k and j doesn't exist.

Hence T1 is schedulable.

For T2:

$$w_2(t) = \begin{cases} 2.5, 0 < t \leq 3 \\ 3.5, 3 < t \leq 5 \end{cases}$$

i = 2, so k = 1, 2

For k=1:

$$j = 1, 2, \dots, \text{floor}(\min(5, 5)/3)) = 1, 2$$

$$t = jp_k = 1 \times 3 = 3, w_2(3) = 2.5 < 3.$$

For k=2:

$$j = 1, 2, \dots, \text{floor}(\min(5, 5)/5)) = 1$$

$$t = jp_k = 1 \times 5 = 5, w_2(5) = 3.5 < 5.$$

Hence T2 is schedulable, since the inequality is satisfied for all the cases.

For T3:

$$w_3(t) = \begin{cases} 3.75, 0 < t \leq 3 \\ 4.75, 3 < t \leq 5 \\ 6.25, 5 < t \leq 6 \\ 7.25, 6 < t \leq 7 \end{cases}$$

i = 3, so k = 1, 2, 3

For k = 1:

$$j = 1, 2, \dots, \text{floor}(\min(7, 7)/3)) = 1, 2$$

$$\text{For } j=1, t = jp_k = 1 \times 3 = 3, w_3(3) = 3.75 > 3.$$

$$\text{For } j=2, t = jp_k = 2 \times 3 = 6, w_3(6) = 6.25 > 6.$$

For k = 2:

$$j = 1, 2, \dots, \text{floor}(\min(7, 7)/5)) = 1$$

$$\text{For } j=1, t = jp_k = 1 \times 5 = 5, w_3(5) = 4.75 < 5.$$

For k = 3:

$$j = 1, 2, \dots, \text{floor}(\min(9, 9)/7)) = 1$$

$$\text{For } j=1, t = jp_k = 1 \times 7 = 7, w_3(7) = 7.25 > 7.$$

Hence T3 is schedulable, since the inequality is satisfied for t = 5.

Time Demand Analysis – Example (contd.)



For T4:

i = 4, so k = 1, 2, 3, 4

For k = 1:

j = 1,2,..., floor(min(9,9)/3)) = 1, 2, 3

For j=1, t = $jp_k = 1 \times 3 = 3$, $w_4(3) = 4.25 > 3$.

For j=2, t = $jp_k = 2 \times 3 = 6$, $w_4(6) = 6.25 > 6$.

For j=3, t = $jp_k = 3 \times 3 = 9$, $w_4(9) = 9$.

For k = 2:

j = 1,2,..., floor(min(9,9)/5)) = 1

For j=1, t = $jp_k = 1 \times 5 = 5$, $w_4(5) = 5.25 > 5$.

For k = 3:

j = 1,2,..., floor(min(9,9)/7)) = 1

For j=1, t = $jp_k = 1 \times 7 = 7$, $w_4(7) = 7.75 > 7$.

For k = 4:

j = 1,2,..., floor(min(9,9)/9)) = 1

For j=1, t = $jp_k = 1 \times 9 = 7$, $w_4(9) = 9$.

T4 is schedulable, since the inequality is satisfied for t = 9.

$$w_4(t) = \begin{cases} 4.25, 0 < t \leq 3 \\ 5.25, 3 < t \leq 5 \\ 6.75, 5 < t \leq 6 \\ 7.75, 6 < t \leq 7 \\ 9.0, 7 < t \leq 9 \end{cases}$$

Busy Intervals



The time interval $(t_0, t]$ is called level Π_i busy interval, if in this interval, the processor is busy all the time executing jobs with priorities Π_i or higher, all the jobs executed in the busy interval are released in the interval, and at the end of the interval there is no backlog of jobs to be executed afterwards.

Thank You.

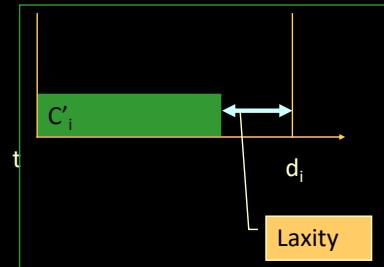
Any Questions?

RMA & EDF Examples

Priority-driven Preemptive Scheduling

Assumptions & Definitions

- Tasks are periodic
- No aperiodic or sporadic tasks
- Job (instance) deadline = end of period
- No resource constraints
- Tasks are preemptable
- Laxity of a Task
- $T_i = d_i - (t + c'_i)$
where d_i : deadline;
 t : current time;
 c'_i : remaining computation time.



RMS (cont.)

Schedule construction (online)

- Task with the smallest period is assigned the highest priority.
- At any time, the highest priority task is executed.

RMS is an optimal preemptive scheduling algorithm with fixed priorities.

Static/fixed priority algorithm assigns the same priority to all the jobs (instances) in each task.

Rate Monotonic Scheduling (RMS)

Schedulability check (off-line)

- A set of n tasks is schedulable on a uniprocessor by the RMS algorithm if the processor utilization (utilization test):

$$\sum_{i=1}^n c_i/p_i \leq n(2^{1/n} - 1).$$

The term $n(2^{1/n} - 1)$ approaches $\ln 2$, (≈ 0.69 as $n \rightarrow \infty$).

- This condition is sufficient, but not necessary.

3

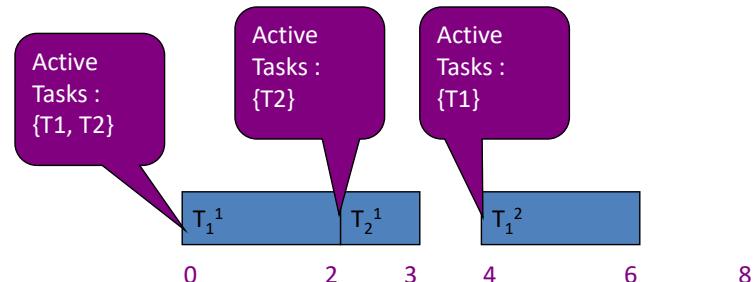
RMS Scheduler -- Example 1

Task set: $T_i = (c_i, p_i)$

$T1 = (2, 4)$ and $T2 = (1, 8)$

Schedulability check:

$$2/4 + 1/8 = 0.5 + 0.125 = 0.625 \leq 2(\sqrt{2} - 1) = 0.82$$



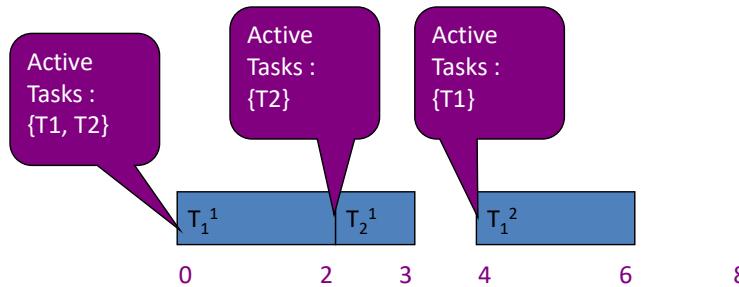
RMS Scheduler -- Example 1

Task set: $T_i = (c_i, p_i)$

$T1 = (2,4)$ and $T2 = (1,8)$

Schedulability check:

$$2/4 + 1/8 = 0.5 + 0.125 = 0.625 \leq 2(\sqrt{2} - 1) = 0.82$$



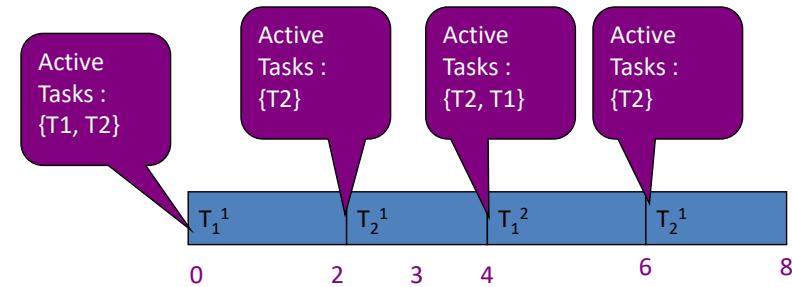
RMS scheduler -- Example-2

Task set: $T_i = (c_i, p_i)$

$T1 = (2,4)$ and $T2 = (4,8)$

Schedulability check:

$$2/4 + 4/8 = 0.5 + 0.5 = 1.0 > 2(\sqrt{2} - 1) = 0.82$$



Some task sets that FAIL the utilization-based schedulability test are also schedulable under RMS ➔ We need exact analysis (necessary & sufficient)

6

7

Earliest Deadline First (EDF)

- **Schedulability check (off-line)**

- A set of n tasks is schedulable on a uniprocessor by the EDF algorithm if the processor utilization.

$$\sum_{i=1}^n c_i/p_i \leq 1$$

- This condition is both necessary and sufficient.
- Least Laxity First (LLF) algorithm has the same schedulability check.

EDF/LLF (cont.)

- **Schedule construction (online)**

- EDF/LLF: Task with the smallest deadline/laxity is assigned the highest priority.
- At any time, the highest priority task is executed.

EDF/LLF is an optimal preemptive scheduling algorithm with dynamic priorities.

Dynamic priority algorithm assigns different priorities to the individual jobs (instances) in each task.

8

9

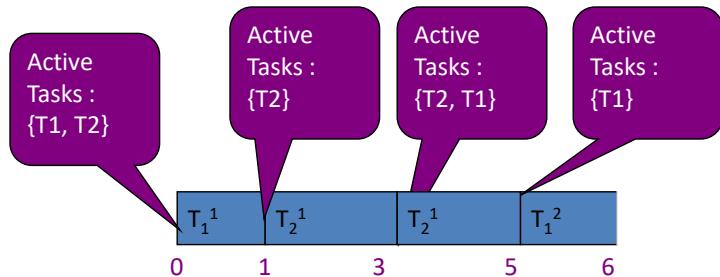
EDF scheduler -- Example

Task set: $T_i = (c_i, p_i, d_i)$

$T_1 = (1,3,3)$ and $T_2 = (4,6,6)$

Schedulability check:

$$1/3 + 4/6 = 0.33 + 0.67 = 1.0$$



Unlike RMS, Only those task sets which pass the schedulability test are schedulable under EDF

RMS vs. EDF/LLF

- RMS is an optimal preemptive scheduling algorithm with fixed priorities.
- EDF/LLF is an optimal preemptive scheduling algorithm with dynamic priorities.
- RMS schedulability properties can be analyzed; rich theory exists and it is widely used in practice.
- EDF/LLF offers higher schedulability than RMS, but it is more difficult to implement.

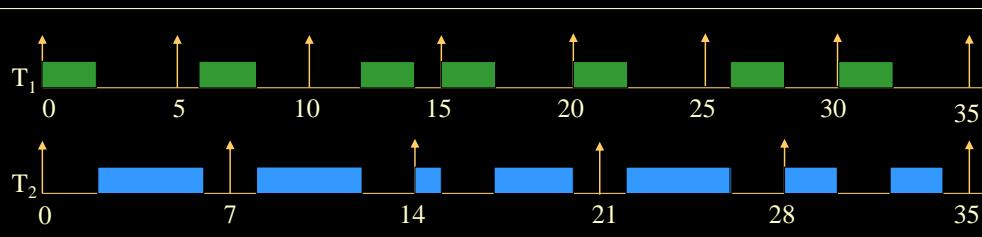
10

11

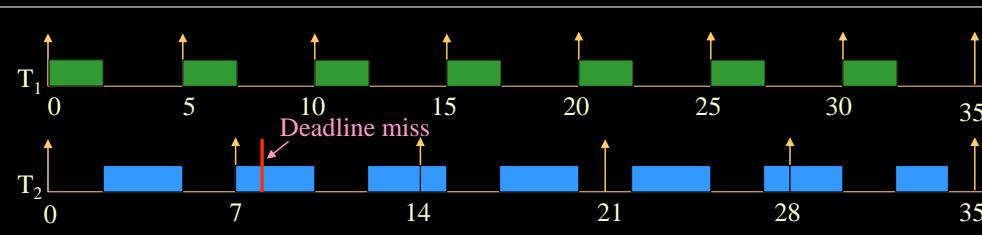
RMS & EDF -- Example

Process	Period, T	WCET, C
T_1	5	2
T_2	7	4

EDF schedule



RMS schedule



Schedulability: UB Test

- Utilization bound (UB) test: a set of n independent periodic tasks scheduled by the rate monotonic algorithm will always meet its deadlines, for all task phasings, if

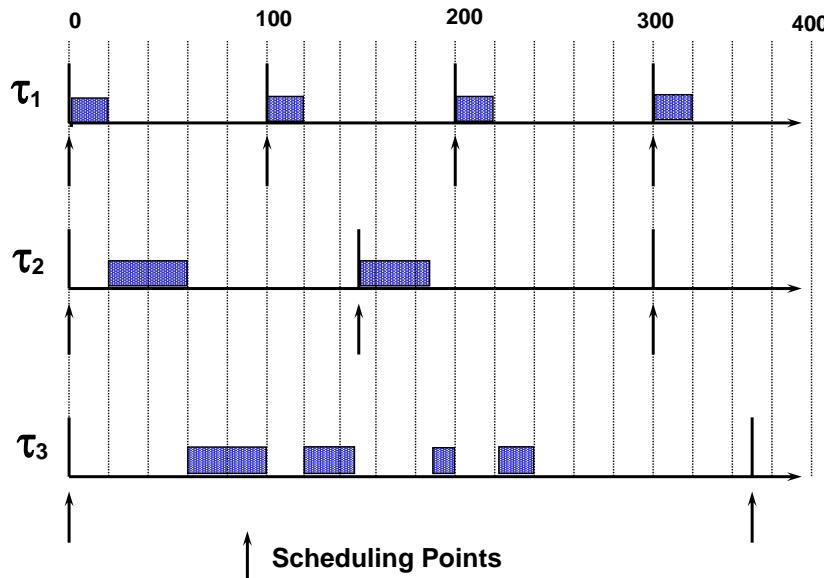
$$\begin{array}{lll} U(1) = 1.0 & U(4) = 0.756 & U(7) = 0.728 \\ U(2) = 0.828 & U(5) = 0.743 & U(8) = 0.724 \\ U(3) = 0.779 & U(6) = 0.734 & U(9) = 0.720 \end{array}$$

- For **harmonic** task sets, the utilization bound is $U(n)=1.00$ for all n .

$$\frac{C_1}{T_1} + \dots + \frac{C_n}{T_n} \leq U(n) = n(2^{1/n} - 1)$$

Introduction to Embedded Systems

Timeline for Sample Problem



Introduction to Embedded Systems

Sample Problem: Applying UB Test

	C	T	U
Task τ_1	20	100	0.200
Task τ_2	40	150	0.267
Task τ_3	100	350	0.286

- Total utilization is $.200 + .267 + .286 = .753 < U(3) = .779$
- The periodic tasks in the sample problem are schedulable according to the UB test

Introduction to Embedded Systems

Exercise: Applying the UB Test

Given:

Task	C	T	U
τ_1	1	4	
τ_2	2	6	
τ_3	1	10	

- What is the total utilization?
- Is the task set schedulable?
- Draw the timeline.
- What is the total utilization if $C_3 = 2$?

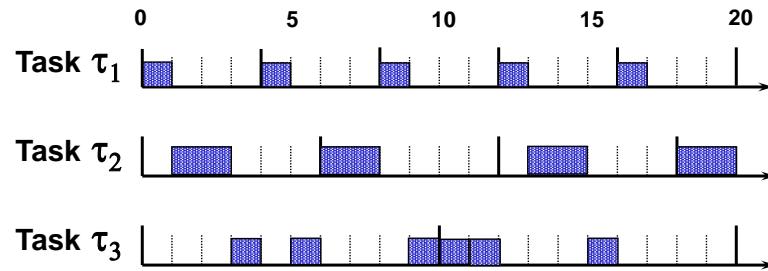
Introduction to Embedded Systems

Solution: Applying the UB Test

a. What is the total utilization? $.25 + .34 + .10 = .69$

b. Is the task set schedulable? Yes: $.69 < U(3) = .779$

c. Draw the timeline.



d. What is the total utilization if $C_3 = 2$?

$$.25 + .34 + .20 = .79 > U(3) = .779$$

Toward a More Precise Test

- UB test has three possible outcomes:

$0 < U < U(n)$	→ Success
$U(n) < U < 1.00$	→ Inconclusive
$1.00 < U$	→ Overload

- UB test is conservative.

- A more precise test can be applied.

Text Book / References



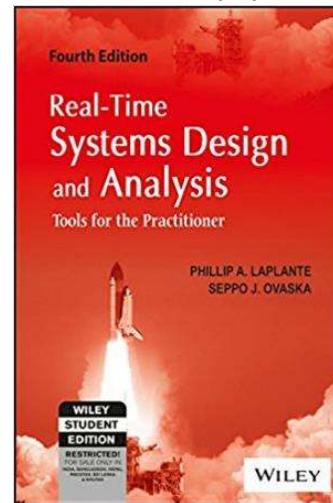
BITS Pilani
Pilani Campus

BITS ZG553: Real Time Systems

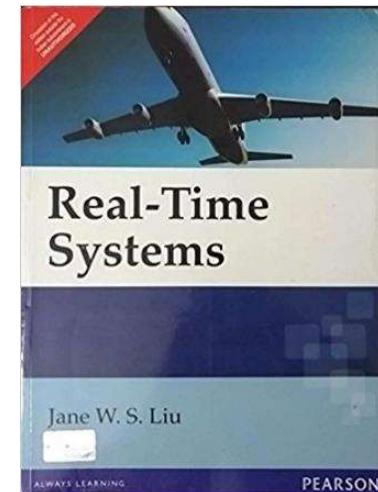
L7 – Scheduling of Aperiodic & Sporadic Jobs

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Reference (R1)



Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Excellent MOOCs Videos

(Coursera, edX,...)



HONOR CODE
CERTIFICATE



This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

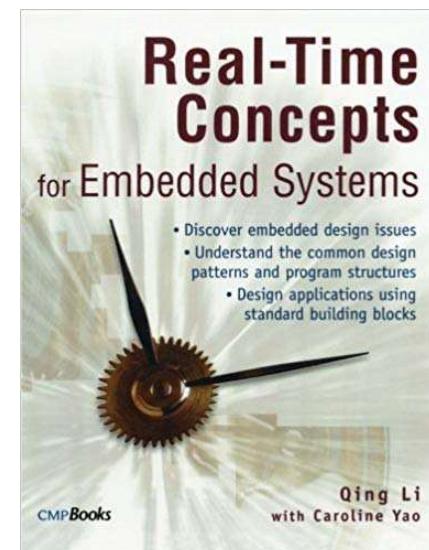
HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44a218f5a1eebf47d4e54

Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

RTS Primer – For Light Reading





L-7: Scheduling Aperiodic & Sporadic Jobs- (Polled Servers, Deferrable Servers, Sporadic Servers)

[Ref: T1/C5]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs, Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

What is a aperiodic task?

- ❑ We say a task is *aperiodic* if the jobs in it have either soft deadlines or no deadlines.
- ❑ We therefore want to optimize the responsiveness of the system for the aperiodic jobs, but never at the expense of hard real-time tasks whose deadlines must be met at all times.

What is a sporadic task?

- ❑ Tasks containing jobs that are released at random time instants and have hard deadlines are *sporadic tasks*.
- ❑ We treat sporadic tasks as *hard real-time tasks*.
- ❑ Our primary concern is to ensure that their *deadlines are met*.
- ❑ Minimizing their response times is of secondary importance.

Assumptions

It is impossible for some sporadic jobs to meet their deadlines no matter what algorithm is used.

So there are two alternatives:

- a) Reject the sporadic jobs that can't complete on time
- b) Accept all sporadic jobs and allow some of them to complete late.

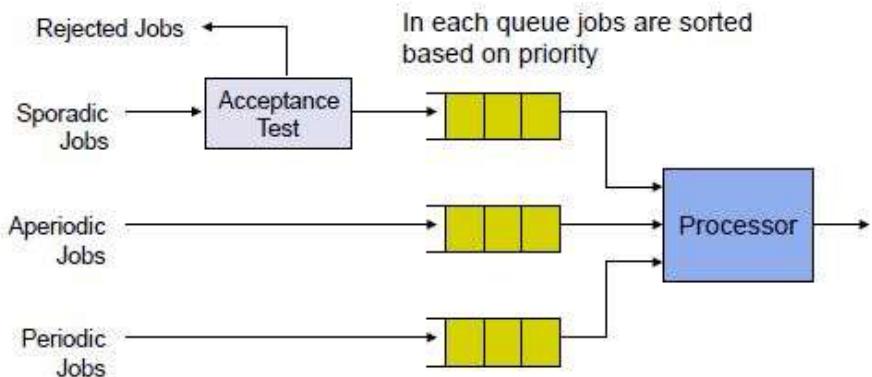
Option (a) will be the focus here.

Objective

Two problems for the Sporadic and Aperiodic job scheduling algorithms:

1. If the scheduler accepts the **sporadic** job, it should schedule the job so that the job completes in time without causing any periodic tasks and previously accepted sporadic jobs to miss their deadlines. The problems are:
 - a) How to do acceptance test
 - b) How to schedule the accepted sporadic job
2. How to complete each **aperiodic** job as soon as possible without causing periodic tasks and accepted sporadic jobs to miss their deadlines

System Model



Optimality

1. An **aperiodic job scheduling algorithm is optimal** if it minimizes either the response time of the aperiodic job at the head of the aperiodic job queue or the average response time of all the aperiodic jobs for the given queuing discipline.
2. A **sporadic job scheduling algorithm is optimal** if it accepts each sporadic job newly offered to the system and schedule the job to complete in time if and only if the new job can be correctly scheduled to complete in time by some means.

Scheduling of Aperiodic tasks - Background Approach

Aperiodic jobs are scheduled and executed only when there is no periodic or sporadic job ready for execution.

Advantages:

- This method always produces correct schedule
- It is simple to implement.

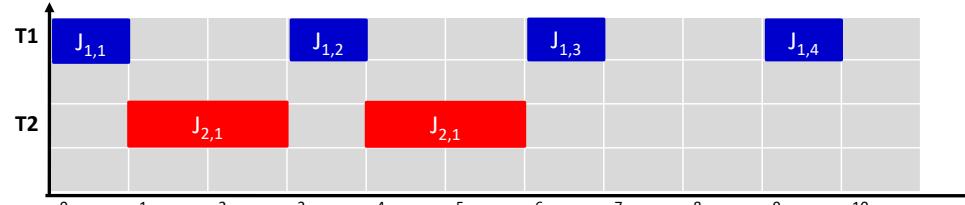
Disadvantage(s):

- Execution of aperiodic jobs are delayed and their response times prolongs unnecessarily.

Scheduling of Aperiodic tasks – Background Approach



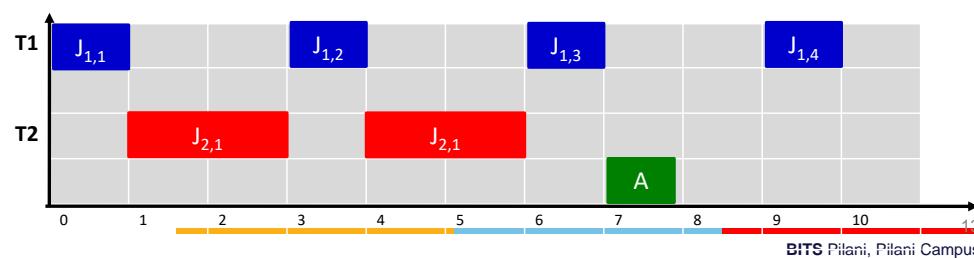
Example: $T_1 = (3, 1)$, $T_2 = (10, 4)$ scheduled as per RM algorithm



Aperiodic job A with execution time 0.8 arrives at time 0.1.

After $J_{1,3}$ is completed, there is free time slot, so A can be scheduled. A completes at 7.8.

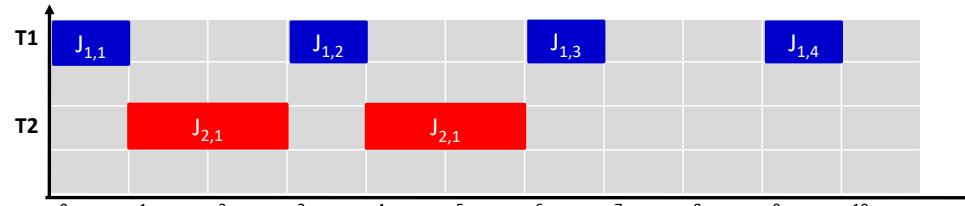
Response time of A = $7.8 - 0.1 = 7.7$



Scheduling of Aperiodic tasks – Interrupt Driven Approach

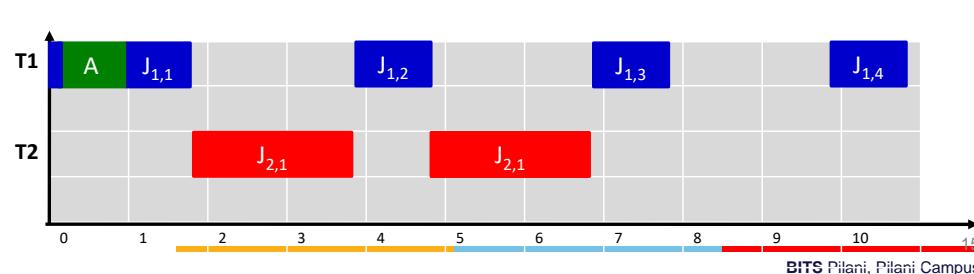


Example: $T_1 = (3, 1)$, $T_2 = (10, 4)$ scheduled as per RM algorithm



Aperiodic job A with execution time 0.8 arrives at time 0.1. So it gets scheduled at time 0.1.

Response time of A = 0.8. But the periodic jobs get delayed by 0.8 times.



Scheduling of Aperiodic tasks – Interrupt Driven Approach



Whenever an aperiodic job arrives, the execution of periodic tasks are interrupted and the aperiodic job is executed.

Advantage:

Aperiodic job will always meet its deadline and will have shortest possible response time.

Disadvantage:

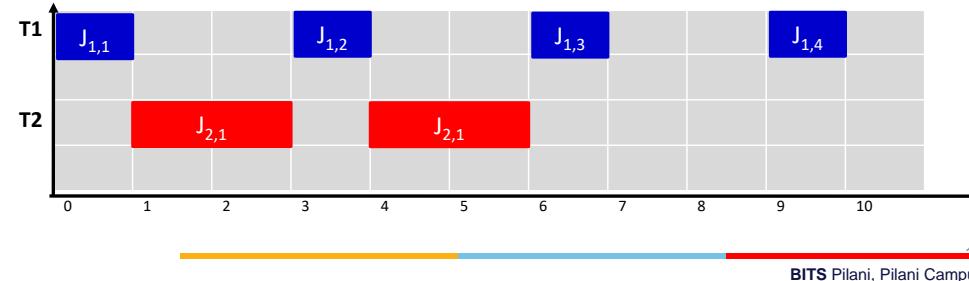
Periodic tasks may miss some deadline.

14
BITS Pilani, Pilani Campus

Scheduling of Aperiodic tasks – Interrupt Driven with Slack Stealing



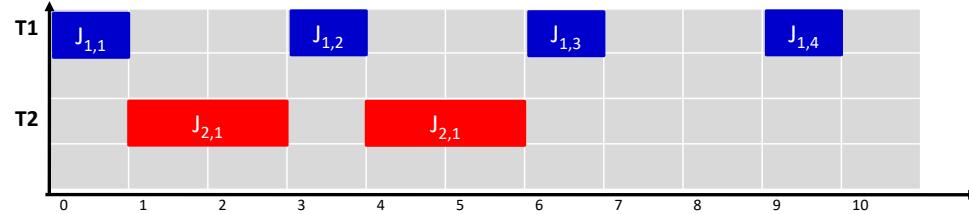
Whenever an aperiodic job arrives, the scheduler checks if there is any slack available. If available, it postpones the periodic job as per the slack availability and executes the aperiodic job during this time. This way periodic job doesn't miss its deadline. At the same time aperiodic job executes as soon as possible.



15
BITS Pilani, Pilani Campus

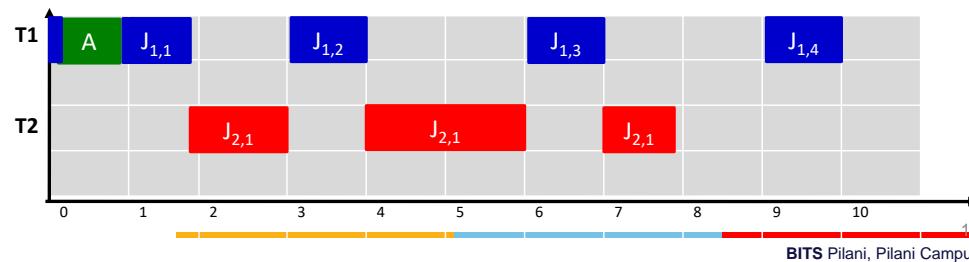
Scheduling of Aperiodic tasks – Interrupt Driven with Slack Stealing

Example: $T_1 = (3,1)$, $T_2 = (10,4)$ scheduled as per RM algorithm



Aperiodic job A with execution time 0.8 arrives at time 0.1.

$J_{2,1}$ and $J_{1,1}$ have slack times, so can be postponed. During which A can be scheduled.



Polled Execution

- ❑ **Periodic Server:** A periodic task which is created to execute aperiodic jobs.
- ❑ A periodic server (p_s, e_s) is defined by its period p_s and execution time e_s
- ❑ The term is e_s called **execution budget** (or simply **budget**)
- ❑ The ratio $u_s = e_s / p_s$ is called the **size** of the server.
- ❑ At the beginning of the period, the budget of the poller is set to e_s . We say that the budget is **replenished**.
- ❑ The time instant when the budget is replenished is called **replenishment time**.
- ❑ We say that the periodic server is **backlogged** whenever the aperiodic job queue is nonempty.
- ❑ The server is **eligible** (i.e. ready) for execution only when it is backlogged and has budget.
- ❑ When the server executes, it consumes its budget at **one per unit time**.
- ❑ The server budget becomes **exhausted** when the budget becomes zero.
- ❑ The budget of a poller becomes **exhausted instantaneously** whenever the poller finds the aperiodic queue empty i.e. itself idle.

Polled Execution

- ❑ A **poller** or **polling server** (p_s, e_s) is a periodic task: p_s is its **period** and e_s is its **execution time**.
- ❑ When executed, it executes an aperiodic job, if the aperiodic job queue is non-empty.
- ❑ Poller suspends execution or is suspended by the scheduler either
 - when it has executed for e_s , or
 - when the aperiodic job queue becomes empty

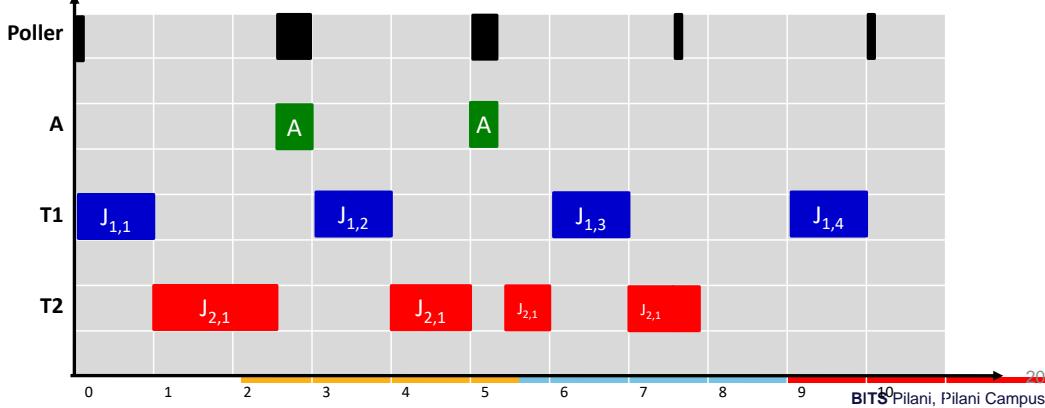
18

BITS Pilani, Pilani Campus

Polling Server Example

Take the old example : $T_1 = (3,1)$, $T_2 = (10,4)$, Aperiodic job A of execution time 0.8 arrives at time 0.1.

- The poller (2.5, 0.5) is considered as the highest priority periodic task.
- At time 0, it wakes up, finds no aperiodic job, so gets suspended.
- At time 0.1, the aperiodic job arrives, but doesn't get scheduled, because poller is not running.
- At time 2.5 the poller gets scheduled again. It finds aperiodic task A, so executes it for 0.5 time units. Then it gets suspended.
- At time 5, the poller gets scheduled again and executes remaining 0.3 time units of the aperiodic job A and then gets suspended.
- Therefore, the job A completes at time 5.3 and its response time is $5.3 - 0.1 = 5.2$.



20

BITS Pilani, Pilani Campus

Bandwidth Preserving Servers

- ❑ A bandwidth-preserving server is a periodic server
- ❑ Compared to polling server **bandwidth preserving servers try to preserve their budget when they are not executed.**
- ❑ Hence there are additional rules for consumption and replenishment for these servers:
 - A backlogged bandwidth-preserving server is ready for execution when it has budget. Scheduler keeps track of the consumption of the server budget. If budget is exhausted server becomes idle. Scheduler moves server back to ready queue, when budget is replenished and server is backlogged
 - **If a new aperiodic job arrives, an idle server becomes backlogged** and is put into the ready queue if it has budget

Deferrable Server Example - 1

Take the old example.

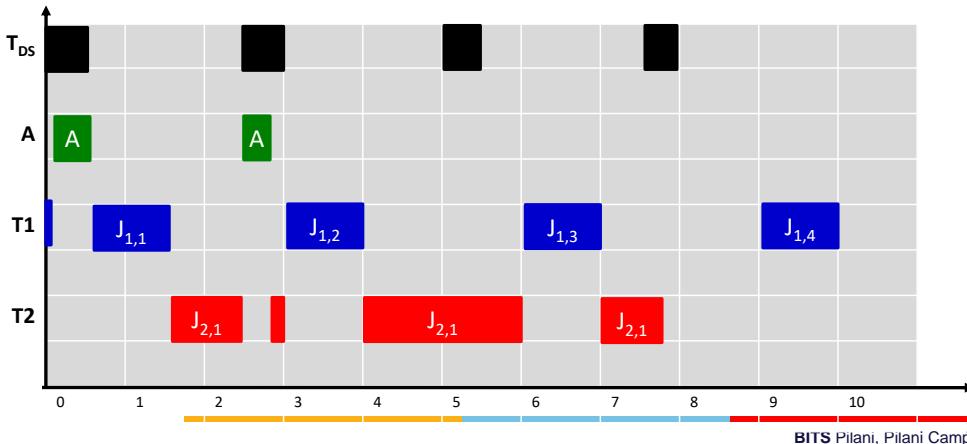
The deferrable server is defined by $T_{DS} = (2.5, 0.5)$ is considered as the highest priority periodic task.

At time 0, it wakes up, finds no aperiodic job, so gets suspended.

At time 0.1, the aperiodic job arrives, it wakes up and scheduled the aperiodic job for 0.5 time units.

At time 2.5 the deferrable server gets scheduled again. It executes remaining 0.3 time units of the aperiodic job.

Therefore, the job A completes at time 2.8 and its response time is $2.8 - 0.1 = 2.7$.



Deferrable Servers

A deferrable server is the simplest bandwidth-preserving server.

Unlike poller, it preserves its budget when there is no aperiodic tasks to execute.

Consumption rule: The execution budget of the server is consumed at the rate of one unit per time whenever the server executes

Replenishment rule: The execution budget of the server is set to k at kpk , $k = 0, 1, 2, \dots$

Server is not allowed to cumulate budget from period to period

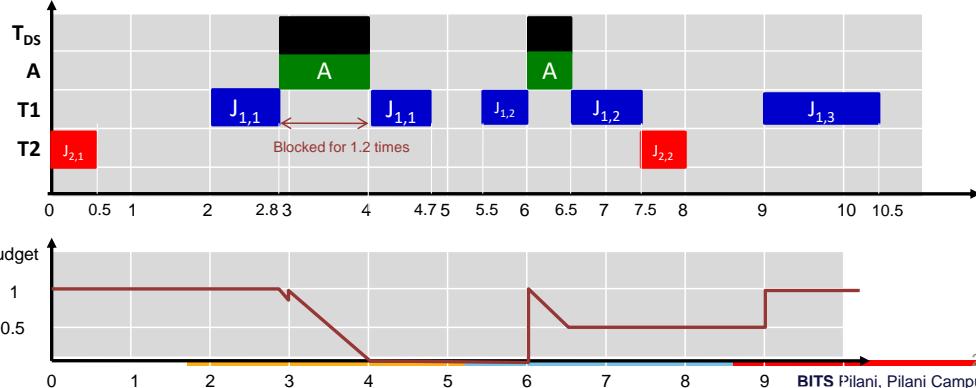
Deferrable Server Example - 2

The deferrable server is defined by $T_{DS} = (3, 1)$ is considered as the highest priority periodic task.

Periodic Tasks T₁ = (2, 3.5, 1.5, 3.5) and T₂ = (6.5, 0.5) are scheduled rate monotonically.

An aperiodic job A with execution time 1.7 arrives at time 2.8.

At time 0, the server budget = 1. No aperiodic task arrives, so server is not executed. At time 1, the server budget = 1. No aperiodic task arrives, so server is not executed. At time 2, the server budget = 1. No aperiodic task arrives, so server is not executed. At time 2.8, the server budget = 1. The aperiodic task arrives, so server is executed. At time 3, the server budget = 0.8. It gets replenished to 1. So it continues execution till time 4. At time 4, the server budget = 0, so the server gets suspended. At time 6, server budget is replenished to 1, so server gets scheduled. This executes the aperiodic task till completion i.e. time 6.5. At time 6.5, server has budget 0.5 remaining. Since there is no aperiodic tasks in the queue, it suspends itself.



Sporadic Servers

- Limitation of deferrable servers – they may delay lower-priority tasks for more time than a periodic task with the same period and execution time

Example: In the last example, T1 was blocked for 1.2 times, even though the budget of the deferrable server is 1.0.

- A sporadic server is designed to eliminate this limitation
 - A different type of periodic server: several different sub-types
 - More complex consumption and replenishment rules ensure that a sporadic server with period p_s and budget e_s never demands more processor time than a periodic task with the same parameters

BITS Pilani, Pilani Campus

Sporadic Server in a Fixed Priority System

Consumption rule

At any time $t \geq t_r$, if the server has budget and if either of the following two conditions is true, the budget is consumed at the rate of 1 per unit time.

- C1: The server is executing
- C2: The server has executed since t_r and $END < t$

When they are not true, the server holds its budget

What does it mean ?

- The server executes for no more time than it has execution budget
- The server retains its budget if a higher-priority job is executing, or It has not executed since t_r (i.e. last replenishment time)
- Otherwise, the budget decreases when the server executes, or if it idles while it has budget

Sporadic Server in a Fixed Priority System

Let us define

- A system, T , of independent preemptable periodic tasks and a sporadic server with parameters (p_s, e_s) and arbitrary priority of Π_s
- T_H : the periodic tasks with higher priority than the server (may be empty)
- t_r : the last time the server budget replenished
- t_f : the first instant after t_r at which the server begins to execute

At any time t define:

- **BEGIN** as the start of the earliest busy interval in the most recent contiguous sequence of busy intervals of T_H starting before t (busy intervals are contiguous if the later one starts immediately the earlier one ends)
- **END** as the end of the latest busy interval in this sequence if this interval ends before t and define $END = \infty$ if the interval ends after t

26
BITS Pilani, Pilani Campus

Sporadic Server in a Fixed Priority System

Replenishment rules

- R1: When system begins executing, and each time budget is replenished, set the budget to e_s and $t_r =$ the current time.

- R2: When server begins to execute (defined as time t_f)
 - if $END = t_f$ then
 $t_e = \max(t_r, BEGIN)$
 - else if $END < t_f$ then
 $t_e = t_f$

The next replenishment time is set to $t_e + p_s$, where t_e = effective replenishment time

- R3: budget replenished at the next replenishment time, except under the following conditions. Under these conditions, replenishment is done at times stated below.
 - If $t_e + p_s$ is earlier than t_f the budget is replenished as soon as it is exhausted
 - If T becomes idle before $t_e + p_s$, and becomes busy again at t_b , the budget is replenished at $\min(t_b, t_e + p_s)$

Sporadic Server in a Fixed Priority System - Example



A sporadic server $(5, 1.5)$ was scheduled with three periodic tasks $T_1 = (3, 0.5)$, $T_2 = (4, 1.0)$, $T_3 = (19, 4.5)$ rate-monotonically. Three aperiodic jobs arrived at times 3, 7 and 15.5, with execution times 1.0, 2.0 and 2.0 respectively.

Solution

As per RM, priorities of the tasks are: $T_1 > T_2, T_3 > T_2$.

T_1 and T_2 are higher priority than T_3 , so create the Gantt chart of them first.

Time 0 to 3: Aperiodic job queue is empty. T_3 is replenished at time 0 with budget 1.5. So $t_r = 0$.

Time 3: Aperiodic job A_1 arrives. T_3 has budget, so A_1 gets scheduled at time 3.5 after T_1 finishes. At this time $t_r = 0$, BEGIN = 3, END = 3.5, $t_e = 3.5$. According to rule R2, effective replenishment time $t_e = \max(0, 3.0) = 3$. So next replenishment time is $3 + 5 = 8$.

Time 4: The server executed A_1 till time 4. At time 4, it gets preempted by T_2 . From time 4 to 5, T_2 executes. Since it is of higher priority than T_3 , the budget of T_3 is not consumed (rule C2).

Time 5: T_2 is ended. So T_3 is scheduled. A_1 completes by 5.5. At time $t = 5.5$, END = 5 since there is no other high priority task in the system (so, END < t). So the budget of the server continues to be consumed 1 unit per time (rule C2). Hence the budget gets exhausted at time 6.

Time 7: Next periodic job A_2 arrives. But the server doesn't have budget. So it waits till next replenishment time 8.

Time 9.5: T_3 schedules A_2 after the high priority jobs are finished. Since the server began to execute for the first time after replenished, now, $t_r = 9.5$. BEGIN = 8, END = 9.5, $t_e = \max(8, 8) = 8$. So next replenishment time is $8 + 5 = 13$.

Time 11: A_2 executes from 9.5 to 11, so budget of T_3 gets exhausted (rule C1). So A_2 gets preempted.

Time 13.5: The budget of the server is replenished at time 13. So A_2 gets scheduled at time 13.5 after high priority tasks are done. It completes at time 14.

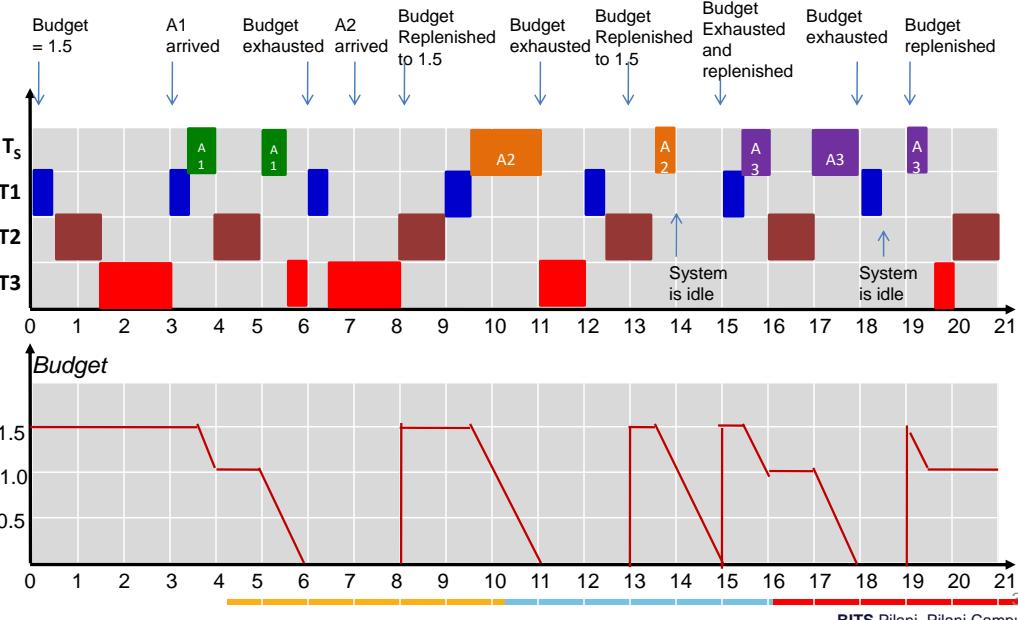
Time 14: Now the system is idle (even T_3 is not there). The system will busy again at time 15. So rule R3 (b) will come into picture. So the server will be replenished at time 15 instead of time 18.

Time 15.5: A_3 arrives, the server has budget. Hence it will be scheduled.

Time 19: System becomes idle again. So Rule R3(b) applies, so the server will be replenished at time 19, in stead of 20.

(T_3 is scheduled in the time slots remained after T_1, T_2 and T_3)

Sporadic Server in a Fixed Priority System - Example



BITs Pilani, Pilani Campus

Schedulability of Sporadic Jobs



- Sporadic jobs are scheduled as done in case of clock-driven scheduling.
- Acceptance tests are performed on sporadic jobs in EDF order.
- Once accepted, Sporadic jobs are ordered among themselves in EDF order.
- In a **fixed-priority system**, they are executed in a bandwidth preserving server.
- In a **deadline-driven system**, they are scheduled with periodic jobs on EDF basis.

BITs Pilani, Pilani Campus

Acceptance Test for Sporadic Jobs in Deadline-Driven Systems



For a sporadic job J_i , that has release time r_i , maximum execution time e_i and absolute deadline d_i ,

$$\text{Density} = e_i / (d_i - r_i)$$

A sporadic job is said to be **active** in its feasible interval $(r_i, d_i]$.

Theorem

A system of independent, preemptable sporadic jobs is schedulable according to EDF algorithm if the total density of all active jobs in the system is no greater than 1 at all times.

BITs Pilani, Pilani Campus

BITs Pilani, Pilani Campus

Acceptance Test for Sporadic Jobs in Deadline-Driven Systems



Let us say, Δ = Total density of all the periodic tasks.

The all accepted sporadic jobs can meet their deadlines as long as the total density of all the active sporadic jobs is no greater than $1 - \Delta$.

Acceptance Test Procedure:

1. Calculate the total density of the periodic jobs Δ .
2. When a Sporadic job $S(t, d, e)$ arrives, divide the time interval into two: time interval I_1 before d and time interval I_2 after d .
3. Let the total densities of active Sporadic jobs in these two intervals are $\Delta_{s,1}$ and $\Delta_{s,2}$ respectively.
4. Accept the sporadic job S , if

$$\frac{e}{d-t} + \Delta_{s,1} \leq 1 - \Delta$$

BITS Pilani, Pilani Campus

Example – Scheduling Sporadic Jobs (contd.)



Two periodic tasks $T1 = (4, 1)$ and $T2 = (6, 1.5)$ are in the system, scheduled based on EDF algorithm.

Four sporadic jobs $S1 = (0, 8, 2)$, $S2 = (2, 7, 0.5)$, $S3 = (4, 14, 1)$ and $S4 = (9, 13, 2)$ are arrived and are to be scheduled.

Solution:

$$\Delta = 0.25 + 0.25 = 0.5.$$

$$\text{So } 1 - \Delta = 1 - 0.5 = 0.5.$$

Time 0: $J_{1,1}$ and $J_{2,1}$ are released. Also sporadic job S_1 is released.

$$\text{Density of } S_1 = 2 / (8-0) = 0.25 < 0.5.$$

So S_1 is accepted. All the 3 jobs are scheduled based on EDF.

S_1 divides the time into two intervals: $(0, 8]$, $(8, \infty)$.

$$\text{Scheduler updates } \Delta_{s,1} = 0.25, \Delta_{s,2} = 0.$$

$J_{1,1}$ has earlier deadline, so scheduled. Then $J_{2,1}$ has earlier deadline, so scheduled at time 1.

Acceptance Test for Sporadic Jobs in Deadline-Driven Systems



Generalising the steps,

Let there are n_s sporadic jobs in the system, then the deadlines of these sporadic jobs will partition the time into $n_s + 1$ disjoint intervals: $I_1, I_2, \dots, I_{n_s+1}$

Let I_j be the time interval during which a new Sporadic job $S(t, d, e)$ arrives.

Accept the sporadic job S , if

$$\frac{e}{d-t} + \Delta_{s,k} \leq 1 - \Delta$$
$$k = 1, 2, 3, \dots, l$$

BITS Pilani, Pilani Campus

Example – Scheduling Sporadic Jobs (contd.)



Time 2: Sporadic job S_2 is released.

$$\text{Density of } S_2 = 0.5 / (7-2) = 0.1$$

$$\text{Total density of sporadic jobs in this interval} = 0.25 + 0.1 = 0.35 < 0.5$$

So S_2 is accepted. All the 2 jobs are scheduled based on EDF.

S_2 divides the first time into two intervals: $(0, 7]$, $(7, 8]$.

$$\text{Scheduler updates } \Delta_{s,1} = 0.35, \Delta_{s,2} = 0.25 \text{ and } \Delta_{s,3} = 0.$$

S_2 has earlier deadline, so scheduled first, and then S_1 is scheduled.

Time 4: $J_{1,2}$ are released, but it has the deadline same as that of S_1 , so S_1 can continue to execute.

Only sporadic job in the system now is S_1 .

S_3 is also released.

$$\text{Density of } S_3 = 1 / (14-4) = 0.1$$

$$\Delta_{s,1} = 0.25. \text{ So total density of sporadic job in this interval} = 0.25 + 0.1 = 0.35 < 0.5$$

So S_3 is accepted. All the 3 jobs are scheduled based on EDF. S_1 having earliest deadline, it continued to execute.

S_3 divides the time into three intervals: $(0, 8]$, $(8, 14]$, $(14, \infty)$.

$$\text{Scheduler updates } \Delta_{s,1} = 0.35, \Delta_{s,2} = 0.1 \text{ and } \Delta_{s,3} = 0$$

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Example – Scheduling Sporadic Jobs (contd.)



Time 5: S1 is done and J_{1,2} is scheduled.

Time 6: J_{2,2} is released. It has earlier deadline than S3, so gets scheduled.

Time 7.5: S3 gets scheduled.

Time 8: J_{1,3} gets released. It has earlier deadline than S3, so S3 gets pre-empted and J_{1,3} gets scheduled.

Time 9: J_{1,3} is completed. S4 gets released.

S4 divides the time into three intervals: (9, 14], (14, ∞)

In the interval (9,14], there are two active sporadic jobs S3 and S4 now.

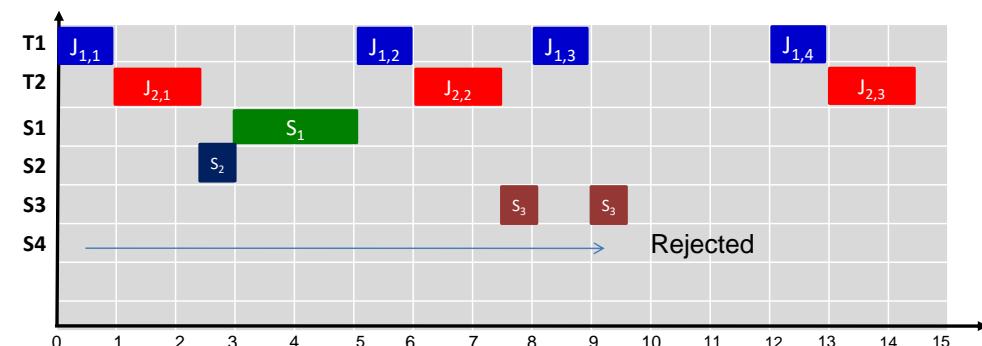
Density of existing sporadic jobs = 0.1

Density of S4 = $2 / (13-9) = 0.5$.

Total Density of sporadic jobs = $0.1 + 0.5 = 0.6 > 0.5$.

Hence S4 is rejected. So S3 gets scheduled at time 9.

Example – Scheduling Spradic Jobs (contd.)



Thank You.

Any Questions?

27 BITS Pilani, Pilani Campus



The slide features the BITS Pilani clock tower in the background against a clear blue sky. In the foreground, the BITS Pilani logo is displayed, consisting of a circular emblem with text and symbols, and the text "BITS Pilani" and "Pilani Campus".

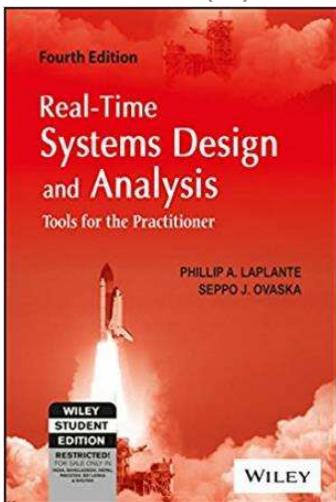
BITS ZG553: Real Time Systems
L8 – Resources & Resource Access Control,
Priority Inversion & Deadlocks

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

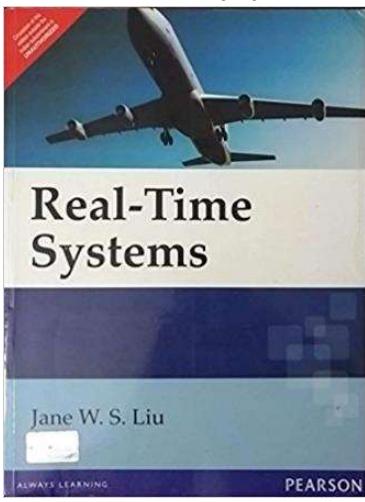
Text Book / References



Reference (R1)



Text Book (T1)



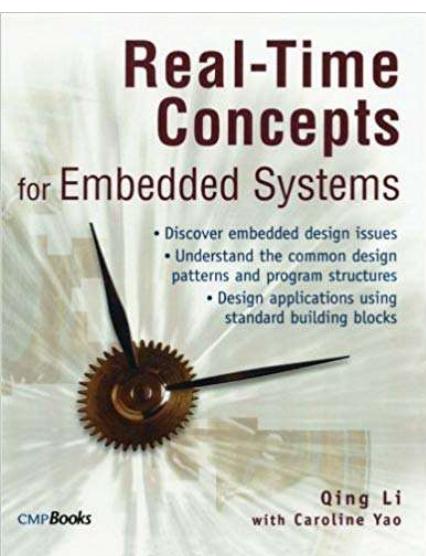
Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



RTS Primer – For Light Reading



4

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Excellent MOOCs Videos (Coursera, edX,...)



HONOR CODE CERTIFICATE



Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIx: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.



VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54

3
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956



BITS Pilani
Pilani|Duba|Gosha|Hyderabad



L-8: Resources & Resource Access Control, Priority Inversion & Deadlocks

[Ref. T1/C8]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

Assumptions

- ❑ System contains only one processor
- ❑ System contains p types of *serially reusable* resources R_1, R_2, \dots, R_p
- ❑ There are v_i *indistinguishable* units of a resource of type R_i
 - Plentiful resources are ignored
 - Binary semaphore has one unit
 - Counting semaphore has n units
 - Only one unit of *write-lock*
 - A system containing 5 printers have 5 units of printer resources
- ❑ Serially reusable resources are allocated to jobs on a *non-preemptive basis and used in a mutually exclusive manner*. It means when a unit of resource is granted to a job, this unit is no longer available to other jobs until the job frees the unit.
- ❑ If a resource can be used by more than one jobs at the same time, this is modelled as a resource with several units, each used in a mutual exclusive manner. Example: A file which can be read by v users at a time is modelled as v mutually exclusive units of file resources

Mutual Exclusive Resource Access

- ❑ A lock-based concurrency control mechanism assumed to be used to enforce mutual exclusive access to resources
- ❑ When a job wants to use η_i units of a resource R_i , it executes a *lock* $L(R_i, \eta_i)$ (η = "eta") to request them
- ❑ When the job no longer needs the resources, it releases them by executing an *unlock* $U(R_i, \eta_i)$
- ❑ When a lock request fails, the requesting job is blocked and loses the processor
- ❑ It stays blocked until the scheduler grants the resources the job is waiting for
- ❑ If a resource has only 1 unit the simpler notations $L(R_i)$ and $U(R_i)$ are used for lock and unlock

Critical Section

- ❑ A segment of a job *that begins with a lock and ends at a matching unlock* is called a *critical section*
- ❑ Resources are released in *last-in-first-out* order
- ❑ A critical section that is not included in other critical sections is called an *outermost critical section*
- ❑ Critical sections are denoted by $[R, \eta; e]$, where
 - R gives the *name* and
 - η the *number of units* of a resource and
 - e the *(maximum) execution time* of the critical section
- ❑ If there is only one unit of a resource the simpler notation $[R; e]$ is used

Nested Critical Sections

- ❑ Nested critical sections are denoted by nested square brackets
 - $[R1; 7 [R2; 3]]$ indicates that the critical section beginning $L(R1)$ from includes another critical section that begins with $L(R2)$
 - Locks are locked and unlocked in the following order: $L(R1), L(R2), U(R2), U(R1)$

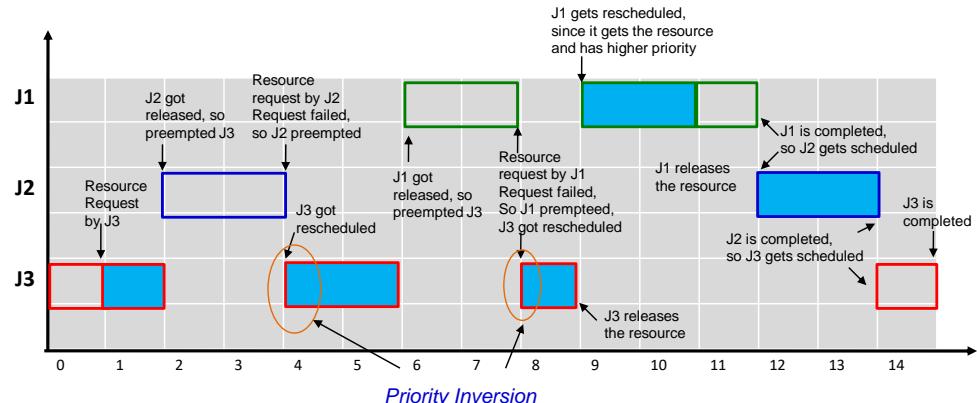
Resource Conflicts

- Two jobs conflict with each other, if some of the resources they require are of the same type
- They contend for a resource when one job requests a resource that the other job already has
- These terms are used interchangeably

10 BITS Pilani, Pilani Campus

Priority Inversion

- Priority inversion occurs, when a low-priority job executes while a ready higher-priority job waits

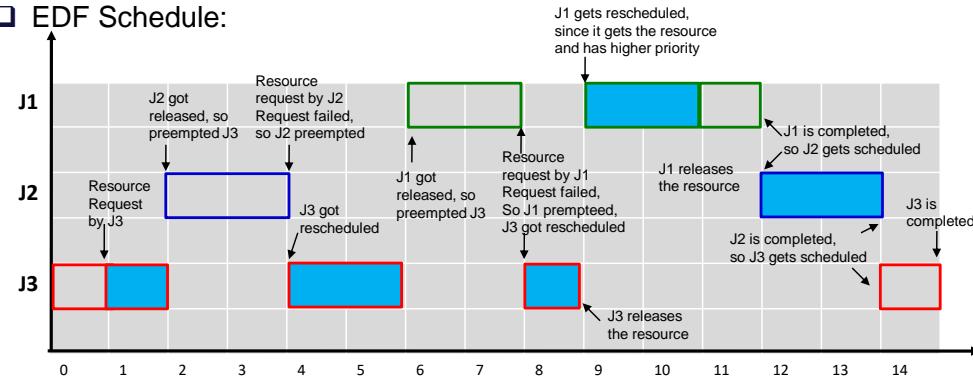


12 BITS Pilani, Pilani Campus

Resource Conflicts - Example

- Jobs **J1**, **J2**, and **J3** with feasible intervals $(6, 12]$, $(2, 14]$, $(0, 15]$
- Critical sections: $[R; 2]$, $[R; 2]$, $[R; 4]$ for jobs 1, 2, 3
- Execution times of these jobs are 5, 4 and 6 time slots respectively.
- Resource requests by the three jobs are after 2, 2 and 1 time units respectively.

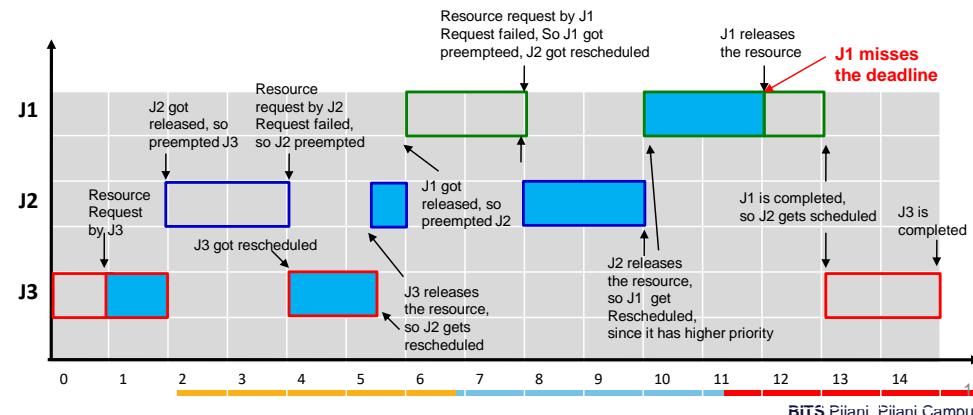
EDF Schedule:



11 BITS Pilani, Pilani Campus

Timing Anomalies

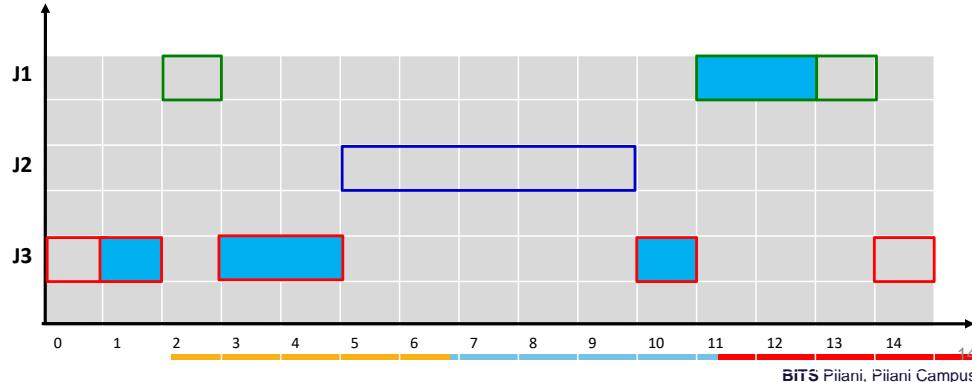
- Timing Anomalies can occur due to priority inversion
- Assume that critical section of job 3 is reduced to $[R; 2.5]$ instead of $[R; 4]$
 - Jobs **J1**, **J2**, and **J3** with feasible intervals $(6, 12]$, $(2, 14]$, $(0, 15]$
 - Critical sections: $[R; 2]$, $[R; 2]$, $[R; 2.5]$ for jobs 1, 2, 3
 - Execution times of these jobs are 5, 4 and 6 time slots respectively.
 - Resource requests by the three jobs are after 2, 2 and 1 time units respectively.



13 BITS Pilani, Pilani Campus

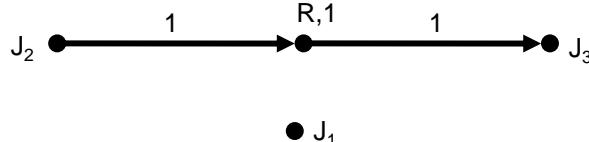
Priority Inversion - Uncontrolled

- ❑ In this example J1 is blocked by J3
- ❑ But J3 is preempted by J2
- ❑ Thus J1 has to wait for the lower-priority job J2 that does not even require the resource that J1 needs.
- ❑ In this case priority inversion is considered *uncontrolled*.
- ❑ There can be arbitrary number of jobs with priorities lower than J1 and higher than J3 released in the meantime. They can further lengthen the duration of [priority inversion].
- ❑ In fact when priority inversion is uncontrolled, then a job can be blocked for an infinitely long time.



Wait-for-graph

- Wait-for-Graphs are used to describe the dynamic-blocking relationship among jobs
- Vertexes represent jobs (J_i) and resources
- Resources are represented by R_i, k ; where R_i is the resource id and k is the number of units of resources available.
- If a job has acquired x units of resources, then it is represented by an edge (line with arrowhead) from the resource to the job with x labelled on it.
- If a job requested y units of resources, but denied earlier, then it is represented by an edge (line with arrowhead) from the job to the resource with y labelled on it.
- **A circular wait-for-graph indicates deadlock.**



- J_3 has one unit of resource R
- J_2 is waiting for one unit of resource R
- J_1 is not waiting for resource R

Deadlock

Deadlock can occur, if jobs block each other from execution

Example:

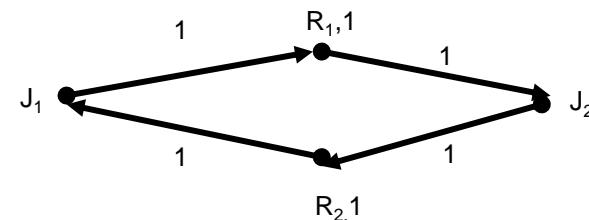
Job 1: $L(R1), L(R2), U(R2), U(R1)$

Job 2: $L(R2), L(R1), U(R1), U(R2)$



Wait-for-graph : Deadlock condition

- **A circular wait-for-graph indicates deadlock.**



- J_2 has one unit of resource R_1
- J_1 is waiting for one unit of resource R_1
- J_1 has one unit of resource R_2
- J_2 is waiting for one unit of resource R_2

So it is a deadlock.



Additional notations

- A periodic task is represented by the tuple $(\Phi_i, p_i, e_i, D_i, [R, x; y])$, where $[R, x; y]$ represents the critical section requiring x units of resource R for y units of time
- When a job requires more than one resources and has more than one critical sections, all of the critical sections are listed in the tuple.
- When the resources are not important, it is represented by the tuple $(\Phi_i, p_i, e_i, D_i, c_i)$, where c_i is the maximum execution time of the longest critical section of the job. (Please note that c_i is included in e_i)
- In general
 - A critical section here means outermost critical section (when inner critical sections are not specified)
 - A critical section of a periodic task means critical section of each job of that periodic task

Blocking time

- In a fixed priority system, let tasks are indexed in the decreasing order of the priority i.e. In a system of tasks $T_1, T_2, T_3, \dots, T_i, T_{i+1}, \dots, T_n$; T_i has higher priority than T_{i+1} .
- A job in task T_i is blocked if a the resource required by T_i is taken by any task of priority lower than T_i i.e any task between T_{i+1} to T_n .
- If c_k is the maximum critical section time of task T_k , then blocking time of Task T_i is

$$b_i(rc) = \max_{i+1 \leq k \leq n} (c_k)$$

where

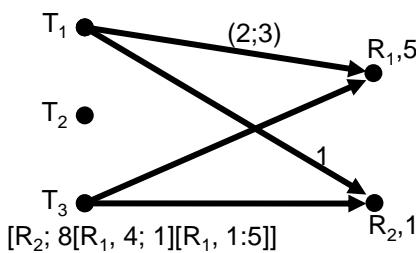
$$i = 0, 1, \dots, n$$

$$k = 1, 2, \dots, n$$

$$n \geq 1$$

Resource requirement graph

- In a resource requirement graph, there is a vertex for each job and resource.
- The integer next to each resource vertex indicates **number of units** of the resource.
- If a job requires a resource, there will be a edge drawn from the job to the resource
- The edge will be labelled by the 2 tuple (x, y) , where the job requires x units of the resources for y units of time in the critical section.
- If there is one unit of resource, then first parameter of the tuple is omitted i.e. Only the time units required for the critical section is mentioned



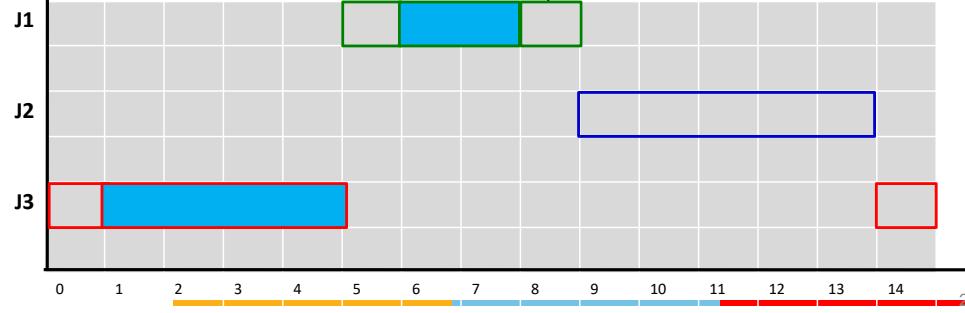
- Each job of T_1 requires 2 units of R_1 for 3 units of time and 1 unit of R_2 for 1 unit of time
- R_1 has 5 units of resources and R_2 has 1 unit of resource.
- T_2 doesn't require any resources
- Each job of T_3 requires R_1 and R_2 as specified by the critical section (rather than indicating over the edges)

Nonpreemptive Critical Section (NPCS) Protocol

- **When a job holds a resource, it executes at a priority higher than the priorities of all jobs.**
- Because no job is ever preempted when it holds a resource, **Deadlock can never occur**.
- **Uncontrolled priority inversion** can't occur

Previous Example

- Priorities of J_1 and J_3 are highest and lowest, J_2 has priority lower than J_1 , but higher than J_3
- When J_3 is released, it acquires the resource and continues till the end of the critical section
- When J_1 is released at time 2, it waits till time 5, when J_3 releases the resource.
- At time 5, J_1 Gets scheduled. At the same time J_2 is released, but it doesn't get scheduled, since J_1 has higher priority.
- At time 6, J_1 acquires the resource and holds it till time 8.
- J_1 completes at time 9. At this time J_2 is scheduled.
- J_2 continues till time 14. Then J_3 is scheduled.
- J_3 completes at time 15.



Nonpreemptive Critical Section (NPCS) Protocol



Advantages

- Simple to implement
- Uncontrolled priority inversion cannot occur
- Good protocol when critical sections are short

Disadvantages

- Every job can be blocked by every lower-priority job even if there is no resource conflict between them

22
BITS Pilani, Pilani Campus

Priority Inheritance Protocol - Rule



Scheduling Rule

- Ready jobs are scheduled on the processor preemptively in a priority-driven manner according to their current priorities.
- At its release time t , the current priority $\pi(t)$ of every job J is equal to its assigned priority
- The job remains at this priority except under the condition stated in the priority-inheritance rule

Allocation Rule

- When a job J requests a resource R at time t ,
 - if R is free, R is allocated to J until J releases the resource, and
 - if R is not free, the request is denied and J is blocked

Priority-Inheritance Rule

- When the requesting job J becomes blocked, the job J , which blocks J inherits the current priority $\pi(t)$ of J .
- The job J executes at its inherited priority $\pi(t)$ until it releases R
- At that time, the priority of J returns to its priority $\pi_i(t')$ at the time t' when it acquired the resource R

24
BITS Pilani, Pilani Campus

Priority Inheritance Protocol



- Like NPCS, it doesn't require any prior knowledge on the resource requirements of the jobs.
- Uncontrolled priority inversion cannot occur
- It does not avoid deadlock. External mechanisms needed to avoid deadlock

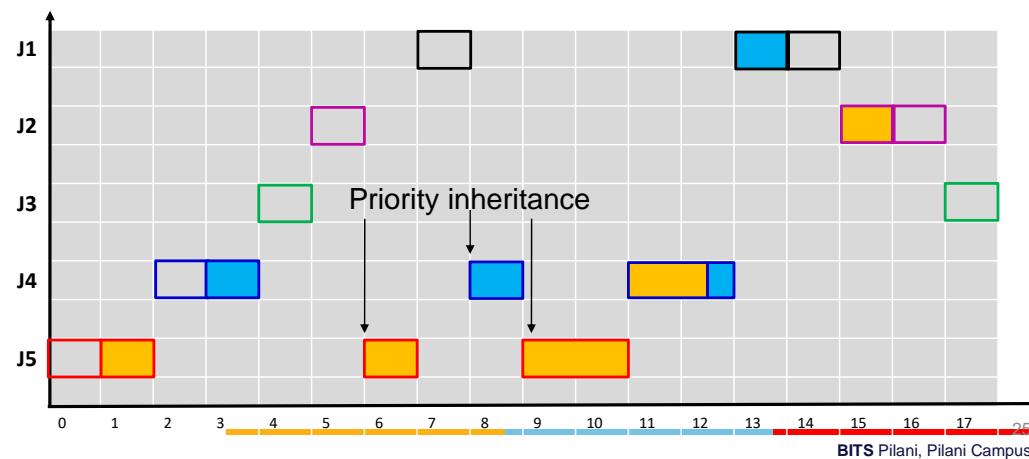
23
BITS Pilani, Pilani Campus

Priority Inheritance Protocol - Example



The jobs require resources as mentioned 1 time slot after start of their execution

Job	r_i	e_i	Π_i	Critical sections
J1	7	3	1	[Light blue; 1]
J2	5	3	2	[Orange; 1]
J3	4	2	3	
J4	2	6	4	[Light blue; 2.5[Orange; 1.5]] (Light blue: 2 slots, Orange: 1.5 slots, Light blue: 0.5 slots)
J5	0	6	5	[Orange; 4]



25
BITS Pilani, Pilani Campus

Priority Inheritance Protocol - Example



At time 0, J5 is released. Since there were no other jobs, it was scheduled.

At time 1, it was granted resource 'orange'

At time 2, CJ4 was released, so J5 was preempted and J4 was executed

At time 3, J4 required light blue resource and was granted

At time 4, J3 was released, so J4 was preempted. J3 didn't require any resource.

At time 5, J2 was released, which preempted J3.

At time 6, J2 required orange resource. But it was held by J5. As per priority inheritance rule, the priority of J5 was made that of J2, hence it was scheduled and J2 was preempted.

At time 7, J1 is released, which is of highest priority. So J5 gets preempted and J1 is scheduled

At time 8, J1 requires light blue resource, which is already acquired by J4. Hence as per priority inheritance rule, priority of J4 is made equal to that of J1. So J1 gets preempted and J4 gets scheduled.

At time 9, J4 requests for resource Orange, which is acquired by J5. Hence J5 inherits the priority of J4 (in turn of J1) and executes.

At time 11, J5 releases the Orange resource, its priority is back to original, so gets preempted. J4 gets scheduled and acquires the orange resource

At time 12.5, J4 completes using orange resource. Then it uses the blue resource and releases it at time 13.

At time 13, J4 releases light blue resource, so its priority becomes the original one. So it gets preempted and J1 executes acquiring the light blue resource

At time 14, J1 releases the light blue resource. But it is the highest priority task hence continues till its completion at time 15.

At time 15, next high priority task J2 executes, acquires the orange resource.

At time 16, J2 releases the orange resource, but being the highest priority task in the system continues till its completion at time 17.

At time 17 J3 gets scheduled and continues till its completion at time 18.

At time 18, J4 gets scheduled and it completes at time 19.

At time 19, J5 gets scheduled and completes at time 20.

Priority Ceiling Protocol



➤ Priority Ceiling Protocol extends Priority Inheritance Protocol to prevent deadlock.

➤ Two key assumptions:

- Assigned priorities of the jobs are fixed.

- Resource required by all jobs are known *a priori* before the execution of any job begins.

➤ Priority Ceiling of any resource R_i is the highest priority of all the jobs that require the resource R_i , denoted by $\Pi(R_i)$.

➤ At any time t , the current priority ceiling (or simply ceiling) $\Pi(t)$ of the system is equal to the highest priority ceiling of the resources that are in use at the time

➤ If all resources are free, $\Pi(t)$ is equal to Ω , a nonexistent priority level that is lower than the lowest priority level of all jobs

Priority Ceiling Protocol - Rule



Scheduling Rule

- Ready jobs are scheduled on the processor preemptively in a priority-driven manner according to their current priorities.
- At its release time t , the current priority $\pi(t)$ of every job J is equal to its assigned priority
- The job remains at this priority except under the condition stated in the priority-inheritance rule

Priority Ceiling Protocol - Rule



Allocation Rule

When a job J requests a resource R at time t ,

- if R is not free, the request is denied and J is blocked
- If R is free,
 - If J 's priority $\pi(t)$ is higher than the current priority ceiling $\Pi(t)$, R is allocated to J .
 - If J 's priority $\pi(t)$ is not higher than the current priority ceiling $\Pi(t)$, R is allocated to J only if J is the job holding the resource(s) whose priority ceiling is equal to $\Pi(t)$, otherwise J 's request is denied and J becomes blocked.

Priority-Inheritance Rule

- When the requesting job J becomes blocked, the job J_i which blocks J inherits the current priority $\pi(t)$ of J .
- The job J_i executes at its inherited priority $\pi(t)$ until it releases every resource whose priority ceiling is equal to or higher than $\pi(t)$ of J .
- At that time, the priority of J_i returns to its priority $\pi_i(t')$ at the time t' when it acquired the resource R

Priority Ceiling Protocol - Rule



Please note:

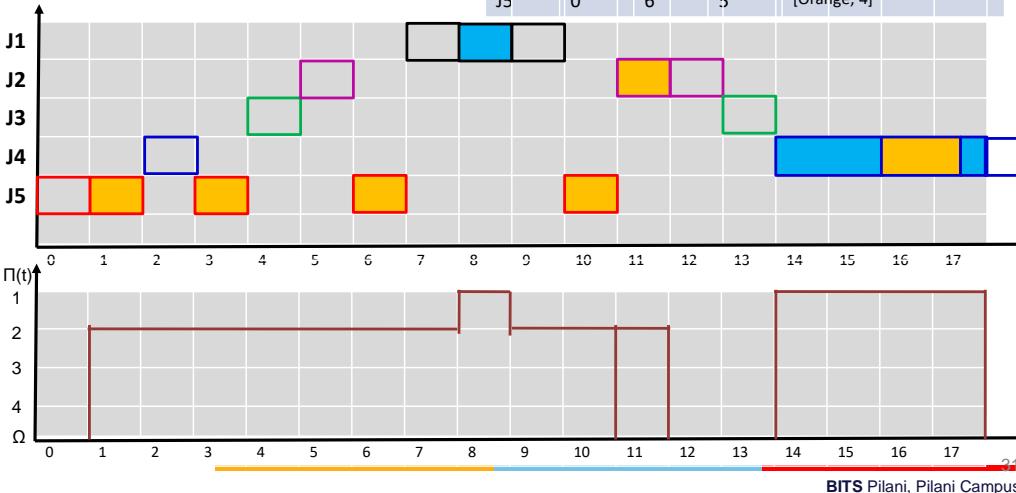
- The Allocation rule assumes that only one job can hold all the resources with priority ceiling equal to $\Pi(t)$.
- The Priority inheritance rule assumes that only one job is responsible for other jobs' request being denied, because it holds either the requested resource or a resource with priority ceiling $\Pi(t)$.

20 BITS Pilani, Pilani Campus

Priority Ceiling Protocol - Example



Job	r_i	e_i	T_i	Critical sections
J1	7	3	1	[Light blue; 1]
J2	5	3	2	[Orange; 1]
J3	4	2	3	
J4	2	6	4	[Light blue; 2.5[Orange; 1.5]
J5	0	6	5	[Orange; 4]



Priority Ceiling Protocol - Example



At time 0, J5 is released. Since there were no other jobs, it was scheduled.

At time 1, it was granted resource 'orange', since the ceiling was Ω , because all resources were free. At this time the ceiling $\Pi(t) = 2$, the priority ceiling of resource 'orange'.

At time 2, J4 was released, so J5 was preempted and J4 was executed

At time 3, J4 required light blue resource. But J4's priority is 4, which is less than the ceiling 2. So J4 is blocked. J5 is the job which blocked J4, since it holds the 'orange' resource which determines the ceiling. So J5 inherits J4's priority i.e. 4 and executes.

At time 4, J3 was released, so J5 was preempted, since J3 has higher priority than the J5's inherited priority. J3 didn't require any resource.

At time 5, J2 was released, which preempted J3.

At time 6, J2 required orange resource. But it was held by J5. As per priority inheritance rule, the priority of J5 was made that of J2, hence it was scheduled and J2 was preempted.

At time 7, J1 is released, which is of highest priority. So J5 gets preempted and J1 is scheduled

At time 8, J1 requires light blue resource. J1 has higher priority than the ceiling. So it is granted the resource. The ceiling $\Pi(t) = 1$. It releases the resource at time 9. Then it continues execution (being the highest priority job) till its completion i.e. Time 10.

At time 10: J5 is scheduled, since it has the next highest priority because of priority inheritance (i.e. 2).

At time 11, J5 releases the Orange resource, its priority is back to original i.e. 5, so gets preempted. The ceiling of the system becomes Ω , since no jobs hold any resource. So highest priority job J2 get scheduled, which acquires the 'orange' resource'. So the ceiling of the system again becomes 2.

At time 12: J2 releases the 'orange' resource. Being the highest priority job in the system, it continues till completion i.e. Time 13.

At time 13: J3 gets scheduled being the highest priority job in the system. It continues till completion i.e. Time 14.

At time 14: J4 gets scheduled. It uses the 'light blue' resource for 2 time slots, then also uses 'orange' resource for 2 time slots and then uses the 'blue' resource again for 0.5 time slots till time 18. After then it executes for one more slot and gets completed.

At time 19, J5 gets scheduled and completes at time 20.

22 BITS Pilani, Pilani Campus

Thank You.

Any Questions?



Text Book / References

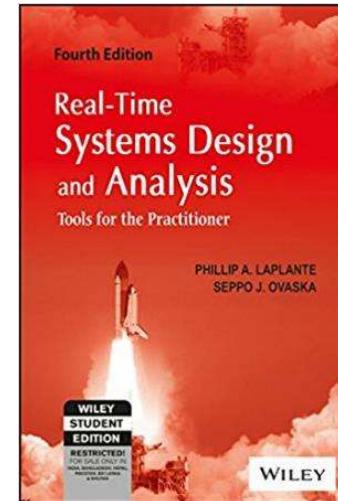


BITS Pilani
Pilani Campus

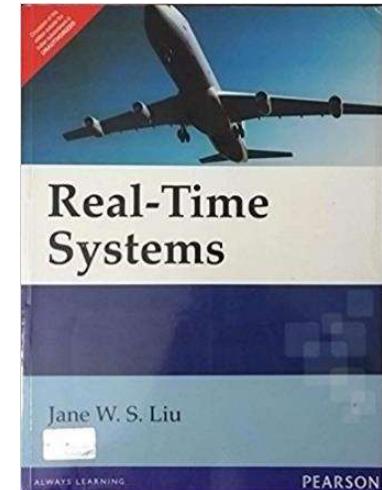
BITS ZG553: Real Time Systems L9 – Hardware for Real Time Systems

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Reference (R1)



Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Excellent MOOCs Videos (Coursera, edX,...)



HONOR CODE
CERTIFICATE



This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

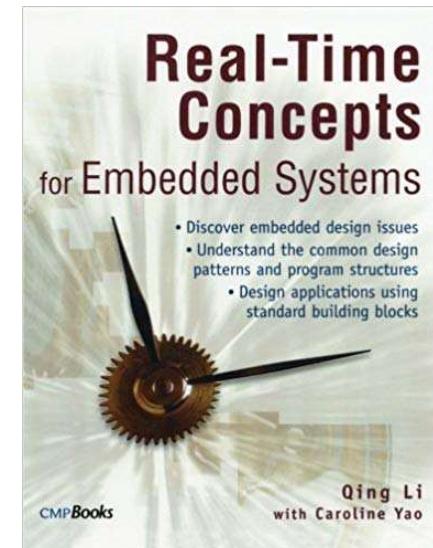
HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44a218f5a1eebf47d4e54

Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

RTS Primer – For Light Reading





L-9: Hardware for Real Time Systems

[Processor Architecture, Memory, Peripheral Interfacing]
 [Ref: Notes/PPT]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

Processor Architectures

- Von Neumann Architecture
- Harvard Architecture

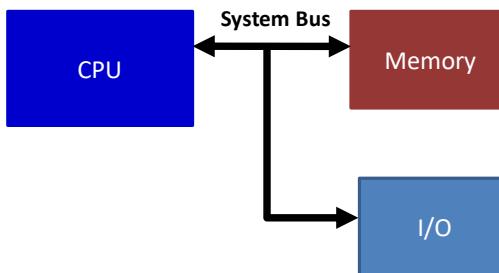
Motivation

- Many real-time systems are **embedded** systems
- Embedded systems interact closely with specialized input and output devices including sensors and actuators
- Hardware structure can affect timing and require special software design techniques
- Real-time systems engineers must have some familiarity with hardware and do some “low level” programming

BITS Pilani, Pilani Campus

Von Neumann Architecture

- Described in 1945 by the mathematician and physicist **John von Neumann** and others in the First Draft of a Report on the EDVAC.
- Also known as **Princeton Architecture**.
- CPU is connected to Memory and I/O through the System Bus.
- In a Von Neumann Architecture, **I/O is Memory Mapped**.



System Bus constitutes:

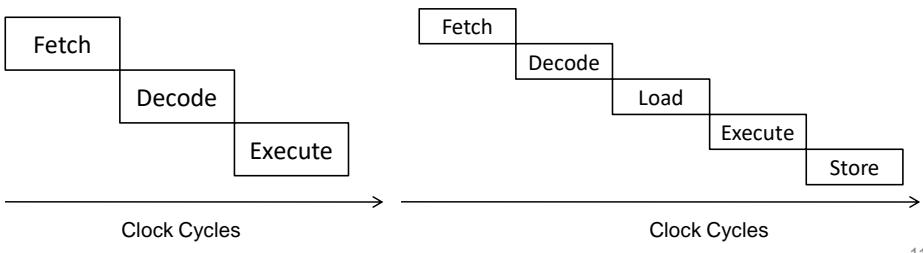
- Unidirectional **Address Bus**
- Bidirectional **Data Bus**
- **Control Bus**, which is a heterogeneous collection of independent lines such as control, status and power lines.

Harvard Architecture

- The name Harvard Architecture comes from the Harvard Mark I relay-based computer.
- Separate data and address bus
- It is possible to access program memory and data memory simultaneously.
- Typically, code (or program) memory is read-only and data memory is read-write. Therefore, it is impossible for program contents to be modified by the program itself.
- But the price paid is increase in chip area (therefore power consumption)
- Many modern CPUs comprises both Harvard and von Neumann characteristics: the separate on-chip instruction and data caches have a Harvard type interface, while the common off-chip cache memory is interfaced through a single system bus.

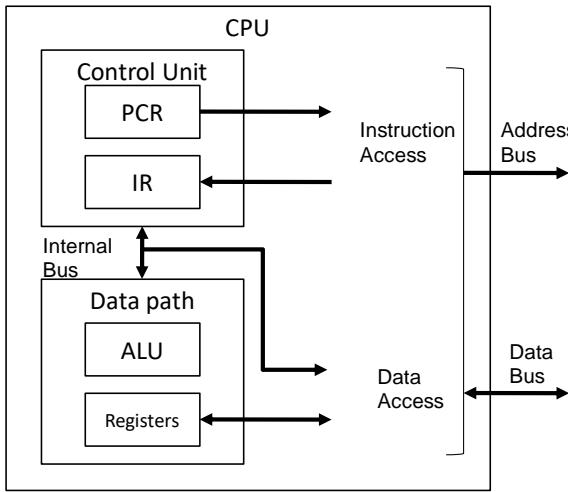
Instruction Processing

- Instruction processing requires multiple consecutive stages.
- Typical instruction processing stages constitute: **Fetch, Decode and Execute**.
- But these stages can be further divided into multiple sub-stages. For example, we can have 5 stages for instruction processing: **Fetch, Decode, Load, Execute and Store**.
- The duration of each instruction may vary, because each of these stages may take different amount of clock cycles.



CPU Internal Architecture

- Instruction processing takes place inside the CPU
- It has 3 parts: Control Unit, Internal Bus and Data Path



- Control Unit
 - Interfaces to the system bus through Program Counter Register (PCR)
 - PC points to the address from which the next instruction to be fetched.
 - Instruction fetched is stored in the Instruction Register (IR)
 - Instruction is then decoded, desired operands are fetched from memory and appropriate command is given to the ALU of the Data Path
- Data Path
 - ALU (Arithmetic and Logic Unit), performs the desired operation.
 - The output of ALU is either stored in Registers of Memory depending upon the Instructions
 - Also various status flags such as Carry/Borrow, Zero etc are updated in the status registers depending upon the state of the CPU. These flags are used for implementing conditional instructions.

Instruction Processing

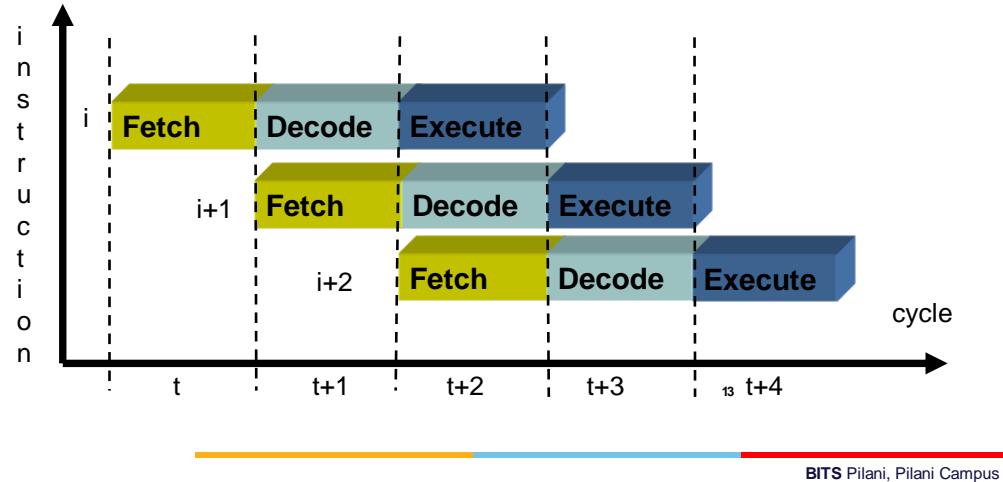
- Each Instruction consists of an **Opcode** and zero, one or more **Operands**.
Example:
MOV AX, BX
Here, MOV is the Opcode and AX and BX are the Operands
- Addressing modes are the way CPU accesses the Operands

Immediate:	ADD AX, 5 ; AX = AX + 5
Register Direct:	SUB AX, BX ; AX = AX - BX
Register Indirect:	ADD AX, [BX]; AX = AX + *BX
.....
- Two principal techniques for implementing the control unit
 - **Microprogramming**
 - Every instruction consists of a sequence of primitive hardware commands, microinstructions for activating the data path.
 - Suitable when the number of instructions is large
 - Instruction processing is slower
 - **Hardwired Logic**
 - Each instruction is mapped to some combinational and sequential logic hardware.
 - Suitable when the number of instructions is small
 - Results in fast instruction processing
- In order to conserve power, modern processors implements '**slowdown modes**'.
 - Specific instructions enable such modes, where the voltage and/or frequency are lowered, reducing the power consumption.

Pipeline

Example of a 3-stage pipeline:

- Fetch, Decode, Execute
- 3 cycle latency



Where do Processors Spend Time ?

Instruction Type	Dynamic Usage
Data movement	43%
Control flow	23%
Arithmetic operations	15%
Comparisons	13%
Logical operations	5%
Others	1%

Source: Steve Furber, ARM System-on-chip Architecture, Second Edition, Pearson, 2007

Pipeline Hazards

- **Structural Hazards**
 - These occur when a single piece of hardware is used in more than one stage of the pipeline, so it's possible for two instructions to need it at the same time.
 - Example: One memory unit being shared by subsequent instructions)

- **Data Hazards**
 - This is when reads and writes of data occur in a different order in the pipeline than in the program code.
 - **RAW (Read After Write):** Occurs when, one instruction reads a location after an earlier instruction writes new data to it, but in the pipeline the write occurs after the read.
 - **WAR (Write After Read):** A write occurs after a read, but the pipeline causes write to happen first.
 - **WAW (Write After Write):** Two writes occur out of order.

In order to resolve these hazards, **pipeline stall** happens.

- **Control Hazards**
 - This is when a decision needs to be made, but the information needed to make the decision is not available yet.
 - Example: Branch instruction in a pipeline.

In some RISC architectures, the instruction following the branch is executed whether or not the branch is taken. It is called **delayed branch**.

BITS Pilani, Pilani Campus

CISC Processor Architecture

- **CISC – Complex Instruction Set Computer**
- Instructions are of variable format and length
- Large number of instructions
- Practically any instruction can reference memory
- Single set of work registers
- Micro-programmed instruction decode logic. Complexity of instruction is handled at this level
- Great variety of addressing modes
- No instruction pipelines
- Programmer's coding effort is simplified
- Example: x86 Architecture

RISC Processor Architecture



- RISC – Reduced Instruction Set Computer
- Fixed 32-bit instruction size
- Load-store architecture – Instructions operating on data operates only on registers.
- Large register bank of thirty-two 32-bit registers
- Hard-wired instruction decode logic (CISC processors use large microcode ROMs)
- Pipelined execution
- Single-cycle execution (most of the instructions)
- Performance is enhanced
- Example: ARM Architecture

Multi-core Processor System



- From HW Perspective:
 - Homogenous: If all CPU cores are of same architecture
 - Heterogeneous: If all CPU cores are of different architecture
- From SW Perspective:
 - AMP: Asymmetric Multi-processing
 - SMP: Symmetric Multi-processing

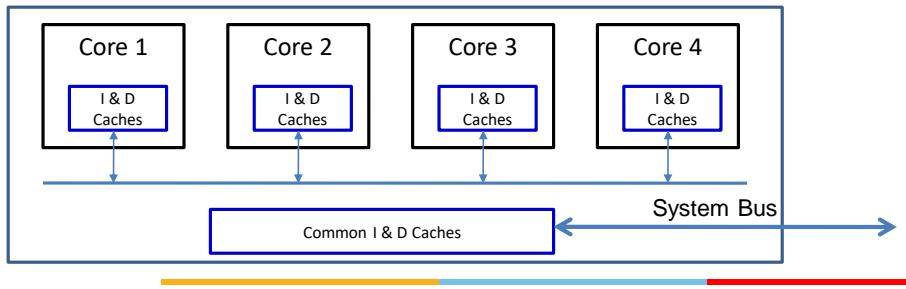
Multi-core Processor

Motivation

- CPU architecture is evolving, with higher frequencies, longer pipelines.
- As clock frequencies increases, power consumption increases.
- Cooling such high speed processors is an costly affair.
- Then why not have processors with multiple cores, operating in parallel ?

Multi-core processors

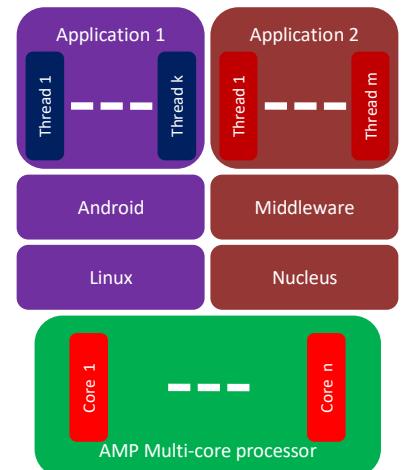
- Have individual cores with private cache memory, separate instruction and data caches.
- These small on-chip caches are interfaced to larger on-chip common cache memories



BITS Pilani, Pilani Campus

Multi-core Processor System - AMP

- CPU cores can be of homogenous or heterogeneous architectures
- Each core has own address space
- Used when different architectures are optimal for specific activities, such as: one is for DSP and other one is for MCU
- There is an opportunity to deploy different O/Ses on different cores

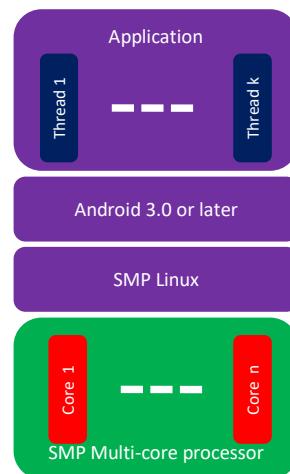


BITS Pilani, Pilani Campus

Multi-core Processor System - SMP



- CPU cores are of homogenous architecture
- CPU cores share the memory space
- A single O/S runs on all the cores
- Some kind of communication facility between the CPUs is provided this is normally through shared memory, but accessed through the API of the OS
- Typically, SMP is used when an embedded application simply needs more CPU power to manage its workload



BITS Pilani, Pilani Campus

Common DSP Processor features



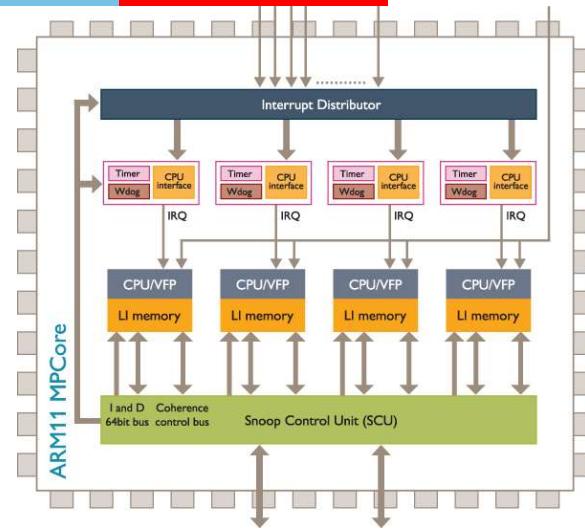
- Harvard architecture
 - Needed for accessing multiple memory locations simultaneously
- Dedicated single-cycle Multiply-Accumulate (MAC) instruction (hardware MAC units)
 - Required for processing arithmetic operations like
- Single-Instruction Multiple Data (SIMD)
- Very Large Instruction Word (VLIW) architecture
- Pipelining
- Saturation arithmetic
- Special Execution Units with
 - Minimum overhead looping
 - Low interrupt and branch latency
 - Fast context switching
 - Register banking / stacking
- Cache
- DMA

$$Y = \sum_{n=1}^N a_n * x_n$$

BITS Pilani, Pilani Campus

ARM11 MPCore Architecture

- **SMP Architecture**
- Shared L1 cache is managed by Snoop Control Unit:
 - Each CPU's snooping unit looks at writes from other processors
 - If a write modifies a location in this CPU's level 1 cache, the snoop unit modifies the locally cached value
- **Interrupt Distributor**
 - Masks and prioritizes the interrupts as in standard interrupt system
 - Identifies the set of CPUs that can handle the interrupts
 - Sends each CPU its highest priority pending interrupt



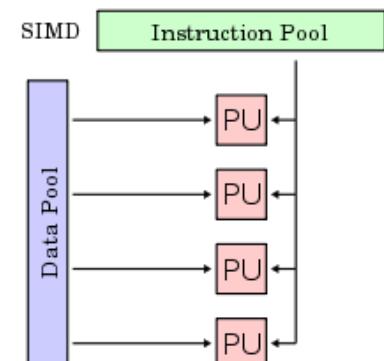
Source: <http://www.arm.com/products/processors/classic/arm11/arm11-mpcore.php>

BITS Pilani, Pilani Campus

SIMD – Single Instruction Multiple Data

One instruction operates on multiple data simultaneously.
Example: Addition of two row vectors:

$$\begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{bmatrix} + \begin{bmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{bmatrix}$$



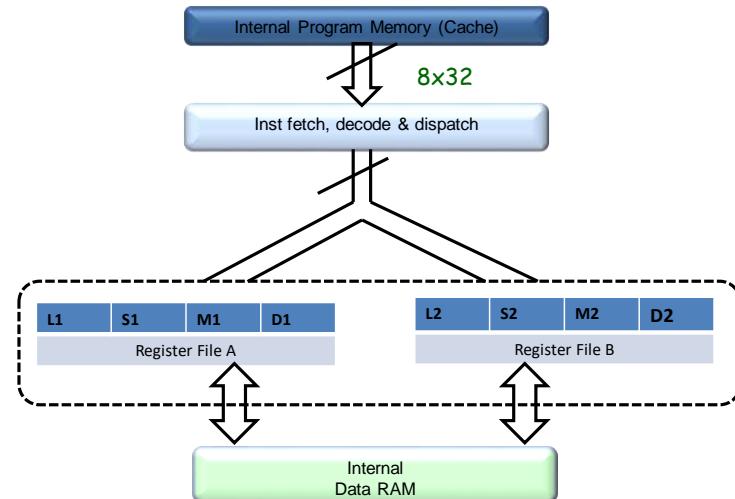
Same ADD operation on 4 separate data in parallel

BITS Pilani, Pilani Campus

VLIW Architecture

- **VLIW - Very Long Instruction Word**
- Exploits instruction level parallelism
- Has **multiple execution units** executing in parallel
- Hence has **multiple register files** to help parallel execution of instructions
- Instruction word is a combination of multiple instructions, hence it is very long
- **Compiler reorders the instructions** in order to achieve maximum parallel execution possible. Hence VLIW compiler is very complex.

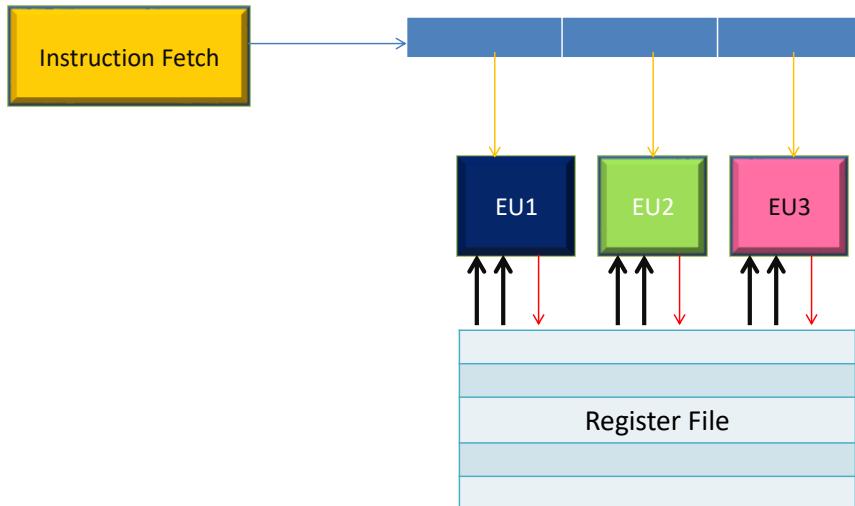
VLIW Architecture



BITS Pilani, Pilani Campus

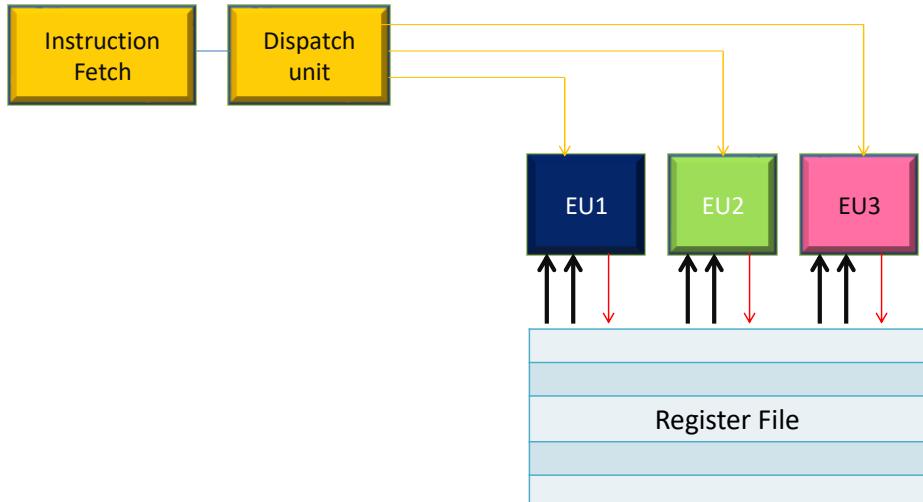
BITS Pilani, Pilani Campus

VLIW Architecture



BITS Pilani, Pilani Campus

Superscalar Architecture



BITS Pilani, Pilani Campus

Difference between VLIW and Superscalar Architectures



VLIW

- Single and very long instruction word containing multiple instructions
- Compiler does instruction reordering (Static instruction scheduling)

Superscalar

- Multiple instructions are issued separately
- Hardware does instruction reordering (Dynamic instruction scheduling)

Any Questions on
Processor Architectures?

Memory Systems



- Memory access has major contribution in latency.
- It becomes important when the utilization factor is in dangerous zone (83% - 99%) (Refer to lecture 2)
- In such systems hierarchical memory architecture may cause aperiodic jobs to miss deadlines with considerable delays.
- So it is necessary to understand memory technologies, so that appropriate design decision can be taken

Different Classes of Memory



- Traditional Distinction between Memories:
 - ROM (Read Only Memory) – Non-volatile
 - RAM (Random Access Memory) – Volatile
- ROMS used to get programmed once at the beginning using a special programming units.
- But at present, EEPROM and Flash memories can be rewritten without any special programming units.
- At present following types of Memories are widely used
 - Non Volatile
 - EEPROM
 - Flash
 - Volatile
 - SRAM
 - DRAM

Non-Volatile Memory

- Erasing is slower than RAMs
- Each device can be rewritten 1000,000 – 1,000,000 times only.
- EEPROMs are usually used for storing nonvolatile programs and parameters.
- Flash memory is usually used for storing both application programs (including operating system) and large data records.
- These are slower to read compared to RAM
- So it is a common practice in many real-time systems to load the application-program (including the operating system) from a flash memory to the RAM and then run it from the RAM.
- In many cases the flash memory used is a movable one (e.g. a SD Card or a USB Memory stick)

Evolution of DRAM

- First page mode (FPM) DRAM
- Extended data output (EDO) DRAM
- Synchronous DRAM (SDRAM)
- Direct Rambus DRAM (DRDRAM)
- Double data rate 3 synchronous DRAM (DDR3 DRAM)

SRAM and DRAM

- SRAM-type memory cell requires typically 6 transistors, while DRAM cell requires a single transistor and a capacitor only.
- So SRAM takes more chip area compared to DRAM.
- Hence SRAMs are costlier than DRAMs.
- But SRAMs are faster than DRAMs.
- DRAM has the issue of leakage in their storage capacitors. So DRAMs need to be refreshed periodically (typical refresh cycle is 3-4 ms).

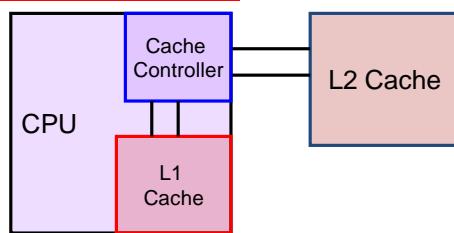
While designing a RAM subsystem of a real-time system, the basic rule of thumb is that if you need a large amount of memory, then use DRAM, otherwise go for SRAM.

Memory Hierarchy

- CPU Clock Rates are increasing at a faster rate compared to DRAM access times.
- Programmers always desire faster and infinite memory.
- Solution is Memory Hierarchy.
- In order of fastest to slowest, memory should be assigned, considering cost as follows:
 - internal CPU memory
 - registers
 - cache
 - main memory
 - memory on board external devices

Cache

- Cache memories rely on the 'locality of reference' principle.
- Locality of reference
 - Refers to the address distance between memory between consecutive code and data access.
 - If the code or data fetched tends to reside close in memory, then the locality of reference is low.



- A cache is a relatively small storage of fast memory where frequently used instructions and data are kept.
- Basic operations of Cache
 - Suppose CPU requests the content of a DRAM.
 - First the cache controller checks the address to see if the particular location is in the cache.
 - If present, the data is immediately retrieved from the cache, which is significantly faster.
 - If the address is not in the cache, then a cache miss occurs. Subsequently the cache content gets written back to the main memory and the new block is loaded into the cache.
- Cache Hierarchy
 - L1 Cache: Nearest cache to CPU
 - L2 Cache: Next level cache between L1 cache and Main Memory

37
BITS Pilani, Pilani Campus

Cache Access Time - Example

Problem:

Assume that the system has a two-level cache. The level-1 cache has a hit rate of 90% and the level-2 cache has a hit rate of 97%. The level 1 cache access time is 4 ns, the level 2 cache access time is 15 ns and access time of main memory is 80 ns. What is the average memory access time?

Solution:

$$\begin{aligned}
 h_1 &= 0.9, \\
 h_2 &= (1-0.9) \times 0.97 = 0.1 \times 0.97 = 0.097, \\
 t_{L1} &= 4 \text{ ns}, t_{L2} = 15 \text{ ns}, t_{main} = 80 \text{ ns}. \\
 t_{av} &= h_1 t_{L1} + h_2 t_{L2} + (1-h_1-h_2) t_{main} \\
 &= 0.9 \times 4 \text{ ns} + 0.097 \times 15 \text{ ns} + (1 - 0.9 - 0.097) \times 80 \text{ ns} \\
 &= 0.9 \times 4 \text{ ns} + 0.097 \times 15 \text{ ns} + 0.003 \times 80 \text{ ns} \\
 &= 3.6 \text{ ns} + 1.455 \text{ ns} + 0.24 \text{ ns} \\
 &= 5.295 \text{ ns}
 \end{aligned}$$

Cache Access Time

- t_{main} – access time for main memory
- t_{L1} – access time for primary cache
- h_1 – hit ratio of primary cache
- t_{L2} – access time for secondary cache
- h_2 – hit ratio of secondary cache but not primary

If system has only L1 cache

$$\text{Average access time, } t_{av} = h_1 t_{L1} + (1-h_1) t_{main}$$

If system has L1 and L2 cache

$$\text{Average access time, } t_{av} = h_1 t_{L1} + h_2 t_{L2} + (1-h_1-h_2) t_{main}$$

38
BITS Pilani, Pilani Campus

Any Questions on
Memory Systems?

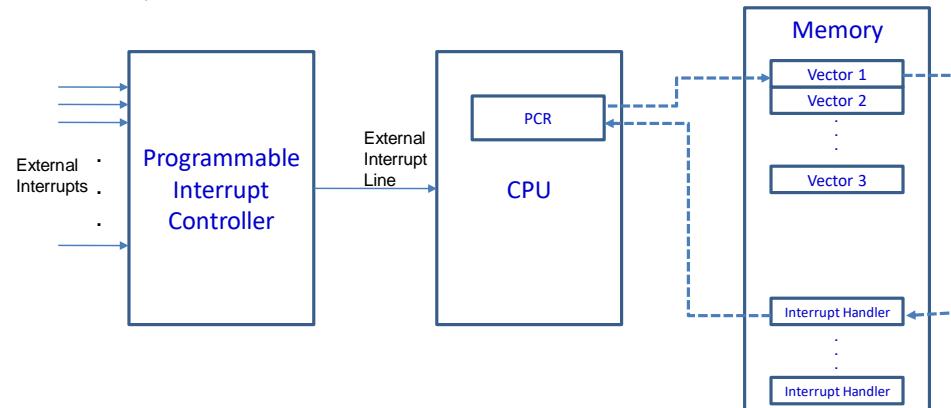
Peripheral Interfacing: Interrupts

- An hardware interrupt is an external hardware signal that initiates an event.
- Each interrupt is associated with an interrupt vector address, which contains the interrupt handler routine.
- Interrupt handler routine is a program for the desired action to be taken when the interrupt occurs.
- There are software interrupts as well, which are generated internally.
- A typical interrupt service process is as follows
 - The interrupt request line is activated
 - The interrupt request is latched by the CPU hardware
 - The processing of ongoing instruction is completed
 - The content of program counter register (PCR) is pushed to the stack
 - The content of status register (SR) is pushed to the stack
 - The PCR is loaded with the interrupt handler's address.
 - The interrupt handler is executed.
 - The original content of the SR is popped back from the stack.
 - The original content of the PCR is popped back from the stack.

41
BITS Pilani, Pilani Campus

Programmable Interrupt Controller (PIC)

- When number of interrupt sources are more than the available external interrupt inputs on a microprocessor/microcontroller, then PIC chips are used.
- Priorities can be assigned among the interrupts.
- Example: 8059 for x86

42
BITS Pilani, Pilani Campus

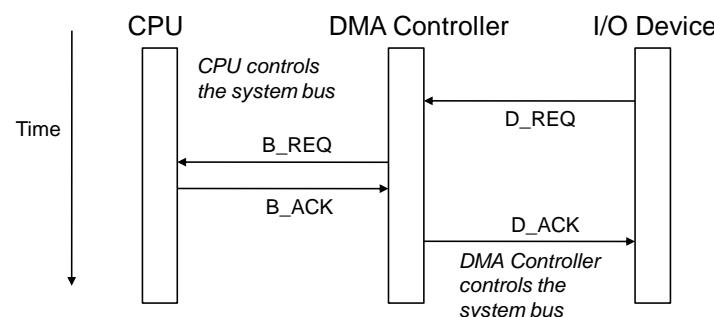
Direct Memory Access (DMA)

- Access to the computer's memory is given to other devices in the system without CPU intervention.
- Information is transferred directly into main memory by the external device.
- DMA controller is required unless the DMA circuitry is integrated into the CPU.
- Because CPU participation is not required, data transfer is fast.
- DMA is the most preferred method for data transfer in real-time systems.

43
BITS Pilani, Pilani Campus

DMA Transfer Process

- I/O Device requests DMA transfer by activating DMA-request signal (D_REQ).
- This cause the DMA Controller issue a bus-request signal (B_REQ) to the CPU.
- Then the CPU finishes its present bus cycle and activates a bus-acknowledgement signal (B_ACK).
- After recognizing B_ACK signal, DMA Controller activates a DMA-acknowledgement signal (D_ACK) to the I/O device.
- Then I/O device performs the data transfer to the memory.
- Once the transfer is done, the DMA controller deactivates the B_REQ signal, giving the bus back to the CPU.

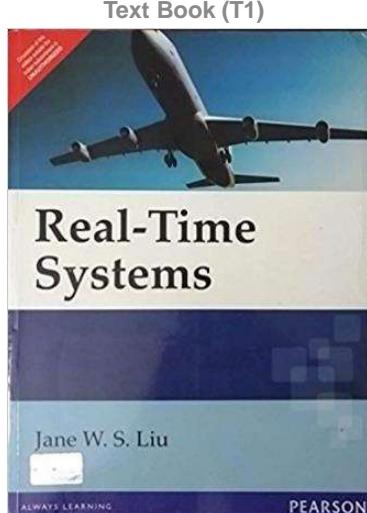
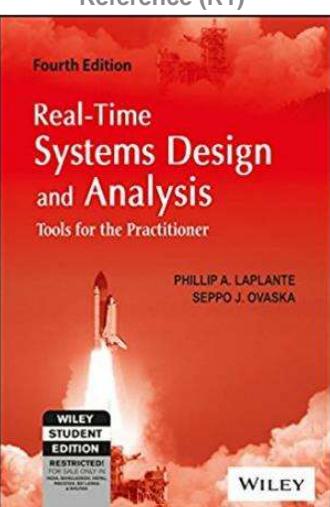
44
BITS Pilani, Pilani Campus

Thank You.

Any Questions?

45
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Text Book / References



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

BITS ZG553: Real Time Systems
L10 – Real Time Operating Systems

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Excellent MOOCs Videos
(Coursera, edX,...)

HONOR CODE
CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

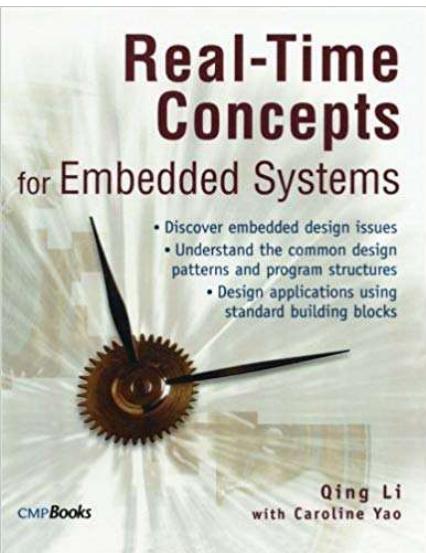
HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54



Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

RTS Primer – For Light Reading



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

4

Operating Systems

The operating system controls resources:

- who gets the CPU;
- when I/O takes place;
- how much memory is allocated.

The most important resource is the CPU itself.

- CPU access controlled by the scheduler.



L-10: Real Time Operating Systems

[Ref: Notes/PPT]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

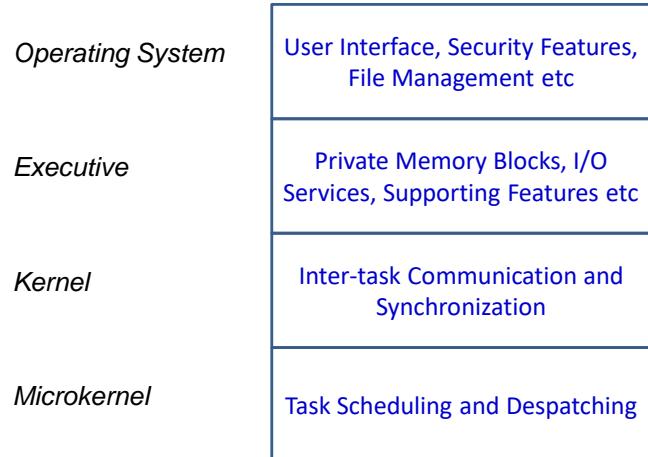
Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

Real Time Operating System

- Real-time operating systems must provide three specific functions with respect to tasks:
 - scheduling
 - dispatching
 - intercommunication and synchronization
- A scheduler determines which task will run next in a multitasking system
- A dispatcher performs the necessary bookkeeping to start that task.
- Inter-task communication and synchronization assures that the tasks cooperate.

Layers of Operating Systems



Pseudokernels



System without O/S and Interrupt.

Polled Loop:

```
while (1)
{ /* do forever */
  if (packet_here) /* check flag */
  {
    process_data(); /* process data */
    packet_here=0; /* reset flag */
  }
}
```

Pros:

- Fast response

Cons:

- Can fail due to burst of events
- Generally not sufficient to handle complex systems
- Waste of CPU time, especially when event being polled occurs infrequently

Pseudokernels

-
- Real-time multitasking in rudimentary form can be achieved **without operating systems and interrupts**.
 - These systems are known as **pseudo-kernels**.
 - In this approach, the resultant systems are often **highly predictable**.
 - Hence these are sometime preferred in **low-end embedded systems**.



Pseudokernels

Polled Loop with Delay:

-
- Used to treat a problematic events exhibiting **Contact Bounce behaviour**.
 - Any electromechanical switch can't change its state instantaneously. It needs some time to settle down the oscillations. It is called Contact Bounce behaviour.
 - So a delay equal to the time little more than the contact bounce behaviour disappears, is introduced in the polled loop

In the following example, contact bounce behaviour disappears after 20 ms. So delay of 20 ms + 1 ms (to be on the safer side) = 21 ms is introduced.

```
while(1)
{ /* do forever */
  if(flag) /* check flag */
  {
    process_event(); /* process event */
    pause(21); /* wait 21 ms */
    flag=0; /* reset flag */
  }
}
```

- The flag is set by an external device.
- Delay (pause functionality) is generated by a timer.

Pros:

- Excellent for high speed data channels, especially when the events occurs at widely dispersed intervals and CPU is dedicated to handle the data channel

Cons:

- Can fail due to burst of events
- Generally not sufficient to handle complex systems
- Waste of CPU time, especially when event being polled occurs infrequently

Pseudokernels

Cyclic Code Structure:

```
while (1)
{ /* do forever */
    Process_1();
    Process_2();
    ...
    Process_N();
}
```

- Tasks are executed in **round-robin fashion**
- Higher priority can be given to a task by repeating it.
- The task list can be managed by a ‘pseudo-operating system’, as the tasks are created and completed.
- Inter-process communication is achieved through global variables
- Suitable for simplest real time systems, but not suitable for complex systems, since there are difficulties in uniformly dividing the tasks and the response times are lengthy.

Pseudokernels

Code driven by FSMs (Finite State Machines)

```
void process_a(void)
{
    for(;;)
    {
        switch(state_a)
        {
            case 1: phase_a1(); break;
            case 2: phase_a2(); break;
            case 3: phase_a3(); break;
            case 4: phase_a4(); break;
            case 5: phase_a5(); break;
        }
    }
}
```

Pseudokernels - Coroutines

- Tasks are separated into multiple phases based on their states in the FSM.
- After executing a phase, the task is suspended and dispatcher is invoked.
- Dispatcher holds the list of the tasks to be executed in a round-robin fashion. It picks up the next task from the list and executes it.
- The states are stored as global variables.
- Inter-task communication and synchronization are managed by global variables.
- Programming languages like Ada have efficient built-in constructs for implementing coroutines.

Pros

- Are easiest type for fairness scheduling.
- Easy to determine the response time, since calls to the dispatcher is given at known intervals.

Cons

- Fairness scheduling is not suitable for real-time system, since the tasks scheduling should be made depending upon importance and urgency.
- Final number of tasks need to be known beforehand.
- Tasks can't be always decomposed uniformly.

Pseudokernels - Coroutines

```
void process_a(void)
{
    for(;;)
    {
        switch(state_a)
        {
            case 1: phase_a1(); break;
            case 2: phase_a2(); break;
            case 3: phase_a3(); break;
            case 4: phase_a4(); break;
            case 5: phase_a5(); break;
        }
    }
}

void process_b(void)
{
    for(;;)
    {
        switch(state_b)
        {
            case 1: phase_b1(); break;
            case 2: phase_b2(); break;
            case 3: phase_b3(); break;
            case 4: phase_b4(); break;
            case 5: phase_b5(); break;
        }
    }
}
```

- Dispatcher is running in the background, which is scheduling `process_a` and `process_b` alternatively
- Dispatcher maintains ‘`state_a`’ and ‘`state_b`’, which are global variables
- Inter-task communication and synchronization are managed by global variables.

Pseudokernels – ‘Interrupt-only Systems



```
void main(void)
{
    init();           /* initialize system, load interrupt handlers */
    while(TRUE);     /* infinite wait loop */
}
void int1 (void)   /* interrupt handler 1 */
{
    save(context); /* save context on stack */
    task1();        /* execute task 1 */
    restore(context);/* restore context from stack */
}
void int2(void)    /* interrupt handler 2 */
{
    save(context); /* save context on stack */
    task2();        /* execute task 2 */
    restore(context);/* restore context from stack */
}
```

- Main program is a jump-to-self instruction (infinite loop)
- Tasks in the systems are scheduled via either H/W or S/W interrupts
- Access to resources is usually controlled by disabling interrupts around any code that reads or writes to the shared resource.
- Snapshot of the current system state (context) for the task must be preserved while switching to other task. It is called **context switching**.

16 BITS Pilani, Pilani Campus

Context Switching



- **Context** switching is the process of saving and restoring sufficient information for a real-time task so that it can be resumed after being interrupted.
- Context is ordinarily saved to a stack data structure.
- Context-switching time is a major contributor to response time.
- The rule for saving context is simple: **save the minimum amount of information necessary to safely restore any process after it has been interrupted**.
- The context data typically includes
 - Contents of general registers
 - Contents of the program counter
 - Contents of coprocessor registers (if present)
 - Contents of memory page register
 - Images of memory-mapped I/O locations (mirror images)

18 BITS Pilani, Pilani Campus

Interrupt Service Routine



- Steps inside an ISR (Interrupt Service Routine)
 - Disables all interrupts
 - Runs code to service the event
 - Clears the interrupt flag that got it called
 - Re-enables interrupts
 - Exits so the processor can go back to its running task
- Interrupt handlers should not take a lot of time to execute, otherwise it will be keeping the external device waiting for longer period of time.
- An ISR (Interrupt Service Routine) is said to be **reentrant** if, while the ISR is handling an interrupt, the same interrupt can occur again and the ISR can process the second occurrence of the interrupt before it has finished processing the first.
- **Reentrant** code is not allowed to
 - Use any data in non-atomic way except when they are saved on the stack
 - Call any other code that is not re-entrant.
 - Use any h/w resources in a non-atomic way.
 - Change its own code.
- An **atomic operation** refers to a group of sub-operations that can be combined to appear as a **single non-interruptable operation**.
- Normally inside the **interrupt handlers**, interrupts are disabled during context switching period.

17 BITS Pilani, Pilani Campus

Preemptive Priority Systems



- A higher-priority task is said to **preempt** a lower-priority task if it interrupts the lower-priority task.
 - The priorities assigned to each interrupt are based on the urgency of the task associated with the interrupt.
 - Prioritized interrupts can be either fixed priority or dynamic priority.
 - **Fixed priority systems** are less flexible, since the task priorities cannot be changed.
 - **Dynamic-priority systems** can allow the priority of tasks to be adjusted at runtime to meet changing process demands.
 - Preemptive-priority schemes can suffer from resource hogging by higher priority tasks. The lower-priority tasks are said to be facing a problem called **starvation**.
 - A special class of fixed-rate preemptive-priority interrupt-driven systems, called **rate-monotonic systems**, comprises those real-time systems where the priorities are assigned so that the higher the execution frequency, the higher the priority.
- Example: In the aircraft navigation system, the task that gathers accelerometer data every 10 milliseconds has the highest priority. The task that collects gyro data, and compensates these data and the accelerometer data every 40 milliseconds, has the second highest priority. Finally, the task that updates the pilot's display every second has lowest priority.

19 BITS Pilani, Pilani Campus

Hybrid Systems

- Hybrid systems include interrupts that occur at both fixed rates and sporadically.
- The sporadic interrupts can be used to handle a critical error that requires immediate attention, and thus have highest priority.
- Another type of hybrid system consists of a combination of round-robin and preemptive priority systems. – generally found in commercial operating systems.

Task Control Block (TCB) Model

- Each task is associated with a data structure – a [task control block \(TCB\)](#).
- TCB contains context. The system stores TCBs in one or more data structures, such as a [linked list](#).
- TCB model can be used in round-robin, preemptive priority or combination systems, although it is generally associated with round-robin systems with a single clock.
- Is the most popular method for implementing commercial, full-featured, real-time operating systems.
- The [operating system manages TCBs](#) by keeping track of the status of each task.
- A task can typically be in any one of the following states:
 - executing
 - ready
 - suspended (or blocked)
 - dormant (or sleeping)
- Every hardware interrupt and every system level call (such as a request on a resource) invokes the real-time operating system.
- The operating system is responsible for maintaining a linked list containing the TCBs of all the ready tasks, and a second linked list of those in the suspended state.
- It also keeps a table of resources and a table of resource requests.

The difference between TCB model and the Interrupt-handler model is that the resources are managed by the operating system in the TCB model, whereas in interrupt-handler model, tasks track their own resources.

Foreground / Background Systems

- Disadvantages of Interrupt driven systems
 - Time is wasted in jump to self loop
 - Difficulty in providing advance services
- Foreground/background systems are '[Interrupt Driven Systems with small variation](#)':
 - Jump-to self routine is replaced by a low priority code that does useful processing.
- It Constitutes
 - a set of interrupt-driven or real-time processes called the foreground
 - a collection of non-interrupt-driven processes called the background
- Background task is fully pre-emptable by any foreground task (hence is the lowest priority task)
- Foreground tasks run in round-robin, pre-emptive priority, or hybrid fashion.
- Background task generally does
 - Initialization
 - House keeping,
 - Low priority tasks such as parameter entry through keypad, low priority display updates etc
- Initialization includes
 - Disable interrupts
 - Setup interrupt vectors and tables
 - Initialize peripherals and other hardware
 - Perform self-tests
 - Perform necessary software initializations
 - Enable interrupts

Task Control Block (TCB) Model

TCB

Task Identifier
Priority
Status
Work Registers
Program Counter
Status Register(s)
Stack Pointer
Pointer to Next TCB

Process and Thread

Process

- A process is an abstraction of a running program and is the logical unit of work scheduled by the operating system.
- Typically represented by a data structure that contains at least a state of execution, an identity (real-time), attributes (e.g., execution time), and the resources associated with it.

Thread

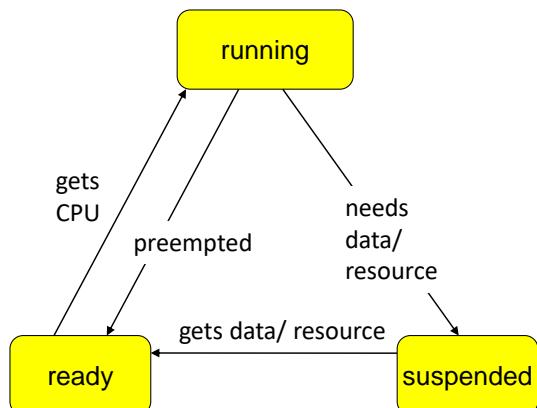
- A thread is a **lightweight process** that shares resources with other processes or threads.
- Each thread must “reside” within some process and make use of the resources of that process.

Timer and Clock Services

- Timing services are essential for Real Time Software
- Example: a diagnostic task checks the health of the system periodically.
- A system call '`delay`' is commonly used to suspend a task.
- The parameter passed to this function is an integer specifying the number of '`ticks`'.
- **A timer circuit is configured to interrupt the CPU at a fixed rate** and the internal system time is incremented at each timer interrupt. The interval of time for which the timer interrupt is programmed is called '`tick`' or time resolution.

Process States

Scheduling Policy defines how processes are selected for **promotion from the ready state to the running state**



Linux Kernel Scheduler

- Linux kernel scheduler supports
 - Nonpreemptible scheduler (initial version of Linux kernel supported this policy only)
 - Preemptible scheduler (introduced from version 2.6) – required for real-time tasks/threads
 - Completely fair scheduler (CFS) (introduced in the version 2.6.23) – Primarily for desktops
- Following function controls the behaviour of process scheduling discipline:

```
int sched_setscheduler (pid_t pid, int policy,
                      const struct sched_param *param);
```

Non-real time policies supported:

SCHED_OTHER: Standard round-robin policy

SCHED_BATCH: Batch style execution – it doesn't preempt processes often

SCHED_IDLE: Runs the task at very low priority background task

Real time policies supported:

SCHED_FIFO: First in first out

(when the process becomes runnable, it will be inserted at the end of the list for its priority)

SCHED_RR: Round robin

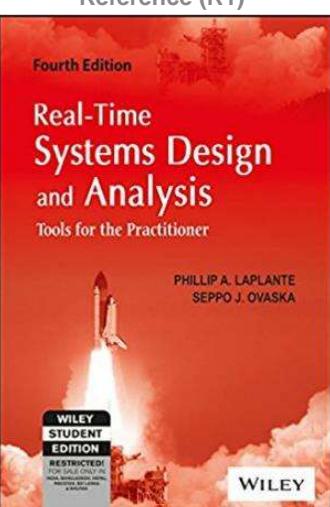
Please note that this function is called per process basis. That means different processes can be scheduled based on different scheduling policies

Thank You.

Any Questions?

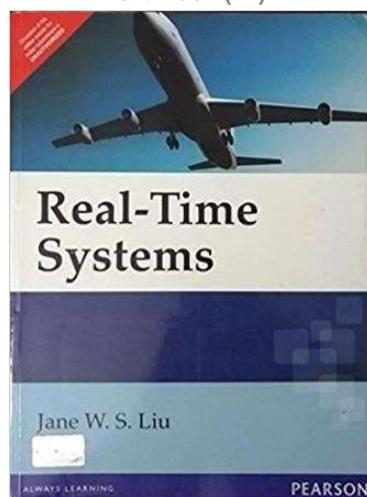
28
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Text Book / References



Reference (R1)

Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

BITS ZG553: Real Time Systems
L11 – Programming Languages for Real Time Systems
K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Excellent MOOCs Videos
(Coursera, edX,...)

HONOR CODE
CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIx: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.



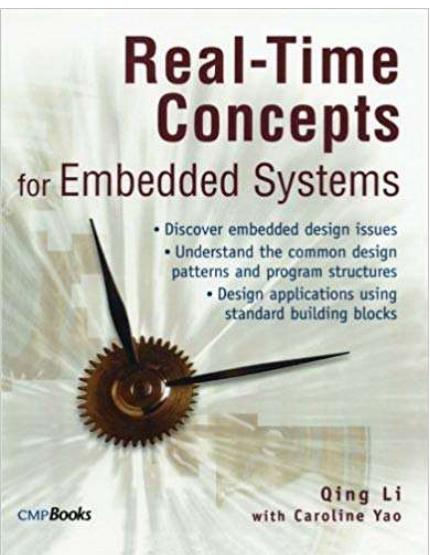
Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54

RTS Primer – For Light Reading



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

4

Language Taxonomy

- Two different ways to partition the world of programming languages:
 - Imperative, functional, logic, other
 - Procedural or declarative
- **Imperative** languages:
 - Involve assignment of variables
 - Most widely used language are imperative, e.g. C, Java, Ada 95, Fortran, BASIC
- **Functional** (applicative) languages employ repeated application of some function (e.g. LISP).
- **Logic programming** involves a formal statement of the problem without any specification of how it is solved (e.g. PROLOG).
- **Procedural** languages are defined by a series of operations (most languages you can think of).
 - “*Procedural*” also describes a style of programming that is not object-oriented.
- **Declarative** (non-procedural) languages involve specification of rules defining the solution to a problem (PROLOG, Spreadsheets).
- **Object-oriented** describes a style of programming. Languages written to support this are called object-oriented.



L-11: Programming Languages for Real Time Systems

[Ref: Notes/PPT]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

Which language to choose ?

- At present do we need pay too much attention on how to select a programming language ?
- For present day real-time systems, **it is almost like to choose between the two: C & C++**.
- If you have large S/W projects where productivity of the programmers and maintainability is an issue, go for C++.
- If response-time is the primary driving factor, go for C.
- Now a days, majority of the applications of real-time systems are written in Java.
- Web-based applications are in HTML5, Javascript, Jquery, Angular JS, Phonegap, PHP, ASP, ...
- So, followings can be generally suggested:
 - System Software: C
 - Middleware and Applications: C++
 - Applications: Java, HTML5 and other web development languages.

Fitness of a Programming Language for Real-Time Applications



Five criteria of Cardelli in selecting a programming language:

- **Economy of execution:** How fast does a program run?
- **Economy of compilation:** How long does it take to go from sources to executables?
- **Economy of small-scale development:** How hard must an individual programmer work?
- **Economy of large-scale development:** How hard must a team of programmers work?
- **Economy of language features:** How hard is it to learn or use a programming language?

Procedural Languages



- Example: Ada, C, Fortran, Visual Basic etc.
- Features advantageous for Real Time Systems:
 - Modularity (design become simpler)
 - Strong typing (prevents data corruption through unwanted type conversion)
 - Abstract data typing (allow user to define custom types)
 - Versatile parameter passing mechanism
 - Dynamic memory allocation
 - Exception handling

Assembly Languages



- In terms of Cardelli's criteria, assembly languages have
 - Excellent economy of execution
 - Excellent economy of compilation
 - Poor economies of small- and large-scale development of languages features
- So Assembly languages are used
 - Where there is **tight timing situations** e.g. interrupt service routines, part of device drivers , sophisticated signal-processing algorithms etc.
 - In **controlling hardware features** that are not supported by compilers
 - Where **predictable performance for the code is extremely difficult or impossible using high level language** because of undesirable programming language-compiler interactions.
 - For writing entire software **for custom-designed CPU** with a small instruction set and without high-level language support.
 - For **debugging hard problems** below the high-level language code and tracing the stream of fetched instructions by logic analyzer.
 - For teaching and learning computer architectures

BITS Pilani, Pilani Campus

Procedural Languages



- Parameter passing mechanisms
- **Call by value:** The value is passed to the function and it manipulates the value passed, not the source
 - **Call by reference:** Address of the parameter is passed, so any manipulation of the parameter inside the function is reflected in the original parameter.
 - **Global variables:**
 - Global variables are within the scope of all code, **so faster than references to variables passed via parameter list**, which require additional memory accesses.
 - But they are **dangerous** because references to them can be made by unauthorized code, potentially introducing faults that can be hard to isolate.
 - **Recursions**
 - A procedure can call itself.
 - It is elegant and sometimes can't be avoidable.
 - Procedure call requires allocation on stack for passing parameters and for storage of local variables.
 - **So precaution needs to be taken for the recursion to terminate early, otherwise it may lead to stack overflows.**
 - Also allocation and deallocation of parameters may be a costly operation with respect to time.

Cardelli's Matrix for Procedural Languages



- Economy of execution
 - High, because strong type-checking improve code generation
- Economy of compilation
 - High, because modules can be compiled independently
- Economy of small-scale development
 - High, since type-checking can catch many errors, reducing testing and debugging efforts
- Economy of large-scale development
 - High, because data abstraction and modularization allow independent team of programmers to develop modules in parallel (provided interfaces between modules are well defined)
- Economy of language features
 - High, because learning is easier due to orthogonality of the language features

Object-Oriented Languages



- Example: C++, C#, Java etc.
- Features:
 - Data Abstraction (e.g. Class)
 - Inheritance (e.g. Derived Class)
 - Polymorphism (e.g. Operator Overloading)
 - Messaging

Cardelli's Matrix for Object-Oriented Languages



- Economy of execution
 - Low, because extra encapsulation (like class and method) causes the execution slower
- Economy of compilation
 - Low, because it may require recompilation of base classes while compiling the inherited classes
- Economy of small-scale development
 - High, since individual programmers can take help of class libraries and frameworks
- Economy of large-scale development
 - Low, sometimes these languages have poor modularity properties. E.g. it is easy to override a method that should not be overridden, or to re-implement a class in a way that causes problems in subclasses.
 - High, since separate teams can develop separate modules, making the development highly parallel
- Economy of language features
 - Low, because learning takes time due to many features

Object-Oriented vs Procedural Languages



- Object oriented languages provide better design.
- But for real time systems, the main disadvantages is that they lead to unpredictable and inefficient system.
- So the decision should be made from case to case basis.
- Usual norms followed is:
 - Procedural languages: For System Software (e.g. C)
 - Object-oriented languages: For Application and Middleware (e.g. C++, Java)

Code Optimization Techniques



- Use of arithmetic identities
 - Addition by 0 or multiplication by 1 can be avoided
- Reduction in Strength
 - Use of fastest machine-language instructions for a given expression.
 - For example
 - $x * 2$ can be replaced by $x << 1$
 - $x / 4$ can be replaced by $x >> 2$
- Repeated calculations of the same sub-expression

Example:

```
x = 6 + a * b;  
y = a * b + z;
```

Can be replaced by

```
t = a * b  
x = 6 + t  
y = t + z
```

Code Optimization Techniques



- Loop induction elimination

```
for(i = 1; i <= 10; i++)  
    a[i+1] = i+1;
```

Can be replaced with

```
for(j = 2; j <= 11; j++)  
    a[j] = j;
```

Here i and j are called loop induction.

- Use of registers and caches
- Dead-code removal

Code Optimization Techniques



- Use of intrinsic functions
 - Intrinsic functions are macros with inline code
 - Use of intrinsic function improves performance, because parameter passing to the functions is eliminated
- Constant folding
 - If a program uses $\pi / 2$, it can be pre-computed and stored as `pi_div_2`.
- Loop invariant removal

Move the computations which can be moved out of the loop.

```
x = 100;  
while (x > 0)  
    x = x -(y + z);
```

Can be replaced with

```
x = 100;  
t = y + z;  
while (x > 0)  
    x = x - t;
```

Code Optimization Techniques



- Flow-of-control optimization

Unnecessary branch to branch instructions are replaced by single branch instructions.

```
goto label_1;  
label_0:    y = 1;  
label_1:    goto label_2;
```

Can be replaced by

```
goto label_2;  
label_0:    y = 1;  
label_1:    goto label_2;
```

Code Optimization Techniques



➤ Constant propagation

```
x = 100;  
y = x;
```

Can be replaced by

```
x = 100;  
y = 100;
```

➤ Dead store elimination

```
t = y + z;  
x = func(t);
```

Can be replaced by

```
x = func(y + z);
```

➤ Dead variable elimination

```
x = y + z;  
x = z;
```

Can be replaced by

```
x = z;
```

Code Optimization Techniques



➤ Loop unrolling

```
for(i = 0; i < 4; i++)  
    a[i] = x[i] * y[i];
```

Can be replaced by

```
a[0] = x[0] * y[0];  
a[1] = x[1] * y[1];  
a[2] = x[2] * y[2];  
a[3] = x[3] * y[3];
```

Such code is very much useful where vector operations are involved.

Code Optimization Techniques



➤ Short circuit Boolean code

```
if(a > 100){  
    if(b < 20) {  
        ...  
    }  
}
```

Is equivalent to

```
if((a > 100) &&(b < 20)) {  
    ...  
}
```

ANSI-C compiler will generate a better code for the second case because it evaluates expressions from left to right and drops out at first FALSE condition, whereas other compilers may generate a better code for the first case.

Code Optimization Techniques



Code Optimization Techniques



➤ Loop jamming (or Loop fusion)

```
for(i = 0; i < 100; i++)  
    x[i] = y[i] * 8;  
for(i = 0; i < 100; i++)  
    a[i] = x[i] * y[i];
```

Can be replaced by

```
for(i = 0; i < 100; i++)  
    a[i] = x[i] * y[i] * 8;
```

Code Optimization Techniques

➤ Cross-branch elimination

If same code appears in multiple cases in a switch-case statement, combine the cases.

Can be replaced by

```
switch(x)
{
    case 0:
        x++;
        break;
    case 1:
        x = x + 2;
        break;
    case 2:
        x++;
        break;
    default:
        break;
}
```

24

Example

Can you optimize the following C function ?

```
int function(int *a, int *b, int count) {
    int c = 0;
    int i = 0;
    int temp1 = 0;
    int temp2 = 0;
    if(a == null)
        return -1;
    if(b == null)
        return -1;
    for(i = 0; i < count; i++) {
        temp1 = a[i] / 2;
        temp2 = b[i] / 4;
        c += temp1 * temp2;
    }
    return c;
}
```

Additional Optimization Considerations

- Arrange entries in the table so that most frequently sought values are the first to be compared
- Replace threshold tests on monotone functions (continuously increasing or continuously decreasing) by tests of their parameters, thereby avoiding evaluation of the function itself

```
if(exp(x) < exp(y)) {
    ...
}
```

can be replaced by

```
if(x < y) {
    ...
}
```

for monotone functions

- Keep the most frequently used procedures nearby in order to take the advantages of locality of reference in a cached or paged system.
- Keep the frequently used data nearby in order to take the advantages of locality of reference in a cached or paged system.
- Aligning memory access consumes lesser time in memory access operation.
- Match data types to the architecture

Example (contd.)

One way could be the following.

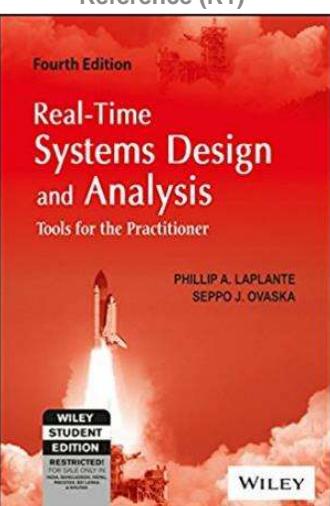
```
int function(int *a, int *b, int count) {
    int c = 0;
    int i = 0;
    if((a == null) || (b == null))
        return -1;
    for(i = 0; i < count; i++) {
        c += (a[i] * b[i]) >> 3; /* Divide by 8 */
    }
    return c;
}
```

Thank You.

Any Questions?

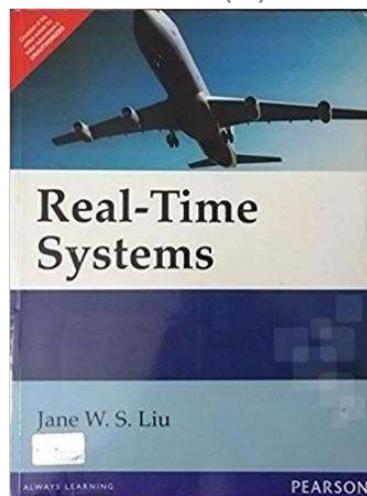
28
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Text Book / References



Reference (R1)

Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

BITS ZG553: Real Time Systems

L12 – Software Design & Development for Real Time Systems

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

BITS Pilani
Pilani Campus



BITS Pilani
Pilani Campus

Excellent MOOCs Videos

(Coursera, edX,...)

HONOR CODE CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

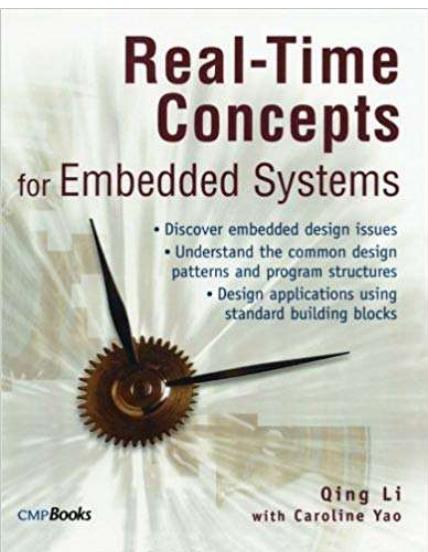


Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

4

Qualities of Real-Time Software

- Reliability
- Correctness
- Performance
- Usability
- Interoperability
- Maintainability
- Portability
- Verifiability



L-12: Software Design & Development for RTS [Quality, Performance Considerations]

[Ref: Notes/PPT]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

Reliability

- A measure of whether a user can depend on the software.
 - The system “stands the test of time.”
 - There is an absence of known catastrophic errors; that is, errors that render the system useless.
 - The system recovers “gracefully” from errors.
 - The software is robust.
- For real-time systems, other informal characterizations of reliability might include
 - Downtime is below a certain threshold.
 - The accuracy of the system is within a certain tolerance.
 - Real-time performance requirements are met consistently.

Traditionally, MTTF (Mean Time To First Failure) and MTBF (Mean Time Between Failures) are used as a measure of Software Reliability.

Software Reliability – Statistical Perspective



Let S = The software system

T = Time instance of failure

Then Reliability of S at any time instance t, is defined by

$$r(t) = P(T > t)$$

$r(t) = 1$ is the ideal case. Practically, $r(t)$ is less than 1.

Example:

In the monitoring system of a nuclear power plant, the failure probability is no more than 10^{-9} per hour.

So, $r(t) = (1 - 10^{-9})^t$, t is in hours.

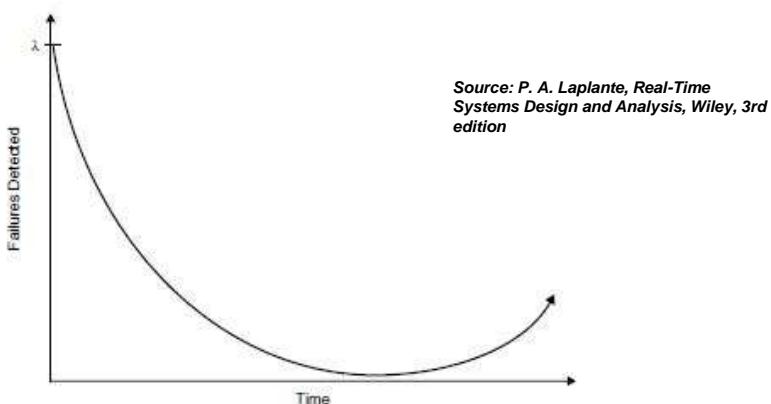
So, as $t \rightarrow \infty$, $r(t) \rightarrow 0$.

BITS Pilani, Pilani Campus

Software Reliability – Statistical Perspective



But the failure function tends to become a *bathtub curve* i.e. failures reduces initially at an exponential rate, but increases towards the end.



10
BITS Pilani, Pilani Campus

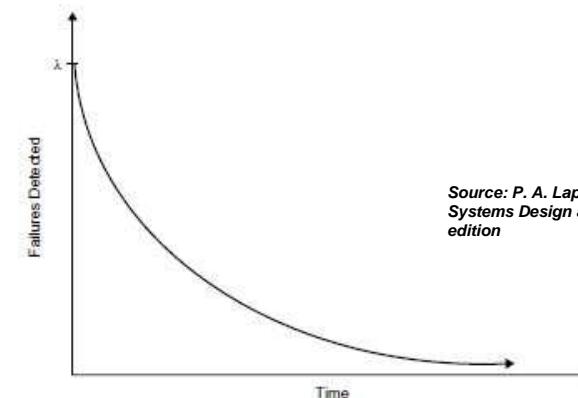
Software Reliability – Statistical Perspective



Another way is to characterize software reliability is in terms of a failure function, which is usually an exponential function,

$$f(t) = \lambda e^{\lambda t}, \quad t \geq 0$$

' λ ' is a system dependent parameter, determined empirically.



Source: P. A. Laplante, Real-Time Systems Design and Analysis, Wiley, 3rd edition

BITS Pilani, Pilani Campus

Software Reliability – Statistical Perspective



Reasons for Bathtub curve for failures:

- Hardware can wear and tear. So after some point of time, the number of failures increases.
- But software can't wear and tear. Then why bathtub curve for the software ?
 - During patch releases, making quick correction without designing them properly.
 - During patch releases, ignoring the effect of the new bug fix on other modules (thereby fixing one issue and introducing more new issues)
 - Inadequate testing of the software during patch releases.
 - Wear and tear of underlying hardware.

11
BITS Pilani, Pilani Campus

Correctness

- Closely related to reliability and the terms are often used interchangeably.
- The main difference is that *minor deviation from the requirements is strictly considered a failure and hence means the software is incorrect*. However, a system may still be deemed reliable if only minor deviations from the requirements are experienced.
- *Correctness is measured in terms of number of failures detected over time.*

Usability

- Often referred to as ease of use, or *user friendliness*.
- Usability is difficult to quantify (but we can easily determine its absence).
- However, informal feedback can be used, as well as user feedback from surveys and problem reports can be used in most cases.

Performance

- A measure of some required behavior.
- For example, an imaging system might be required to display a filtered image at a rate of 30 frames per second.
- A photo reproduction system might be required to digitize, clean and output color copies at a rate of 1 every two seconds.
- *Can be measured via mathematical or algorithmic complexity, direct timing, or simulation.*

Interoperability

- Refers to the ability of the software system to coexist and cooperate with other systems.
 - In imaging systems the software must be able to communicate with various devices using standard bus structures and protocols.
- *Open systems and standards foster interoperability.*
- *Interoperability can be measured in terms of compliance with open system standards.*

Maintainability

- A software system in which *changes are relatively easy to make*, has a high level of maintainability.
- Two contributing properties – evolvability and reparability.
 - Evolvability is a measure of how easily the system can be changed to accommodate new features or modification of existing features.
 - Software is repairable if it allows for the fixing of defects.

Verifiability

- A software system is verifiable if its properties can be verified easily.
- Can be achieved through the insertion of code that is intended to *monitor various qualities* such as performance or correctness.
- **Modular design, rigorous software engineering practices** and the effective use of an appropriate programming language can also contribute to verifiability.

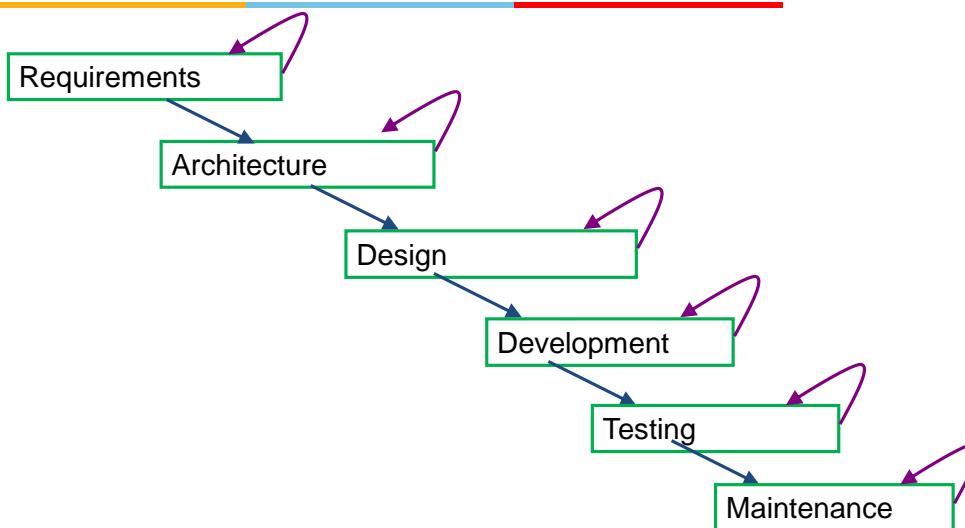
Portability

- Software is portable if it can easily run in different environments.
- Portability is achieved through a deliberate design strategy in which hardware dependent code is confined to the fewest code units as possible.
- For making the upper-layer software independent of the underlying hardware architecture, usually a **Hardware Abstraction Layer (HAL)** is implemented.
- Can be achieved using either object-oriented or procedural programming languages and through object-oriented or structured approaches.
- Also using standardized APIs (such as **POSIX** calls), improves the portability
- *Person months required to perform the port are the standard measure of this property*

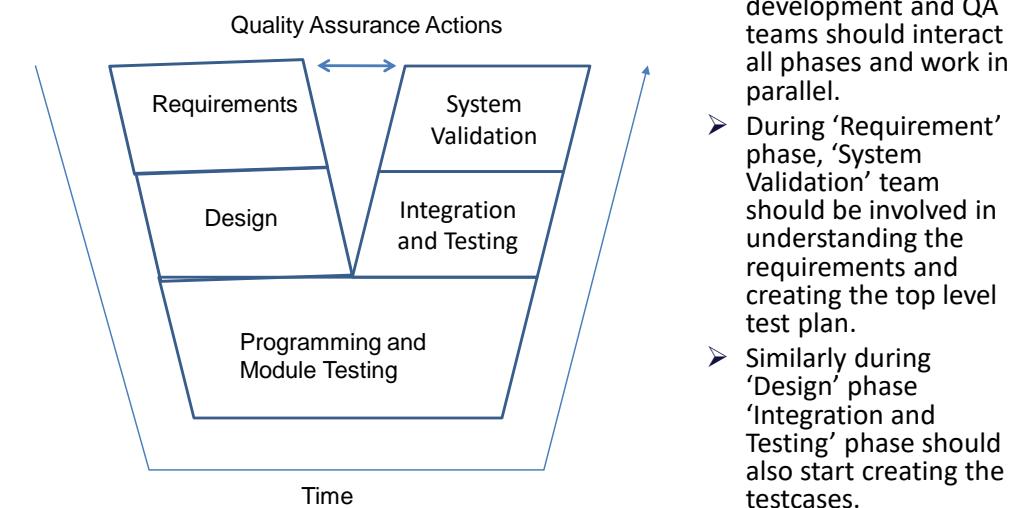
Software Design Life Cycle (SDLC) Models

- Waterfall Model
- V Model
- Spiral Model
- Agile Methodologies

Waterfall Model



V Model



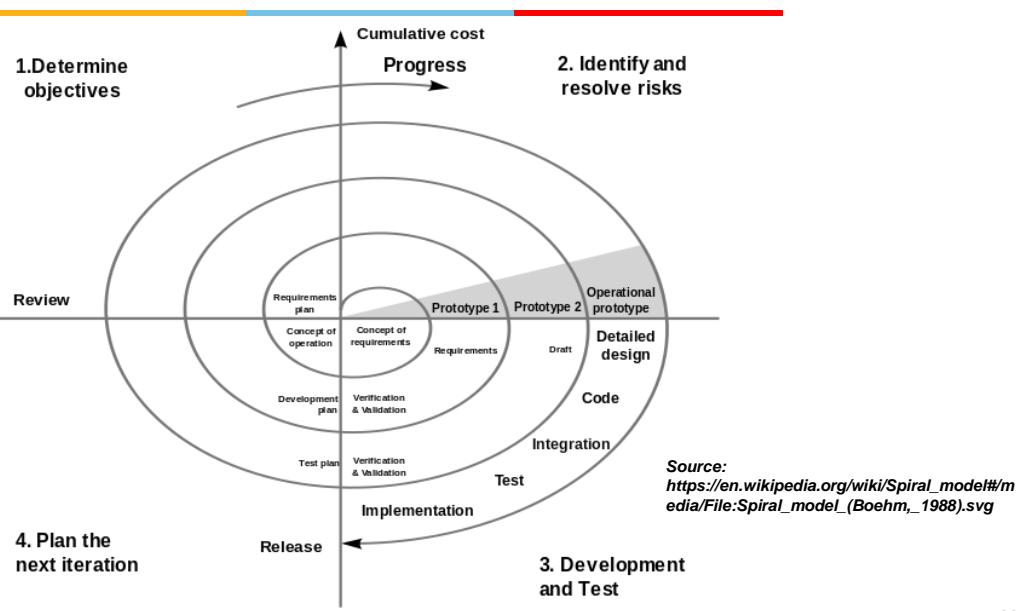
- As per the model, the development and QA teams should interact all phases and work in parallel.
- During 'Requirement' phase, 'System Validation' team should be involved in understanding the requirements and creating the top level test plan.
- Similarly during 'Design' phase 'Integration and Testing' phase should also start creating the testcases.

20 BITS Pilani, Pilani Campus

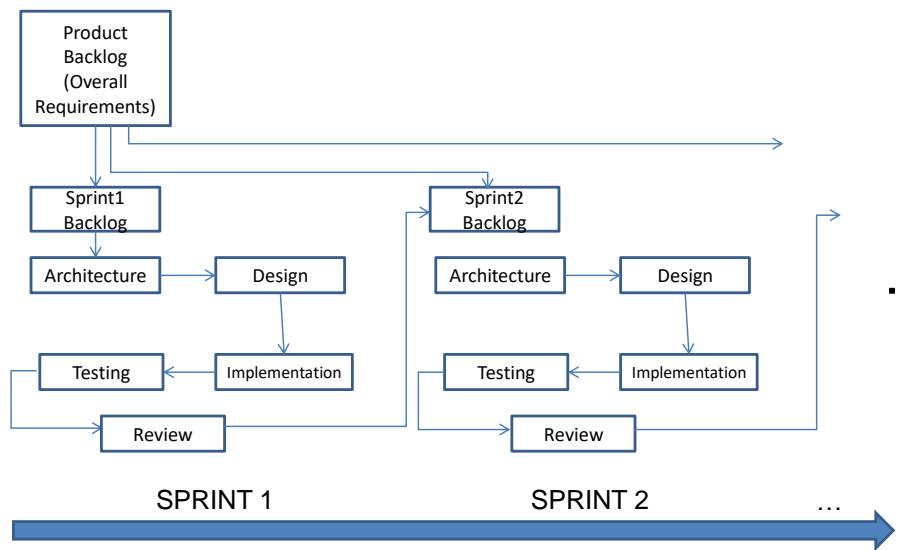
21

BITS Pilani, Pilani Campus

Spiral Model



Agile Scrum



22 BITS Pilani, Pilani Campus

23

BITS Pilani, Pilani Campus

Software Design Approaches

- Procedural Design Approach
 - Parnas Partitioning
 - Structured Design
- Object Oriented Design Approach
 - UML (Unified Modeling Language) based Design

Software Design Approaches

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

25

Parnas Partitioning

- Cohesion indicates the dependencies between modules. i.e.
 - High cohesion between two modules indicates tight dependencies between them.
 - Low cohesion between two modules indicates lesser dependencies between them.
- Parnas Partitioning involve the software into multiple modules, which have **low external coupling and high internal cohesion**.
 - If any change has to be made to a module, then there is high probability that the changes are localized to the specific module.
- Individual modules then hide their internal implementations to other modules and expose only required interfaces meant for communication between the modules. It is the principle of '**information hiding**'.

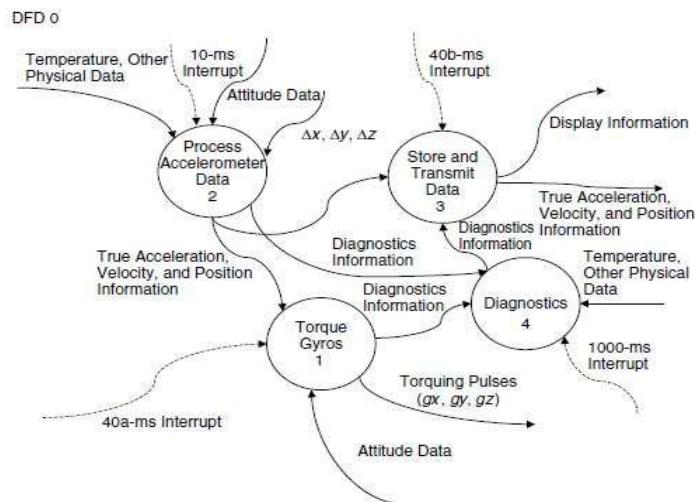
Structured Design (SD)

- Companion methodology to **Structured Analysis (SA)** – Both are combined known as **SA/SD (SD follows SA)**
- Transition from SA to SD is manual.
- One of the very important diagram is **Data Flow Diagram (DFD)**
- DFD evolves from the Context Diagram (CD)
- DFDs are used to partition the system.
 - Level 0 DFD is the first DFD, which illustrates the highest level of system abstraction.
 - Then these DFDs are subdivided into lower and lower levels until they are ready for detailed design.
 - Modules are represented by circles (and labeled by verb phrases)
 - The data flow is indicated by solid lines with arrowhead indicating the direction of the flow (and labeled by noun phrases)
 - Control flow is indicated by dashed lines with arrowhead indicating the direction of the flow

Data Flow Diagrams (DFD)

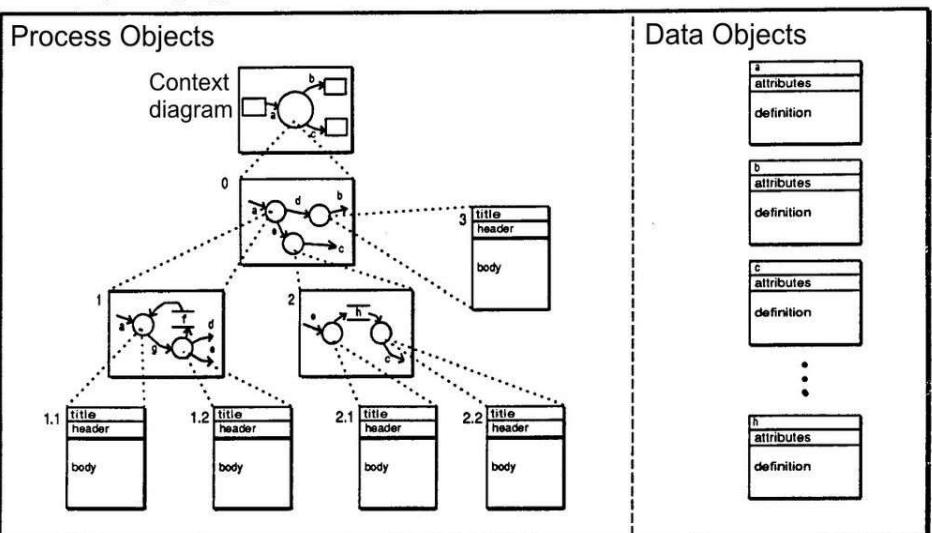
- One of the very important diagram is **Data Flow Diagram (DFD)**
- DFD evolves from the Context Diagram (CD)
- DFDs indicate flow of the data between different modules in the system.
- Level 0 DFD is the first DFD, which illustrates the highest level of system abstraction.
- Then these DFDs are subdivided into lower and lower levels until they are ready for detailed design.
- **Modules** are represented by **circles** (and labeled by verb phrases)
- The data flow is indicated by **solid lines** with arrowhead indicating the direction of the flow (and labeled by noun phrases)
- Control flow is indicated by **dashed lines** with arrowhead indicating the direction of the flow

DFD Example



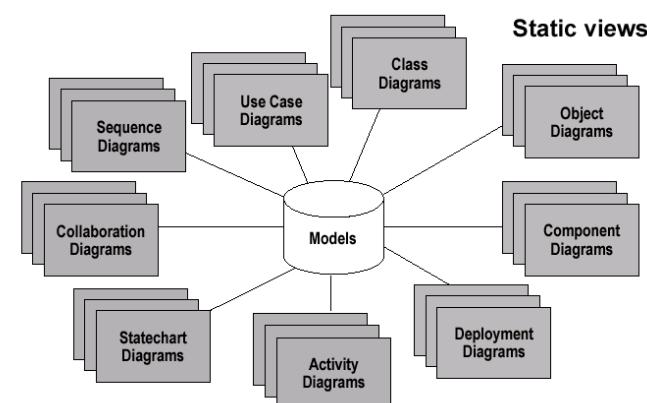
Level 0 DFD for an inertial measurement system. Dashed lines indicate control flow.
Source: P. A. Laplante, *Real-Time Systems Design and Analysis*, Wiley, 3rd edition

Evolution of DFD from CD to Lower Level Design Specification



Unified Modeling Language (UML)

- It is the de facto standard for Object Oriented Design.
- Graphical language based on the concept that any system can be composed of communities of interacting entities.



Dynamic views

Class Diagram

- Gives an overview of a system by showing its classes and the relationships among them.
 - Class diagrams are static
 - they display what interacts but not what happens when they do interact
- Also shows attributes and operations of each class
- Good way to describe the overall architecture of system components

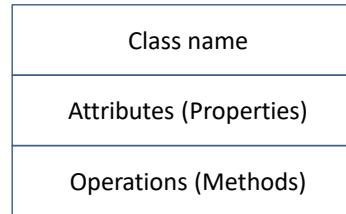
BITS Pilani, Pilani Campus

Relationship Between Classes

- Association**
 - A relationship between instances of two classes, where one class must know about the other to do its work, e.g. client communicates to server
 - Each classes have their own life-cycle (i.e. their creation and destruction are not dependent on each other.
 - indicated by a straight line or arrow
- Aggregation**
 - An association where one class belongs to a collection
 - But the child object doesn't die if parent object dies
 - Indicated by an empty diamond on the side of the collection
- Composition**
 - Strong form of Aggregation
 - Lifetime control; components cannot exist without the aggregate
 - When the parent object dies, the child object also dies
 - Indicated by a solid diamond on the side of the collection
- Inheritance**
 - An inheritance link indicating one class a superclass relationship, e.g. human is a mammal
 - Indicated by triangle pointing to superclass

BITS Pilani, Pilani Campus

Class Representation



Example:

Employee
-Name: string
+ID: int
#Salary: int
+GetName(): string
+SetName(string sName)
-CalcSalary(int ID)

- A class is a rectangle divided into three parts

- Class name
- Class attributes (i.e. data members, variables)
- Class operations (i.e. methods)

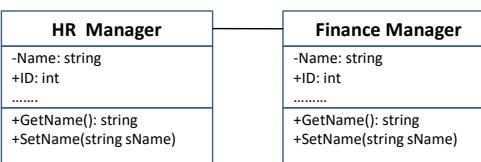
Modifiers

- Private: -
- Public: +
- Protected: #
- Static: Underlined (i.e. shared among all members of the class)

- Abstract class: Name in italics

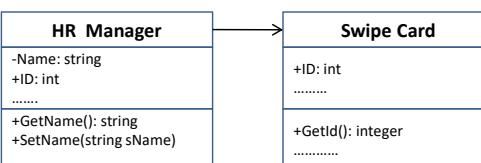
BITS Pilani, Pilani Campus

Association



Binary Association (Both Classes Know About Each Other)

Indicated by a straight line.



Unary Association (Only One Class Know About the Other, Other Doesn't)

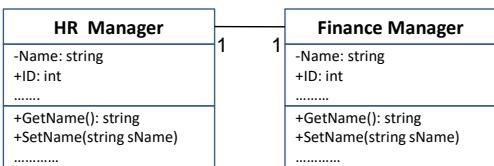
Indicated by an arrow starting from the class that has knowledge about the other class and pointing to the class that doesn't have knowledge about the previous class.

```

public Class HRManager {
    SwipeCard swipeCard;
    .....
}
  
```

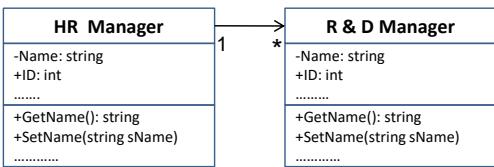
BITS Pilani, Pilani Campus

Association: Multiplicity

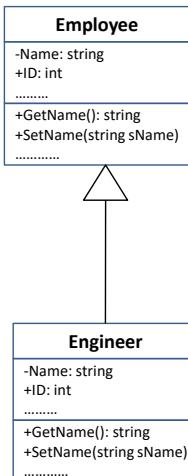


Multiplicity

Symbol	Meaning
1	One and only one
0..1	Zero or one
M..N	From M to N
* Or 0..*	From zero to any positive integer
1..*	From one to any positive integer



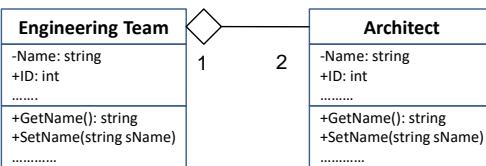
Inheritance



Inheritance

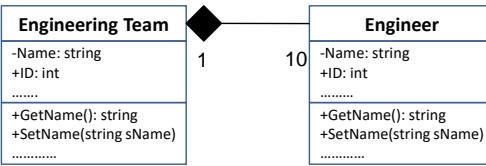
- ❑ The child class inherits the properties and operations of the base class
- ❑ Adds its own specific functionalities as well

Aggregation & Composition



Aggregation

- ❑ A bigger entity containing smaller entities or components
- ❑ A component can be part of an aggregation
- ❑ Example: Engineering Team can have architects, who are shared between multiple team.



Composition

- A special case of aggregation, where the component can't exist without the aggregated entity

```

public Class EngineeringTeam {
    Engineer engineer[10];
    .....
}
  
```



Metrics in Software Engineering

- Lines of Source Code
- Cyclomatic Complexity
- Halstead's Metrics
- Function Points
- Feature Points

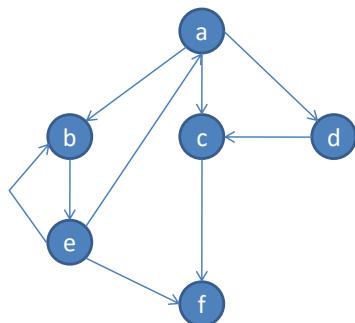
Lines of Source Code

- Measured in thousands of lines of code (**KLOC**).
- It doesn't include comments.
- But it doesn't indicate the complexity of the code (we can't say that a program having 1000 lines of code is more complex and can have more errors than a program having 100 lines of code).
- Another major disadvantage is that KLOC can be calculated after the code is written.
- Another related metric is '**Delta KLOC**', indicating the amount of code added incrementally.
- As the project matures, '**Delta KLOC**' should reduce.

Cyclomatic Complexity - Example

For the graph shown, $n = 6$, $e = 9$.

So, cyclomatic complexity $C = e - n + 2 = 9 - 6 + 2 = 5$



Cyclomatic Complexity

- Introduced by McCabe
- The concept fits well with Procedural Programming, but not quite well with Object-Oriented Programming
- It is based on determining the number of linearly independent paths in a program.
- Consider a graph, where the nodes represent the program segment and the edges represent independent paths.
- Let the number of nodes = n and number of edges = e .
- Then the cyclomatic complexity,

$$C = e - n + 2$$

- Commercial tools are available to calculate cyclomatic complexity.

Halstead's Metrics

- One drawback of Cyclomatic complexity is that it measures complexity as a function of control flow, but complexity can exist the way programming language is used.
- Halstead's metrics measures how intensively the programming language is used.

Let

n_1 = The number of distinct syntactic begin-end pairs called 'operators' - **operator carries out an action**

n_2 = The number of distinct statements called 'operands' - **operand participates in such an action**

N_1 = Total number of occurrences of n_1 in the program

N_2 = Total number of occurrences of n_2 in the program

Halstead's Metrics

- Program *Vocabulary*, $n = n_1 + n_2$
- Program *Length*, $N = N_1 + N_2$
- Program *Volume*, $V = N \log_2 n$
- Calculated program length, $N^* = n_1 \log_2 n_1 + n_2 \log_2 n_2$
- *Difficulty* Level, $D = (n_1 / 2)X(N_2 / n_2)$
- Programming *Effort*, $E = D \times V$
- The difficulty measure D is related to the *difficulty of the program to write or understand*
- The effort measure translates into actual coding time using the following relation.
- Time required to program, $T = E / 18 \text{ seconds}$
- Number of delivered bugs, $B = V / 3000$ (it was $E^{2/3}/3000$ earlier)

Halstead's Metrics - Example

(Source: https://en.wikipedia.org/wiki/Halstead_complexity_measures)

Program Vocabulary, $n = n_1 + n_2 = 10 + 7 = 17$

Program Length, $N = N_1 + N_2 = 16 + 15 = 31$

Program Volume, $V = N \log_2 n = 31 \log_2 17 = 31 \times 4.087 = 126.697$

Calculated program length,

$N^* = n_1 \log_2 n_1 + n_2 \log_2 n_2 = 10 \log_2 10 + 7 \log_2 7 = 10 \times 3.32 + 7 \times 2.81 = 33.2 + 19.67 = 52.87$

Difficulty Level,

$D = (n_1 / 2)X(N_2 / n_2) = (10 / 2) \times (15 / 7) = 5 \times 2.142 = 10.714$

Programming effort,

$E = D \times V = 10.714 \times 126.697 = 1357.468$

Time required to program,

$T = E / 18 \text{ seconds} = 1357.468 / 18 \text{ seconds} = 75.41 \text{ seconds}$

Number of delivered bugs,

$B = V / 3000 = 126.697 / 3000 = 0.04223$

(With the earlier formula, $B = E^{2/3}/3000 = (1357.468)^{2/3} / 3000 = 122.5991 / 3000 = 0.04$)

Halstead's Metrics - Example

(Source: https://en.wikipedia.org/wiki/Halstead_complexity_measures)

Let us consider the following C program.

```
void main()
{
    int a, b, c, avg;
    scanf("%d %d %d", &a, &b, &c);
    avg = (a + b + c) / 3;
    printf("avg = %d", avg);
}
```

Here,

The unique operators are: main, (), {}, int, scanf, &, =, +, /, printf

The unique operands are: a, b, c, avg, "%d %d %d", 3, "avg = %d"

So,

$$n_1 = 10, n_2 = 7,$$

$$N_1 = 16, N_2 = 15$$

Function Points

- The philosophy here is to measure the functionality of the software via the number of interfaces between modules and subsystems in programs or entire systems.
- Generally used for business information systems, but is also used for large scale real-time systems like multimedia and communication systems.
- 5 software characteristics for each module or subsystem:
 - A1 = Number of inputs to the application (I)
 - A2 = Number of outputs (O)
 - A3 = Number of user inquiries (Q)
 - A4 = Number of files used (F)
 - A5 = Number of external interfaces (X)

Function Points

Then the function point,

$$FP = \left[\sum_i A_i W_i \right] [0.65 + 0.01 \sum_j F_j]$$

Where, W_i = Weighting factors

A_i = Item counts

F_j = Complexity Adjustment Factors

The weighting factors and complexity adjustment factors are adjusted experimentally.

Following is the way to calculate the complexity adjustment factors.

Function Points – Complexity Adjustment Factors

The 14 Questions to be asked for a real-time system:

- A. Does the system require reliable backup and recovery?
- B. Are data communications required?
- C. Are there distributed processing functions?
- D. Is performance critical?
- E. Will the system run in an existing, heavily utilized operational environment?
- F. Does the system require on-line data entry?
- G. Does the on-line data entry require the input transactions to be built over multiple screens or operations?
- H. Are the master files updated on-line?
- I. Are the inputs, outputs, files, or inquiries complex?
- J. Is the internal processing complex?
- K. Is the code designed to be reusable?
- L. Are the conversion and installation included in the design?
- M. Is the system designed for multiple installations in different organizations?
- N. Is the application designed to facilitate change and ease of use by the user?

Function Points – Complexity Adjustment Factors

A set of 14 questions that are answered on a scale from 0 to 5 where:

- | | |
|---|--------------|
| 0 | no influence |
| 1 | incidental |
| 2 | moderate |
| 3 | average |
| 4 | significant |
| 5 | essential |

Function Points – Example

For an embedded system, the answers to the 14 questions came like the following.

A : 4	B : 5	C : 5	D : 5	E : 5	F : 4	G : 4
H : 5	I : 4	J : 4	K : 4	L : 5	M : 5	N : 5

$$\text{Then } \sum F_j = 4 \times 6 + 5 \times 8 = 64$$

Let the item numbers and corresponding weights are as follows

A1 = I = 5	W1 = 4
A2 = U = 7	W2 = 4
A3 = Q = 8	W3 = 5
A4 = F = 5	W4 = 10
A5 = X = 5	W5 = 7

Then the function point,

$$\begin{aligned} FP &= \left[\sum_i A_i W_i \right] [0.65 + 0.01 \sum_j F_j] \\ &= (5 \times 4 + 7 \times 4 + 8 \times 5 + 5 \times 10 + 5 \times 7) \times (0.65 + 0.01 \times 64) \\ &= (20 + 28 + 40 + 50 + 35) \times 1.29 \approx 223 \end{aligned}$$

Lines of code per function points



Programming Language	Lines of code / Function point
Assembly	320
C	128
C++	64

Object oriented languages take lesser code to implement a feature.

Source: P. A. Laplante & S. J. Ovaska, *Real-Time Systems Design and Analysis: Tools for the Practitioner*, Wiley, 4th edition

52
BITS Pilani, Pilani Campus

Feature Points- Example



In the previous example, let us have

$$A6 = 10, W6 = 7.$$

Then feature point,

$$\begin{aligned} FP &= \left[\sum_i A_i W_i \right] \left[0.65 + 0.01 \sum_j F_j \right] \\ &= (5 \times 4 + 7 \times 4 + 8 \times 5 + 5 \times 10 + 5 \times 7 + 10 \times 7) \times (0.65 + 0.01 \times 64) \\ &= (20 + 28 + 40 + 50 + 35 + 70) \times 1.29 \approx 294 \end{aligned}$$

So, if the language used is C,
then approx lines of code for this module

$$= 128 \times 294 = 37.632 \text{ KLOC}$$

Feature Points

- It is an extension of Function Points.
- In order to make it more suitable for real-time systems like mobile communication systems, process control systems etc, one more item is included – that is '[Algorithm](#)'
- 6 software characteristics for each module or subsystem:
 - A1 = Number of inputs to the application (*I*)
 - A2 = Number of outputs (*O*)
 - A3 = Number of user inquiries (*Q*)
 - A4 = Number of files used (*F*)
 - A5 = Number of external interfaces (*X*)
 - A6 = Item count for algorithms (*A*)

53
Thank You.

Any Questions?

54
BITS Pilani, Pilani Campus

55
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Text Book / References



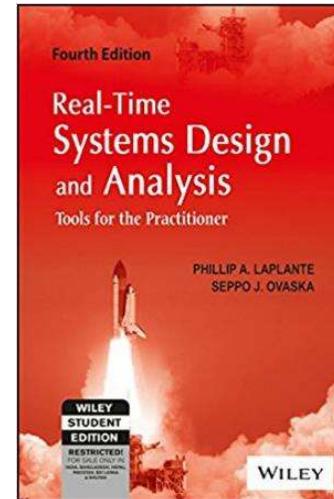
BITS Pilani
Pilani Campus

BITS ZG553: Real Time Systems

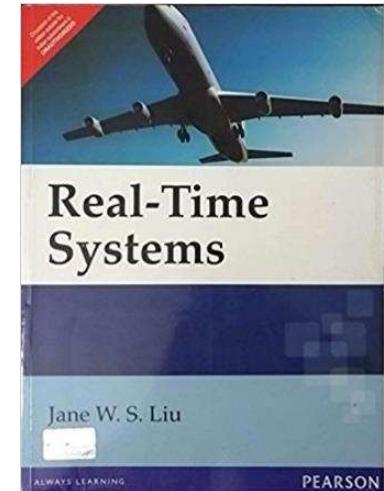
L13 – Requirements Engineering for Real Time Systems

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

Reference (R1)



Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Excellent MOOCs Videos

(Coursera, edX,...)



HONOR CODE
CERTIFICATE



This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

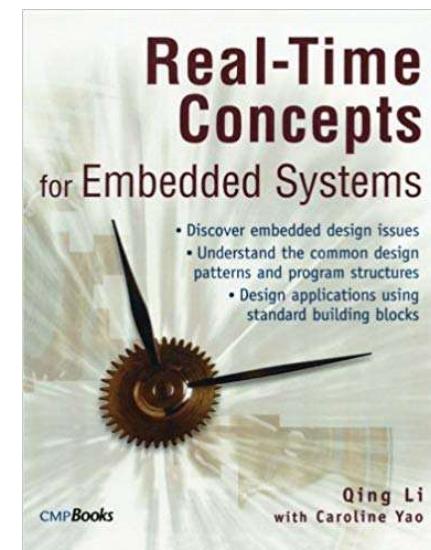
RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44a218f5a1eebf47d4e54

RTS Primer – For Light Reading





L-13: Requirements Engineering for RTS

[Ref: Notes/PPT]

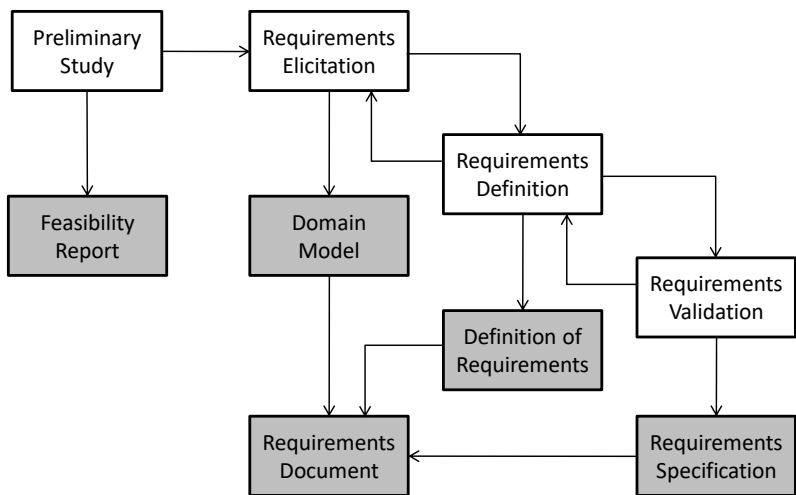
Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

Requirement Engineering Process



Adapted from Sommerville (2000)

Introduction

- Requirements engineering is the subdiscipline of software engineering
- Concerned with determining the goals, functions, and constraints of software systems.
- The goal is to create a requirements specification that is complete, correct, and understandable to both customers and developers.
- While the development of the solution is considered increasingly as commoditized activities (if requirements are made clear), hence can be outsourced, requirement engineering is a crucial activity and therefore should be conducted by the development organisation – together with an appropriate group of customer representatives.

Types of Requirements



As per IEEE, a general software requirement for a real-time software are of following types:

- C1. **Functional** : Fundamental actions of features
- C2. External Interfaces: Inputs and Outputs
- C3. Performance : Static and Dynamic numerical requirements
- C4. Logical database: Logical requirements related to database
- C5. Design constraints: Restrictions
- C6. Software system attributes: Various quantifiable attributes

Types C2-C6 are categorized under '**Non-Functional**' Requirements.

Functional Requirements

- All system inputs
- Exact sequence of operations and responses (outputs) to normal and abnormal situations for every input possibility
- May use case-by-case description or other general form of description (e.g. using universal quantification, use cases, user stories)

External Interface Requirements

- name of item
- description of purpose
- source of input or destination of output
- valid range, accuracy, and/or tolerance
- units of measure
- timing
- relationships to other inputs/outputs
- screen formats/organization
- window formats/organization
- data formats
- command formats

Performance Requirements

- Static and dynamic requirements placed on the software or on human interaction with the software as a whole .
- Might include:
 - the number of simultaneous users to be supported
 - the numbers of transactions and tasks the amount of data to be processed within certain time periods for both normal and peak workload conditions

...
...

Logical database requirements

- Types of information used by various functions such as
 - Frequency of use
 - Accessing capabilities
 - Data entities and their relationships
 - Integrity constraints
 - Data retention requirements

Design constraint requirements



Related to :

- Standards compliance
- Hardware limitations

Software system attribute requirements



- Reliability
 - Availability
 - Security
 - Maintainability
 - Portability
 - Any other miscellaneous requirements you can think of
- ...

Requirements specification for real-time systems- Approaches



There is no one way to do it. Various ways include:

- Top-down process decomposition or structured analysis
- Object-oriented approaches
- Program description languages (PDL) or pseudo-code
- High-level functional specifications that are not further decomposed
- Ad hoc techniques, including simply natural language and mathematical description, and are always included in virtually every system specification

3 Types of Specification Techniques



- **Formal:** Involve rigorous mathematical or logical modeling
- **Informal:** Can't be translated into rigorous mathematical notations and associated rules (e.g. specifications involving flowcharts where the decision blocks involve natural languages)
- **Semiformal:** Methods which don't fall into the above two categories e.g. models involving UML.

Boundaries between these are sometime unclear.

Formal Methods in System Specification

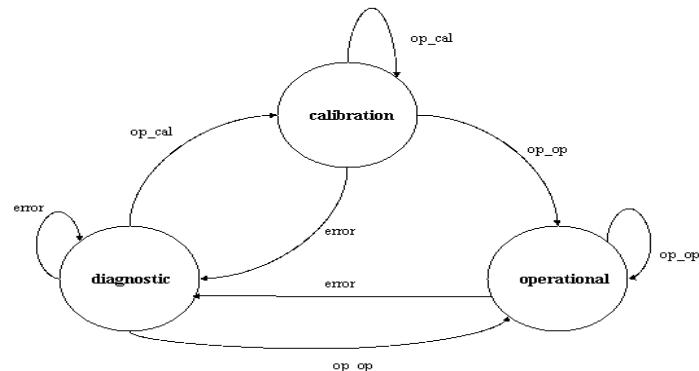


- Contribute significantly to requirements formulation and validation by use of extensive mathematical techniques
- By nature, real-time systems usually contain some formalism of mathematical expression of interaction with the operating environment.
- Three typical use of formal methods
 - **Consistency checking:** System behavioral requirements are described using a mathematically-based notation.
 - **Model checking:** Uses state machines to verify whether a given property is satisfied under all conditions.
 - **Theorem proving:** Axioms of system behavior are used to derive a proof that a system will behave in a given way.

17

BITS Pilani, Pilani Campus

Finite State Machine - Example



$S = \{\text{calibration, diagnostic, operational}\}$, $i = \text{calibration}$, $T = S$ and $\Sigma = \{\text{op_op, op_cal, error}\}$. The transition function δ can be described by a set of triples of the form (state, signal, next_state).

Source: P. A. Laplante, Real-Time Systems Design and Analysis, Wiley, 3rd edition

Finite State Machine (FSM)



A finite state machine can be described mathematically by a five tuple, $M = \{S, i, T, \Sigma, \delta\}$

- S is a finite, non-empty set of states
- i is the initial state (i is a member of S),
- T is the set of terminal states,
- Σ is an alphabet of symbols or events used to mark transitions,
- δ is a transition function that describes the next state of the machine given the current state, and a symbol from the alphabet (an event).

18

BITS Pilani, Pilani Campus

State Charts



Statechart = FSM + Depth + Orthogonality + Broadcast Communication

- FSM is a finite state machine,
- Depth represents levels of detail
- Orthogonality represents the existence of separate tasks
- Broadcast communication is a method for allowing different orthogonal processes to react to the same event.

The statechart resembles a finite state machine where each state can contain its own FSM that describes its behaviour.

Statechart is part of UML.

19

BITS Pilani, Pilani Campus

20

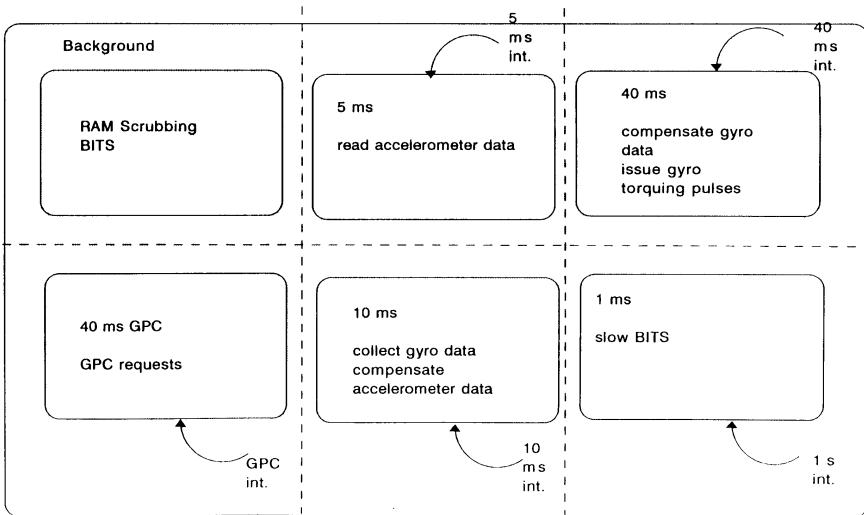
BITS Pilani, Pilani Campus

State Charts

Various components of the statechart:

- The FSM is represented in the usual way, with capital letters or descriptive phrases used to label the states.
- Depth is represented by the insideness of states.
- Broadcast communications are represented by labelled arrows, in the same way as FSMs.
- Orthogonality is represented by dashed lines separating states.
- Symbols a, b, \dots, z represent events that trigger transitions, in the same way that transitions are represented in FSMs.
- Small letters within parentheses represent conditions that must be true for the transitions to occur.

State Chart - Example



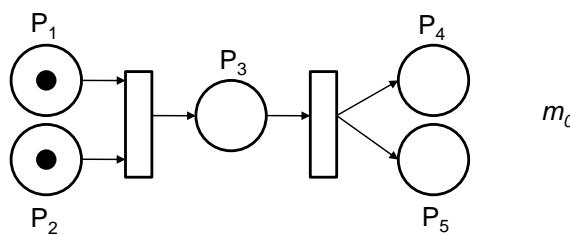
Source: P. A. Laplante & S. J. Ovaska, Real-Time Systems Design and Analysis: Tools for the Practitioner, Wiley, 4th edition

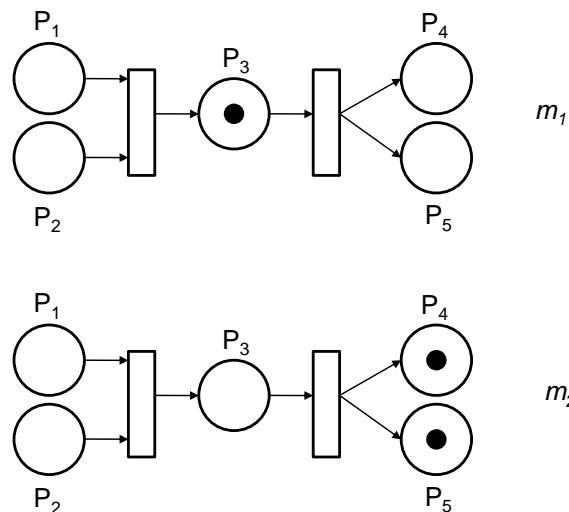
Petri Nets

- First introduced by Carl Adam Petri in 1962.
- A diagrammatic tool to model concurrency and synchronization in distributed systems.
- Based on strong mathematical foundation.
- Consists of three parts:
 - Circles: Called Places (means States)
 - Rectangles: Transitions (or events)
 - Arcs (Arrows): Connects Places to Transitions
- Each Arc is labeled with a number. If there is no label on an arrow, it is considered as '1'.
- Tokens are represented as black dots, which move from one place to other at each transition.
- Change of transition at an event is called **firing of event**.
- **Markings (m_0, m_1, \dots)** represents initial data count (represented by tokens) in all the places
- At each transition, tokens move from one place to other based on the direction of the arrow.

Petri Nets

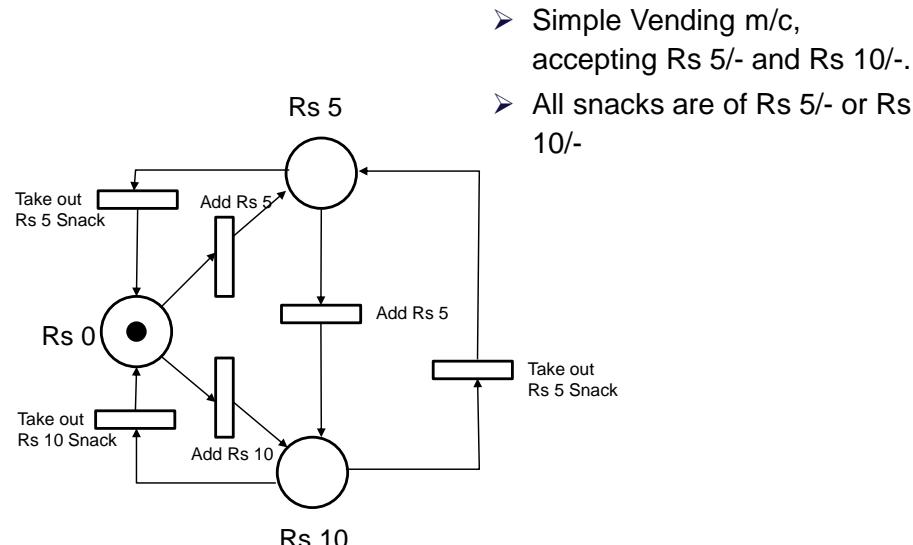
Firing Table					
	P_1	P_2	P_3	P_4	P_5
m_0	1	1	0	0	0
m_1	0	0	1	0	0
m_2	0	0	0	1	1





25 BITS Pilani, Pilani Campus

Petri Nets – Example – Vending Machine



26 BITS Pilani, Pilani Campus

Semi-Formal Methods

Two wide-spread approaches:

- Structured Analysis and Structured Design (SA/SD)
 - Closely associated with procedural programming languages like C
- Unified Modeling Language (UML)
 - Closely associated with object oriented programming languages like C++ and Java

UML is covered extensively in the subject like 'Software for Embedded Systems', 'Object Oriented Analysis and Design' etc.

Structured Analysis and Structured Design (SA/SD)

Three complementary models to describe a real-time system.

- M1. Environmental Model
 - Context Diagram
 - Event List
 - Natural Language
- M2. Behavioural Model
 - Data Flow Diagram
 - Control Flow Diagram
 - Entity Relationship Diagram
 - Data Dictionary
 - Process Specification
 - Control Specification
 - State Transition Diagram
 - Natural Language
- M3. Implementation Model
 - Structure Charts
 - Pseudocode
 - Temporal Logic
 - Natural Language

27 BITS Pilani, Pilani Campus

28 BITS Pilani, Pilani Campus

Structured Analysis and Structured Design (SA/SD)

- The purpose of the **environmental model** is to model the system at **high** level of abstraction.
- Behavioral model embodies the design aspect of SA/SD as a series of data flow diagrams, control flow diagrams, entity-relationship diagrams, process and control specifications, state transition diagrams and data dictionary.
- At the **lowest level of abstraction**, implementation models like **pseudocode** come into picture.

Requirement Validation

The purpose of requirement validation is to ask this question –
“Am I building the right software ?”

Requirement Validation involves checking the followings:

- **Validity** - Does the system provide the functions that best support the customer's needs?
- **Consistency** -Are there any requirements conflicts?
- **Completeness** -Are all functions required by the customer included?
- **Realism** -Can the requirements be implemented given available budget and technology?
- **Verifiability** -Can the requirements be checked?

The Requirement Document

SRS (Software Requirement Specification) can be viewed as a binding contract among the designers, developers, testers and customers.

Recommended table of contents for the SRS as per IEEE Std 830-1998:

1. Introduction
 - 1.1 Purpose
 - 1.2 Scope
 - 1.3 Definitions and Acronyms
 - 1.4 References
 - 1.5 Overview
2. Overall Description
 - 2.1 Product Perspective
 - 2.2 Product Functions
 - 2.3 User Characteristics
 - 2.4 Constraints
 - 2.5 Assumptions and Dependencies
3. Specific Requirements
- Appendices
- Index

Requirement Validation

Many ways of checking the software requirements specification for conformance to IEEE standard's best practices:

1. Automated consistency analysis
2. Checking the consistency of a structured requirements description.
3. Comparing the requirements to those for a similar system
4. Developing tests for requirements to check testability.
5. Prototyping.
6. Requirements reviews.
7. Systematic manual analysis of the requirements.
8. Test-case generation.
9. Using an executable model of the system to check requirements.

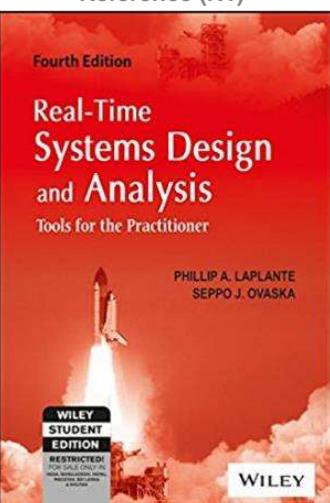
Thank You.

Any Questions?

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

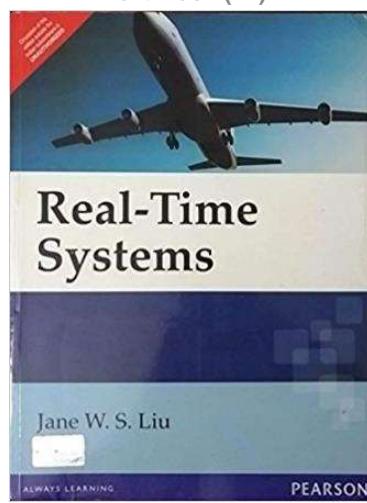
33

Text Book / References



Reference (R1)

Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

**BITS ZG553: Real Time Systems
L14 – Performance Analysis**

K G Krishna
WILP Division, BITS-Pilani, Hyderabad



BITS Pilani
Pilani Campus

Excellent MOOCs Videos
(Coursera, edX,...)



**HONOR CODE
CERTIFICATE**

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54



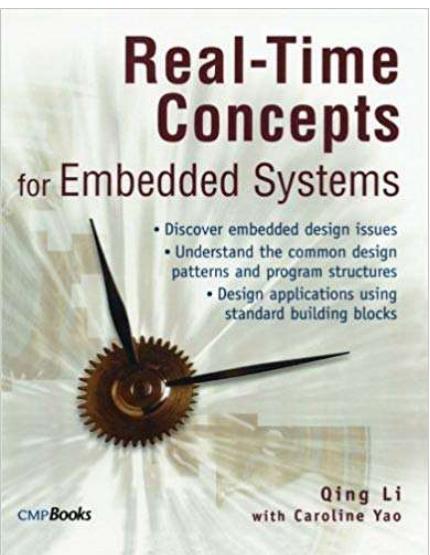
Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

3

RTS Primer – For Light Reading



BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

4

Introduction

- Performance analysis activities occurs in all phases of development.
- In testing phase, it is practical to measure performance in a real working environment.
- But there are needs to analyze the performance at design phase itself.
- Such performance measures are usually predicted and estimated.
- It requires theoretical performance analysis techniques.



L-14: Performance Analysis of RTS

[Ref: Notes/PPT]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

Execution Time Estimation

- Estimation of the execution time e , is the most important task in modeling.
- The most accurate method is using 'Logic Analyzer'. But it can be done, when physical system is ready i.e. at the time of testing and characterization.
- When the whole system is not ready, the execution time can be estimated by examining the compiler output, either manually or through a tool.
- Another possible method is to use system timer
 - Note the system time before start of the execution.
 - Note the system time after end of the execution.
 - The difference in the time is the execution time.
 - Do this exercise for all possible code execution path and take the worst case time.

Analysis of round-robin systems

Assume

Number of tasks in the system = n

Time slice (time quantum) = q

Maximum execution time of each task $T_i = c_i$

Maximum execution time of all tasks $c = \max(c_1, c_2, \dots, c_n)$

Then amount of time each task has to wait in one round = $(n-1)q$

Maximum number of time slices each task require to complete = $\lceil c/q \rceil$

Then amount of time (worst case) each task has to wait in total = $(n-1)q \lceil c/q \rceil$

The worst case time for a task to complete from readiness to completion

$$T = c + (n-1)q \lceil c/q \rceil$$

If there is a context-switching overhead of 'o', then for switching 'n' tasks in a cycle, the context switching overhead will be ' $n.o$ '.

In that case, the worst case time for a task to complete from readiness to completion

$$T = c + [(n-1)q + n.o] \lceil c/q \rceil$$

10 BITS Pilani, Pilani Campus

Fixed Priority System

From time-demand analysis, the worst case response time is given by

$$t = e_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{p_k} \right\rceil e_k, 0 < t \leq p_i$$

But due to the ceiling function, this equation can't be solved in a straight forward method. So let us do it in iterative method. We can rewrite this equation as,

$$t_{n+1} = e_i + \sum_{k=1}^{i-1} \left\lceil \frac{t_n}{p_k} \right\rceil e_k, 0 < t \leq p_i$$

where 'n' is the number of iteration.

Start the recursion with $t_0 = e_i$. The recursion stops, when $t_{n+1} = t_n$

If the recursion doesn't converge, then the processor is overloaded.

Round Robin System - Example

In a round-robin system, there are 6 equally important tasks, each with a maximum execution time of 600 ms and the time quantum is 40 ms, and every context switch costs 2 ms, then find out the response time of each task.

Solution:

Here, $n = 6$, $q = 40$ ms, $c = 600$ ms, $o = 2$ ms.

$$\begin{aligned} \text{So, } T &= c + [(n-1)q + n.o] \lceil c/q \rceil \\ &= 600ms + [(6-1) \times 40 + 6 \times 2] \lceil 600/40 \rceil ms \\ &= 600ms + (200 + 12) \times 15ms \\ &= (600 + 212 \times 15)ms \\ &= (600 + 3180)ms \\ &= 3780ms \end{aligned}$$

11 BITS Pilani, Pilani Campus

Fixed Priority System - Example

Find out the worst case response times of the tasks $T_1 = (9, 3)$, $T_2 = (12, 4)$ and $T_3 = (18, 2)$, when they are scheduled rate-monotonically.

Solution:

As per RM algorithm, the priorities of the tasks are: $T_1 > T_2 > T_3$.

First check the total utilization.

Total utilization $U = 3/9 + 4/12 + 2/18 = 0.72 < 1$.

Hence processor is not overloaded.

Since T_1 is the highest priority task, its response time is equal to its execution time i.e. 3.

Fixed Priority System - Example

For T2,
 $t_{n+1} = 4 + \left\lceil \frac{t_n}{9} \right\rceil \times 3, 0 < t_n \leq 12$

$$t_0 = 4 : t_1 = 4 + \left\lceil \frac{4}{9} \right\rceil \times 3 = 4 + 3 = 7$$

$$t_1 = 7 : t_2 = 4 + \left\lceil \frac{7}{9} \right\rceil \times 3 = 4 + 3 = 7$$

So for T2, worst case response time is 7.

For T3,
 $t_{n+1} = 2 + \left\lceil \frac{t_n}{9} \right\rceil \times 3 + \left\lceil \frac{t_n}{12} \right\rceil \times 4, 0 < t_n \leq 18$

$$t_0 = 2 : t_1 = 2 + \left\lceil \frac{2}{9} \right\rceil \times 3 + \left\lceil \frac{2}{12} \right\rceil \times 4 = 2 + 3 + 4 = 9$$

$$t_1 = 9 : t_2 = 2 + \left\lceil \frac{9}{9} \right\rceil \times 3 + \left\lceil \frac{9}{12} \right\rceil \times 4 = 2 + 3 + 4 = 9$$

So for T3, worst case response time is 9.

Analysis of Non-periodic Systems

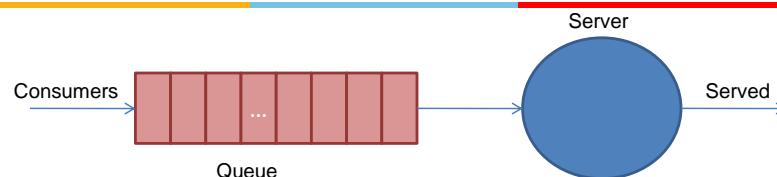
- A real-time system having one or more aperiodic or sporadic job can be modeled (as covered in the previous lectures)
- But practically, they can be modeled rate monotonically as having periods equal to worst case inter-arrival times.
- Such rough approximation leads to unacceptably high utilization.
- So some heuristic approach should be used instead.
- Queuing theory could also be helpful.

Sources of indeterministic behaviour

- Interrupt Latency
- DMA
- Caches
- Instruction Pipelines
- Floating point instructions

All of the above are to improve the CPU performance, but they destroy the determinism. This leads to go for a probabilistic performance model.

Single Server Queue Model



- Known as M/M/1 Queue
 - First M: Probability distribution of the time arrival of consumers : [Here interrupt requests](#)
 - Second M: Probability distribution of the time needed to service each customer: [Here interrupt servicing](#)
 - Third 1: Number of servers
 - The probability distribution of arrival is [Poisson Distribution](#) (inter-arrival times are exponentially distributed)
 - The service or processing time is also exponentially distributed.

Single Server Queue Model

Let

$1/\lambda$ = Mean inter-arrival time of the customers (interrupt requests)

$1/\mu$ = Mean service (or interrupt processing) time

$1/\lambda > 1/\mu$

N = Number of customers (interrupt requests) in the queue

Then, the expected number (mean) of customers (interrupt requests or jobs) in the queue,

$$\bar{N} = \frac{\rho}{1-\rho}, \text{ where } \rho = \frac{\lambda}{\mu}$$

The corresponding variance,

$$\sigma_N^2 = \frac{\rho}{(1-\rho)^2}$$

Example 1

In an interrupt driven system, the mean processing time of the interrupt is 5 ms. What should be the mean inter-arrival time of the interrupts, so that there is 98% confident that CPU will not be overloaded?

Solution:

Mean interrupt processing time, $1/\mu = 5$ ms.

Probability of overload = $100\% - 98\% = 2\%$

$$\Rightarrow \Pr[\geq 2] = \rho^2 \leq 0.02$$

$$\Rightarrow \rho \leq \sqrt{0.02} = 0.141$$

$$\Rightarrow \frac{\lambda}{\mu} \leq 0.141$$

$$\Rightarrow \frac{\mu}{\lambda} \geq \frac{1}{0.141}$$

$$\Rightarrow \frac{1}{\lambda} \geq \frac{1}{0.141} \times \frac{1}{\mu} = \frac{5}{0.141} \text{ ms} = 35.46 \text{ ms}$$

Hence mean interrupt inter-arrival time should be more than 35.46 ms.

Single Server Queue Model

The mean time a customer (job) spends in the system,

$$\bar{T} = \frac{1/\mu}{1-\rho}$$

If Y is a random variable for the time spent in the system, which has an exponential probability distribution,

$$s(y) = \mu(1-\rho)e^{-\mu(1-\rho)y}$$

Probability that at least k customers (jobs) are in the queue simultaneously,

$$\Pr[\geq k] = \rho^k$$

When there are more than one customer waiting in the system, it is a **time-overloaded condition**. From this equation, it is obvious that the probability of exceeding a certain number of customers (jobs) in the system **decreases geometrically (since $\rho \leq 1$)**. So it is pointless to determine the probability when there are more than two customer in the systems, since it will be negligible. Hence for determining the probability of time overloading, **it is sufficient to consider the probability of two customers (jobs) in the system i.e. $\Pr[2]$** .

Example 2

In an interrupt driven system, the mean inter-arrival time of the interrupt is 10 ms. What should be the mean processing time of each interrupt, so that there is 98% confident that CPU will not be overloaded?

Solution:

Mean inter-arrival time, $1/\lambda = 10$ ms.

Probability of overload = $100\% - 98\% = 2\%$

$$\Rightarrow \Pr[\geq 2] = \rho^2 \leq 0.02$$

$$\Rightarrow \rho \leq \sqrt{0.02} = 0.141$$

$$\Rightarrow \frac{\lambda}{\mu} \leq 0.141$$

$$\Rightarrow \frac{1}{\mu} \leq \frac{0.141}{\lambda} \text{ ms} = 0.141 \times 10 \text{ ms} = 1.41 \text{ ms}$$

Hence mean interrupt processing time should be less than 1.41 ms.

Example 3

A producer task is known to process data at rate that is exponentially distributed with average service time of 3 ms per task. What is the maximum allowable average data rate if the probability of collision is to be no more than 1%. Assume that the data arrive at intervals that are exponentially distributed.

Solution:

Mean task processing time, $1/\mu = 3$ ms.

Probability of collision(overload) = 1%

$$\begin{aligned} \Rightarrow \Pr[\geq 2] &= \rho^2 \leq 0.01 \\ \Rightarrow \rho &\leq \sqrt{0.01} = 0.1 \\ \Rightarrow \frac{\lambda}{\mu} &\leq 0.1 \\ \Rightarrow \frac{\mu}{\lambda} &\geq \frac{1}{0.1} \\ \Rightarrow \frac{1}{\lambda} &\geq \frac{1}{0.1} \times \frac{1}{\mu} = \frac{3}{0.1} \text{ ms} = 30 \text{ ms} \end{aligned}$$

Hence maximum allowable data rate = $1 / (30 \text{ ms}) = 1000 / 30 \text{ per sec} = 33.33 \text{ per sec.}$

Analysis of Memory Requirements

A typical memory map constitutes

- Program
- Stack
- Data
- Parameters

Total Memory Utilization,

$$M_T = M_{PG} \cdot P_{PG} + M_{ST} \cdot P_{ST} + M_{DT} \cdot P_{DT} + M_{PM} \cdot P_{PM}$$

$$M_A = \frac{U_A}{T_A}, A \in \{T, PG, ST, DT, PM\}$$

Here M_{PG} , M_{ST} , M_{DT} , M_{PM} represents **memory utilization** of program memory, stack memory, data memory and parameter memory respectively. Similarly P_{PG} , P_{ST} , P_{DT} , P_{PM} represents the **fraction of the total memory used** for these types of memories respectively.

Example 4

In an interrupt driven system, the mean inter-arrival time of the interrupt is 5 ms and mean interrupt processing time is 3 ms. What should be the mean response time of each interrupt? What should be the mean and variance of the buffer required to store the incoming interrupt requests?

Solution:

Mean inter-arrival time, $1/\lambda = 5$ ms.

Mean interrupt processing time, $1/\mu = 3$ ms.

Average response time,

$$\bar{T} = \frac{1/\mu}{1-\rho} = \frac{1/\mu}{1-\lambda/\mu} = \frac{3}{1-3/5} \text{ ms} = 3/0.4 \text{ ms} = 7.5 \text{ ms}$$

Average length of the buffer required to store the interrupt requests,

$$\bar{N} = \frac{\rho}{1-\rho} = \frac{\lambda/\mu}{1-\lambda/\mu} = \frac{3/5}{1-3/5} = 0.6/0.4 = 1.5$$

Variance of the buffer length,

$$\sigma_N = \sqrt{\frac{\rho}{(1-\rho)^2}} = \sqrt{\frac{\lambda/\mu}{(1-\lambda/\mu)^2}} = \sqrt{\frac{3/5}{(1-3/5)^2}} = \sqrt{0.6/(0.4)^2} = \sqrt{0.6/0.16} = \sqrt{3.75} = 1.936$$

Example - Memory Requirements

A system has 64 M bytes of program memory, 16 M bytes of data memory and 8 M bytes of stack area. A program requires 75% of program memory, 25% of data memory and 50% of stack area, which are the worst case scenario. The system doesn't have separate parameter memory area. Calculate the memory utilization of the program.

Solution:

$$\text{Total Memory} = 64 + 16 + 8 \text{ M bytes} = 88 \text{ M bytes}$$

Total memory utilization,

$$\begin{aligned} M_T &= 0.75 \times \frac{64}{88} + 0.25 \times \frac{16}{88} + 0.50 \times \frac{8}{88} \\ &= 0.5454 + 0.0454 + 0.0454 = 63.63\% \end{aligned}$$

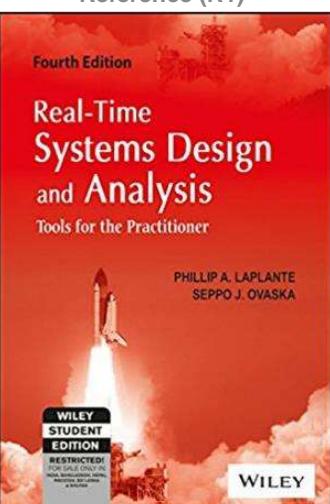
Thank You.

Any Questions?

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

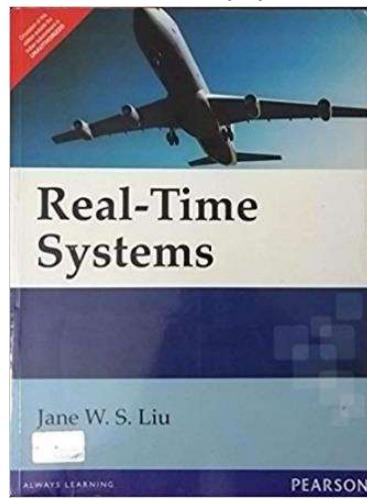
26

Text Book / References



Reference (R1)

Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

BITS ZG553: Real Time Systems
L15 – Verification & Validation Techniques

K G Krishna
WILP Division, BITS-Pilani, Hyderabad

HONOR CODE CERTIFICATE

This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIx: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54

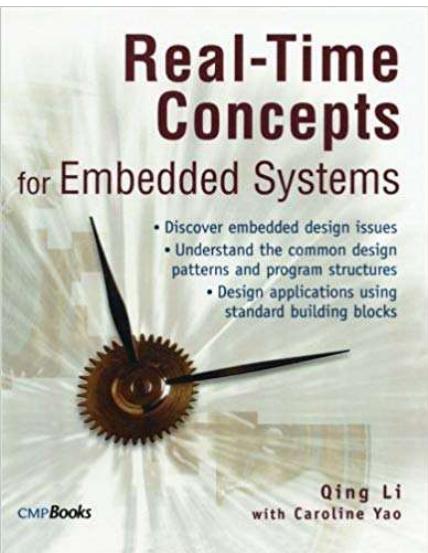


Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University
Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

3

RTS Primer – For Light Reading

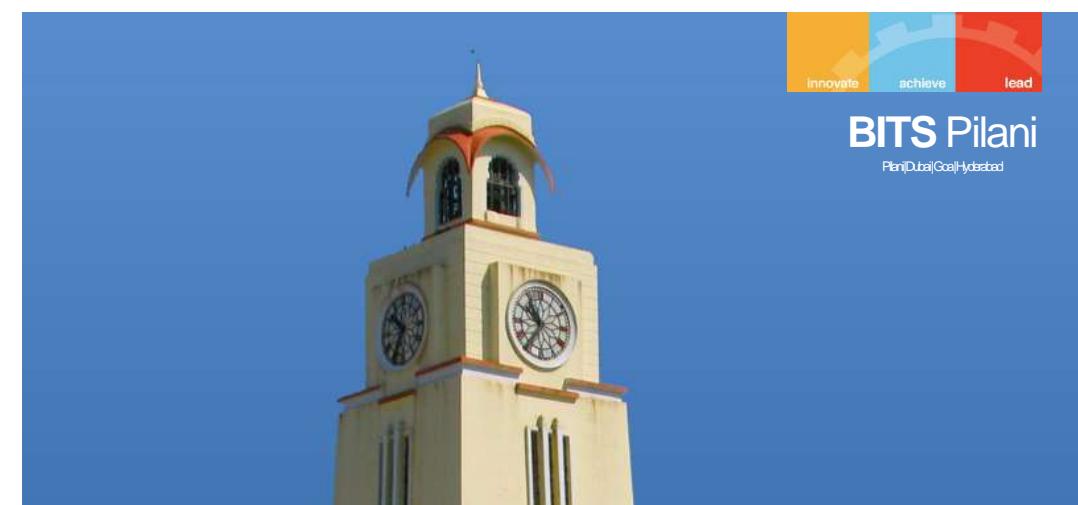


BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

4

The role of testing

- Testing is the process of exercising a program with the specific intent of finding errors prior to delivery to the end user.
- A good test case is one that has a high probability of finding an error.
- A successful test is one that uncovers an error.
- The purpose of testing is also to address various aspects of qualities like
 - Errors
 - Performance
 - Reliability etc



L-15: Verification & Validation Techniques

[Ref: R1/C7 & C8]

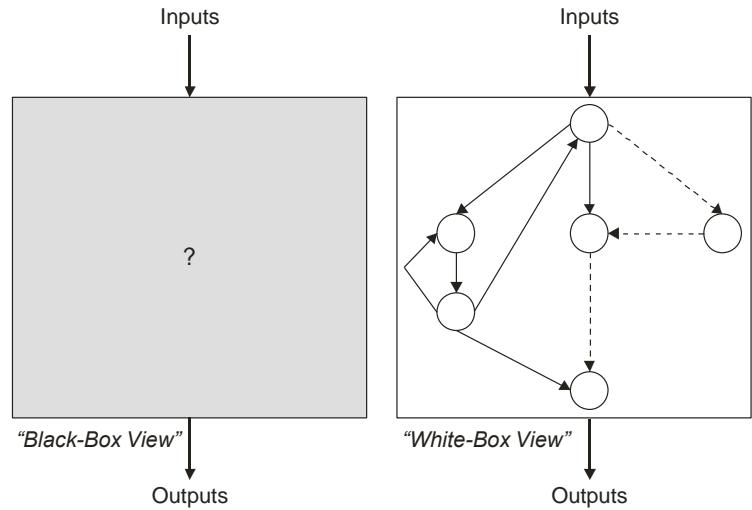
Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

5

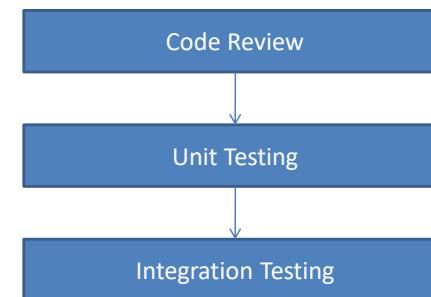
Testing real-time software



Testing principles

- All tests should be traceable to requirements
- Tests should be planned long before testing begins
- Testing should begin “in the small” and progress toward testing “in the large”
- Test to also improve quality, not just demonstrate correctness

Quality Assurance Activity Sequence



Code Review

- The purpose of code review is to
 - Find bugs
 - Suggest code restructuring for better maintainability, readability etc
 - Improve performance
 -
- It is performed before Integration / Testing
- More eyes may catch more issues
- Input to Code Review
 - Coding Guidelines
 - Willingness to participate in code review
 - Incorporating time for code review in the plan

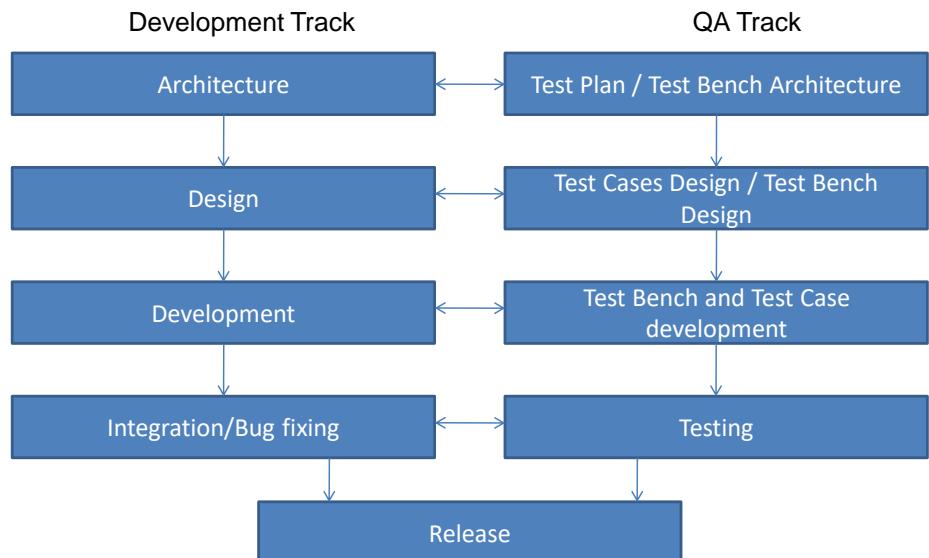
Unit level testing

- Several methods can be used to test individual modules or units
- Include black and white box techniques
- Black box techniques include
 - exhaustive testing
 - boundary value testing (boundary value analysis)
 - random test generation
 - worst case testing
- White box techniques include
 - judicious test selection
 - formal program proving
 - code inspections

System-level testing

- Top-down testing
 - Start with high-level system and integrate from the top-down replacing individual components by stubs where appropriate
- Bottom-up testing
 - Integrate individual components in levels until the complete system is created
- In practice, most integration involves a combination of these strategies

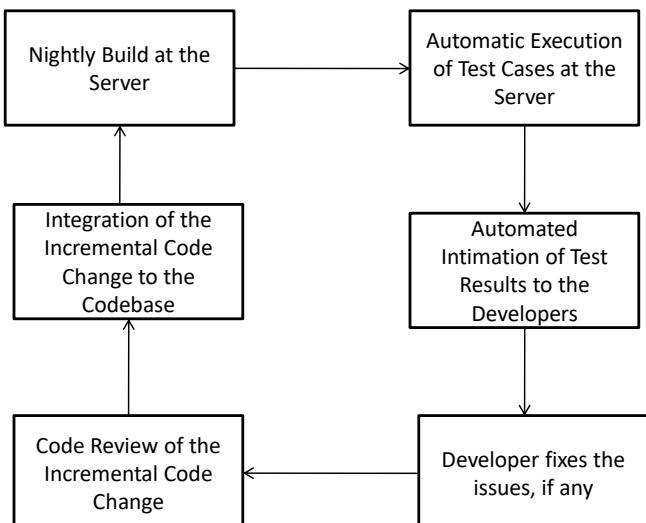
Development and QA Activities should be in parallel



12 BITS Pilani, Pilani Campus

13

Continuous integration



14 BITS Pilani, Pilani Campus

Thank You.

Any Questions?



15

Text Book / References



BITS Pilani
Pilani Campus

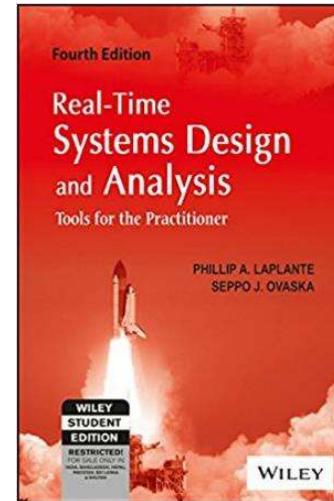
BITS ZG553: Real Time Systems

L16 – Uncertainties in RTS & Fault-Tolerance Techniques

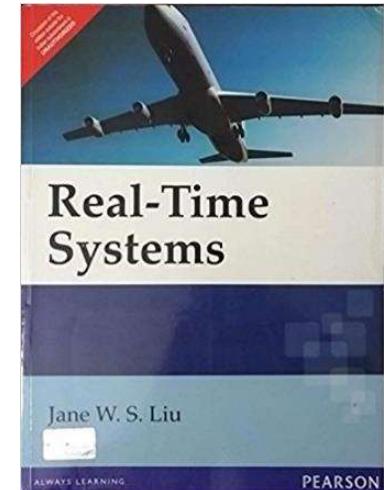
K G Krishna
WILP Division, BITS-Pilani, Hyderabad



Reference (R1)



Text Book (T1)



Note: As the above two books focus on theoretical treatment of the subject, Students are strongly advised to refer to web sources / MOOCs videos / library within their own organizations for more practical understanding of the topics.

2
BITS Pilani, Deemed to be University under Section 3 of UGC Act, 1956

Excellent MOOCs Videos

(Coursera, edX,...)



HONOR CODE
CERTIFICATE



This is to certify that

GOPALA KRISHNA KONERU

successfully completed and received a passing grade in

RTSIX: Introduction to Real-Time Systems

a course of study offered by IEEEEx, an online learning initiative of Institute of Electrical and Electronics Engineers through edX.

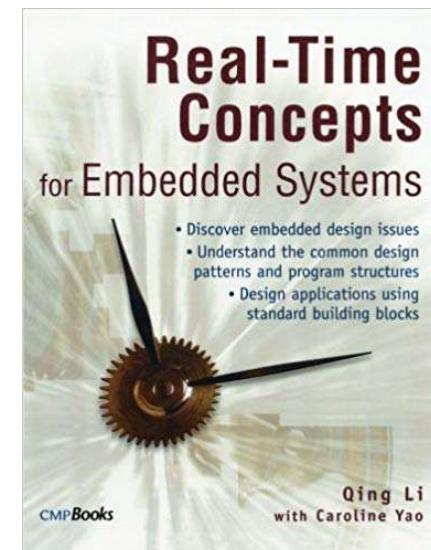
HONOR CODE CERTIFICATE
Issued November 28, 2015

VALID CERTIFICATE ID
06ae78da4df44e218f5a1eebf47d4e54

Dr. Philip Laplante
Professor of Software and Systems Engineering
The Pennsylvania State University

Dr. Saurabh Sinha
Vice President
IEEE Educational Activities Board

RTS Primer – For Light Reading





L-16: Uncertainties in RTS & Fault-Tolerance Techniques

[Ref: Notes/PPT]

Note: Students are requested to NOT to rely on PPTs/Recorded sessions as their only source of knowledge, explore sources within your own organization or web for any specific topic; attend classes regularly and involve in discussions;

PLEASE DO NOT PRINT PPTs. Save the Environment!

Source PPT Courtesy: Some of the contents of this PPT is sourced from Presentations of Prof K R Anupa / Prof B Mishra, BITS-Pilani WILP Division

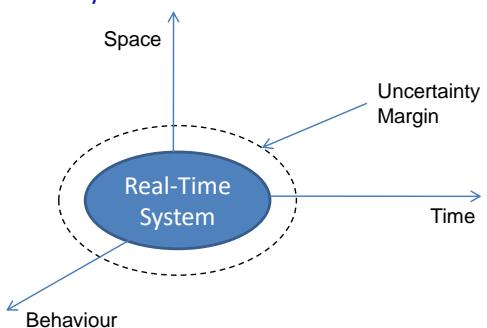
5

Introduction

Heisenberg Uncertainty Principle:

The precise position and momentum of a particle can't be known precisely; trying to be more certain about one comes at the expense of increased uncertainty of others.

Similarly, in case of real-time systems, trying to bring more certainty in one dimension, bring uncertainty to other dimensions.



Introduction

- Focus of real-time system design is shifting from meeting mere performance goals to *designing for uncertainty*
- Uncertainty exists in embedded systems along 3 dimensions – Time, Space and Behaviour.
- **Time:**
 - Main of the challenge in design of embedded systems is unpredictability of response time.
 - RM Algorithm produces a deterministic schedule only when sporadic tasks, aperiodic tasks, mutual exclusion and resource contention are excluded.
- **Space:**
 - Space dimension deals with the limitation posed by physical resources like memory.
- **Behaviour**
 - Component behaviour, environmental behaviour (such as inputs, outputs), state, uncertainty of requirements, uncertainty in testing etc.

Sources of Uncertainties

Kind of Uncertainties	Signs	Possible Causes	Possible Solutions
System under control	Bizarre inputs; Bizarre outputs;	Nature of application; Faulty hardware	Use aggressive fault tolerant design
Operating environment	Bizarre inputs; Bizarre outputs;	SEUs, sand, salt, shock, etc.	Use aggressive fault tolerant design
Requirements	Sparse requirements; Numerous "TBDs"	Incomplete requirements; Inconsistent requirements	Goals based requirements analysis; formal consistency checking
Testing	System that passes testing fails in the field	Poor testing regimen; Incomplete coverage	Improve testing process
Input	Strange behavior; Explicating comments	Unstable input sources; Defective hardware	Kalman filters, data fusion, median filters or averaging, roll back and recovery blocks.
Output	Strange behavior; Explicating comments	Defective hardware; Corrupted inputs from SUC	Kalman filters, data fusion, median filters or averaging, roll back and recovery blocks.
State	Strange behavior; Explicating comments	Program code jumping; Pointer errors	Model checking; Software black boxes

Sources of Uncertainties

Kind of Uncertainties	Signs	Possible Causes	Possible Solutions
Timing and schedulability	Dubious constraints; Missed deadlines; Explicating comments	Delays as loops; Speculative generality; Failed off-the-shelf components	Use OS provided timer facilities; Model checking; Fault injection
Programming Language & Compiler	Explicating comments;	Compiler induced errors	Test the compiler; Improve coding techniques
Off-the-shelf components	Missed deadlines; Inexplicable failure;	Badly tested software or hardware; Falsely advertised claims	Fault injection; Software black boxes
Execution Time	Missed deadlines	Gradual build up of various truncation and round-off errors	Rejuvenation (stopping and restarting the system regularly – should be used cautiously)

Fault Tolerance Techniques

Faults and Their Sources

Fault is an erroneous state of software or hardware resulting from failures of its components

Fault Sources

- Mechanical --“wears out”
 - Deterioration: wear, fatigue, corrosion
 - Shock: fractures, overload, etc.
- Electronic Hardware --“bad fabrication; wears out”
 - Latent manufacturing defects
 - Operating environment: noise, heat, ESD, electro-migration
 - Design defects (Pentium F-DIV bug)
- Software --“bad design”
 - Design defects
 - “Code rot” --accumulated run-time faults
- System Misuse – “people”

Faults Classifications

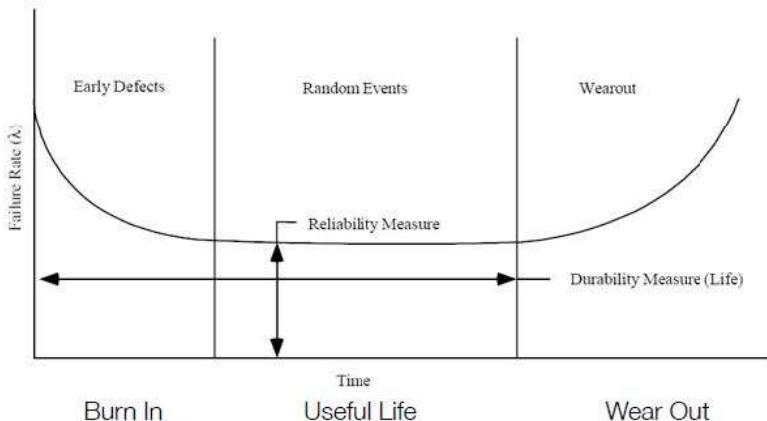
- Failure:
 - Component does not provide service
- Fault:
 - A defect within a system
- Error:
 - A deviation from the required operation of the system or subsystem

Reliability

- A measure of whether a user can depend on the system.
 - The system “stands the test of time.”
 - There is an absence of known catastrophic errors; that is, errors that render the system useless.
 - The system recovers “gracefully” from errors.
 - The software is robust.
- For real-time systems, other informal characterizations of reliability might include
 - Downtime is below a certain threshold.
 - The accuracy of the system is within a certain tolerance.
 - Real-time performance requirements are met consistently.

Traditionally, MTFF (Mean Time To First Failure) and MTBF (Mean Time Between Failures) are used as a measure of Software Reliability.

Component Reliability Model



Reliability – Statistical Perspective

Let S = The system

T = Time instance of fault

Then Reliability of S at any time instance t , is defined by

$$R(t) = P(T > t) \text{ i.e. Probability of not having any fault till time } t.$$

$R(t) = 1$ is the ideal case. Practically, $r(t)$ is less than 1.

Example:

In the monitoring system of a nuclear power plant, the failure probability is no more than 10^{-9} per hour.

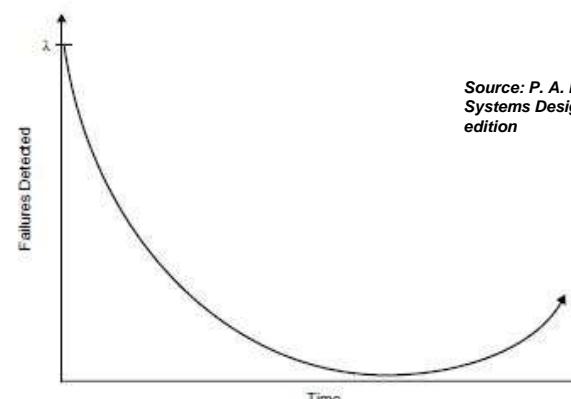
So, $R(t) = (1 - 10^{-9})^t$, t is in hours.

So, as $t \rightarrow \infty$, $r(t) \rightarrow 0$.

Software Reliability – Statistical Perspective



Failure function tends to become a *bathtub curve* i.e. failures reduces initially at an exponential rate, but increases towards the end.



Source: P. A. Laplante, Real-Time Systems Design and Analysis, Wiley, 3rd edition

Software Reliability – Statistical Perspective



Reasons for Bathtub curve for failures:

- Hardware can wear and tear. So after some point of time, the number of failures increases.
- But software can't wear and tear. Then why bathtub curve for the software ?
 - During patch releases, making quick correction without designing them properly.
 - During patch releases, ignoring the effect of the new bug fix on other modules (thereby fixing one issue and introducing more new issues)
 - Inadequate testing of the software during patch releases.
 - Wear and tear of underlying hardware.

17
BITS Pilani, Pilani Campus

What is Fault Tolerance ?



Fault Tolerance in real-time system is the tendency to continue functioning in the presence of hardware or software failures.

Two Types:

- **Spatial Fault Tolerance:** Includes methods involving redundant hardware and/or software solutions
- **Temporal Fault Tolerance:** Involves miscellaneous techniques that allow for tolerating missed deadlines

19
BITS Pilani, Pilani Campus

Serial and Parallel System Reliability



Serial Component Probability

Let there are n serially connected components, with reliability of $R_k(t)$ for kth component

Then total system reliability, $R(t) = \prod_{i=1}^n R_i(t)$

Parallel Component Probability

If there n parallel connected components , with reliability of $R_k(t)$ for kth component

Then total failure probability, $Q(t) = \prod_{i=1}^n Q_i(t) = \prod_{i=1}^n (1 - R_i(t))$

So total reliability, $R(t) = 1 - Q(t) = 1 - \prod_{i=1}^n (1 - R_i(t))$

So, **reliability increases with parallel components**

18
BITS Pilani, Pilani Campus

Methods to increase fault tolerance

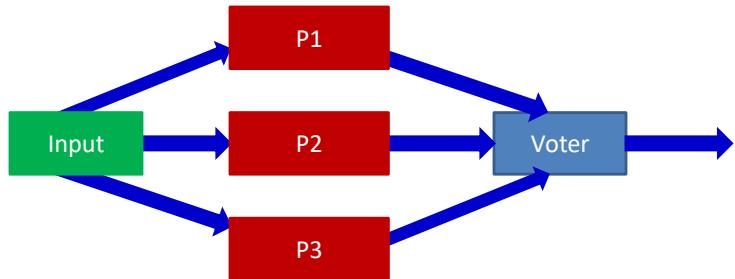


- Voting schemes
- Software Check Points
- Software Black Boxes
- Built-in-Test Software
- CRC check
- Hamming code error detection and correction
- Handling Spurious and Missed Interrupts

20
BITS Pilani, Pilani Campus

Voting Schemes

- Combine input from two or more process blocks (hardware or software) and combine in some manner
- Voting schemes use majority rule or average of block outputs to obtain final values



- ❖ P1, P2 and P3 processors execute different versions of the code for the same application (usually developed by different teams).
- ❖ Voter compares the results and forward the majority vote of results (two out of three).

21 BITS Pilani, Pilani Campus

Software Check Points

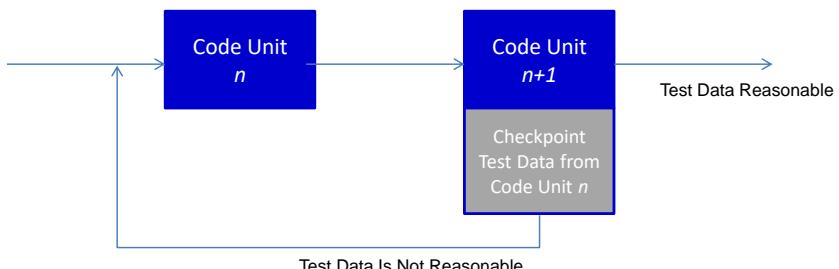
- Intermediate results are dumped at fixed locations for diagnostic purpose
- Can be used during verification or regular operation
- If used during testing only then this additional code is called a **test probe**

22

BITS Pilani, Pilani Campus

Recovery Blocks

- Fault tolerance is increased by using checkpoints with predetermined reset points in the software.
- These reset points mark **recovery blocks**
- At the end of each recovery block, the corresponding check point is tested.
- If the results are not reasonable, then the processing resumes with a prior recovery block.



BITs Pilani, Pilani Campus

N-Version Programming

- Used to increase software fault tolerance
- Generally redundant processors are used.
- Same software developed by different teams using different algorithms runs on these redundant processors.
- In case one such software locks-up, then the system can ignore this software and go ahead with the other software which is functioning correctly on different processors.
- This is called N-Version programming.
- Generally this method is used mission-critical systems like the space shuttle's general purpose computer.

24

BITs Pilani, Pilani Campus

Built-in-Test Software (BITS)



Can increase fault-tolerance in presence of hardware errors

- CPU Tests
- Memory Tests (ROM and RAM)
- Other Devices

25
BITS Pilani, Pilani Campus

BITS – Memory test



- All kind of memory can be corrupted via electrostatic discharge, power surging, vibration or other physical means
- Corruption of memory devices by randomly encountered charge particles is a particular problem in space.

27
BITS Pilani, Pilani Campus

BITS – CPU test

- Perform carefully constructed set of tests
- Must be performed with interrupts disabled
- Too slow to perform in foreground

26
BITS Pilani, Pilani Campus

ROM check – using CRC



- CRC – Cyclic Redundancy Code
- CRC value of each memory is stored and checked when data from the memory is used
- CRC Calculation
 - Treats memory as a long binary polynomial
 - Divide polynomial by generator polynomial
 - $X^{16}+X^{12}+X^5+1$ (CCITT)
 - $X^{16}+X^{15}+X^2+1$ (CRC-16)
 - Remainder is CRC

28
BITS Pilani, Pilani Campus

3-bits CRC Example

Generator polynomial: x^3+x+1

$$= 1x^3+0x^2+1x+1$$

So generator polynomial is represented as 1011

Let the input polynomial is 11010011101100.

Pad the input polynomial by 3 bits (3 bits CRC) : 11010011101100000

Divide 11010011101100000 by 1011:

Essentially x-or the 4 bits of input polynomial with 4 bits of generator polynomial starting from left bit. Then move the generator polynomial to the next 1 of the result.

3-bits CRC Example (contd.)

11010011101100 000 <--- input right padded by 3 bits

1011 <--- divisor

01100011101100 000 <--- result (note the first four bits are the XOR with the divisor beneath, the rest of the bits are unchanged)

1011 <--- divisor ...

00111011101100 000

1011

00010111101100 000

1011

00000001101100 000 <--- note that the divisor moves over to align with the next 1 in the dividend (since quotient for that step was zero)

(in other words, it doesn't necessarily move one bit per iteration)

1011

00000000110100 000

1011

00000000011000 000

1011

00000000001110 000

1011

00000000000101 000

101 1

0000000000000000 100 <--- remainder (3 bits). Division algorithm stops here as quotient is equal to zero.

So, CRC is 100

RAM Check – using EDC

- CRC is not suitable for RAM, since RAM content is volatile
- A Hardware circuit implementing some **Hamming code error detection and correction (EDC)** is used
- During normal data access, the data must pass through the EDC hardware before into or out of the memory
- The hardware compares the data against the check bits and makes necessary corrections if necessary.

Spurious and Missed Interrupts

Spurious Interrupts

- Spurious (or *phantom*) Interrupts are the extraneous and unwanted interrupts.
- Caused by noisy hardware, power surges, electrostatic discharges etc.
- They can destroy algorithmic integrity, stack overflows, compromise the deadlines etc.
- Such interrupt can be tolerated by
 - Having redundant interrupt hardware
 - The device issuing the interrupt can issue a redundant check like
 - Using DMA to send a confirming flag.
 - Then the device receiving the interrupt can check if this flag upon receiving it.

Missed Interrupts

- Missed interrupts are more difficult to deal with.
- They can be tolerated by robust algorithms

Watchdog Timer

- A **watchdog timer** is used to automatically detect software anomalies and reset the processor if any such anomaly occur.
- It is a timing device such that it is set for a preset time interval and an event must occur during that interval else the device will generate the timeout signal on failure to get that event in the watched time interval
- Watchdog timer can be
 - A Hardware
 - A Software task
- A simple watchdog timer example:
 - Watchdog timer is based on a counter that counts down from some initial value to zero.
 - The embedded software selects the counter's initial value and periodically restarts it.
 - If the counter ever reaches zero before the software restarts it, the software is presumed to be malfunctioning and the processor's reset signal is asserted.
- Another Example:
 - Assume that we anticipate that a set of tasks must finish in 100 ms interval.
 - The watchdog timer is disabled and stopped by the program instruction in case the tasks finish within 100 ms interval.
 - In case task does not finish (not disabled by the program instruction), watchdog timer generates interrupts after 100 ms and executes a routine, which is programmed to run because there is failure of finishing the task in anticipated interval.

Thank You.

Any Questions?